



Open Core Protocol Specification

Release 2.2

OCP-IP Confidential



Open Core Protocol Specification

Document Revision - 1.0

© 2006 OCP-IP Association, All Rights Reserved.

Open Core Protocol Specification 2.2
Document Revision 1.0

This document, including all software described in it, is furnished under the terms of the Open Core Protocol Specification License Agreement (the "License") and may only be used or copied in accordance with the terms of the License. The information in this document is a work in progress, jointly developed by the members of OCP-IP Association ("OCP-IP") and is furnished for informational use only.

The technology disclosed herein may be protected by one or more patents, copyrights, trademarks and/or trade secrets owned by or licensed to OCP-IP. OCP-IP reserves all rights with respect to such technology and related materials. Any use of the protected technology and related material beyond the terms of the License without the prior written consent of OCP-IP is prohibited.

This document contains material that is confidential to OCP-IP and its members and licensors. The user should assume that all materials contained and/or referenced in this document are confidential and proprietary unless otherwise indicated or apparent from the nature of such materials (for example, references to publicly available forms or documents). Disclosure or use of this document or any material contained herein, other than as expressly permitted, is prohibited without the prior written consent of OCP-IP or such other party that may grant permission to use its proprietary material.

The trademarks, logos, and service marks displayed in this document are the registered and unregistered trademarks of OCP-IP, its members and its licensors. The following trademarks of Sonics, Inc. have been licensed to OCP-IP: FastForward, SonicsIA, CoreCreator, SiliconBackplane, SiliconBackplane Agent, InitiatorAgent Module, TargetAgent Module, ServiceAgent Module, SOCCreator, and Open Core Protocol.

The copyright and trademarks owned by OCP-IP, whether registered or unregistered, may not be used in connection with any product or service that is not owned, approved or distributed by OCP-IP, and may not be used in any manner that is likely to cause customer confusion or that disparages OCP-IP. Nothing contained in this document should be construed as granting by implication, estoppel, or otherwise, any license or right to use any copyright without the express written consent of OCP-IP, its licensors or a third party owner of any such trademark.

Printed in the United States of America.

Part number: 161-000125-0003

EXCEPT AS OTHERWISE EXPRESSLY PROVIDED, THE OPEN CORE PROTOCOL (OCP) SPECIFICATION IS PROVIDED BY OCP-IP TO MEMBERS "AS IS" WITHOUT WARRANTY OF ANY KIND, EXPRESS, IMPLIED OR STATUTORY, INCLUDING BUT NOT LIMITED TO ANY IMPLIED WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT OF THIRD PARTY RIGHTS.

OCP-IP SHALL NOT BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL OR CONSEQUENTIAL DAMAGES OF ANY KIND OR NATURE WHATSOEVER (INCLUDING, WITHOUT LIMITATION, ANY DAMAGES ARISING FROM LOSS OF USE OR LOST BUSINESS, REVENUE, PROFITS, DATA OR GOODWILL) ARISING IN CONNECTION WITH ANY INFRINGEMENT CLAIMS BY THIRD PARTIES OR THE SPECIFICATION, WHETHER IN AN ACTION IN CONTRACT, TORT, STRICT LIABILITY, NEGLIGENCE, OR ANY OTHER THEORY, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

Contents

1 Overview	1
OCP Characteristics	2
Compliance	3
Part I Specification	5
2 Theory of Operation	7
3 Signals and Encoding	13
Dataflow Signals	14
Basic Signals	14
Simple Extensions.	16
Burst Extensions	19
Tag Extensions	22
Thread Extensions.	23
Sideband Signals	25
Reset, Interrupt, Error, and Core-Specific Flag Signals	26
Control and Status Signals	27
Test Signals	27
Scan Interface	28
Clock Control Interface	28
Debug and Test Interface	29
Signal Configuration	29
Signal Directions	33
4 Protocol Semantics	35
Signal Groups	36
Combinational Dependencies	37
Signal Timing and Protocol Phases	38
OCP Clock	38
Dataflow Signals	39
Sideband and Test Signals	44
Transfer Effects	45
Partial Word Transfers	47
Posting Semantics.	47

Endianness	47
Burst Definition	48
Burst Address Sequences	49
Burst Length, Precise and Imprecise Bursts	50
Constant Fields in Bursts	51
Atomicity	51
Single Request / Multiple Data Bursts (Packets)	51
MReqLast, MDataLast, SRespLast	52
MReqRowLast, MDataRowLast, SRespRowLast	52
Tags	53
Ordering Restrictions	53
Threads and Connections	54
OCP Configuration	55
Protocol Options	55
Phase Options	59
Signal Options.	60
Minimum Implementation	60
OCP Interface Interoperability	60
Configuration Parameter Defaults	63
5 Interface Configuration File	69
Lexical Grammar	69
Syntax	70
6 Core RTL Configuration File	75
Syntax	75
Components	76
Sample RTL Configuration File	84
7 Core Timing	87
Timing Parameters	88
Minimum Parameters	88
Hold-time Parameters	88
Technology Variables.	89
Connecting Two OCP Cores	90
Core Synthesis Configuration File	92
Syntax Conventions	92
Version Section	94

Clock Section	94
Area Section.	94
Port Constraints Section.	95
Max Delay Constraints	100
False Path Constraints.	100
Sample Core Synthesis Configuration File	101
Part II Guidelines	103
8 Timing Diagrams	105
Simple Write and Read Transfer	105
Request Handshake	107
Request Handshake and Separate Response	108
Write with Response	109
Non-Posted Write	110
Burst Write	111
Non-Pipelined Read	113
Pipelined Request and Response	114
Response Accept	115
Incrementing Precise Burst Read	116
Incrementing Imprecise Burst Read	118
Wrapping Burst Read	120
Incrementing Burst Read with IDLE Request Cycle	121
Incrementing Burst Read with NULL Response Cycle	123
Single Request Burst Read	124
Datahandshake Extension	126
Burst Write with Combined Request and Data	127
2-Dimensional Block Read	129
Tagged Reads	130
Tagged Bursts	132
Threaded Read	134
Threaded Read with Thread Busy	136
Threaded Read with Thread Busy Exact	137
Threaded Read with Pipelined Thread Busy	138
Reset	140
Reset with Clock Enable	141
Basic Read with Clock Enable	142

9	Developers Guidelines	143
	Basic OCP	143
	Divided Clocks	144
	Signal Timing	145
	State Machine Examples	147
	OCP Subsets.	152
	Simple OCP Extensions	153
	Byte Enables	153
	Multiple Address Spaces	154
	In-Band Information	155
	Burst Extensions	156
	OCP Burst Capabilities	156
	Compatibility with the OCP 1.0 Burst Model	160
	Tags	162
	Threads and Connections	163
	Threads	163
	Connections	168
	OCP Specific Features	170
	Write Semantics.	170
	Lazy Synchronization	171
	OCP and Endianness	174
	Security	175
	Sideband Signals	176
	Reset Handling	176
	Debug and Test Interface	178
	Scan Control.	179
	Clock Control	179
10	Timing Guidelines	181
	Level0 Timing	182
	Level1 Timing	182
	Level2 Timing	182
11	OCP Profiles	185
	Profile Types	187
	Native OCP Profiles	188
	Block Data Flow Profile	188
	Sequential Undefined Length Data Flow Profile	190

Register Access Profile	192
Bridging Profiles	194
Simple H-bus Profile	194
X-Bus Packet Write Profile	196
X-Bus Packet Read Profile	198
Layered Profiles	201
Security	201
12 Core Performance	203
Report Instructions	203
Sample Report	206
Performance Report Template	208
Part III Protocol Compliance	211
13 Compliance	213
Configuration Compliance	213
Interface Configuration	213
Configuration Parameter Extraction	214
Protocol Compliance	214
Select the Relevant Checks	214
Verification Techniques	215
Dynamic Verification	215
Static Verification	216
14 Protocol Compliance Checks	219
Activation Tables	220
Compliance Checks	227
Dataflow Signals Checks	227
DataFlow Phase Checks	230
Dataflow Burst Checks	242
DataFlow Transfer Checks	254
DataFlow ReadEx Checks	256
Sideband Checks	258
15 Configuration Compliance Checks	263
Request Group	264
Datahandshake Group	271
Response Group	276

Sideband Group	279
Test Group	281
Interface Interoperability	281
16 Functional Coverage	289
Signal Level	290
Transfer Level.	293
Transaction Level	293
Sideband Signals Coverage	298
Naming Conventions	298
Index	301

Introduction

The Open Core Protocol™ (OCP™) delivers the only non-proprietary, openly licensed, core-centric protocol that comprehensively describes the system-level integration requirements of intellectual property (IP) cores.

While other bus and component interfaces address only the data flow aspects of core communications, the OCP unifies all inter-core communications, including sideband control and test harness signals. OCP's synchronous unidirectional signaling produces simplified core implementation, integration, and timing analysis.

OCP eliminates the task of repeatedly defining, verifying, documenting and supporting proprietary interface protocols. The OCP readily adapts to support new core capabilities while limiting test suite modifications for core upgrades.

Clearly delineated design boundaries enable cores to be designed independently of other system cores yielding definitive, reusable IP cores with reusable verification and test suites.

Any on-chip interconnect can be interfaced to the OCP rendering it appropriate for many forms of on-chip communications:

- Dedicated peer-to-peer communications, as in many pipelined signal processing applications such as MPEG2 decoding.
- Simple slave-only applications such as slow peripheral interfaces.
- High-performance, latency-sensitive, multi-threaded applications, such as multi-bank DRAM architectures.

The OCP supports very high performance data transfer models ranging from simple request-grants through pipelined and multi-threaded objects. Higher complexity SOC communication models are supported using thread identifiers to manage out-of-order completion of multiple concurrent transfer sequences.

The CoreCreator™ tool automates the tasks of building, simulating, verifying and packaging OCP-compatible cores. IP core products can be fully "componentized" by consolidating core models, timing parameters, synthesis scripts, verification suites and test vectors in accordance with the *OCP Specification*. CoreCreator does not constrain the user to either a specific methodology or design tool.

Support

The *OCP Specification* is maintained by the Open Core Protocol International Partnership (OCP-IP™), a trade organization solely dedicated to OCP, supporting products and services. For all technical support inquiries, please contact techsupport@ocpip.org. For any other information or comments, please contact admin@ocpip.org.

Changes for Revision 2.2

The following changes and additions have been implemented for this release.

- The EnableClk signal and the accompanying enableclk parameter have been added to indicate which rising edges of Clk are to be considered rising edges of the OCP clock. The addition of EnableClk allows the system to control the effective clock frequency of the interface dynamically without introducing extra outputs from PLLs, or requiring additional low-skew clock distribution networks when divided frequency clocks are used.
- The assertion edge of SReset_n and MReset_n may now be asynchronous to the OCP clock, although de-assertion must still be synchronous to the OCP clock.
- Support has been added for two-dimensional block burst sequences by introducing a new encoding for the MBurstSeq signal and a new parameter, burstseq_blk_enable in addition to height (MBlockHeight) and row-to-row address offset (MBlockStride) signals. The MReqRowLast, MDataRowLast, and SRespRowLast signals have been added to enable end of row framing. The new reqrowlast, datarowlast, and resprowlast parameters serve to configure the signals.
- New non-blocking (exact) flow control options have been added that move the *ThreadBusy signal assertion into the cycle before the associated phase assertion. These options facilitate fully synchronous design styles (for example, with no combinational dependencies at the interface), but new combinational dependencies are also allowed. The mthreadbusy_pipelined, sdatathreadbusy_pipelined, and sthreadbusy_pipelined parameters can be used to set the relative protocol timing of the MThreadBusy, SDataThreadBusy, and SThreadBusy signals with respect to their phases. These new flow control pipelining options are only supported for exact flow control (for instance, situations where the associated *threadbusy_exact parameter is set).

- A security profile has been added to describe the consistent mapping of security-related information into MReqInfo to indicate access permissions associated with requests.
- A new part dedicated to verification provides compliance, configuration and functional coverage checks for the OCP interface that can help you create checking solutions in the language and tool of your choice.

Acknowledgements

The following companies were instrumental in the development of the *Open Core Protocol Specification*, Release 2.2.

All OCP-IP Specification Working Group members, including participants from:

- MIPS Technologies, Inc.
- Nokia Technology Platforms
- Sonics, Inc.
- Texas Instruments Incorporated
- Toshiba Corporation
- Yogitech SPA

Other member companies providing thoughtful comments including TransEDA and all additional reviewers who participated in the general member review.

The Functional Verification Working Group with particular thanks to Steven McMaster of Synopsys Inc., Anshul Bhargava of Mips Technologies and Ravi Venugopalan of Sonics Inc.

1 Overview

The Open Core Protocol™ (OCP) defines a high-performance, bus-independent interface between IP cores that reduces design time, design risk, and manufacturing costs for SOC designs.

An IP core can be a simple peripheral core, a high-performance microprocessor, or an on-chip communication subsystem such as a wrapped on-chip bus. The Open Core Protocol:

- Achieves the goal of IP design reuse. The OCP transforms IP cores making them independent of the architecture and design of the systems in which they are used
- Optimizes die area by configuring into the OCP only those features needed by the communicating cores
- Simplifies system verification and testing by providing a firm boundary around each IP core that can be observed, controlled, and validated

The approach adopted by the Virtual Socket Interface Alliance's (VSIA) Design Working Group on On-Chip Buses (DWGOB) is to specify a bus wrapper to provide a bus-independent Transaction Protocol-level interface to IP cores.

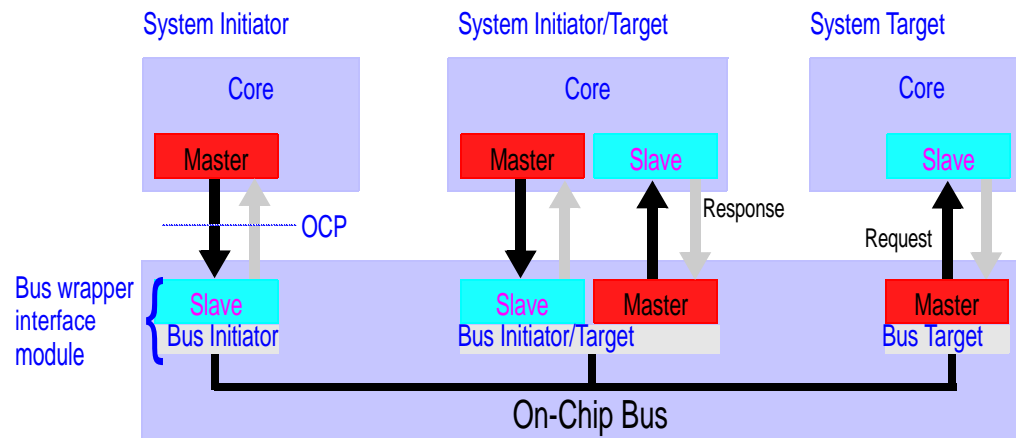
The OCP is equivalent to VSIA's Virtual Component Interface (VCI). While the VCI addresses only data flow aspects of core communications, the OCP is a superset of VCI additionally supporting configurable sideband control signaling and test harness signals. The OCP is the only standard that defines protocols to unify all of the inter-core communication.

OCP Characteristics

The OCP defines a point-to-point interface between two communicating entities such as IP cores and bus interface modules (bus wrappers). One entity acts as the master of the OCP instance, and the other as the slave. Only the master can present commands and is the controlling entity. The slave responds to commands presented to it, either by accepting data from the master, or presenting data to the master. For two entities to communicate in a peer-to-peer fashion, there need to be two instances of the OCP connecting them - one where the first entity is a master, and one where the first entity is a slave.

Figure 1 shows a simple system containing a wrapped bus and three IP core entities: one that is a system target, one that is a system initiator, and an entity that is both.

Figure 1 System Showing Wrapped Bus and OCP Instances



The characteristics of the IP core determine whether the core needs master, slave, or both sides of the OCP; the wrapper interface modules must act as the complementary side of the OCP for each connected entity. A transfer across this system occurs as follows. A system initiator (as the OCP master) presents command, control, and possibly data to its connected slave (a bus wrapper interface module). The interface module plays the request across the on-chip bus system. The OCP does not specify the embedded bus functionality. Instead, the interface designer converts the OCP request into an embedded bus transfer. The receiving bus wrapper interface module (as the OCP master) converts the embedded bus operation into a legal OCP command. The system target (OCP slave) receives the command and takes the requested action.

Each instance of the OCP is configured (by choosing signals or bit widths of a particular signal) based on the requirements of the connected entities and is independent of the others. For instance, system initiators may require more

address bits in their OCP instances than do the system targets; the extra address bits might be used by the embedded bus to select which bus target is addressed by the system initiator.

The OCP is flexible. There are several useful models for how existing IP cores communicate with one another. Some employ pipelining to improve bandwidth and latency characteristics. Others use multiple-cycle access models, where signals are held static for several clock cycles to simplify timing analysis and reduce implementation area. Support for this wide range of behavior is possible through the use of synchronous handshaking signals that allow both the master and slave to control when signals are allowed to change.

Compliance

For a core to be considered OCP compliant, it must satisfy the following conditions:

1. The core must include at least one OCP interface.
2. The core and OCP interfaces must be described using an RTL configuration file with the syntax specified in Chapter 6 on page 75.
3. Each OCP interface on the core must:
 - Comply with all aspects of the OCP interface specification
 - Have its timing described using a synthesis configuration file following the syntax specified in Chapter 7 on page 87.
4. The following practices are recommended but not required:
 - a. Each non-OCP interface on the core should:
 - Be described using an interface configuration file with the syntax specified in Chapter 5 on page 69.
 - Have its timing described using a synthesis configuration file with the syntax specified in Chapter 7 on page 87.
 - b. A performance report as specified in Chapter 12 on page 203 (or an equivalent report) should be included for the core.

Part I *Specification*

2 *Theory of Operation*

The Open Core Protocol interface addresses communications between the functional units (or IP cores) that comprise a system on a chip. The OCP provides independence from bus protocols without having to sacrifice high-performance access to on-chip interconnects. By designing to the interface boundary defined by the OCP, you can develop reusable IP cores without regard for the ultimate target system.

Given the wide range of IP core functionality, performance and interface requirements, a fixed definition interface protocol cannot address the full spectrum of requirements. The need to support verification and test requirements adds an even higher level of complexity to the interface. To address this spectrum of interface definitions, the OCP defines a highly configurable interface. The OCP's structured methodology includes all of the signals required to describe an IP cores' communications including data flow, control, and verification and test signals.

This chapter provides an overview of the concepts behind the Open Core Protocol, introduces the terminology used to describe the interface and offers a high-level view of the protocol.

Point-to-Point Synchronous Interface

To simplify timing analysis, physical design, and general comprehension, the OCP is composed of uni-directional signals driven with respect to, and sampled by the rising edge of the OCP clock. The OCP is fully synchronous (with the exception of reset) and contains no multi-cycle timing paths with respect to the OCP clock. All signals other than the clock signal are strictly point-to-point.

Bus Independence

A core utilizing the OCP can be interfaced to any bus. A test of any bus-independent interface is to connect a master to a slave without an intervening on-chip bus. This test not only drives the specification towards a fully symmetric interface but helps to clarify other issues. For instance, device selection techniques vary greatly among on-chip buses. Some use address decoders. Others generate independent device select signals (analogous to a board level chip select). This complexity should be hidden from IP cores, especially since in the directly-connected case there is no decode/selection logic. OCP-compliant slaves receive device selection information integrated into the basic command field.

Arbitration schemes vary widely. Since there is virtually no arbitration in the directly-connected case, arbitration for any shared resource is the sole responsibility of the logic on the bus side of the OCP. This permits OCP-compliant masters to pass a command field across the OCP that the bus interface logic converts into an arbitration request sequence.

Commands

There are two basic commands, Read and Write and five command extensions. The WriteNonPost and Broadcast commands have semantics that are similar to the Write command. A WriteNonPost explicitly instructs the slave not to post a write. For the Broadcast command, the master indicates that it is attempting to write to several or all remote target devices that are connected on the other side of the slave. As such, Broadcast is typically useful only for slaves that are in turn a master on another communication medium (such as an attached bus).

The other command extensions, ReadExclusive, ReadLinked and WriteConditional, are used for synchronization between system initiators. ReadExclusive is paired with Write or WriteNonPost, and has blocking semantics. ReadLinked, used in conjunction with WriteConditional has non-blocking (lazy) semantics. These synchronization primitives correspond to those available natively in the instruction sets of different processors.

Address/Data

Wide widths, characteristic of shared on-chip address and data buses, make tuning the OCP address and data widths essential for area-efficient implementation. Only those address bits that are significant to the IP core should cross the OCP to the slave. The OCP address space is flat and composed of 8-bit bytes (octets).

To increase transfer efficiencies, many IP cores have data field widths significantly greater than an octet. The OCP supports a configurable data width to allow multiple bytes to be transferred simultaneously. The OCP refers to the chosen data field width as the *word size* of the OCP. The term *word* is used in the traditional computer system context; that is, a *word* is the natural transfer unit of the block. OCP supports word sizes of power-of-two and non-power-of-two as would be needed for a 12-bit DSP core. The OCP address is a byte address that is word aligned.

Transfers of less than a full word of data are supported by providing byte enable information that specifies which octets are to be transferred. Byte enables are linked to specific data bits (byte lanes). Byte lanes are not associated with particular byte addresses. This makes the OCP endian-neutral, able to support both big and little-endian cores.

Pipelining

The OCP allows pipelining of transfers. To support this feature, the return of read data and the provision of write data may be delayed after the presentation of the associated request.

Response

The OCP separates requests from responses. A slave can accept a command request from a master on one cycle and respond in a later cycle. The division of request from response permits pipelining. The OCP provides the option of having responses for Write commands, or completing them immediately without an explicit response.

Burst

To provide high transfer efficiency, burst support is essential for many IP cores. The extended OCP supports annotation of transfers with burst information. Bursts can either include addressing information for each successive command (which simplifies the requirements for address sequencing/burst count processing in the slave), or include addressing information only once for the entire burst.

In-band Information

Cores can pass core-specific information in-band in company with the other information being exchanged. In-band extensions exist for requests and responses, as well as read and write data. A typical use of in-band extensions is to pass cacheable information or data parity.

Tags

Tags are available in the OCP interface to control the ordering of responses. Without tags, a slave must return responses in the order that the requests were issued by the master. Similarly, writes must be committed in order. With the addition of tags, responses can be returned out-of-order, and write data can be committed out-of-order with respect to requests, as long as the transactions target different addresses. The tag links the response back to the original request.

Tagging is useful when a master core such as a processor can handle out-of-order return, because it allows a slave core such as a DRAM controller to service requests in the order that is most convenient, rather than the order in which requests were sent by the master.

Out-of-order request and response delivery can also be enabled using multiple threads. The major differences between threads and tags are that threads can have independent flow control for each thread and have no ordering rules for transactions on different threads. Tags, on the other hand, exist within a single thread and are restricted to shared flow control. Tagged transactions cannot be re-ordered with respect to overlapping addresses. Implementing independent flow control requires independent buffering for each thread, leading to more complex implementations. Tags enable lower overhead implementations for out-of-order return of responses at the expense of some concurrency.

Threads and Connections

To support concurrency and out-of-order processing of transfers, the extended OCP supports the notion of multiple threads. Transactions within different threads have no ordering requirements, and independent flow control from one another. Within a single thread of data flow, all OCP transfers must remain ordered unless tags are in use. Transfers within a single thread must remain ordered unless tags are in use. The concepts of threads and tags are hierarchical: each thread has its own flow control, and ordering within a thread either follows the request order strictly, or is governed by tags.

While the notion of a thread is a local concept between a master and a slave communicating over an OCP, it is possible to globally pass thread information from initiator to target using connection identifiers. Connection information helps to identify the initiator and determine priorities or access permissions at the target.

Interrupts, Errors, and other Sideband Signaling

While moving data between devices is a central requirement of on-chip communication systems, other types of communications are also important. Different types of control signaling are required to coordinate data transfers (for instance, high-level flow control) or signal system events (such as interrupts). Dedicated point-to-point data communication is sometimes required. Many devices also require the ability to notify the system of errors that may be unrelated to address/data transfers.

The OCP refers to all such communication as *sideband* (or out-of-band) signaling, since it is not directly related to the protocol state machines of the dataflow portion of the OCP. The OCP provides support for such signals through sideband signaling extensions.

Errors are reported across the OCP using two mechanisms. The error response code in the response field describes errors resulting from OCP transfers that provide responses. Write-type commands without responses cannot use the in-band reporting mechanism. The second method for reporting errors across the OCP uses out-of band error fields. These signals report more generic sideband errors, including those associated with posted write commands.

3 *Signals and Encoding*

OCP interface signals are grouped into dataflow, sideband, and test signals. The dataflow signals are divided into basic signals, simple extensions, burst extensions, tag extensions, and thread extensions. A small set of the signals from the basic dataflow group is required in all OCP configurations. Optional signals can be configured to support additional core communication requirements. All sideband and test signals are optional.

The OCP is a synchronous interface with a single clock signal. All OCP signals, other than the clock and reset are driven with respect to, and sampled by, the rising edge of the OCP clock. Except for clock, OCP signals are strictly point-to-point and uni-directional. The complete set of OCP signals is shown in Figure 4 on page 34.

Dataflow Signals

The dataflow signals consist of a small set of required signals and a number of optional signals that can be configured to support additional core communication requirements. The dataflow signals are grouped into basic signals, simple extensions (options such as byte enables and in-band information), burst extensions (support for bursting), tag extensions (re-ordering support), and thread extensions (multi-threading support).

The naming conventions for dataflow signals use the prefix M for signals driven by the OCP master and S for signals driven by the OCP slave.

Basic Signals

Table 1 lists the basic OCP signals. Only Clk and MCmd are required. The remaining OCP signals are optional.

Table 1 Basic OCP Signals

Name	Width	Driver	Function
Clk	1	varies	Clock input
EnableClk	1	varies	Enable OCP clock
MAddr	configurable	master	Transfer address
MCmd	3	master	Transfer command
MData	configurable	master	Write data
MDataValid	1	master	Write data valid
MRespAccept	1	master	Master accepts response
SCmdAccept	1	slave	Slave accepts transfer
SData	configurable	slave	Read data
SDataAccept	1	slave	Slave accepts write data
SResp	2	slave	Transfer response

Clk

Input clock signal for the OCP clock. The rising edge of the OCP clock is defined as a rising edge of Clk that samples the asserted EnableClk. Falling edges of Clk and any rising edge of Clk that does not sample EnableClk asserted do not constitute rising edges of the OCP clock.

EnableClk

EnableClk indicates which rising edges of Clk are the rising edges of the OCP clock, that is, which rising edges of Clk should sample and advance interface state. Use the `enableclk` parameter to configure this signal. EnableClk is driven by a third entity and serves as an input to both the master and the slave.

When `enableclk` is set to 0 (the default), the signal is not present and the OCP behaves as if `EnableClk` is constantly asserted. In that case all rising edges of `Clk` are rising edges of the OCP clock.

MAddr

The Transfer address, `MAddr` specifies the slave-dependent address of the resource targeted by the current transfer. To configure this field into the OCP, use the `addr` parameter. To configure the width of this field, use the `addr_width` parameter.

`MAddr` is a byte address that must be aligned to the OCP word size (`data_width`). `data_width` defines a minimum `addr_width` value that is based on the data bus byte width, and is defined as:

$$\text{min_addr_width} = \max(1, \text{ceil}(\log_2(\text{data_width}))-3)$$

If the OCP word size is larger than a single byte, the aggregate is addressed at the OCP word-aligned address and the lowest order address bits are hardwired to 0. If the OCP word size is not a power-of-2, the address is the same as it would be for an OCP interface with a word size equal to the next larger power-of-2.

MCmd

Transfer command. This signal indicates the type of OCP transfer the master is requesting. Each non-idle command is either a read or write type request, depending on the direction of data flow. Commands are encoded as follows.

Table 2 Command Encoding

MCmd[2:0]	Command	Mnemonic	Request Type
0 0 0	Idle	IDLE	(none)
0 0 1	Write	WR	write
0 1 0	Read	RD	read
0 1 1	ReadEx	RDEX	read
1 0 0	ReadLinked	RDL	read
1 0 1	WriteNonPost	WRNP	write
1 1 0	WriteConditional	WRC	write
1 1 1	Broadcast	BCST	write

The set of allowable commands can be limited using the `write_enable`, `read_enable`, `readex_enable`, `writenonpost_enable`, `rdlwrc_enable`, and `broadcast_enable` parameters as described in “Protocol Options” on page 55.

MData

Write data. This field carries the write data from the master to the slave. The field is configured into the OCP using the `mdata` parameter and its width is configured using the `data_width` parameter. The width is not restricted to multiples of 8.

MDataValid

Write data valid. When set to 1, this bit indicates that the data on the MData field is valid. Use the `datahandshake` parameter to configure this field into the OCP.

MRespAccept

Master response accept. The master indicates that it accepts the current response from the slave with a value of 1 on the MRespAccept signal. Use the `respaccept` parameter to enable this field into the OCP.

SCmdAccept

Slave accepts transfer. A value of 1 on the SCmdAccept signal indicates that the slave accepts the master's transfer request. To configure this field into the OCP, use the `cmdaccept` parameter.

SData

This field carries the requested read data from the slave to the master. The field is configured into the OCP using the `sdata` parameter and its width is configured using the `data_width` parameter. The width is not restricted to multiples of 8.

SDataAccept

Slave accepts write data. The slave indicates that it accepts pipelined write data from the master with a value of 1 on SDataAccept. This signal is meaningful only when `datahandshake` is in use. Use the `dataaccept` parameter to configure this field into the OCP.

SResp

Response field from the slave to a transfer request from the master. The field is configured into the OCP using the `resp` parameter. Response encoding is as follows.

Table 3 Response Encoding

SResp[1:0]		Response	Mnemonic
0	0	No response	NULL
0	1	Data valid / accept	DVA
1	0	Request failed	FAIL
1	1	Response error	ERR

Simple Extensions

Table 4 lists the simple OCP extensions. The extensions add to the OCP interface address spaces, byte enables, and additional core-specific information for each phase.

Table 4 Simple OCP Extensions

Name	Width	Driver	Function
MAddrSpace	configurable	master	Address space
MByteEn	configurable	master	Request phase byte enables
MDataByteEn	configurable	master	Datahandshake phase write byte enables
MDataInfo	configurable	master	Additional information transferred with the write data
MReqInfo	configurable	master	Additional information transferred with the request
SDataInfo	configurable	slave	Additional information transferred with the read data
SRespInfo	configurable	slave	Additional information transferred with the response

MAddrSpace

This field specifies the address space and is an extension of the MAddr field that is used to indicate the address region of a transfer. Examples of address regions are the register space versus the regular memory space of a slave or the user versus supervisor space for a master.

The MAddrSpace field is configured into the OCP using the `addrspace` parameter. The width of the MAddrSpace field is configured with the `addrspace_width` parameter. While the encoding of the MAddrSpace field is core-specific, it is recommended that slaves use 0 to indicate the internal register space.

MByteEn

Byte enables. This field indicates which bytes within the OCP word are part of the current transfer. See “Partial Word Transfers” on page 47 for more detail on request and datahandshake phase byte enables and their relationship. There is one bit in MByteEn for each byte in the OCP word. Setting MByteEn[n] to 1 indicates that the byte associated with data wires [(8n + 7):8n] should be transferred. The MByteEn field is configured into the OCP using the `byteen` parameter and is allowed only if `data_width` is a multiple of 8 (that is, the data width is an integer number of bytes).

The allowable patterns on MByteEn can be limited using the `force_aligned` parameter as described on page 56.

MDataByteEn

Write byte enables. This field indicates which bytes within the OCP word are part of the current write transfer. See “Partial Word Transfers” on page 47 for more detail on request and datahandshake phase byte enables and their relationship. There is one bit in MDataByteEn for each byte in the OCP word. Setting MDataByteEn[n] to 1 indicates that the byte associated with MData wires [(8n + 7):8n] should be transferred. The MDataByteEn field is configured into the OCP using the `mdatabyteen`

parameter. Setting `mdatabyteen` to 1 is only allowed if `datahandshake` is 1, and only if `data_width` is a multiple of 8 (that is, the data width is an integer number of bytes).

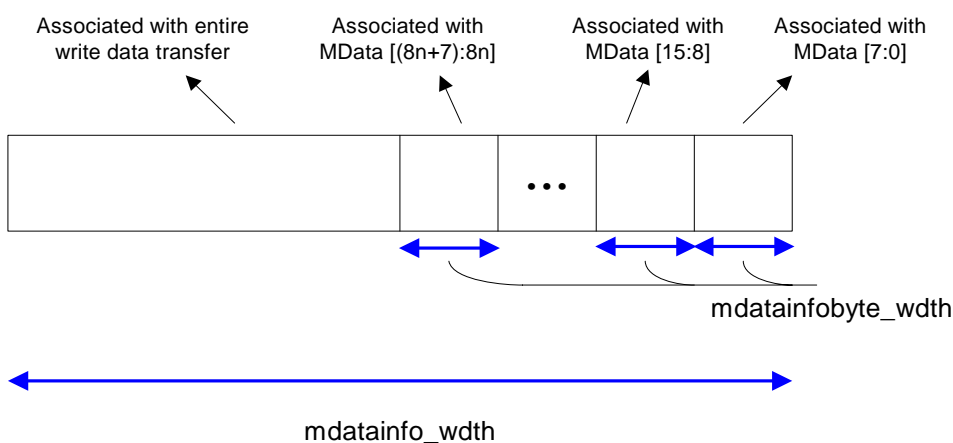
The allowable patterns on `MDataByteEn` can be limited using the `force_aligned` parameter as described on page 56.

MDataInfo

Extra information sent with the write data. The master uses this field to send additional information sequenced with the write data. The encoding of the information is core-specific. To be interoperable with masters that do not provide this signal, design slaves to be operable in a normal mode when the signal is tied off to its default tie-off value as specified in Table 13 on page 30. Sample uses are data byte parity or error correction code values. Use the `mdatainfo` parameter to configure this field into the OCP, and the `mdatainfo_width` parameter to configure its width.

This field is divided in two: the low-order bits are associated with each data byte, while the high-order bits are associated with the entire write data transfer. The number of bits to associate with each data byte is configured using the `mdatainfobyte_width` parameter. The low-order `mdatainfobyte_width` bits of `MDataInfo` are associated with the `MData[7:0]` byte, and so on.

Figure 2 MDataInfo Field



MReqInfo

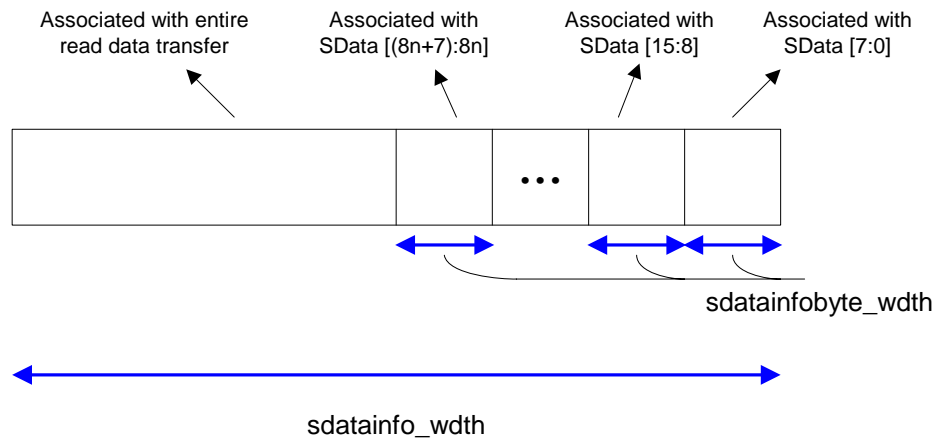
Extra information sent with the request. The master uses this field to send additional information sequenced with the request. The encoding of the information is core-specific. To be interoperable with masters that do not provide this signal, design slaves to be operable in a normal mode when the signal is tied off to its default tie-off value as specified in Table 13 on page 30. Sample uses are cacheable storage attributes or other access mode information. Use the `reqinfo` parameter to configure this field into the OCP, and the `reqinfo_width` parameter to configure its width.

SDataInfo

Extra information sent with the read data. The slave uses this field to send additional information sequenced with the read data. The encoding of the information is core-specific. To be interoperable with slaves that do not provide this signal, design masters to be operable in a normal mode when the signal is tied off to its default tie-off value as specified in Table 13 on page 30. Sample uses are data byte parity or error correction code values. Use the `sdatainfo` parameter to configure this field into the OCP, and the `sdatainfo_width` parameter to configure its width.

This field is divided into two pieces: the low-order bits are associated with each data byte, while the high-order bits are associated with the entire read data transfer. The number of bits to associate with each data byte is configured using the `sdatainfobyte_width` parameter. The low-order `sdatainfobyte_width` bits of `SDataInfo` are associated with the `SData[7:0]` byte, and so on.

Figure 3 *SDataInfo Field*



SRespInfo

Extra information sent with the response. The slave uses this field to send additional information sequenced with the response. The encoding of the information is core-specific. To be interoperable with slaves that do not provide this signal, design masters to be operable in a normal mode when the signal is tied off to its default tie-off value as specified in Table 13 on page 30. Sample uses are status or error information such as FIFO full or empty indications. Use the `respinfo` parameter to configure this field into the OCP, and the `respinfo_width` parameter to configure its width.

Burst Extensions

Table 5 lists the OCP burst extensions. The burst extensions allow the grouping of multiple transfers that have a defined address relationship. The burst extensions are enabled only when `MBurstLength` is included in the interface, or tied off to a value other than one.

Table 5 OCP Burst Extensions

Name	Width	Driver	Function
MAtomicLength	configurable	master	Length of atomic burst
MBlockHeight	configurable	master	Height of 2D block burst
MBlockStride	configurable	master	Address offset between 2D block rows
MBurstLength	configurable	master	Burst length
MBurstPrecise	1	master	Given burst length is precise
MBurstSeq	3	master	Address sequence of burst
MBurstSingleReq	1	master	Burst uses single request/ multiple data protocol
MDataLast	1	master	Last write data in burst
MDataRowLast	1	master	Last write data in row
MReqLast	1	master	Last request in burst
MReqRowLast	1	master	Last request in row
SRespLast	1	slave	Last response in burst
SRespRowLast	1	slave	Last response in row

MAtomicLength

This field indicates the minimum number of transfers within a burst that are to be kept together as an atomic unit when interleaving requests from different initiators onto a single thread at the target. To configure this field into the OCP, use the `atomiclength` parameter. To configure the width of this field, use the `atomiclength_width` parameter. A binary encoding of the number of transfers is used. A value of 0 is not a legal encoding for MAtomicLength.

MBlockHeight

This field indicates the number of rows of data to be transferred in a 2-dimensional block burst (the height of the block of data). A binary encoding of the height is used. To configure this field into the OCP, use the `blockheight` parameter. To configure the width of this field, use the `blockheight_width` parameter.

MBlockStride

This field indicates the address difference between the first data word in each consecutive row in a 2-dimensional block burst. The stride value is a binary encoded byte address offset and must be aligned to the OCP word size (`data_width`). To configure this field into the OCP, use the `blockstride` parameter. To configure the width of this field, use the `blockstride_width` parameter.

MBurstLength

The number of transfers in a burst. For precise bursts, the value indicates the total number of transfers in the burst, and is constant throughout the burst. For imprecise bursts, the value indicates the best guess of the

number of transfers remaining (including the current request), and may change with every request. To configure this field into the OCP, use the `burstlength` parameter. To configure the width of this field, use the `burstlength_width` parameter. A binary encoding of the number of transfers is used. A value of 0 is not a legal encoding for `MBurstLength`.

MBurstPrecise

This field indicates whether the precise length of a burst is known at the start of the burst or not. When set to 1, `MBurstLength` indicates the precise length of the burst during the first request of the burst. To configure this field into the OCP, use the `burstprecise` parameter. When set to 0, `MBurstLength` for each request is a hint of the remaining burst length.

MBurstSeq

This field indicates the sequence of addresses for requests in a burst. To configure this field into the OCP, use the `burstseq` parameter. The encodings of the `MBurstSeq` field are shown in Table 6. The precise definition of the address sequences is described in “Burst Address Sequences” on page 49.

Table 6 *MBurstSeq Encoding*

MBurstSeq[2:0]			Burst Sequence	Mnemonic
0	0	0	Incrementing	INCR
0	0	1	Custom (packed)	DFLT1
0	1	0	Wrapping	WRAP
0	1	1	Custom (not packed)	DFLT2
1	0	0	Exclusive OR	XOR
1	0	1	Streaming	STRM
1	1	0	Unknown	UNKN
1	1	1	2-dimensional Block	BLCK

MBurstSingleReq

The burst has a single request with multiple data transfers. This field indicates whether the burst has a request per data transfer, or a single request for all data transfers. To configure this field into the OCP, use the `burstsinglereq` parameter. When this field is set to 0, there is a one-to-one association of requests to data transfers; when set to 1, there is a single request for all data transfers in the burst.

MDataLast

Last write data in a burst. This field indicates whether the current write data transfer is the last in a burst. To configure this field into the OCP, use the `datalast` parameter with `datahandshake` set to 1. When this field is set to 0, more write data transfers are coming for the burst; when set to 1, the current write data transfer is the last in the burst.

MDataRowLast

Last write data in a row. This field identifies the last transfer in a row. The last data transfer in a burst is always considered the last in a row, and BLCK burst sequences also have a last in a row transfer after every MBurstLength transfers. To configure this field into the OCP, use the `datarowlast` parameter. If this field is set to 0, additional write data transfers can be expected for the current row; when set to 1, the current write data transfer is the last in the row.

MReqLast

Last request in a burst. This field indicates whether the current request is the last in this burst. To configure this field into the OCP, use the `reqlast` parameter. When this field is set to 0, more requests are coming for this burst; when set to 1, the current request is the last in the burst.

MReqRowLast

Last request in a row. This field identifies the last request in a row. The last request in a burst is always considered the last in a row, and BLCK burst sequences also have a last in a row request after every MBurstLength requests. To configure this field into the OCP, use the `reqrowlast` parameter. When this field is set to 0, more requests can be expected for the current row; when set to 1, the current request is the last in the row.

SRespLast

Last response in a burst. This field indicates whether the current response is the last in this burst. To configure this field into the OCP, use the `resplast` parameter. When the field is set to 0, more responses are coming for this burst; when set to 1, the current response is the last in the burst.

SRespRowLast

Last response in a row. This field identifies the last response in a row. The last response in a burst is always considered the last in a row, and BLCK burst sequences also have a last in a row response after every MBurstLength responses. To configure this field into the OCP, use the `resprolast` parameter. When this field is set to 0, more can be expected for the current row; when set to 1, the current response is the last in the row.

Tag Extensions

Table 7 lists OCP tag extensions. The extensions add support for tagging OCP transfers to enable out-of-order responses and write data commit.

Table 7 OCP Tag Extensions

Name	Width	Driver	Function
MDataTagID	configurable	master	Ordering tag for write data
MTagID	configurable	master	Ordering tag for request

MTagInOrder	1	master	Do not reorder this request
STagID	configurable	slave	Ordering tag for response
STagInOrder	1	slave	This response is not reordered

MDataTagID

Write data tag. This variable-width field provides the tag associated with the current write data. The field carries the binary-encoded value of the tag.

MDataTagID is required if `tags` is greater than 1 and the `datahandshake` parameter is set to 1. The field width is the next whole integer of $\log_2(\text{tags})$.

MTagID

Request tag. This variable-width field provides the tag associated with the current transfer request. If `tags` is greater than 1, this field is enabled. The field width is the next whole integer of $\log_2(\text{tags})$.

MTagInOrder

Assertion of this single-bit field indicates that the current request should not be reordered with respect to other requests that had this field asserted. This field is enabled by the `taginorder` parameter. Both MTagInOrder and STagInOrder are present in the interface, or else neither may be present.

STagID

Response tag. This variable-width field provides the tag associated with the current transfer response. This field is enabled if `tags` is greater than 1, and the `resp` parameter is set to 1. The field width is the next whole integer of $\log_2(\text{tags})$.

STagInOrder

Assertion of this single-bit field indicates that the current response is associated with an in-order request and was not reordered with respect to other requests that had MTagInOrder asserted. This field is enabled if both the `taginorder` parameter is set to 1 and the `resp` parameter is set to 1.

Thread Extensions

Table 8 shows a list of OCP thread extensions. The extensions add support for multi-threading of the OCP interface.

Table 8 OCP Thread Extensions

Name	Width	Driver	Function
MConnID	configurable	master	Connection identifier
MDataThreadID	configurable	master	Write data thread identifier
MThreadBusy	configurable	master	Master thread busy
MThreadID	configurable	master	Request thread identifier

SDataThreadBusy	configurable	slave	Slave write data thread busy
SThreadBusy	configurable	slave	Slave request thread busy
SThreadID	configurable	slave	Response thread identifier

MConnID

Connection identifier. This variable-width field provides the binary encoded connection identifier associated with the current transfer request.

To configure this field use the `connid` parameter. The field width is configured with the `connid_width` parameter.

MDataThreadID

Write data thread identifier. This variable-width field provides the thread identifier associated with the current write data. The field carries the binary-encoded value of the thread identifier.

MDataThreadID is required if `threads` is greater than 1 and the `datahandshake` parameter is set to 1. MDataThreadID has the same width as MThreadID and SThreadID.

MThreadBusy

Master thread busy. The master notifies the slave that it cannot accept any responses associated with certain threads. The MThreadBusy field is a vector (one bit per thread). A value of 1 on any given bit indicates that the thread associated with that bit is busy. Bit 0 corresponds to thread 0, and so on. The width of the field is set using the `threads` parameter. It is legal to enable a one-bit MThreadBusy interface for a single-threaded OCP. To configure this field, use the `mthreadbusy` parameter. See “Ungrouped Signals” on page 42 for a precise description of the flow control options associated with MThreadBusy.

MThreadID

Request thread identifier. This variable-width field provides the thread identifier associated with the current transfer request. If `threads` is greater than 1, this field is enabled. The field width is the next whole integer of $\log_2(\text{threads})$.

SDataThreadBusy

Slave write data thread busy. The slave notifies the master that it cannot accept any new datahandshake phases associated with certain threads. The SDataThreadBusy field is a vector, one bit per thread. A value of 1 on any given bit indicates that the thread associated with that bit is busy. Bit 0 corresponds to thread 0, and so on.

The width of the field is set using the `threads` parameter. It is legal to enable a one-bit SDataThreadBusy interface for a single-threaded OCP. To configure this field, use the `sdatathreadbusy` parameter. See “Ungrouped Signals” on page 42 for a precise description of the flow control options associated with SDataThreadBusy.

SThreadID

Response thread identifier. This variable-width field provides the thread identifier associated with the current transfer response. This field is enabled if `threads` is greater than 1 and the `resp` parameter is set to 1. The field width is the next whole integer of $\log_2(\text{threads})$.

SThreadBusy

Slave thread busy. The slave notifies the master that it cannot accept any new requests associated with certain threads. The `SThreadBusy` field is a vector, one bit per thread. A value of 1 on any given bit indicates that the thread associated with that bit is busy. Bit 0 corresponds to thread 0, and so on. The width of the field is set using the `threads` parameter. It is legal to enable a one-bit `SThreadBusy` interface for a single-threaded OCP. To configure this field, use the `sthrdbusy` parameter. See “Ungrouped Signals” on page 42 for a precise description of the flow control options associated with `SThreadBusy`.

Sideband Signals

Sideband signals are OCP signals that are not part of the dataflow phases, and so can change asynchronously with the request/response flow but are generally synchronous to the rising edge of the OCP clock. Sideband signals convey control information such as reset, interrupt, error, and core-specific flags. They also exchange control and status information between a core and an attached system. All sideband signals are optional except for reset. Either the `MReset_n` or the `SReset_n` signal must be present.

Table 9 shows a list of the sideband extensions to the OCP.

Table 9 Sideband OCP Signals

Name	Width	Driver	Function
MError	1	master	Master Error
MFlag	configurable	master	Master flags
MReset_n	1	master	Master reset
SError	1	slave	Slave error
SFlag	configurable	slave	Slave flags
SInterrupt	1	slave	Slave interrupt
SReset_n	1	slave	Slave reset
Control	configurable	system	Core control information
ControlBusy	1	core	Hold control information
ControlWr	1	system	Control information has been written
Status	configurable	core	Core status information
StatusBusy	1	core	Status information is not consistent
StatusRd	1	system	Status information has been read

Reset, Interrupt, Error, and Core-Specific Flag Signals

MError

Master error. When the MError signal is set to 1, the master notifies the slave of an error condition. The MError field is configured with the `merror` parameter.

MFlag

Master flags. This variable-width set of signals allows the master to communicate out-of-band information to the slave. Encoding is completely core-specific.

To configure this field into the OCP, use the `mflag` parameter. To configure the width of this field, use the `mflag_width` parameter.

MReset_n

Master reset. The MReset_n signal is active low, as shown in Table 10. The MReset_n field is enabled by the `mreset` parameter.

Table 10 MReset Signal

MReset_n	Function
0	Reset Active
1	Reset Inactive

SError

Slave error. With a value of 1 on the SError signal the slave indicates an error condition to the master. The SError field is configured with the `seerror` parameter.

SFlag

Slave flags. This variable-width set of signals allows the slave to communicate out-of-band information to the master. Encoding is completely core-specific.

To configure this field into the OCP, use the `sflag` parameter. To configure the width of this field, use the `sflag_width` parameter.

SInterrupt

Slave interrupt. The slave may generate an interrupt with a value of 1 on the SInterrupt signal. The SInterrupt field is configured with the `interrupt` parameter.

SReset_n

Slave reset. The SReset_n signal is active low, as shown in Table 11. The SReset_n field is enabled by the `sreset` parameter.

Table 11 *SReset Signal*

SReset_n	Function
0	Reset Active
1	Reset Inactive

Control and Status Signals

The remaining sideband signals are designed for the exchange of control and status information between an IP core and the rest of the system. They make sense only in this environment, regardless of whether the IP core acts as a master or slave across the OCP interface.

Control

Core control information. This variable-width field allows the system to drive control information into the IP core. Encoding is core-specific.

Use the `control` parameter to configure this field into the OCP. Use the `control_width` parameter to configure the width of this field.

ControlBusy

Core control busy. When this signal is set to 1, the core tells the system to hold the control field value constant. Use the `controlbusy` parameter to configure this field into the OCP.

ControlWr

Core control event. This signal is set to 1 by the system to indicate that control information is written by the system. Use the `controlwr` parameter to configure this field into the OCP.

Status

Core status information. This variable-width field allows the IP core to report status information to the system. Encoding is core-specific.

Use the `status` parameter to configure this field into the OCP. Use the `status_width` parameter to configure the width of this field.

StatusBusy

Core status busy. When this signal is set to 1, the core tells the system to disregard the status field because it may be inconsistent. Use the `statusbusy` parameter to configure this field into the OCP.

StatusRd

Core status event. This signal is set to 1 by the system to indicate that status information is read by the system. To configure this field into the OCP, use the `statusrd` parameter.

Test Signals

The test signals add support for scan, clock control, and IEEE 1149.1 (JTAG). All test signals are optional.

Table 12 Test OCP Signals

Name	Width	Driver	Function
Scanctrl	configurable	system	Scan control signals
Scanin	configurable	system	Scan data in
Scanout	configurable	core	Scan data out
ClkByp	1	system	Enable clock bypass mode
TestClk	1	system	Test clock
TCK	1	system	Test clock
TDI	1	system	Test data in
TDO	1	core	Test data out
TMS	1	system	Test mode select
TRST_N	1	system	Test reset

Scan Interface

The Scanctrl, Scanin, and Scanout signals together form a scan interface into a given IP core.

Scanctrl

Scan mode control signals. Use this variable width field to control the scan mode of the core. Set scanport to 1 to configure this field into the OCP interface. Use the scanctrl_width parameter to configure the width of this field.

Scanin

Scan data in for a core's scan chains. Use the scanport parameter, to configure this field into the OCP interface and scanport_width to control its width.

Scanout

Scan data out from the core's scan chains. Use the scanport parameter to configure this field into the OCP interface and scanport_width to control its width.

Clock Control Interface

The ClkByp and TestClk signals together form the clock control interface into a given IP core. This interface is used to control the core's clocks during scan operation.

ClkByp

Enable clock bypass signal. When set to 1, this signal instructs the core to bypass the external clock source and use TestClk instead. Use the clkctrl_enable parameter to configure this signal into the OCP interface.

TestClk

A gated test clock. This clock becomes the source clock when ClkByp is asserted during scan operations. Use the `clkctrl_enable` parameter to configure this signal into the OCP interface.

Debug and Test Interface

The TCK, TDI, TDO, TMS, and TRST_N signals together form an IEEE 1149 debug and test interface for the OCP.

TCK

Test clock as defined by IEEE 1149.1. Use the `jtag_enable` parameter to add this signal to the OCP interface.

TDI

Test data in as defined by IEEE 1149.1. Use the `jtag_enable` parameter to add this signal to the OCP interface.

TDO

Test data out as defined by IEEE 1149.1. Use the `jtag_enable` parameter to add this signal to the OCP interface.

TMS

Test mode select as defined by IEEE 1149.1. Use the `jtag_enable` parameter to add this signal to the OCP interface.

TRST_N

Test logic reset signal. This is an asynchronous active low reset as defined by IEEE 1149.1. Use the `jtagtrst_enable` parameter to add this signal to the OCP interface.

Signal Configuration

The set of signals making up the OCP interface is configurable to match the characteristics of the IP core. Throughout this chapter, configuration parameters were mentioned that control the existence and width of various OCP fields. Table 13 on page 30 summarizes the configuration parameters, sorted by interface signal group. For each signal, the table also specifies the default constant tie-off, which is the inferred value of the signal if it is not configured into the OCP interface and if no other constant tie-off is specified.

For the ControlBusy, EnableClk, MRespAccept, SCmdAccept, SDataAccept, MReset_n, SReset_n, and StatusBusy signals, the default tie-off is also the only legal tie-off.

Table 13 OCP Signal Configuration Parameters

Group	Signal	Parameter to add signal to interface	Parameter to control width	Default Tie-off
Basic	Clk	Required	Fixed	n/a
	EnableClk	enableclk	Fixed	1
	MAddr	addr	addr_width	0
	MCmd	Required	Fixed	n/a
	MData	mdata	data_width	0
	MDataValid	datahandshake	Fixed	n/a
	MRespAccept ¹	respaccept	Fixed	1
	SCmdAccept	cmdaccept	Fixed	1
	SData ¹	sdata	data_width	0
	SDataAccept ²	dataaccept	Fixed	1
	SResp	resp	Fixed	Null
Simple	MAddrSpace	addrspace	addrspace_width	0
	MByteEn ³	byteen	data_width	all 1s
	MDataByteEn ⁴	mdatabyteen	data_width	all 1s
	MDataInfo	mdatainfo	mdatainfo_width ⁵	0
	MReqInfo	reqinfo	reqinfo_width	0
	SDataInfo ¹	sdatainfo	sdatainfo_width ⁶	0
	SRespInfo ¹	respinfo	respinfo_width	0
Burst	MAtomicLength ¹⁹	atomiclength	atomiclength_width	1
	MBlockHeight ^{19, 21}	blockheight	blockheight_width ²²	1
	MBlockStride ^{19, 21}	blockstride	blockstride_width	0
	MBurstLength	burstlength	burstlength_width ²⁰	1
	MBurstPrecise ^{19, 23}	burstprecise	Fixed	1
	MBurstSeq ¹⁹	burstseq	Fixed	INCR
	MBurstSingleReq ^{7, 19}	burstsinglereq	Fixed	0
	MDataLast ^{8, 19}	datalast	Fixed	n/a
	MDataRowLast ^{8, 19, 21}	datarowlast	Fixed	n/a
	MReqLast ¹⁹	reqlast	Fixed	n/a
	MReqRowLast ^{19, 21, 24}	reqrowlast	Fixed	n/a
	SRespLast ^{1, 19}	resplast	Fixed	n/a
	SRespRowLast ^{19, 21, 26}	resprowlast	Fixed	n/a

Group	Signal	Parameter to add signal to interface	Parameter to control width	Default Tie-off
Tag	MDataTagID ⁹	tags>1 and datahandshake	tags	0
	MTagID	tags>1	tags	0
	MTagInOrder ¹⁰	taginorder	Fixed	0
	STagID	tags>1 and resp	tags	0
	STagInOrder ¹¹	taginorder and resp	Fixed	0
Thread	MConnID	connid	connid_width	0
	MDataThreadID	threads>1 and datahandshake	threads	0
	MThreadBusy ^{1, 12}	mthreadbusy	threads	0
	MThreadID	threads>1	threads	0
	SDataThreadBusy ¹³	sdatathreadbusy	threads	0
	SThreadBusy ¹⁴	stthreadbusy	threads	0
	SThreadID	threads>1 and resp	threads	0
Sideband	Control	control	control_width	0
	ControlBusy ¹⁵	controlbusy	Fixed	0
	ControlWr ¹⁶	controlwr	Fixed	n/a
	MError	merror	Fixed	0
	MFlag	mflag	mflag_width	0
	MReset_n	mreset	Fixed	1
	SError	serror	Fixed	0
	SFlag	sflag	sflag_width	0
	SInterrupt	interrupt	Fixed	0
	SReset_n	sreset	Fixed	1
	Status	status	status_width	0
	StatusBusy ¹⁷	statusbusy	Fixed	0
	StatusRd ¹⁸	statusrd	Fixed	n/a

Group	Signal	Parameter to add signal to interface	Parameter to control width	Default Tie-off
Test	ClkByp	clkctrl_enable	Fixed	n/a
	Scanctrl	scanport	scanctrl_width	n/a
	Scanin	scanport	scanport_width	n/a
	Scanout	scanport	scanport_width	n/a
	TCK	jtag_enable	Fixed	n/a
	TDI	jtag_enable	Fixed	n/a
	TDO	jtag_enable	Fixed	n/a
	TestClk	clkctrl_enable	Fixed	n/a
	TMS	jtag_enable	Fixed	n/a
	TRST_N ²⁷	jtagtrst_enable	Fixed	n/a

- 1 MRespAccept, MThreadBusy, SData, SDataInfo, SRespInfo, SRespLast, and SRespRowLast may be included only if the resp parameter is set to 1.
- 2 SDataAccept can be included only if datahandshake is set to 1.
- 3 MByteEn has a width of data_width/8 and can only be included when either mdata or sdata is set to 1 and data_width is an integer multiple of 8.
- 4 MDataByteEn has a width of data_width/8 and can only be included when mdata is set to 1, datahandshake is set to 1, and data_width is an integer multiple of 8.
- 5 mdatainfo_width must be \geq mdatainfobyte_width * data_width/8 and can be used only if data_width is a multiple of 8. mdatainfobyte_width specifies the partitioning of MDataInfo into transfer-specific and per-byte fields.
- 6 sdatainfo_width must be \geq sdatainfobyte_width * data_width/8 and can be used only if data_width is a multiple of 8. sdatainfobyte_width specifies the partitioning of SDataInfo into transfer-specific and per-byte fields.
- 7 If any write-type commands are enabled, MBurstSingleReq can only be included when datahandshake is set to 1. If the only enabled burst address sequence is UNKN, MBurstSingleReq cannot be included or tied off to a non-default value.
- 8 MDataLast and MDataRowLast can be included only if the datahandshake parameter is set to 1.
- 9 MDataTagID is included if tags is greater than 1 and the datahandshake parameter is set to 1.
- 10 MTagInOrder can only be included if tags is greater than 1.
- 11 STagInOrder can only be included if tags is greater than 1.
- 12 MThreadBusy has a width equal to threads. It may be included for single-threaded OCP interfaces.
- 13 SDataThreadBusy has a width equal to threads. It may be included for single-threaded OCP interfaces and may only be included if datahandshake is 1.
- 14 SThreadBusy has a width equal to threads. It may be included for single-threaded OCP interfaces.
- 15 ControlBusy can only be included if both Control and ControlWr exist.
- 16 ControlWr can only be included if Control exists.
- 17 StatusBusy can only be included if Status exists.
- 18 StatusRd can only be included if Status exists.
- 19 MAtomicLength, MBlockHeight, MBlockStride, MBurstPrecise, MBurstSeq, MBurstSingleReq MDataLast, MDataRowLast, MReqLast, MReqRowLast, SRespLast, and SRespRowLast may be included in the interface or tied off to non-default values only if MBurstLength is included or tied off to a value other than 1.
- 20 burstlength_width can never be 1.
- 21 MBlockHeight, MBlockStride, MDataRowLast, MReqRowLast, and SRespRowLast may be included or tied off to non-default values only if burstseq_blk_enable is set to 1 and MBurstPrecise is included or tied off to a value of 1.
- 22 blockheight_width can never be 1.
- 23 If no sequences other than WRAP, XOR, or BLCK are enabled, MBurstPrecise must be tied off to 1.
- 24 MDataRowLast can only be included if MDataLast is included.
- 25 MReqRowLast can only be included if MReqLast is included.
- 26 SRespRowLast can only be included if SRespLast is included.
- 27 TRST_N can only be included if jtag_enable is set to 1.

Signal Directions

Figure 4 on page 34 summarizes all OCP signals. The direction of some signals (for example, MCmd) depends on whether the module acts as a master or slave, while the direction of other signals (for example, Control) depends on whether the module acts as a system or a core. The combination of these two choices provides four possible module configurations as shown in Table 14.

Table 14 Module Configuration Based on Signal Directions


	Acts as Core or has No System Signals	Acts as System
Acts as OCP Master	Master	System master
Acts as OCP Slave	Slave	System slave

For example, if a module acts as OCP master and also as system, it is designated a system master. There is also a monitor type that observes all OCP signals. The permitted connectivity between interface types is shown in Table 15.

Table 15 Interface Types

Type	Connect To	Cannot Connect To
Master	System slave Slave	Master System master
Slave	System master Master	Slave System slave
System master	Slave System slave	Master System master
System slave	Master System master	Slave System slave
Monitor	Any except monitor	Monitor


The Clk and EnableClk signals are special in that they are supplied by a third (external) entity that is neither of the two modules connected through the OCP interface.



4 *Protocol Semantics*

This chapter defines the semantics of the OCP protocol by assigning meanings to the signal encodings described in the preceding chapter. Figure 5 provides a graphic view of the hierarchy of elements that compose the OCP.

Figure 5 *Hierarchy of Elements*



Signal Groups

Some OCP fields are grouped together because they must be active at the same time. The data flow signals are divided into three signal groups: request signals, response signals, and datahandshake signals. A list of the signals that belong to each group is shown in Table 16.

Table 16 OCP Signal Groups

Group	Signal	Condition
Request	MAddr	always
	MAddrSpace	always
	MAtomicLength	always
	MBlockHeight	always
	MBlockStride	always
	MBurstLength	always
	MBurstPrecise	always
	MBurstSeq	always
	MBurstSingleReq	always
	MByteEn	always
	MCmd	always
	MConnID	always
	MData*	datahandshake = 0
	MDataInfo*	datahandshake = 0
	MReqInfo	always
	MReqLast	always
	MReqRowLast	always
	MTagID	always
	MTagInOrder	always
	MThreadID	always

Group	Signal	Condition
Response	SData	always
	SDataInfo	always
	SResp	always
	SRespInfo	always
	SRespLast	always
	SRespRowLast	always
	STagID	always
	STagInOrder	always
	SThreadID	always
Datahandshake	MData*	datahandshake = 1
	MDataByteEn	always
	MDataInfo*	datahandshake = 1
	MDataLast	always
	MDataRowLast	always
	MDataTagID	always
	MDataThreadID	always
	MDataValid	always

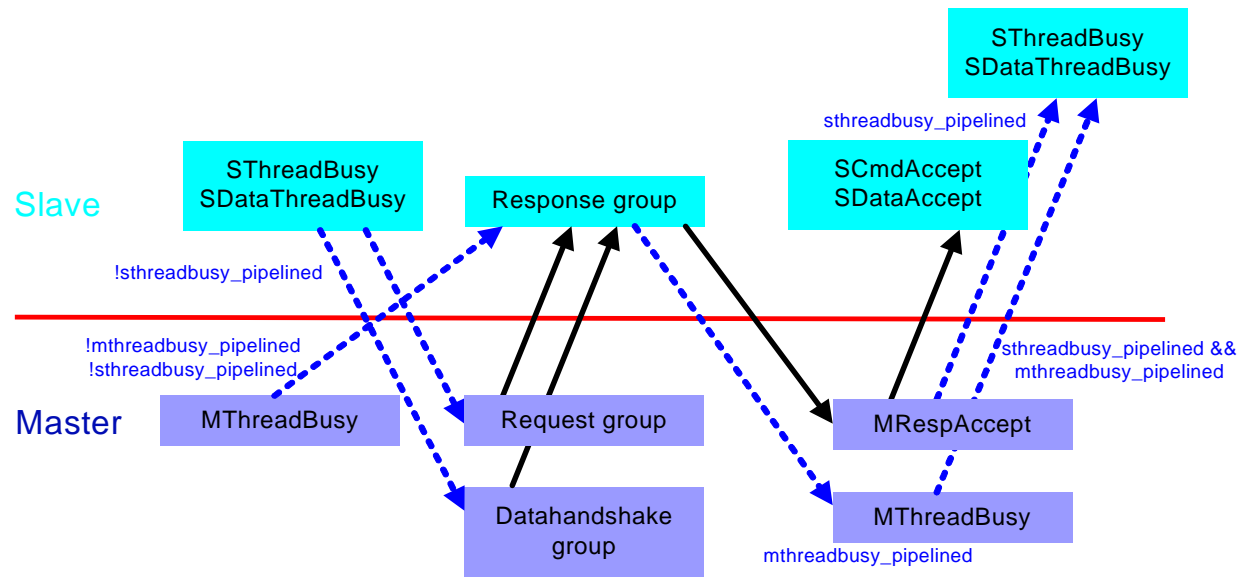


*MData and MDataInfo belong to the request group, unless the datahandshake configuration parameter is enabled. In that case they belong to the datahandshake group.

Combinational Dependencies

It is legal for some signal or signal group outputs to be derived from inputs without an intervening latch point, that is combinationally. To avoid combinational loops, other outputs cannot be derived in this manner. Figure 6 describes a partial order of combinational dependency. For any arrow shown, the signal or signal group can be derived combinationally from the signal at the point of origin of the arrow or another signal earlier in the dependency chain. No other combinational dependencies are allowed.

Figure 6 Legal Combinational Dependencies Between Signals and Signal Groups



MThreadBusy, SDataThreadBusy, and SThreadBusy each appear twice in Figure 6. The two appearances of each signal are mutually exclusive based upon the setting of the `mthreadbusy_pipelined`, `sdatathreadbusy_pipelined`, and `sthreadbusy_pipelined` parameters. Refer to “Ungrouped Signals” on page 42 for more information about these parameters.

Combinational paths are not allowed within the sideband and test signals, or between those signals and the data flow signals. The only legal combinational dependencies are within the data flow signals. Data flow signals, however, may be combinationaly derived from `MReset_n` and `SReset_n`.

For timing purposes, some of the allowed combinational paths are designated as preferred paths and are described in Table 33 on page 183.

Signal Timing and Protocol Phases

This section specifies when a signal can or must be valid.

OCP Clock

The rising edge of the OCP clock signal is used to sample other OCP signals to advance the state of the interface. When the `EnableClk` signal is not present, the OCP clock is simply the `Clk` signal. When the `EnableClk` signal is present (`enableclk` is 1), only rising edges of `Clk` that sample `EnableClk` asserted are considered rising edges of the OCP clock. Therefore, when `EnableClk` is 0, rising edges of `Clk` are not rising edges of the OCP clock and OCP state is not advanced.

This restriction applies to all signals in the OCP interface. In particular, the requirements for reset assertion described on page 44, are measured in OCP clock cycles.

Dataflow Signals

Signals in a signal group must all be valid at the same time.

- The request group is valid whenever a command other than Idle is presented on the MCmd field.
- The response group is valid whenever a response other than Null is presented on the SResp field.
- The datahandshake group is valid whenever a 1 is presented on the MDataValid field.

The accept signal associated with a signal group is valid only when that group is valid.

- The SCmdAccept signal is valid whenever a command other than Idle is presented on the MCmd field.
- The MRespAccept signal is valid whenever a response other than Null is presented on the SResp field.
- The SDataAccept signal is valid whenever a 1 is presented on the MDataValid field.

The signal groups map on a one-to-one basis to protocol phases. All signals in the group must be held steady from the beginning of a protocol phase until the end of that phase. Outside of a protocol phase, all signals in the corresponding group (except for the signal that defines the beginning of the phase) are “don’t care”.

In addition, the MData and MDataInfo fields are a “don’t care” during read-type requests, and the SData and SDataInfo fields are a “don’t care” for responses to write-type requests. Non-enabled data bytes in MData and bits in MDataInfo as well as non-enabled bytes in SData and bits in SDataInfo are a “don’t care”. The MDataByteEn field is a “don’t care” during read-type transfers. If MDataByteEn is present, the MByteEn field is a “don’t care” during write-type transfers. MTagID is a “don’t care” if MTagInOrder is asserted and MDataTagID is a “don’t care” for the corresponding datahandshake phase. Similarly, STagID is a “don’t care” if STagInOrder is asserted.

- A request phase begins whenever the request group becomes active. It ends when the SCmdAccept signal is sampled by the rising edge of the OCP clock as 1 during a request phase.
- A response phase begins whenever the response group becomes active. It ends when the MRespAccept signal is sampled by the rising edge of the OCP clock as 1 during a response phase.

If MRespAccept is not configured into the OCP interface (respaccept = 0) then MRespAccept is assumed to be on; that is the response phase is exactly one cycle long.

- A datahandshake phase begins whenever the datahandshake signal group becomes active. It ends when the SDataAccept signal is sampled by the rising edge of the OCP clock as 1 during a datahandshake phase.

For all phases, it is legal to assert the corresponding accept signal in the cycle that the phase begins, allowing the phase to complete in a single cycle.

Phases in a Transfer

An OCP transfer consists of several phases as shown in Table 17. Every transfer has a request phase. Read-type requests always have a response phase. For write-type requests, the OCP can be configured with or without responses or datahandshake. Without a response, a write-type request completes upon completion of the request phase (posted write model).

Table 17 Phases in each Transfer for MBurstSingleReq set to 0

MCmd	Phases	Condition
Read, ReadEx, ReadLinked	Request, response	always
Write, Broadcast	Request	datahandshake = 0 writeresp_enable = 0
Write, WriteNonPost, WriteConditional, Broadcast	Request, response	datahandshake = 0 writeresp_enable = 1
Write, Broadcast	Request, datahandshake	datahandshake = 1 writeresp_enable = 0
Write, WriteNonPost, WriteConditional, Broadcast	Request, datahandshake, response	datahandshake = 1 writeresp_enable = 1

Single request / multiple data bursts, described in “Single Request / Multiple Data Bursts (Packets)” on page 51, have a single request phase and multiple data transfer phases as shown in Table 18.

Table 18 Phases in a Transaction for MBurstSingleReq set to 1

MCmd	Phases	Condition
Read	Request, H*L* response	always
Write, Broadcast	Request, H*L* datahandshake	datahandshake = 1 writeresp_enable = 0
Write, WriteNonPost, Broadcast	Request, H*L* datahandshake, response	datahandshake = 1 writeresp_enable = 1

- * H refers to the burst height (MBlockHeight), and is 1 for all burst sequences other than BLCK
- * L refers to the burst length (MBurstLength)

Phase Ordering Within a Transfer

The OCP is causal: within each transfer a request phase must precede the associated datahandshake phase which in turn, must precede the associated response phase. The specific constraints are:

- A datahandshake phase cannot begin before the associated request phase begins, but can begin in the same OCP clock cycle.
- A datahandshake phase cannot end before the associated request phase ends, but can end in the same OCP clock cycle.
- A response phase cannot begin before the associated request phase begins, but can begin in the same OCP clock cycle.
- A response phase cannot end before the associated request phase ends, but can end in the same OCP clock cycle.
- If there is a datahandshake phase and a response phase, the response phase cannot begin before the associated datahandshake phase (or last datahandshake phase for single request/multiple data bursts) begins, but can begin in the same OCP clock cycle.
- If there is a datahandshake phase and a response phase, the response phase cannot end before the associated datahandshake phase (or last datahandshake phase for single request/multiple data bursts) ends, but can end in the same OCP clock cycle.

Phase Ordering Between Transfers

If tags are not in use, the ordering of transfers is determined by the ordering of their request phases. Also, the following rules apply.

- Since two phases of the same type belonging to different transfers both use the same signal wires, the phase of a subsequent transfer cannot begin before the phase of the previous transfer has ended. If the first phase ends in cycle x , the second phase can begin in cycle $x + 1$.
- The ordering of datahandshake phases must follow the order set by the request phases including multiple datahandshake phases for single request / multiple data bursts.
- The ordering of response phases must follow the order set by the request phases including multiple response phases for single request / multiple data bursts.
- For single request / multiple data bursts, a request or response phase is shared between multiple transfers. Each individual transfer must obey the ordering rules described in “Phase Ordering Within a Transfer” on page 41, even when a phase is shared with another transfer.

- Where no phase ordering is specified, by the previous rules, the effect of two transfers that are addressing the same location (as specified by MAddr, MAddrSpace, and MByteEn) must be the same as if the two transfers were executed in the same order but without any phase overlap. This ensures that read/write hazards yield predictable results.

For example, on an OCP interface with datahandshake enabled, a read following a write to the same location cannot start its response phase until the write has started its datahandshake phase, otherwise the latest write data cannot be returned for the read.

If tags are in use, the preceding rules are partially relaxed. Refer to “Ordering Restrictions” on page 53 for a more detailed explanation.

Ungrouped Signals

Signals not covered in the description of signal groups and phases are MThreadBusy, SDataThreadBusy, and SThreadBusy. Without further protocol restriction, the cycle timing of the transition of each bit that makes up each of these three fields is not specified relative to the other dataflow signals. This means that there is no specific time for an OCP master or slave to drive these signals, nor a specific time for the signals to have the desired flow-control effect. Without further restriction, MThreadBusy, SDataThreadBusy, and SThreadBusy can only be treated as a hint.

For stricter semantics use the protocol configuration parameters `mthreadbusy_exact`, `sdathreadbusy_exact`, and `sthreadbusy_exact`. These parameters can be set to 1 only when the corresponding signal has been enabled.

The parameters `mthreadbusy_pipelined`, `sdathreadbusy_pipelined`, and `sthreadbusy_pipelined` can be used to set the relative protocol timing of the MThreadBusy, SDataThreadBusy, and SThreadBusy signals with respect to their phases. The `_pipelined` parameters can only be enabled when the corresponding `_exact` parameter is enabled.

`_exact` parameters define strict protocol semantics for the corresponding phase. The receiver of the phase identifies (through the corresponding ThreadBusy signals) to the sender of the phase which threads can accept an asserted phase. The sender will not assert a phase on a busy thread, and the receiver accepts any phases asserted on threads that are not busy. To avoid ambiguity, the corresponding phase Accept signal may not be present on the interface, and is considered tied off to 1. The resulting phase has exact flow control and is non-blocking. See “Flow Control Options” on page 57 for configuration restrictions associated with ThreadBusy-related parameters.

`_pipelined` parameters control the cycle relationship between the ThreadBusy signal and the corresponding phase assertion. When the `_pipelined` parameter is disabled (the default), the ThreadBusy signal defines the flow control for the current cycle, so phase assertion depends upon that cycle's ThreadBusy values. This mode corresponds to the timing arcs in Figure 6 where the ThreadBusy generation appears at the beginning of the OCP cycle. When a `_pipelined` parameter is enabled, the ThreadBusy signal defines the

flow control for the next cycle enabling fully sequential interface behavior, where non-blocking flow control can be accomplished without combinational paths crossing the interface twice in a single cycle.

Combinational paths may still be present to enable the phase receiver to consider interface activity in the current cycle before signaling the ThreadBusy signal that affects the next cycle. This corresponds to the timing arcs in Figure 6 where ThreadBusy appears at the end of the OCP cycle. When a `_pipelined` parameter is enabled, exact flow control is not possible for the first cycle after reset is de-asserted, since the associated ThreadBusy would have to be provided while reset was asserted. When `sthreadbusy_pipelined` is enabled the master may not assert the request phase in the first cycle after reset.

The effect of these parameters is as follows:

- If `mthreadbusy_exact` is enabled, `mthreadbusy_pipelined` is disabled, and the master cannot accept a response on a thread, it must set the MThreadBusy bit for that thread to 1 in that cycle. The slave must not present a response on a thread when the corresponding MThreadBusy bit is set to 1. Any response presented by the slave on a thread that is not busy is accepted by the master in that cycle.
- If `mthreadbusy_exact` and `mthreadbusy_pipelined` are enabled and the master cannot accept a response on a thread in the next cycle, it must set the MThreadBusy bit for that thread to 1 in the current cycle. If an MThreadBusy bit was set to 1 in the prior cycle, the slave cannot present a response on a thread in the current cycle. Any response presented by the slave on a thread that was not busy in the prior cycle is accepted by the master in that cycle.
- If `sdatathreadbusy_exact` is enabled, `sdatathreadbusy_pipelined` is disabled, and the slave cannot accept a datahandshake phase on a thread, the slave must set the SDataThreadBusy bit for that thread to 1 in that cycle. The master must not present a datahandshake phase on a thread when the corresponding SDataThreadBusy bit is set to 1. Any datahandshake phase presented by the master on a thread that is not busy is accepted by the slave in that cycle.
- If `sdatathreadbusy_exact` and `sdatathreadbusy_pipelined` are enabled and the slave cannot accept a datahandshake phase on a thread in the next cycle, the slave must set the SDataThreadBusy bit for that thread to 1 in the current cycle. If an SDataThreadBusy bit was set to 1 in the prior cycle, the master cannot present a datahandshake on the corresponding thread in the current cycle. Any datahandshake presented by the master on a thread that was not busy in the prior cycle is accepted by the slave in that cycle.
- If `sthreadbusy_exact` is enabled, `sthreadbusy_pipelined` is disabled, and the slave cannot accept a command on a thread, the slave must set the SThreadBusy bit for that thread to 1 in that cycle. The master must not present a request on a thread when the corresponding SThreadBusy bit is set to 1. Any request presented by the master on a thread that is not busy is accepted by the slave in that cycle.

- If `sthreadbusy_exact` and `sthreadbusy_pipelined` are enabled and the slave cannot accept a request on a thread in the next cycle, the slave must set the `SThreadBusy` bit for that thread to 1 in the current cycle. If an `SThreadBusy` bit was set to 1 in the prior cycle, the master cannot present a request on the corresponding thread in the current cycle. Any request presented by the master on a thread that was not busy in the prior cycle is accepted by the slave in that cycle.

Sideband and Test Signals

Reset

The OCP interface provides an interface reset signal for each master and slave. At least one of these signals must be present. If both signals are present, the composite reset state of the interface is derived as the logical AND of the two signals (that is, the interface is in reset as long as one of the two resets is asserted).

Treat OCP reset signals as fully synchronous to the OCP clock, where the receiver samples the incoming reset using the rising edge of the clock and deassertion of the reset meets the receiver's timing requirements with respect to the clock. An exception to this rule exists when the assertion edge of an OCP reset signal is asynchronous to the OCP clock. This behavior handles the practice of forcing all reset signals to be combinationaly asserted for power-on reset or other hardware reset conditions without waiting for a clock edge.

Once a reset signal is sampled asserted by the rising edge of the OCP clock, all incomplete transactions, transfers and phases are terminated and both master and slave must transition to a state where there are no pending OCP requests or responses. When a reset signal is asserted asynchronously, there may be ambiguity about transactions that completed, or were aborted due to timing differences between the arrival of the OCP reset and the OCP clock.

For systems requiring precision use synchronous reset assertion, or only apply reset asynchronously if the interface is either quiescent or hung. `MReset_n` and `SReset_n` must be asserted for at least 16 cycles of the OCP clock to ensure that the master and slave reach a consistent internal state. When one or both of the reset signals are asserted in a given cycle, all other OCP signals must be ignored in that cycle. The master and slave must each be able to reach their reset state regardless of the values presented on the OCP signals. If the master or slave require more than 16 cycles of reset assertion, the requirement must be documented in the IP core specifications.

At the clock edge that all reset signals present are sampled deasserted, all OCP interface signals must be valid. In particular, it is legal for the master to begin its first request phase in the same clock cycle that reset is deasserted.

Interrupt, Error, and Core Flags

There is no specific timing associated with `SInterrupt`, `SError`, `MFlag`, `MError`, and `SFlag`. The timing of these signals is core-specific.

Status and Control

The following rules assure that control and status information can be exchanged across the OCP without any combinational paths from inputs to outputs and at the pace of a slow core.

- Control must be held steady for a full cycle after the cycle in which it has transitioned, which means it cannot transition more frequently than every other cycle. If ControlBusy was sampled active at the end of the previous cycle, Control can not transition in the current cycle. In addition, Control must be held steady for the first two cycles after reset is deasserted.
- If Control transitions in a cycle, ControlWr (if present) must be driven active for that cycle. ControlWr following the rules for Control, cannot be asserted in two consecutive cycles.
- ControlBusy allows a core to force the system to hold Control steady. ControlBusy may only start to be asserted immediately after reset, or in the cycle after ControlWr is asserted, but can be left asserted for any number of cycles.
- While StatusBusy is active, Status is a “don’t care”. StatusBusy enables a core to prevent the system from reading the current status information. While StatusBusy is active the core may not read Status. StatusBusy can be asserted at any time and be left asserted for any number of cycles.
- StatusRd is active for a single cycle every time the status register is read by the system. If StatusRd was asserted in the previous cycle, it must not be asserted in the current cycle, so it cannot transition more frequently than every other cycle.

Test Signals

Scanin and Scanout are “don’t care” while Scanctrl is inactive (but the encoding of inactive for Scanctrl is core-specific).

TestClk is “don’t care” while ClkByp is 0.

The timing of TRST_N, TCK, TMS, TDI, and TDO is specified in the IEEE 1149 standard.

Transfer Effects

A successful transfer is one that completes without error. For write-type requests without responses, there is no in-band error indication. For all other requests, a non-ERR response (that is, a DVA or FAIL response) indicates a successful transfer. The FAIL response is legal only for WriteConditional commands. This section defines the effect that a successful transfer has on a slave. The request acts on the addressed location, where the term address refers to the combination of MAddr, MAddrSpace, and MByteEnable (or MDataByteEn, if applicable). Two addresses are said to match if they are identical in all components. Two addresses are said to conflict, if the mutual exclusion (lock or monitor) logic is built to alias the two addresses into the same mutual exclusion unit. The transfer effects of each command are:

Idle

None.

Read

Returns the latest value of the addressed location on the SData field.

ReadEx

Returns the latest value of the addressed location on the SData field. Sets a lock for the initiating thread on that location. The next request on the thread that issued a ReadEx must be a Write or WriteNonPost to the matching address. Requests from other threads to a conflicting address that is locked are blocked from proceeding until the lock is released. If the ReadEx request returns an ERR response, it is slave-specific whether the lock is actually set or not.

ReadLinked

Returns the latest value of the addressed location on the SData field. Sets a reservation in a monitor for the corresponding thread on at least that location. Requests of any type from any thread to a conflicting address that is reserved are not blocked from proceeding, but may clear the reservation.

Write/WriteNonPost

Places the value on the MData field in the addressed location. Unlocks access to the matched address if locked by a ReadEx. Clears the reservations on any conflicting addresses set by other threads.

WriteConditional

If a reservation is set for the matching address and for the corresponding thread, the write is performed as it would be for a Write or WriteNonPost. Simultaneously, the reservation is cleared for all threads on any conflicting address. If no reservation is set for the corresponding thread, the write is not performed, a FAIL response is returned, and no reservations are cleared.

Broadcast

Places the value on the MData field in the addressed location that may map to more than one slave in a system-dependent way. Broadcast clears the reservations on any conflicting addresses set by other threads.

If a transfer is unsuccessful, the effect of the transfer is unspecified. Higher-level protocols must determine what happened and handle any clean-up.

The synchronization commands ReadEx / Write, ReadEx / WriteNonPost, and ReadLinked / WriteConditional have special restrictions with regard to data width conversion and partial words. In systems where these commands are sent through a bridge or interconnect that performs wide-to-narrow data width conversion between two OCP interfaces, the initiator must issue only commands within the subset of partial words that can be expressed as a single word of the narrow OCP interface. For maximum portability, single-byte synchronization operations are recommended.

Partial Word Transfers

An OCP interface may be configured to include partial word transfers by using either the MByteEn field, or the MDataByteEn field, or both.

- If neither field is present, then only whole word transfers are possible.
- If only MByteEn is present, then the partial word is specified by this field for both read type transfers and write type transfers as part of the request phase.
- If only MDataByteEn is present, then the partial word is specified by this field for write type transfers as part of the datahandshake phase, and partial word reads are not supported.
- If both MByteEn and MDataByteEn are present, then MByteEn specifies partial words for read transfers as part of the request phase, and MDataByteEn specifies partial words for write transfers as part of the datahandshake phase.

It is legal to use a request with all byte enables deasserted. Such requests must follow all the protocol rules, except that they are treated as no-ops by the slave: the response phase signals SData and SDataInfo are "don't care" for read-type commands, and nothing is written for write-type commands.

Posting Semantics

The OCP includes a basic Write command. Typically, a system designer decides what write completion model to assign to a core's write commands. If the OCP is configured to not respond to write-type commands (writeresp_enable set to 0), a posted write completion model is assumed. It is also possible to achieve a non-posted model by not accepting the command until the write completes, but this de-pipelines the OCP interface and is not recommended.

If the OCP is configured to respond to write-type commands (writeresp_enable set to 1), either a posted or a non-posted write completion model can be used. For cores that need to determine on a per-transfer basis whether a write is to be treated as posted or non-posted, the OCP provides the WriteNonPost command.

Endianness

An OCP interface by itself is inherently endian-neutral. Data widths must match between master and slave, addressing is on an OCP word granularity, and byte enables are tied to byte lanes (data bits) without tying the byte lanes to specific byte addresses.

The issue of endianness arises in the context of multiple OCP interfaces, where the data widths of the initiator of a request and the final target of that request do not match. Examples are a bridge or a more general interconnect used to connect OCP-based cores.

When the OCP interfaces differ in data width, the interconnect must associate an endianness with each transfer. It does so by associating byte lanes and byte enables of the wider OCP with least-significant word address bits of the narrower OCP. Packing rules, described in “Packing” on page 50 must also be obeyed for bursts.

OCP interfaces can be designated as little, big, both, or neutral with respect to endianness. This is specified using the protocol parameter `endian` described in “Endianness” on page 58. A core that is designated as both typically represents a device that can change endianness based upon either an internal configuration register or an external input. A core that is designated as neutral typically represents a device that has no inherent endianness. This indicates that either the association of an endianness is arbitrary (as with a memory, which traditionally has no inherent endianness) or that the device only works with full-word quantities (when `byteen` and `mdatabyteen` are set to 0).

When all cores have the same endianness, an interconnect should match the endianness of the attached cores. The details of any conversion between cores of different endianness is implementation-specific.

Burst Definition

A *burst* is a set of transfers that are linked together into a transaction having a defined address sequence and number of transfers. There are three general categories of bursts:

Imprecise bursts

Request information is given for each transfer. Length information may change during the burst.

Precise bursts

Request information is given for each transfer, but length information is constant throughout the burst.

Single request / multiple data bursts (also known as packets)

Also a precise burst, but request information is given only once for the entire burst.

To express bursts on the OCP interface, at least the address sequence and length of the burst must be communicated, either directly using the `MBurstSeq` and `MBurstLength` signals, or indirectly through an explicit constant tie-off as described in “Signal Mismatch Tie-off Rules” on page 62.

A single (non-burst) request on an OCP interface with burst support is encoded as a request with any legal burst address sequence and a burst length of 1.

The `ReadEx`, `ReadLinked`, and `WriteConditional` commands can not be used as part of a burst. The unlocking `Write` or `WriteNonPost` command associated with a `ReadEx` command also can not be used as part of a burst.

Burst Address Sequences

The relationship of the MBurstSeq encodings and corresponding address sequences are shown in Table 19. The table also indicates whether a burst sequence type is packing or not, a concept discussed on page 50.

Table 19 Burst Address Sequences

Mnemonic	Name	Address Sequence	Packing
BLCK	2D block	see below for precise definition	yes
DFLT1	custom (packed)	user-specified	yes
DFLT2	custom (not packed)	user-specified	no
INCR	incrementing	incremented by OCP word size each transfer*	yes
STRM	streaming	constant each transfer	no
UNKN	unknown	none specified	implementation specific
WRAP	wrapping	like INCR, except wrap at address boundary aligned with MBurstLength * OCP word size	yes
XOR	exclusive OR	see below for precise definition	yes

*Bursts must not wrap around the OCP address size.

The address sequence for 2-dimensional block bursts is as follows. The address sequence begins at the provided address and proceeds through a set of MBlockHeight subsequences, each of which follows the normal INCR address sequence for MBurstLength transfers. The starting address for each following subsequence is the starting address of the prior subsequence plus MBurstStride.

The address sequence for exclusive OR bursts is as follows. Let BASE be the lowest byte address in the burst, which must be aligned with the total burst size. Let FIRST_OFFSET be the byte offset (from BASE) of the first transfer in the burst. Let CURRENT_COUNT be the count of the current transfer in the burst, starting at 0. Let WORD_SHIFT be the log2 of the OCP word size in bytes. Then the current address of the transfer is $\text{BASE} \mid (\text{FIRST_OFFSET} \wedge (\text{CURRENT_COUNT} \ll \text{WORD_SHIFT}))$.

The burst address sequence UNKN is used if the address sequence is not statically known for the burst. Single request/multiple data bursts (described on page 51) with a burst address sequence of UNKN are illegal. In contrast, the DFLT1 and DFLT2 address sequences are known, but are core or system specific.

The burst address sequences BLCK, WRAP, and XOR can only be used for precise bursts. Additionally, the burst sequences WRAP and XOR can only have a power-of-two burst length and a data width that is a power-of-two number of bytes.

Not all masters and slaves need to support all burst sequences. A separate protocol parameter described in “Optional Burst Sequences” on page 55 is provided for each burst sequence to indicate support for that burst sequence.

Byte Enable Restrictions

Burst address sequences STRM and DFLT2 must have at least one byte enable asserted for each transfer in the burst. Bursts with the STRM address sequence must have the same byte enable pattern for each transfer in the burst.

Packing

Packing allows the system to make use of the burst attributes to improve the overall data transfer efficiency in the face of multiple OCP interfaces of different data widths. For example, if a bridge is translating a narrow OCP to a wide OCP, it can aggregate (or pack) the incoming narrow transfers into a smaller number of outgoing wide transfers. Burst address sequences are classified as either packing, or not packing.

For burst address sequences that are packing, the conversion between different OCP data widths is achieved through aggregation or splitting. Narrow OCP words are collected together to form a wide OCP word. A wide OCP word is split into several narrow OCP words. The byte-specific portion of MDataInfo and SDataInfo is aggregated or split with the data. The transfer-specific portion of MDataInfo and SDataInfo is unaffected. The packing and unpacking order depends on endianness as described on page 47.

For burst address sequences that are not packing, conversion between different OCP data widths is achieved using padding and stripping. A narrow OCP word is padded to form a wide OCP word with only the relevant byte enables turned on. A wide OCP word is stripped to form a narrow OCP word. The byte-specific portion of MDataInfo and SDataInfo is zero-padded or stripped with the data. The transfer-specific portion of MDataInfo and SDataInfo is unaffected. Width conversion can be performed reliably only if the wide OCP interface has byte enables associated with it. For wide to narrow conversion the byte enables are restricted to a subset that can be expressed within a single word of the narrow OCP interface.

Since the address sequence of DFLT1 is user-specified, the behavior of DFLT1 bursts through data width conversion is implementation-specific.

Burst Length, Precise and Imprecise Bursts

The MBurstLength field indicates the number of transfers in the burst.

Precise bursts (MBurstPrecise set to 1)

MBurstLength must be held constant throughout the burst, so the exact burst length can be obtained from the first transfer. A precise burst is completed by the transfer of the correct number of OCP words. Precise bursts are recommended over imprecise bursts because they allow for increased hardware optimization.

Imprecise bursts (MBurstPrecise set to 0)

MBurstLength can change throughout the burst, and indicates the current best guess of the number of transfers left in the burst (including the current one). An imprecise burst is completed by an MBurstLength of 1.

Constant Fields in Bursts

MCmd, MAddrSpace, MConnID, MBurstPrecise, MBurstSingleReq, MBurstSeq, MAtomicLength, MBlockHeight, MBlockStride, and MReqInfo must all be held steady by the master for every transfer in a burst, regardless of whether the burst is precise or imprecise. If possible, slaves should hold SRespInfo steady for every transfer in a burst.

Atomicity

When interleaving requests from different initiators on the way to or at the target, the master uses MAtomicLength to indicate the number of OCP words within a burst that must be kept together as an atomic quantity. If MAtomicLength is greater than the actual length of the burst, the atomicity requirement ends with the end of the burst. Specifying atomicity requirements explicitly is especially useful when multiple OCP interfaces are involved that have different data widths.

For master cores, it is best to make the atomic size as small as required and, if possible, to keep the groups of atomic words address-aligned with the group size.

Single Request / Multiple Data Bursts (Packets)

MBurstSingleReq specifies whether a burst can be communicated using a single request / multiple data protocol. When MBurstSingleReq is 0, each request has a single data word associated with it. When MBurstSingleReq is 1, each request may have multiple data words associated with it, according to the values of MBurstLength and MBlockHeight. MBurstSingleReq may be set to 1 only if MBurstPrecise is set to 1. In addition, if any write-type commands are enabled, datahandshake must be set to 1.

When MBurstSingleReq is set to 1, write type transfers have MBurstLength * height datahandshake phases and a single response phase (if writeresp_enable=1) per request; while read-type transfers have MBurstLength * height response phases per request as shown in Table 18 on page 40. The height is MBlockHeight for BLCK address sequences, and 1 for all others.

For write type transfers when MBurstSingleReq is set to 1 and the MDataByteEn field is present, that field in each data transfer phase specifies the partial word pattern for the phase. When MBurstSingleReq is set to 1 and the MDataByteEn field is not present, the MByteEn pattern of the request phase applies to all data transfer phases.

For read type transfers when MBurstSingleReq is set to 1, the MByteEn field specifies the byte enable pattern that is applied to all data transfers in the burst.

MReqLast, MDataLast, SRespLast

Optional signals MReqLast, MDataLast, and SRespLast provide redundant information that indicates the last request, datahandshake, and response phase in a burst, respectively. These signals are provided as a convenience to the recipient of the signal. To avoid separate counting mechanisms to track bursts, cores that have the information available internally are encouraged to provide it at the OCP interface.

MReqLast is 0 for all request phases in a burst except the last one. MReqLast is 1 for the last request phase in a burst, for single request / multiple data bursts, and for single requests.

MDataLast is 0 for all datahandshake phases in a burst except the last one. MDataLast is 1 for the last datahandshake phase in a burst and for the only datahandshake phase of a single request.

SRespLast is 0 for all response phases in a burst except the last one. SRespLast is 1 for the last response phase in a burst, for the response to a write-type single request / multiple data burst, and for the response to a single request.

MReqRowLast, MDataRowLast, SRespRowLast

For the BLCK burst address sequence, the optional signals MReqRowLast, MDataRowLast, and SRespRowLast identify the last request, datahandshake, and response phase in a row. The last phase in a burst is always considered the last phase in a row, and BLCK burst sequences reach the end of a row every MBurstLength phases (at the end of each INCR sub-sequence, see page 64). To avoid separate counting mechanisms needed to track BLCK burst sequences, cores that have the end of row information available should provide it at the OCP interface.

For all request phases in a non-BLCK burst except the last one, MReqRowLast is 0. MReqRowLast is 0 for every request phase in a BLCK burst sequence that is not an integer multiple of MBurstLength. MReqRowLast is 1 for:

- The last request phase in a burst including:
 - The only request phase in a single request/multiple data burst
 - The only request phase in a single word request
- Every request phase in a BLCK burst sequence that is an integer multiple of MBurstLength

For all datahandshake phases in a BLCK burst except the last one, MDataRowLast is 0. MDataRowLast is 0 for every datahandshake phase in a BLCK burst sequence that is not an integer multiple of MBurstLength. MDataRowLast is 1 for:

- The last datahandshake phase in a burst including the only datahandshake phase of a single word request
- Every datahandshake phase in a BLCK burst sequence that is an integer multiple of MBurstLength

For all response phases in a BLCK burst except the last one, SRespRowLast is 0. SRespRowLast is 0 for every response phase in a BLCK burst sequence that is not an integer multiple of MBurstLength. SRespRowLast is 1 for:

- The last response phase in a burst including:
 - The only response phase in a write-type single request/multiple data burst
 - The only response phase in a single word request
- Every response phase in a BLCK burst sequence that is an integer multiple of MBurstLength

Tags

Tags allow out-of-order return of responses and out-of-order commit of write data.

A master drives a tag on MTagID during the request phase. The value of the tag is determined by the master and may or may not convey meaning beyond ordering to the slave. For write transactions with data handshake enabled, the master repeats the same tag on MDataTagID during the datahandshake phase. For read transactions and writes with responses the slave returns the tag of the corresponding request on STagID while supplying the response. The same tag must be used for an entire transaction.

Ordering Restrictions

The sequence of requests by the master determines the initial ordering of tagged transactions. For tagged write transactions with datahandshake enabled, the datahandshake phase must observe the same order as the request phase. The master cannot interleave requests or datahandshake phases with different tags within a transaction.

Tag values can be re-used for multiple outstanding transactions. Slaves are responsible for committing write data and sending responses for multiple transactions that have the same tag, in order.

Responses with different tags can be returned in any order for all commands that have responses. Responses with the same tag must remain in order with respect to one another. Responses that are part of the same transaction must stay together, up to the `tag_interleave_size` (see “Burst Interleaving with Tags” on page 58). Beyond the `tag_interleave_size`, responses with different tags can be interleaved. This allows for blocks of responses corresponding to `tag_interleave_size` from one burst to be interleaved with blocks of responses from other bursts.

Responses to requests that target overlapping addresses (as determined by MAddrSpace, MAddr, and MByteEn) are never re-ordered with respect to one another. Similarly, responses to requests that are issued with MTagInOrder asserted are also never reordered with respect to one another. The value returned on STagInOrder with the slave's response must match the value provided on MTagInOrder with the master's request.

When tags are in use and writes are not posted, write response order indicates the order that writes are committed by the slave. When writes are posted, the response to a tagged read can only be guaranteed to be ordered with respect to writes to overlapping addresses. This is different than non-tagged versions of OCP interfaces, where read responses are sometimes used to ensure completion of all earlier posted writes.

Threads and Connections

Using multiple threads, it is possible to support concurrent activity, and out-of-order completion of transfers. All transfers within a given thread must either remain strictly ordered or follow the tag ordering rules, but there are no ordering rules for transfers that are in different threads. Mapping of individual requests and responses to threads is handled through the MThreadID and SThreadID fields respectively. If datahandshake has been enabled when multiple threads are present, there must also be an MDataThreadID field to annotate the datahandshake phase. If datahandshake is set to 1 and sdatathreadbusy is set to 0, the order of datahandshake phases must follow the order of request phases across all threads. If sdatathreadbusy is set to 1, the request order and datahandshake order are independent across threads.

The use of thread IDs allows two entities that are communicating over an OCP interface to assign transfers to particular threads. If one of the communicating entities is itself a bridge to another OCP interface, the information about which transfers are part of which thread must be maintained by the bridge, but the actual assignment of thread IDs is done on a per-OCP-interface basis. There is no way for a slave on the far side of a bridge to extract the original thread ID unless the slave design comprehends the characteristics of the bridge.

Use connections whenever source thread information about a request must be sent end-to-end from master to slave. Any bridges in the path between the end-to-end partners preserve the connection ID, even as thread IDs are re-assigned on each OCP interface in the path. The MConnID field transfers the connection ID during the request phase. Since this establishes the mapping onto a thread ID, the other phases do not require a connection ID but are unambiguous with only a thread ID.

The SThreadBusy, SDataThreadbusy, and MThreadBusy signals are used to indicate that a particular thread is busy. The protocol parameters sthreadbusy_exact, sdatathreadbusy_exact, and mthreadbusy_exact can be used to force precise semantics for these signals and assure that a multi-threaded OCP interface never blocks. For more information, see “Ungrouped Signals” on page 42.

OCP Configuration

This section describes configuration options that control interface capabilities.

Protocol Options

Optional Commands

Not all devices support all commands. Each command in Table 20 has an enabling parameter to indicate if that command is supported.

Table 20 Command Enabling Parameters

Command	Parameter
Broadcast	broadcast_enable
Read	read_enable
ReadEx	readex_enable
ReadLinked and WriteConditional	rdlwrc_enable
Write	write_enable
WriteNonPost	writenonpost_enable

The following conditions apply to command support:

- A master with one of these options set to 0 must not generate the corresponding command.
- A slave with one of these options set to 0 cannot service the corresponding command.
- At least one of the command enables must be set to 1.
- `writenonpost_enable` can be set to 1 only if `writeresp_enable` is set to 1.
- `rdlwrc_enable` can be set to 1 only if `writeresp_enable` is set to 1.
- If any read-type command is enabled or `writeresp_enable` is set to 1, `resp` must be set to 1.
- If `readex_enable` is set to 1, `write_enable` or `writenonpost_enable` must be set to 1.

Optional Burst Sequences

Not all masters and slaves need to support all burst address sequences. Table 21 lists the protocol parameter for each burst sequence. A master that has the parameter set to 1 may generate the corresponding burst sequence. A slave

that has the parameter set to 1 can service the corresponding burst sequence. If MBurstSeq is disabled and tied off to a constant value, the corresponding burst sequence parameter must be enabled and all others disabled. If MBurstSeq is enabled at least one of the burst sequence parameters must be enabled.

Table 21 Burst Sequence Parameters

Burst Sequence	Parameter
BLCK	burstseq_blk_enable
DFLT1	burstseq_dflt1_enable
DFLT2	burstseq_dflt2_enable
INCR	burstseq_incr_enable
STRM	burstseq_strm_enable
UNKN	burstseq_unkn_enable
WRAP	burstseq_wrap_enable
XOR	burstseq_xor_enable

The BLCK burst sequence can only be enabled if both MBlockHeight and MBlockStride are included in the interface or tied off to non-default values. For additional information describing bursts, see “Burst Definition” on page 48.

Byte Enable Patterns

Not all devices support all allowable byte enable patterns. The `force_aligned` parameter limits byte enable patterns on MByteEn and MDataByteEn to be power-of-two in size and aligned to that size. The byte enable pattern of all 0s is explicitly included in the legal `force_aligned` patterns.

- A master with this option set to 1 must not generate any byte enable patterns that are not force aligned.
- A slave with this option set to 1 cannot handle any byte enable patterns that are not force aligned.

`force_aligned` can be set to 1 only if `data_width` is set to a power-of-two value.

Burst Alignment

The `burst_aligned` parameter provides information about the length and alignment of INCR bursts issued by a master and can be used to optimize the system. Setting `burst_aligned` to 1 requires all INCR bursts to:

- Have an exact power-of-two number of transfers
- Have their starting address aligned with their total burst size

- Be issued as precise bursts.

The `burst_aligned` parameter does not apply to the INCR subsequences within BLCK burst sequences.

Flow Control Options

To permit the SThreadBusy and MThreadBusy signals to guarantee a non-blocking, multi-threaded OCP interface, the `sthreadbusy_exact` and `mthreadbusy_exact` parameters require strict semantics. See “Ungrouped Signals” on page 42 for a precise definition of these parameters. Table 22 describes the legal combinations of phase handshake signals.

Table 22 Request Phase Without Datahandshake

cmdaccept	sthreadbusy	sthreadbusy_exact	Explanation
0	0	0	Legal: no flow control
0	0	1	Illegal: sthreadbusy_exact must be 0 when sthreadbusy is 0
0	1	0	Illegal: no real flow control
0	1	1	Legal: non-blocking flow control
1	0	0	Legal: blocking flow control
1	0	1	Illegal: sthreadbusy_exact must be 0 when sthreadbusy is 0
1	1	0	Legal: blocking flow control with hints
1	1	1	Illegal: since SCmdAccept is present flow control cannot be exact

When `datahandshake` is set to 1, the preceding rules for `cmdaccept`, `sthreadbusy`, and `sthreadbusy_exact` also apply to `dataaccept`, `sdatathreadbusy`, and `sdatathreadbusy_exact`. In addition, blocking and non-blocking flow control must not be mixed for the request and datahandshake phase. A phase using no flow control can be mixed with phases using either blocking or non-blocking type flow control. The legal combinations are shown in Table 23.

Table 23 Request Phase with Datahandshake

		Datahandshake Phase Flow Control		
		None	Blocking	Non-blocking
Request Phase Flow Control	None	Legal	Legal ¹	Legal
	Blocking	Legal ²	Legal	Illegal
	Non-blocking	Legal	Illegal	Legal ³

¹ Only legal if `reqdata_together` is set to 0.² Only legal if `reqdata_together` is set to 0. In addition the master must not assert the `datahandshake` phase until after the associated request phase has been accepted.³ Only legal if `sthreadbusy_pipelined` and `sdatathreadbusy_pipelined` are both set to the same value.

The preceding rules for the request phase using `cmdaccept`, `sthreadbusy`, and `sthreadbusy_exact` also apply to the response phase for `respaccept`, `mthreadbusy`, and `mthreadbusy_exact`.

Endianness

The `endian` parameter specifies the endianness of a core. The behavior of each endianness choice is summarized in Table 24.

Table 24 Endianness

Endianness	Description
little	core is little-endian
big	core is big-endian
both	core can be either big or little endian, depending on its static or dynamic configuration (e.g. CPUs)
neutral	core has no inherent endianness (e.g. memories, cores that deal only in OCP words)

As far as OCP is concerned, little endian means that lower addresses are associated with lower numbered data bits (byte lanes), while big endian means that higher addresses are associated with lower numbered data bits (byte lanes). This becomes significant when packing is concerned (see “Packing” on page 50). In addition, for non-power-of-2 data widths, tie-off padding is always added at the most significant end of the OCP word. See “Endianness” on page 47 for additional information.

Burst Interleaving with Tags

When `tags > 1`, the `tag_interleave_size` parameter limits the interleaving permitted for responses associated with packing burst sequences.

`tag_interleave_size` indicates the size of a power-of-two, aligned data block (in OCP words) within which there can be no interleaving of responses from packing bursts with different tags. Interleaving of non-packing burst sequence responses is not limited by `tag_interleave_size`, and interleaving of packing burst responses is allowed whenever the next response would cross the data block boundary, regardless of whether a full data block of responses has been returned.

Restricting interleaving opportunities for packing burst responses reduces the storage required for width conversion when multiple tags are present. For slaves, enabling the parameter restricts the aligned boundary within which the slave interleaves responses with different tags. For masters, the parameter gives the minimum aligned boundary at which the master can tolerate interleaving of responses with different tags.

When `tag_interleave_size` is set to 1, interleaving is permitted at the OCP word level and is unrestricted. When `tag_interleave_size` is set to the maximum power-of-two burst size permitted on `MBurstLength`, no interleaving between packing burst sequence responses with different tags is permitted.

Phase Options

The `datahandshake` parameter allows write data to have a handshake interface separate from the request group.

Datahandshake

If `datahandshake` is set to 1, the `MDataValid` and optionally the `SDataAccept` signals are added to the OCP interface, a separate `datahandshake` phase is added, and the `MData` and `MDataInfo` fields are moved from the request group to the `datahandshake` group. `Datahandshake` can be set to 1 only if at least one write-type command is enabled.

Request and Data Together

While `datahandshake` is required for OCP interfaces that are capable of communicating single request / multiple data bursts, a fully separated `datahandshake` may be overkill for some cores. The parameter `reqdata_together` is used to specify that the request and `datahandshake` phases of the first transfer in a single request / multiple data write-type burst begin and end together.

A master with `reqdata_together` set to 1 must present the request and first write data word in the same cycle and can expect that the slave will accept them together. If `sthreadbusy_exact` and `sdatathreadbusy_exact` are both set to 1, this implies that a request and first write data can be presented only when both `SThreadBusy` and `SDataThreadBusy` for the corresponding thread are 0. A slave with `reqdata_together` set to 1 must accept the request and first write data word in the same cycle and can expect that they will be presented together.

The parameter `reqdata_together` can only be set to 1 if `burstsinglereq` is set to 1, or `burstsinglereq` is set to 0 and `MBurstSingleReq` is tied off to 1.

If both `reqdata_together` and `burstsinglereq` are set to 1, the master must present the request and associated write data word together for each transfer in any multiple request / multiple data writes it issues. The slave must accept both request and write data together for all such transfers.

Write Responses

To configure the OCP to include response phases for write-type requests, use the `writeresp_enable` parameter.

- If responses are not enabled on writes (`writeresp_enable` set to 0), then all write commands complete on command acceptance, and the `WriteNonPost`, `WriteConditional`, and `ReadLinked` commands are not allowed.

- If responses are enabled (`writeresp_enable` set to 1), writes may follow either a posted or non-posted model depending on when the response is returned. For this case, `resp` must be set to 1.

Signal Options

The configuration parameters described in “Signal Configuration” on page 29, not only configure the corresponding signal into the OCP interface, but also enable the function. For example, if the `burstseq` and `burstlength` parameters are enabled the `MBurstSeq` and `MBurstLength` fields are added and the interface also supports burst extensions as described in “Burst Definition” on page 48.

Minimum Implementation

A minimal OCP implementation must support at least the basic OCP dataflow signals. OCP-interoperable masters and slaves must support the command type `Idle` and at least one other command type.

If the `SResp` field is present in the OCP interface, OCP-interoperable masters and slaves must support response types `NULL` and `DVA`. The `ERR` response type is optional and should only be included if the OCP-interoperable slave has the ability to report errors. All OCP masters must be able to accept the `ERR` response. If `rdlwrc_enable` is set to 1, the `FAIL` response type must be supported by OCP masters and slaves.

OCP Interface Interoperability

Two devices connected together each have their own OCP configuration. The two interfaces are only interoperable (allowing the two devices to be connected together and communicate using the OCP protocol semantics) if they are interoperable at the core, protocol, phase, and signal levels.

1. At the core level:
 - One interface must act as master and the other as slave.
 - If system signals are present, one interface must act as core and the other as system.
2. At the protocol level, the following conditions determine interface interoperability:
 - If the slave has `read_enable` set to 0, the master must have `read_enable` set to 0, or it must not issue `Read` commands.
 - If the slave has `readex_enable` set to 0, the master must have `readex_enable` set to 0, or it must not issue `ReadEx` commands.
 - If the slave has `rdlwrc_enable` set to 0, the master must have `rdlwrc_enable` set to 0, or it must not issue either `ReadLinked` or `WriteConditional` commands.

- If the slave has `write_enable` set to 0, the master must have `write_enable` set to 0, or it must not issue Write commands.
- If the slave has `writenonpost_enable` set to 0, the master must have `writenonpost_enable` set to 0, or it must not issue WriteNonPost commands.
- If the slave has `broadcast_enable` set to 0, the master must have `broadcast_enable` set to 0, or it must not issue Broadcast commands.
- If the slave has `burstseq_blk_enable` set to 0, the master must have `burstseq_blk_enable` set to 0, or it must not issue BLCK bursts.
- If the slave has `burstseq_incr_enable` set to 0, the master must have `burstseq_incr_enable` set to 0, or it must not issue INCR bursts.
- If the slave has `burstseq_strm_enable` set to 0, the master must have `burstseq_strm_enable` set to 0, or it must not issue STRM bursts.
- If the slave has `burstseq_dflt1_enable` set to 0, the master must have `burstseq_dflt1_enable` set to 0, or it must not issue DFLT1 bursts.
- If the slave has `burstseq_dflt2_enable` set to 0, the master must have `burstseq_dflt2_enable` set to 0, or it must not issue DFLT2 bursts.
- If the slave has `burstseq_wrap_enable` set to 0, the master must have `burstseq_wrap_enable` set to 0, or it must not issue WRAP bursts.
- If the slave has `burstseq_xor_enable` set to 0, the master must have `burstseq_xor_enable` set to 0, or it must not issue XOR bursts.
- If the slave has `burstseq_unkn_enable` set to 0, the master must have `burstseq_unkn_enable` set to 0, or it must not issue UNKN bursts.
- If the slave has `force_aligned`, the master has `force_aligned` or it must limit itself to aligned byte enable patterns.
- Configuration of the `mdatabyteen` parameter is identical between master and slave.
- If the slave has `burst_aligned`, the master has `burst_aligned` or it must limit itself to issue all INCR bursts using `burst_aligned` rules.
- If the interface includes `SThreadBusy`, the `sthreadbusy_exact` and `sthreadbusy_pipelined` parameters are identical between master and slave.

- If the interface includes MThreadBusy, the `mthreadbusy_exact` and `mthreadbusy_pipelined` parameter are identical between master and slave.
 - If the interface includes SDataThreadBusy, the `sdatathreadbusy_exact` and `sdatathreadbusy_pipelined` parameters are identical between master and slave.
 - All combinations of the `endian` parameter between master and slave are interoperable as far as the OCP interface is concerned. There may be core-specific issues if the endianness is mismatched.
 - If `tags > 1`, the master's `tag_interleave_size` is smaller than or equal to the slave's `tag_interleave_size`.
3. At the phase level the two interfaces are interoperable if:
- Configuration of the `datahandshake` parameter is identical between master and slave.
 - Configuration of the `writeresp_enable` parameter is identical between master and slave.
 - Configuration of the `reqdata_together` parameter is identical between master and slave.
4. At the signal level, two interfaces are interoperable if:
- `data_width` is identical for master and slave, or if one or both `data_width` configurations are not a power-of-two, if that `data_width` rounded up to the next power-of-two is identical for master and slave.
 - The master and slave both have `mreset` or `sreset` set to 1.
 - If the master has `mreset` set to 1, the slave has `mreset` set to 1.
 - If the slave has `sreset` set to 1, the master has `sreset` set to 1.
 - Both master and slave have `tags` set to ≥ 1 or if only one core's `tags` parameter is set to 1, the other core behaves as though MTagInOrder were asserted for every request.
 - The tie-off rules, described in the next section are observed for any mismatch at the signal level for fields other than MData and SData.

Signal Mismatch Tie-off Rules

There are two types of signal mismatches: both interfaces may have configured the signal, but to different widths or only one interface may have configured the signal.

Width mismatch for all fields other than MData and SData is handled through a set of signal tie-off rules. The rules state whether a master and slave that are mismatched in a particular field width configuration are interoperable, and if so how to connect them by tying off the mismatched signals.

If there is a width mismatch between master and slave for a particular signal configuration the following rules apply:

- If there are more outputs than inputs (the driver of the field has a wider configuration than the receiver of the field) the low-order output bits are connected to the input bits, and the high-order output bits are lost. The interfaces are interoperable if the sender of the field explicitly limits itself to encodings that only make use of the bits that are within the configuration of the receiver of the field.
- If there are more inputs than outputs (the driver of field has a narrower configuration than the receiver of the field) the low-order input bits are connected to the output bits, and the high-order input bits are tied to logical 0. The interfaces are always interoperable, but only a portion of the legal encodings are used on that field.

If one of the cores has a signal configured and the other does not, the following rules apply:

- If the core that would be the driver of the field does not have the field configured, the input is tied off to the constant specified in the driving core's configuration, or if no constant tie-off is specified, to the default tie-off constant (see Table 13 on page 30). The interfaces are interoperable if the encodings supported by the receiver's configuration of the field include the tie-off constant.
- If the core that would be the receiver of the field does not have the field configured, the output is lost. The receiver of the signal must behave as though in every phase it were receiving the tie-off constant specified in its configuration, or lacking a constant tie-off, the default tie-off constant (see Table 13 on page 30). The interfaces are interoperable if the driver of the signal can limit itself to only driving the tie-off constant of the receiver.
- If only one core has the EnableClk signal configured, the interfaces are interoperable only when the EnableClk signal is asserted, matching the tie-off value of the core that has `enableclk=0`.

If neither core has a signal configured, the interfaces are interoperable if both cores have the same tie-off constant, where the tie-off constant is either explicitly specified, or if no constant tie-off is specified explicitly, is the default tie-off (see Table 13 on page 30).

While the tie-off rules allow two mismatched cores to be connected, this may not be enough to guarantee meaningful communication, especially when core-specific encodings are used for signals such as `MReqInfo`.

As the previous rules suggest, specifying core specific tie-off constants that are different than the default tie-offs for a signal (see Table 13 on page 30) makes it less likely that the core will be interoperable with other cores.

Configuration Parameter Defaults

To assure OCP interface interoperability between a master and a slave requires complete knowledge of the OCP interface configuration of both master and slave. This is achieved by a combination of (a) requiring some parameters to be explicitly specified for each core, and (b) defining defaults that are used when a parameter is not explicitly specified for a core.

Table 25 lists all configuration parameters. For parameters that do not need to be specified, a default value is listed, which is used whenever an explicit parameter value is not specified. Certain parameters are always required in certain configurations, and for these no default is specified.

Table 25 Configuration Parameter Defaults

Type	Parameter	Default
Protocol	broadcast_enable	0
	burst_aligned	0
	burstseq_blk_enable	0
	burstseq_dfft1_enable	0
	burstseq_dfft2_enable	0
	burstseq_incr_enable	1
	burstseq_strm_enable	0
	burstseq_unkn_enable	0
	burstseq_wrap_enable	0
	burstseq_xor_enable	0
	endian	little
	force_aligned	0
	mthreadbusy_exact	0
	rdlwrc_enable	0
	read_enable	1
	readex_enable	0
	sdatathreadbusy_exact	0
	sthreadbusy_exact	0
	tag_interleave_size	1
	write_enable	1
	writenonpost_enable	0
Phase	datahandshake	0
	reqdata_together	0
	writeresp_enable	0

Type	Parameter	Default
Signal (Dataflow)	addr	1
	addr_width	No default - must be explicitly specified if addr is set to 1
	addrspace	0
	addrspace_width	No default - must be explicitly specified if addrspace is set to 1
	atomiclength	0
	atomiclength_width	No default - must be explicitly specified if atomiclength is set to 1
	blockheight	0
	blockheight_width	No default - must be explicitly specified if blockheight is set to 1
	blockstride	0
	blockstride_width	No default - must be explicitly specified if blockstride is set to 1
	burstlength	0
	burstlength_width	No default - must be explicitly specified if burstlength is set to 1
	burstprecise	0
	burstseq	0
	burstsinglereq	0
	byteen	0
	cmdaccept	1
	connid	0
	connid_width	No default - must be explicitly specified if connid is set to 1
	dataaccept	0
	datalast	0
	datrowalast	0
	data_width	No default - must be explicitly specified if mdata or sdata is set to 1
	enableclk	0
	mdata	1
	mdatabyteen	0
	mdatainfo	0

Type	Parameter	Default
Signal (Dataflow)	mdatainfo_width	No default - must be explicitly specified if mdatainfo is set to 1
	mdatainfobyte_width	
	mthreadbusy	0
	mthreadbusy_pipelined	0
	reqinfo	0
	reqinfo_width	No default - must be explicitly specified if reqinfo is set to 1
	reqlast	
	reqrowlast	0
	resp	1
	respaccept	0
	respinfo	0
	respinfo_width	No default - must be explicitly specified if respinfo is set to 1
	resplast	
	resprowlast	0
	sdata	1
	sdatainfo	0
	sdatainfo_width	No default - must be explicitly specified if sdatainfo is set to 1
	sdatainfobyte_width	
	sdatathreadbusy	0
	sdatathreadbusy_pipelined	0
	stthreadbusy	0
	stthreadbusy_pipelined	0
	tags	1
	taginorder	0
	threads	1

Type	Parameter	Default
Signal (Sideband)	control	0
	controlbusy	0
	control_width	No default - must be explicitly specified if control is set to 1
	controlwr	0
	interrupt	0
	merror	0
	mflag	0
	mflag_width	No default - must be explicitly specified if mflag is set to 1
	mreset	No default - must be explicitly specified
	serror	0
	sflag	0
	sflag_width	No default - must be explicitly specified if sflag is set to 1
	sreset	No default - must be explicitly specified
	status	0
	statusbusy	0
	statusrd	0
	status_width	No default - must be explicitly specified if status is set to 1
Signal (Test)	clkctrl_enable	0
	jtag_enable	0
	jtagrst_enable	0
	scanctrl_width	0
	scanport	0
	scanport_width	No default - must be explicitly specified if scanport is set to 1

5 *Interface Configuration File*

The interface configuration file describes a group of signals, called a bundle. For OCP interfaces, the bundle is pre-defined, and no interface configuration file is required. If you are using an interface other than OCP in your core RTL configuration file, the interface configuration file is required.

Name the file <bundle-name>_intfc.conf where bundle-name is the name given to the bundle that is being defined in the file.

Lexical Grammar

The lexical conventions used in the interface configuration file are:

<name> : (<letter> | '_') (<letter> | '_' | <digit>)*

<letter> : 'a' .. 'z' | 'A' .. 'Z'

<digit> : '0' .. '9'

<number> : <integer> | <float>

<integer> : <digit>+

<float> : <mantissa> [<exponent>]

<mantissa>: (<integer> '.') | ('.' <integer>) | (<integer> '.' <integer>)

<exponent>: ('e' | 'E') ['+' | '-'] <integer>

Syntax

The interface configuration file is written using standard Tcl syntax. Syntax is described using the following conventions:

- [] optional construct
- | or, alternate constructs
- * zero or more repetitions
- + one or more repetitions
- < > enclose names of syntactic units
- () are used for grouping
- { } are part of the format and are required. An open brace must always appear on the same line as the statement
- # comments

The syntax of the interface configuration file is:

```
version <version_string>
bundle <bundle_name> [revision <revision_string>] {<bundle_stmt>+}
```

where:

```
<bundle_stmt>:
    interface_types <interface_type-name>+
    net <net_name> {<net_stmt>*}
    proprietary <vendor_code> <organization_name>
        {<proprietary_statements>}
<net_stmt>:
    direction (input|output|inout)+
    width <number-of-bits>
    vhdl_type <type-string>
    type <net-type>
    proprietary <vendor_code> <organization_name>
        {<proprietary_statements>}
```

The file must contain a single version statement followed by a single bundle statement. The bundle statement must contain exactly one interface_types statement, and one or more net statements. Each net statement must contain exactly one direction statement and may contain additional statements of other types.

version

The version statement identifies the version of the interface configuration file format. The version string consists of major and minor version numbers separated by a decimal. The current version is 4.5.

bundle

This statement is required and indicates that a bundle is being defined instead of a core or a chip. Make the bundle-name the same name as the one used in the interface configuration file name. Use the bundle_name ocp for OCP 1.0 bundles, and ocp2 for OCP 2.x bundles. The optional revision_string identifies a specific revision for the bundle. If not provided, the revision_string defaults to 0. The pre-defined ocp and ocp2 bundles

both use the default `revision_string` to refer to the 1.0 and 2.0 versions of the *OCP Specification*. Set `revision_string` to 2 to refer to the 2.2 version of the *OCP Specification*.

`interface_types`

The `interface_types` statement lists the legal values for the interface types associated with the bundle. Interface types are used by the toolset in conjunction with the `direction` statement to determine whether an interface uses a net as an input or output signal. This statement is required and must have at least one type defined.

Predefined interface types for OCP bundles are `slave`, `master`, `system_slave`, `system_master`, and `monitor`. These are explained in Table 15 on page 33.

`net`

The `net` statement defines the signals that comprise the bundle. There should be one `net` statement for each signal that is part of the bundle. A net can also represent a bus of signals. In this case the net width is specified using the `width` statement. If no `width` statement is provided, the net width defaults to one. A bundle is required to contain at least one net. The net-name field is the same as the one used in the net-name field of the port statements in the core rtl file described in Chapter 6.

`proprietary`

For a description, see “Proprietary Statement” on page 83.

`direction`

The `direction` statement indicates whether the net is an input, output, or inout. This field is required and must have as many direction-values as there are interface types. The order of the values must duplicate the order of the interface types in the `interface_types` statement. The legal values are `input`, `output`, and `inout`.

`vhdl_type`

By default VHDL signals and ports are assumed to be `std_logic` and `std_logic_vector`, but if you have ports on a core that are of a different type, the `vhdl_type` command can be used on a net. This type will be used only when `soccomp` is run with the `design_top=vhdl` option to produce a VHDL top-level netlist.

`type`

The `type` statement specifies that a net has special handling needs for downstream tools such as synthesis and layout. Table 23 shows the allowed `<net-type>` options. If no `<net-type>` is specified, `normal` is assumed.

Table 26 *net-type Options*

<net-type>	Description
<code>clock</code>	clock net
<code>clock_sample</code>	clock sample net
<code>jtag_tck</code>	JTAG test clock
<code>jtag_tdi</code>	JTAG test data in

<net-type>	Description
jtag_tdo	JTAG test data out
jtag_tms	JTAG test mode select
jtag_trstn	JTAG test logic reset
normal	default for nets without special handling needs
reset	reset net
scan_enable	scan enable net, serves as mode control between functional and scan data inputs
scan_in	scan input net
scan_out	scan output net
test_mode	test mode net, puts logic into a special mode for use during production testing

proprietary

For a description, see “Proprietary Statement” on page 83.

The following example defines an SRAM interface. The bundle being defined is called sram16.

```
bundle "sram16" {  
  
    # Two interface types are defined, one is labeled  
    # "controller" and the other is labeled "memory"  
    interface_types controller memory  
  
    # A net named Address is defined to be part of this bundle.  
    net "Address" {  
  
        # The direction of the "Address" net is defined to be  
        # "output" for interfaces of type "controller" and "input"  
        # for interfaces of type "memory".  
        direction output input  
  
        # The width statement indicates that there are 14 bits in  
        # the "Address" net.  
        width 14  
    }  
    net "WData" {  
        direction output input  
        width 16  
    }  
    net "RData" {  
        # The direction of the "RData" net is defined to be  
        # "input" for bundle of type "controller" and "output" for  
        # bundles of type "memory".  
        direction input output  
        width 16  
    }  
}
```



```
    net "We_n" {  
        direction output input  
    }  
    net "Oe_n" {  
        direction output input  
    }  
    net "Reset" {  
        direction output input  
        type reset  
    }  
# close the bundle  
}
```


6 Core RTL Configuration File

The required core RTL configuration file provides a description of the core and its interfaces. The name of the file needs to be <corename>_rtl.conf, where corename is the name of the module to be used. For example, the file defining a core named uart must be called uart_rtl.conf.

For a description of the lexical grammar, see page 69.

Syntax

The core RTL configuration file is written using standard Tcl syntax. Syntax is described using the following conventions:

- [] optional construct
- | or, alternate constructs
- * zero or more repetitions
- + one or more repetitions
- <> enclose names of syntactic units
- () are used for grouping
- { } are part of the format and are required. An open brace must always appear on the same line as the statement
- # comments

The syntax for the core RTL configuration file is:

```
version <version_string>
module <core_name> {<core_stmt>+}
```

core_name is the name of the core being described and:

```
<core_stmt>:
| icon <file_name>
| core_id <vendor_code> <core_code> <revision_code>
```

```
[<description>]
| interface <interface_name> bundle <bundle_name> [revision
| <revision_string>]
|   [{<interface_body>*}]
|   addr_region <name> {<addr_region_body>*}
|   proprietary <vendor_code> <organization_name>
|     {<proprietary_statements>}
```

The file must contain a single version statement followed by a single module statement. The module statement contains multiple core statements. One `core_id` must be included. At least one interface statement must be included. One icon statement and one or more `addr_region` and `proprietary` statements may also be included.

Components

This section describes the core RTL configuration file components.

Version Statement

The version statement identifies the version of the core RTL configuration file format. The version string consists of major and minor version numbers separated by a period. The current version of the file is 4.5.

Icon Statement

This statement specifies the icon to display on a core. The syntax is:

```
icon <file_name>
```

`file_name` is the name of the graphic file, without any directory names. Store the file in the design directory of the core. For example:

```
icon "myCore.ppm"
```

The supported graphic formats are GIF, PPM, and PGM. Graphics should be no larger than 80x80 pixels. Since the text used for the core is white, use a dark background for your icon, otherwise it will be difficult to read.

Core_id Statement

The `core_id` statement provides identifying information to the tools about the core. This information is required. Syntax of the `core_id` statement is:

```
core_id <vendor_code> <core_code> <revision_code> [<description>]
```

where:

vendor_code	An OCP-IP-assigned vendor code that uniquely identifies the core developer. OCP-IP maintains a registry of assigned vendor codes. The allowed range is 0x0000 - 0xFFFF. Use 0x5555 to denote an anonymous vendor. For a list of codes check www.ocpip.org .
core_code	A developer-assigned core code that (in combination with the vendor code) uniquely identifies the core. OCP-IP provides suggested values for common cores. Refer to “Defined Core Code Values”. The allowed range is 0x000 - 0xFFF.
revision_code	A developer-assigned revision code that can provide core revision information. The allowed range is 0x0 - 0xF.
description	An optional Tcl string that provides a short description of the core.

Defined Core Code Values

0x000 - 0x7FF: Pre-defined

0x000 - 0x0FF: Memory

Sum values from following choices:

ROM:

0x0: None

0x1: ROM/EPROM

0x2: Flash (writable)

0x3: Reserved

SRAM:

0x0: None

0x4: Non-pipelined SRAM

0x8: Pipelined SRAM

0xC: Reserved

DRAM:

0x00: None

0x10: DRAM (trad., page mode, EDO, etc.)

0x20: SDRAM (all flavors)

0x30: RDRAM (all flavors)

0x40: Several

0x50: Reserved

0x60: Reserved

0x70: Reserved

Built-in DMA:

0x00: No

0x80: Yes

Values from 0x000 - 0x0FF are defined/reserved

Example: Memory controller supporting only SDRAM & Flash
would have <cc> = 0x2 + 0x20 = 0x022

0x100 - 0x1FF: General-purpose processors

Sum values from following choices plus offset 0x100:

Word size:

0x0: 8-bit

0x1: 16-bit
0x2: 32-bit
0x3: 64-bit
0x4 - 0x7: Reserved
Embedded cache:
0x0: No cache
0x8: Cache (Instruction, Data, combined, or both)
Processor Type:
0x00: CPU
0x10: DSP
0x20: Hybrid
0x30: Reserved

Only values from 0x100 - 0x13F are defined/reserved

Example: 32-bit CPU with embedded cache

would have <cc> = 0x100 + 0x2 + 0x8 + 0x00 = 0x10A

0x200 - 0x2FF: Bridges

Sum values from following choices plus offset 0x200:

Domain:

0x00 - 0x7F: Computing
 0x00 - 0x3F: PC's
 0x00: ISA (inc. EISA)
 0x01 - 0x0F: Reserved
 0x10: PCI (33MHz/32b)
 0x11: PCI (66MHz/32b)
 0x12: PCI (33MHz/64b)
 0x13: PCI (66MHz/64b)
 0x14 - 0x1F: AGP, etc.
 0x40 - 0x7F: Reserved
0x80 - 0xBF: Telecom
 0xA0 - 0xAF: ATM
 0xA0: Utopia Level 1
 0xA1: Utopia Level 2
 ...
0xC0 - 0xFF: Datacom

0x300 - 0x3FF: Reserved

0x400 - 0x5FF: Other processors

(enumerate types: MPEG audio, MPEG video, 2D Graphics,
3D Graphics, packet, cell, QAM, Vitterbi, Huffman,
QPSK, etc.)

0x600 - 0x7FF: I/O

(enumerate types: Serial UART, Parallel, keyboard, mouse,
gameport, USB, 1394, Ethernet 10/100/1000, ATM PHY,
NTSC, audio in/out, A/D, D/A, I2C, PCI, AGP, ISA,
etc.)

0x800 - 0xFFF: Vendor-defined
(explicitly left up to vendor)

Interface Statement

The interface statement defines and names the interfaces of a core. The interface name is required so that cores with multiple interfaces can specify to which interface a particular connection should be made. Syntax for the interface statement is:

```
interface <interface_name> bundle <bundle_name> [revision
<revision_string>]
[{{<interface_body>*}}]
```

Parameters lacking a default must be specified using a param statement. For a list of the required parameters, see “Configuration Parameter Defaults” on page 63. All other interface body statements are optional

The <bundle_name> must be a defined bundle such as ocp or ocp2 or a bundle specified in an interface configuration file as described on page 69. The optional <revision_string> must match that of the referenced bundle. Different interfaces can refer to different revisions of the same bundle. The pre-defined ocp and ocp2 bundles both use the default revision_string to refer to the 1.0 and 2.0 versions of the *OCP Specification*. Set revision_string to 2 to refer to the 2.2 version of the *OCP Specification*.

In the following example, an interface named xyz is defined as an OCP 2 bundle. The quotation marks around xyz are not required but help to distinguish the format.

```
interface "xyz" bundle ocp2 revision 2

<interface_body>:
| interface_type <type_name>
| port <port_name> net <net_name>
| reference_port <interface_name>.<port_name> net <net_name>
| prefix <name>
| param <name> <value> [{(<attribute> <value>)*}]
| subnet <net_name> <bit_range_list> <subnet_name>
| location (n|e|w|s|) <number>
| proprietary <vendor_code> <organization_name>
| {<proprietary_statements>}
```

Ports on a core interface may have names that are different than the nets defined in the bundle type for the interface. In this case, each port in the interface must be mapped to the net in the bundle with which it is associated. Mapping links the module port <prefix><port_name> with the bundle <net_name>.

The default rules for mapping are that the port_name is the same as the net_name and the prefix is the name of the interface. These rules can be overridden using the Port and Prefix statements.

Interface_type Statement

The `interface_type` statement defines characteristics of the bundle. Typically, the different types specify whether the core drives or receives a particular signal within the bundle. Syntax for the `interface_type` statement is:

```
interface_type <type_name>
```

The `type_name` must be a type defined in the bundle definition. If the bundle is OCP, the allowed types are: `master`, `system_master`, `slave`, `system_slave`, and `monitor` as described in Table 15 on page 33. To define a type, specify it in the interface configuration file (described on page 69).

Port Statement

Use the port statement to map a single port corresponding to a signal that is defined in the bundle. Syntax for the port statement is:

```
port <port_name> net <net_name>
```

The module port named `<prefix><port name>` implements the `<net_name>` function of the bundle. The legal `net_name` values are defined in the bundle definition. For OCP bundles, the net names are defined in “Signals and Encoding” on page 13.

Reference_port Statement

The `reference_port` statement re-directs a net to another bundle. Syntax for the port statement is:

```
reference_port <interface_name>.<port_name> net <net_name>
```

The interface (in which the `reference_port` is declared) does not have the reference port and the bundle does not have the reference net. The `reference_port` statement declares that the net is internally connected to the given port of the referenced interface. For example, consider the following two interfaces:

```
interface tp bundle ocp {
    reference_port ip.Clk_i net Clk
    reference_port ip.SReset_ni net MReset_n
    reference_port ip.EnableClk_i net EnableClk
    port Control_i net Control
    port MCmd_i net MCmd
}

interface ip bundle ocp {
    port Clk_i net Clk
    port SReset_ni net SReset_n
    port EnableClk_i net EnableClk
    port Control_i net Control
    port MCmd_o net MCmd
}
```


The diagram illustrates the internal use of signals between the **ip** (input processor) and **tp** (transfer processor) blocks. The **ip** block has four ports: **Clk_i** (input), **SReset_n** (input), **Control_i** (input), and **MCmd_o** (output). The **tp** block has four ports: **Clk** (input), **MReset_n** (input), **Control** (input), and **MCmd_i** (input). The **ip** block is connected to the **tp** block via a signal line. The **tp** block also has a **reference port** (indicated by a dashed line and a circle) for **Clk**. A legend indicates that a blue circle represents a **port** and a white circle represents a **reference port**.

Prefix Command

```
prefix <name>
```

If the prefix command is omitted, the interface name will be inserted as the default prefix. To omit the prefix from the port name, specify it as an empty string, that is `prefix ""`.

For configurable interfaces, parameters specify configurations. The specific parameters for OCP are described in Chapters 3 and 4 and summarized in Table 25 on page 64. The syntax for setting a parameter is:

If the parameter is used to configure a signal, the attribute list can be used to attach additional values to that signal. The supported attributes are the tie-off (if the signal is configured out of the interface) and the signal width (if the

signal is configured into the interface). Specifying the signal width using an attribute attached to the signal parameter is equivalent to using the corresponding signal width parameter but the attribute syntax is preferred. The width of the signals MData, SData, MByteEn, and MDataByteEn are derived from the single data_width parameter, so cannot have their width specified using an attribute.

For example, an OCP might be configured to include an interrupt signal as follows.

```
param interrupt 1
```

The following example shows the MBurstLength field tied off to a constant value of 4.

```
param burstlength 0 {tie_off 4}
```

The following code shows two equivalent ways of setting the address width to 16 bits though the second method is preferred.

```
param addr_width 16
```

```
param addr 1 {width 16}
```

Subnet Statement

The subnet statement assigns names to bits or contiguous bit-fields within a net. Syntax for the subnet statement is:

```
subnet <net_name> <bit_range_list> <subnet_name>
<bit_range_list>: <bit_range>[,<bit_range>]*
<bit_range>: <bit_number>[:<bit_number>]
```

The subnet_name is assigned to the bit_range within the given net_name. Bit_range can be either a single bit or a range. Subnet_name is a Tcl string.

For example bit 3 of the MReqInfo net may be assigned the name “cacheable” as follows:

```
subnet MReqInfo 3 cacheable
```

Location Statement

The location statement provides a way for the core to indicate where to place this interface when a schematic symbol for the core is drawn. The location is specified as a compass direction of north(n), south(s), east(e), west(w) and a number. The number indicates a percentage from the top or left edge of the block. Syntax for the location statement is:

```
location (n|e|w|s) <number>
```

To place an interface on the bottom (south-side) in the middle (50% from the left edge) of the block, for example, use this definition:

location s 50

Address Region Statement

The address region statement specifies address regions within the complete address space of a core. It allows you to give a symbolic name to a region, and to specify its base, size, and behavior.

```
addr_region <name> {<addr_region_body>*
```

where:

```
<addr_region_body>: addr_base <integer> | addr_size <integer>
                    | addr_space <integer>
                    | proprietary <vendor_code> <organization_name>
                    {<proprietary_statements>}
```

- The `addr_base` statement specifies the base address of the region being defined and is specified as an integer.
- The `addr_size` statement similarly specifies the size of the region.
- The `addr_space` statement specifies to which OCP address space the region belongs. If the `addr_space` statement is omitted, the region belongs to all address spaces.

Proprietary Statement

The proprietary statement enables proprietary extensions of the core RTL configuration file syntax. Standard parsers must be able to ignore the extensions, while proprietary parsers can extract additional information about the core. Syntax for the proprietary statement is:

```
proprietary <vendor_code> <organization_name>
            {<proprietary_statements>}
```

The `vendor_code` uniquely identifies the vendor associated with the proprietary extensions and is described in more detail on page 77.

The `organization_name` specifies the name of the organization that specified the extensions. Any number of proprietary statements can be included between the braces but must follow legal Tcl syntax.

The proprietary statement can be included at multiple levels of the syntax hierarchy, allowing it to use scoping to imply context. If multiple proprietary statements are included in a single scope, the parser must process these in an additive fashion.

Sample RTL Configuration File

The format for a core RTL configuration file for a core is shown in Example 1.

Example 1 Sample flashctrl_rtl.conf File

```
# define the module
version 4.5

module flashctrl {
  core_id 0xB BBB 0x001 0x1 "Flash/Rom Controller"

  # Use the Vista icon
  icon "vista.ppm"

  addr_region "FLASHCTRL0" {
    addr_base 0x0
    addr_size 0x100000
  }

  # one of the interfaces is an OCP slave using the pre-defined ocp2 bundle
  # Revision is "1", indicating compliance with OCP 2.1
  interface tp bundle ocp2 revision 1 {

    # this is a slave type ocp interface
    interface_type slave

    # this OCP is a basic interface with byteen support plus a named SFlag
    # and MReset_n
    param mreset 1
    param sreset 0
    param byteen 1
    param sflag 1 {width 1}
    param addr 1 {width 32}
    param mdata 1 {width 64}
    param sdata 1 {width 64}

    prefix tp
    # since the signal names do not exactly match the signal
    # names within the bundle, they must be explicitly linked
    port Reset_ni net MReset_n
    port Clk_i net Clk
    port TMCmd_i net MCmd
    port TMAAddr_i net MAddr
    port TMByteEn_i net MByteEn
    port TMData_i net MData
    port TCCmdAccept_o net SCmdAccept
    port TCResp_o net SResp
    port TCData_o net SData
    port TCErr_o net SFlag
  }
}
```

```

    # name SFlag[0] access_error
    subnet SFlag 0 access_error

    # stick this interface in the middle of the top of the module
    location n 50

} # close interface tp defininition

# The other interface is to the flash device defined in an interface file
# Define the interface for the Flash control
interface emem bundle flash {

    # the type indicates direction and drive of the control signals
    interface_type controller

    # since this module has direction indication on some of the signals
    # ('_o','_b') and is missing assertion level indicators '_n' on
    # some of the signals, the names must again be directly linked to
    # the signal names within the bundle
    port Addr_o    net addr
    port Data_b    net data
    port OE        net oe_n
    port WE        net we_n
    port RP        net rp_n
    port WP        net wp_n

    # all of the signals on this port have the prefix 'emem_'
    prefix "emem_"

    # stick this interface in the middle of the bottom of the module
    location s 50

} # close interface emem defininition

} # close module definition

```

The flash bundle is defined in the following interface configuration file. See “Interface Configuration File” on page 69 for the syntax definition of the interface configuration file.

```

bundle flash {
    #types of flash interfaces
    #controller: flash controller; flash: flash device itself.
    interface_types controller flash
    net addr {
        #Address to the Flash device
        direction output input
        width 19
    }
}

```

```
net data {
    #Read or Write Data
    direction inout inout
    width 16
}
net oe_n {
    # Output Enable, active low.
    direction output input
}
net we_n {
    # Write Enable, active low.
    direction output input
}
net rp_n {
    # Reset, active low.
    direction output input
}
net wp_n {
    # Write protect bit, Active low.
    direction output input
}
}
```

7 *Core Timing*

To connect two entities together, allowing communication over an OCP interface, the protocols, signals, and pin-level timing must be compatible. This chapter describes how to define interface timing for a core. This process can be applied to OCP and non-OCP interfaces.

Use the core synthesis configuration file to set timing constraints for ports in the core. The file consists of any of the constraint sections: port, max delay, and false path. If the core has additional non-OCP clocks, the file should contain their definitions.

When implementing IP cores in a technology independent manner it is difficult to specify only one timing number for the interface signals, since timing is dependent on technology, library and design tools. The methodology specified in this chapter allows the timing of interface signals to be specified in a technology independent way.

To make your core description technology independent use the technology variables defined in the *Core Preparation Guide*. The technology variables range from describing the default setup and clock-to-output times for a port to defining a high drive cell in the library.

Timing Parameters

There is a set of minimum timing parameters that must be specified for a core interface. Additional optional parameters supply more information to help the system designer integrate the core. Hold-time parameters allow hold time checking. Physical-design parameters provide details on the assumptions used for deriving pin-level timing.


Minimum Parameters

At a minimum, the timing of an OCP interface is specified in terms of two parameters:

- `setuptime` is the latest time an input signal is allowed to change before the rising edge of the clock.
- `c2qtime` is the latest time an output signal is guaranteed to become stable after the rising edge of the clock.

Figure 8 shows the definition of `setuptime` and `c2qtime`. See “Port Constraint Keywords” on page 95 for a description of these parameters.

Figure 8 OCP Timing Parameters



Hold-time Parameters


Hold-time parameters are needed to allow the system integrator to check hold time requirements. On the output side, `c2qtimemin` specifies the minimum time for a signal to propagate from a flip-flop to the given output pin. On the input side, `holdtime` specifies the minimum time for a signal to propagate from the input pin to a flip-flop.

Technology Variables

To give meaning to the timing values, timing requirements on input and output pins must be accompanied by information on the assumed environment for which these numbers are determined. This information also adds detail on the expected connection of the pin.


For an input signal, the parameter `drivingcellpin` indicates the cell library name for a cell representative of the strength of the driver that needs to be used to drive the signal. This is shown in Figure 9.

Figure 9 Driver Strength



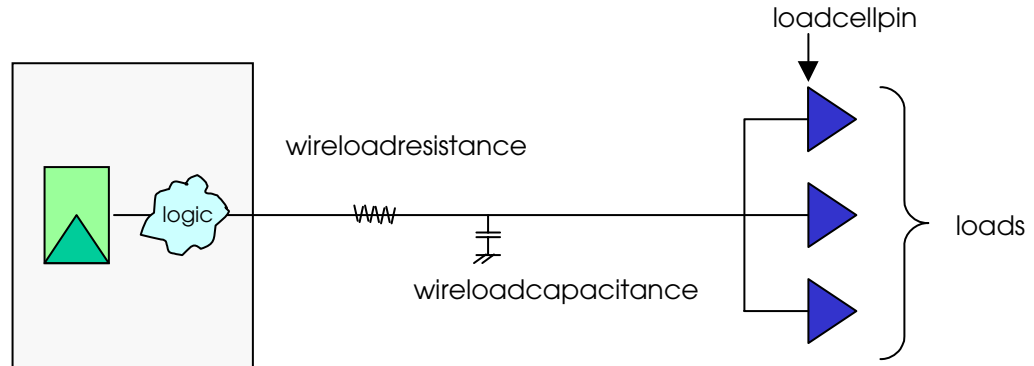
For an output signal, the variable `loadcellpin` indicates the input load of the gate that the signal is expected to drive. The variable `loads` indicates how many `loadcellpins` the signal is expected to drive. Additionally, information on the capacitive load of the wire must be included. There are two options. Either the variable `wireloaddelay` can be specified, as shown in Figure 10. Or, the combination `wireloadcapacitance/wireloadresistance` must be specified, as shown in Figure 11.

Figure 10 Variable Loads - `wireloaddelay`



For instructions on calculating a delay, refer to the *Synopsys Design Compiler Reference*.

Figure 11 Variable Loads - wireloadresistance/wireloadcapacitance



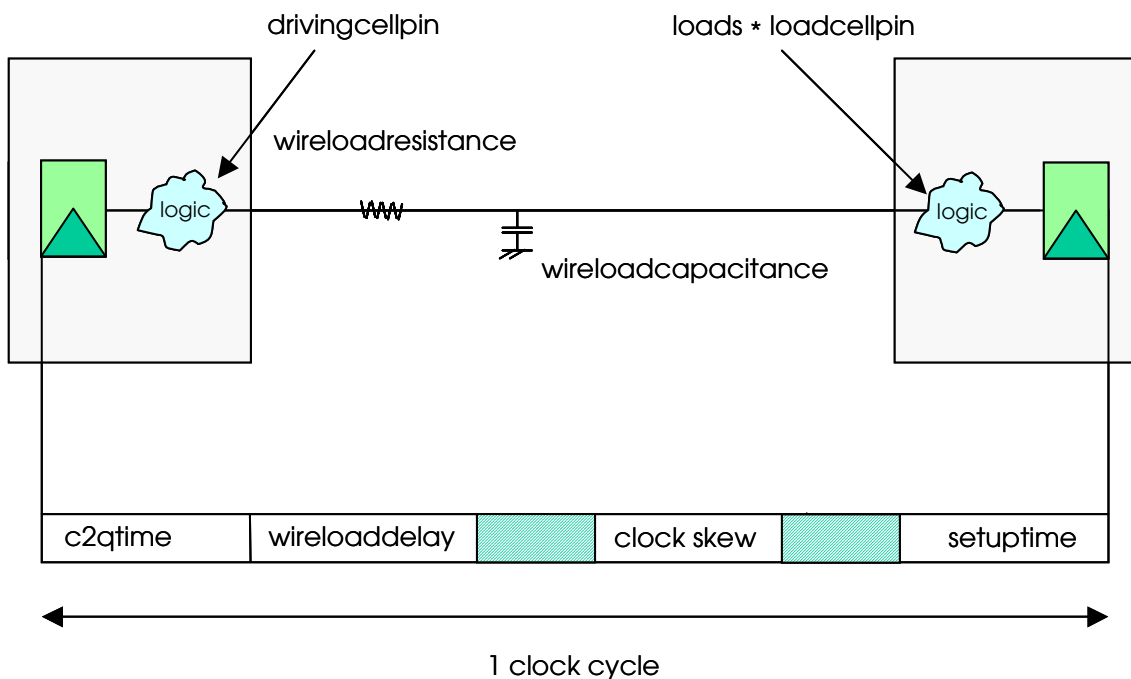
Connecting Two OCP Cores

Figure 12 shows the timing model for interconnecting two OCP compliant cores.

The sum of `setuptime`, `c2qtime` and wire delay must be less than the clock period or cycle time minus the clock-skew. Similarly, the minimum clock-cycle for two cores to interoperate is determined by the maximum of the sum of `c2qtime`, `setuptime`, wire delay and clock-skew over all interface signals.

The wireload delay is defined by either the variable `wireloaddelay` or the set `wireloadcapacitance/wireloadresistance`.


Figure 12 Connecting Two OCP Compliant Cores



Max Delay

In addition to the setup and c2qtime paths for a core, there may also be combinational paths between input and output ports. Use `maxdelay` to specify the timing for these paths.

Figure 13 Max Delay Timing



False Paths

It is possible to identify a path between two ports as being logically impossible. Such paths can be specified using the `falsepath` constraint syntax.

For instructions on specifying the core's timing parameters, see "False Path Constraints" on page 100.

Core Synthesis Configuration File

The core synthesis configuration file contains the following sections:

Version

Specifies the current version of the synthesis configuration file format.
The current version is 1.3.

Clock

Describes clocks brought into the core.

Area

Defines the area in gates of the core.

Port

Defines the timing of IP block ports.

Max Delay

Specifies the delay between two ports on a combinational path.

False Path

Specifies that a path between input and output ports is logically impossible.

Syntax Conventions

Observe the following syntax conventions:

- Enclose all `expr` statements within braces `{ }`, to differentiate between expressions that are to be evaluated while the file is being parsed (without braces) and those that are to be evaluated during synthesis constraint file generation (with braces).
- Although not required by Tcl, enclose strings within quotation marks `""`, to show that they are different than keywords.
- Specify keywords using lower case.

Parameter values are specified using Tcl syntax. Expressions can use any of the technology or environment variables, and any of the following variables:

clockperiod

This variable should only be used in calculations of timing values for ports. When evaluating expressions that use `$clockperiod`, the program will determine which clock the port is relative to, determine its period (in nanoseconds), and apply that value to the equation. For example:

```
port "in" {  
    setuptime {[expr $clockperiod * .5]}  
}
```

rootclockperiod

This variable is set to the period of the main system clock, usually referred to as the root clock. It is typically used when a value needs to be a multiple of the root clock, such as for non-OCP clocks. For example:

```
clock "myClock" {
    period {[expr $rootclockperiod * 4]}
}
```

The design_syn.conf file can also use conditional settings of the parameters in the design as outlined by the following arrays. These variables are only used at the time the file is read into the tools.

param

This array is indexed by the configuration parameters that can be found on a particular instance. Only use param for core_syn.conf files since it is only applicable to the instance being processed. For example:

```
if { $param("dma_fd") == 1 } {
    port "T12_ipReset_no" {
        c2qtime {[expr $clockperiod * 0.7]}
    }
}
```

chipparam

This array is indexed by the configuration parameters that are defined at the chip or design level. These variables can be used in both the design_syn.conf and core_syn.conf files as they are more global in nature than those specified by param. For example:

```
if { $chipparam("full") == 1 } {
    instance "bigcore" {
        port "in" {
            setuptime {[expr $clockperiod * 0.7]}
        }
    }
}
```

interfaceparam

This array is indexed by the interface name and the configuration parameters that are on an interface. It should only be used for core_syn.conf files since it is only applicable to the interfaces on the instance being processed. In the following example the interface name is ip.

```
if { $interfaceparam("ip_respaccept") == 1 } {
    port "ipMRespAccept_o" {
        c2qtime {[expr $clockperiod * 21/25]}
    }
}
```

Version Section

The version of the core synthesis configuration file is required. Specify the version with the version command, for example: `version 1.3`

Clock Section

If you have non-OCF clocks for an IP block or want to specify the `worstcasedelay` of any clock (including OCF clocks) used in the core, specify the names of the clocks in the core synthesis configuration file. Use the following syntax to specify the name of the clock and its `worstcasedelay`:

```
clock <clockName> {  
    worstcasedelay <delay Value>  
}
```

`clockName` refers to the name of the port that brings the clock into the core for the core synthesis configuration file. For example:

```
clock "myClock"
```

`worstcasedelay`

The worst case delay value is the longest path through the core or instance for a particular clock. The value is used to check that the core can meet the timing requirements of the current design. To help make this value more portable, you may want to use the technology variable `gatedelay`. For example:

```
clock "myClock" {  
    worstcasedelay {[10.5 * $gatedelay]}  
}  
  
clock "otherClock" {  
    worstcasedelay 5  
}
```

Constant values are specified in nanoseconds. For consistency, use expressions that can be interpreted in nanoseconds.

Area Section

The area is the size in gates of the core or instance. By specifying the size in gates the area can be calculated based on the size of a typical two input nand gate in a particular synthesis library. For example:

```
area {[expr 20500 / $gatesize]}  
area 5000
```

Constant values are specified in two input nand gate equivalents. For consistency, use expression that can be interpreted in gates.

Port Constraints Section

Use the port constraints section to specify the timing parameters. Input port information that can be specified includes the setup time, related clock (non-OCP ports), and driving cell. For output ports, the clock to output times, related clock (non-OCP ports), and the loading information must be supplied.

Port Constraint Keywords

The keywords that can be used to specify information about port constraints are:

c2qtime

The c2q (clock to q or clock to output) time is the longest path using worst case timing from a starting point in the core (register or input port) to the output port. This includes the c2qtime of the register. To maintain portability, most cores specify this as a percentage of the fastest clock period used while synthesizing the core. For example:

```
c2qtime {[expr $timescale * 3500]}
c2qtime {[expr $clockperiod * 0.25]}
```

Constant values are specified in nanoseconds. For consistency, use expressions that can be interpreted in nanoseconds.

c2qtimemin

The c2q (clock to q or clock to output) time min is the shortest path using best case timing from a starting point in the core (register or input port) to the output port. This includes the c2qtime of the register. Most cores use the default from the technology section, defaultc2qtimemin. For example:

```
c2qtimemin {[expr $timescale * 100]}
c2qtimemin {$defaultc2qtimemin}
```

Constant values are specified in nanoseconds. For consistency, use expressions that can be interpreted in nanoseconds.

clockname

This is an optional field for all OCP ports and is a string specifying the associated clock portname. For input ports, input delays use this clock as the reference clock. For output ports, output delays use this clock as the reference clock. For example:

```
clockname "myClock"
```

drivingcellpin

This variable describes which cell in the synthesis library is expected to be driving the input. To maintain portability set this variable to use one of the technology values of high/medium/lowdrivegatepin.

Values are a string that specifies the logical name of the synthesis library, the cell from the library, and the pin that will be driving an input for the core. The pin is optional. For example:

```
drivingcellpin {$mediumdrivegatepin}  
drivingcellpin "pt25u/nand2/O"
```

holdtime

The hold time is the shortest path using best case timing from an input port to any endpoint in the core. Most cores use the default from the technology section, `defaultholdtime`. For example:

```
holdtime {[expr $timescale * 100]}  
holdtime {$defaultholdtime}
```

Constant values are specified in nanoseconds. For consistency, use expressions that can be interpreted in nanoseconds.

loadcellpin

The name of the load library/cell/pin that this output port is expected to drive. The value is specified to the synthesis tool as the gate to use (along with the number of loads) in its load calculations for output ports of a module. For portability use the default.

Values are a string that specifies the logical name of the synthesis library, the cell from the library, and the pin that the load calculation is derived from. The pin is optional. For example:

```
loadcellpin "pt25u/nand2/I1"  
loadcellpin {$defaultloadcellpin}
```

loads

The number of loadcellpins that this output port is expected to drive. The value is communicated to the synthesis tool as the number of loads to use in load calculations for output ports of a module. The typical setting for this is the technology value of `defaultloads`. Values are an expression that evaluates to an integer. For example:

```
loads 5  
loads {$defaultloads}
```

maxfanout

This keyword limits the fanout of an input port to a specified number of fanouts. To maintain portability set this variable in terms of the technology variable `defaultfanoutload`. Constant values are specified in library units. For example:

```
maxfanout {[expr $defaultfanoutload * 1]}
```

setuptime

The longest path using worst case timing from an input port to any endpoint in the core. To maintain portability, most cores specify this as a percentage of the fastest clock period used during synthesis of the core. For example:

```
setuptime {[expr $timescale * 2500]}  
setuptime {[expr $clockperiod * 0.25]}
```


Constant values are specified in nanoseconds. For consistency, use expressions that can be interpreted in nanoseconds.

wireloaddelay

Replaces capacitance/resistance as a way to specify expected delays caused by the interconnect. To maintain portability set this variable to use a technology value of long/medium/shortnetdelay.

The resulting values get added to the worst case clock-to-output times of the ports to anticipate net delays of connections to these ports. To improve the accuracy of the delay calculation cores should use the resistance and capacitance settings.

You cannot specify both wireloaddelay and wireloadresistance/capacitance for the same port. For example:

```
wireloaddelay {[expr $clockperiod * .25]}
wireloaddelay {$mediumnetdelay}
```

Constant values are specified in nanoseconds. For consistency, use expressions that can be interpreted in nanoseconds.

wireloadresistance

wireloadcapacitance

Specify expected loading and resistance caused by the interconnect. If available, specify both resistance and capacitance. To maintain portability set this variable to use one of the technology values of long/medium/shortnetrcresistance/capacitance.

If these constraints are specified they show up as additional loads and resistances on output ports of a module. You cannot use both wireloaddelay and wireloadresistance/capacitance for the same port.

Specify constant values as expressions that result in kOhms for resistance and picofarads (pf) for capacitance. For example:

```
wireloadresistance {[expr $resistancescale * .09]}
wireloadcapacitance {[expr $capacitancescale * .12]}
wireloadresistance {$mediumnetrcresistance}
wireloadcapacitance {$mediumnetrccapacitance}
```

Input Port Syntax

For input and inout ports (*inout* ports have both an input and an output definition) use the following syntax:

```
port <portName> {
    clockname <clockName>
    drivingcellpin <drivingCellName>
    setuptime <Value>
    holdtime <Value>
    maxfanout <Value>
}
```

Examples

In the following example, the clock is not specified since this is an OCP port and is known to be controlled by the OCP clock. If a clock were specified as something other than the OCP clock, an error would result.

```
port "MCmd_i" {
    drivingcellpin {$mediumdrivegatepin}
    setuptime {[expr $clockperiod * 0.2]}
}
```

In the following example, the setup time is required to be 2ns. Time constants are assumed to be in nanoseconds. Use the `timescale` variable to convert library units to nanoseconds.

```
port "MAddr_i" {
    drivingcellpin {$mediumdrivegatepin}
    setuptime 2
}
```

The following example shows how to associate a non OCP clock to a port. The example uses `maxfanout` to limit the fanout of `myInPort` to 1. If the logic for `myInPort` required it to fanout to more than one connection, the synthesis tool would add a buffer to satisfy the `maxfanout` requirement.

```
port "myInPort" {
    clockname "myClock"
    drivingcellpin {$mediumdrivegatepin}
    setuptime 2
    maxfanout {[expr $defaultfanoutload * 1]}
}
```

Output Port Syntax

For output and inout ports (inout ports have both an input and an output definition) use the following syntax:

```
port <portName> {
    clockname <clockName>
    loadcellpin <loadCellPinName>
    loads <Value>
    wireloadresistance <Value>
    wireloadcapacitance <Value>
    wireloaddelay <Value>
    c2qtime <Value>
    c2qtimemin <Value>
}
```

You cannot specify both `wireloaddelay` and `wireloadresistance/capacitance` for the same port.

Examples

In the following example, the clock is not specified since this is an OCP port and is known to be controlled by the OCP clock.

```
port "SCmdaccept_o"
  loadcellpin {$defaultloadcellpin}
  loads {$defaultloads}
  wireloadresistance {$mediumnetrcresistance}
  wireloadcapacitance {$mediumnetrccapacitance}
  c2qtime {[expr $clockperiod * 0.2]}
}
```

In the following example, the clock to output time is required to be 2 ns. Time constants are assumed to be in nanoseconds. Use the `timescale` variable to convert library units to nanoseconds.

```
port "SResp_o"
  loadcellpin {$defaultloadcellpin}
  loads {$defaultloads}
  wireloadresistance {$mediumnetrcresistance}
  wireloadcapacitance {$mediumnetrccapacitance}
  c2qtime 2
}
```

The following example shows how to associate a clock to an output port.

```
port "myOutPort"
  clockname "myClock"
  loadcellpin {$defaultloadcellpin}
  loads 10
  wireloaddelay {$longnetdelay}
  c2qtime {[expr $clockperiod * .2]}
}
```

InOut Port Example

```
port "Signal_io"
  drivingcellpin {$mediumdrivegatepin}
  setuptime {[expr $clockperiod * 0.2]}
}
port "Signal_io"
  loadcellpin {$defaultloadcellpin}
  loads {$defaultloads}
  wireloadresistance {$mediumnetrcresistance}
  wireloadresistance {$mediumnetrccapacitance}
  c2qtime {[expr $clockperiod * 0.2]}
}
```

Max Delay Constraints

Using the max delay constraints you can specify the delay between two ports on a combinational path. This is useful when synthesizing two communicating OCP interfaces. The syntax for `maxdelay` is:

```
maxdelay {  
    delay <delayValue> fromport <portName> toport <portName>  
    .  
    .  
    .  
}
```

where: `<delayValue>` can be a constant or a Tcl expression.

In the following example, a `maxdelay` of 3 ns is specified for the combinational path between `myInPort1` and `myOutPort1`. A `maxdelay` of 50% of the clockperiod is specified for the path between `myInPort2` and `myOutPort2`. The braces around the expression delay evaluation until the expression is used by the mapping program.

```
maxdelay {  
    delay 3 fromport "myInPort1" toport "myOutPort1"  
    delay {[expr $clockperiod *.5]} fromport "myInPort2" toport "myOutPort2"  
}
```

False Path Constraints

Using the false path constraints you can specify that a path between certain input and output ports is logically impossible.

The syntax for `falsepath` is:

```
falsepath{  
    fromport <portName> toport <portName>  
    .  
    .  
    .  
}
```

In the following example, a `falsepath` is set up between `myInPort1` and `myOutPort1` as well as `myInPort2` and `myOutPort2`. This tells the synthesis tool that the path is not logically possible and so it will not try to optimize this path to meet timing.

```
falsepath {  
    fromport "myInPort1" toport "myOutPort1"  
    fromport "myInPort2" toport "myOutPort2"  
}
```

Sample Core Synthesis Configuration File

The following example shows a complete core synthesis configuration file.

```

version 1.3
port "Reset_ni" {
    drivingcellpin {$mediumgatedrivepin}
    setuptime {[expr $clockperiod * .5]}
}
port "MCmd_i" {
    drivingcellpin {$mediumgatedrivepin}
    setuptime {[expr $clockperiod * .9]}
}
port "MAddr_i" {
    drivingcellpin {$mediumgatedrivepin}
    setuptime {[expr $clockperiod * .5]}
}
port "MWidth_i" {
    drivingcellpin {$mediumgatedrivepin}
    setuptime {[expr $clockperiod * .5]}
}
port "MData_i" {
    drivingcellpin {$mediumgatedrivepin}
    setuptime {[expr $clockperiod * .5]}
}
port "SCmdAccept_o" {
    loadcellpin {$defaultloadcellpin}
    loads {$defaultloads}
    wireloaddelay {$mediumnetdelay}
    c2qtime {[expr $clockperiod * .9]}
}
port "SResp_o" {
    loadcellpin {$defaultloadcellpin}
    loads {$defaultloads}
    wireloaddelay {$mediumnetdelay}
    c2qtime {[expr $clockperiod * .8]}
}
port "SData_o" {
    loadcellpin {$defaultloadcellpin}
    loads {$defaultloads}
    wireloaddelay {$mediumnetdelay}
    c2qtime {[expr $clockperiod * .8]}
}
maxdelay {
    delay 2 fromport "MData_i" toport
    "SResp_o"
}
falsepath {
    fromport "MData_i" toport "SData_o"
}

```


Part II *Guidelines*

8 *Timing Diagrams*

The timing diagrams within this chapter look at signals at strategic points and are not intended to provide full explanations but rather, highlight specific areas of interest. The diagrams are provided solely as examples. For related information about phases, see “Signal Timing and Protocol Phases” on page 38.

Most of the timing diagrams in this chapter are based upon simple OCP clocking, where the OCP clock is completely determined by the Clk signal. A few diagrams are repeated to show the impact of the EnableClk signal. Most fields are unspecified whenever their corresponding phase is not asserted. This is indicated by the striped pattern in the waveforms. For example, when MCmd is IDLE the request phase is not asserted, so the values of MAddr, MData, and SCmdAccept are unspecified.

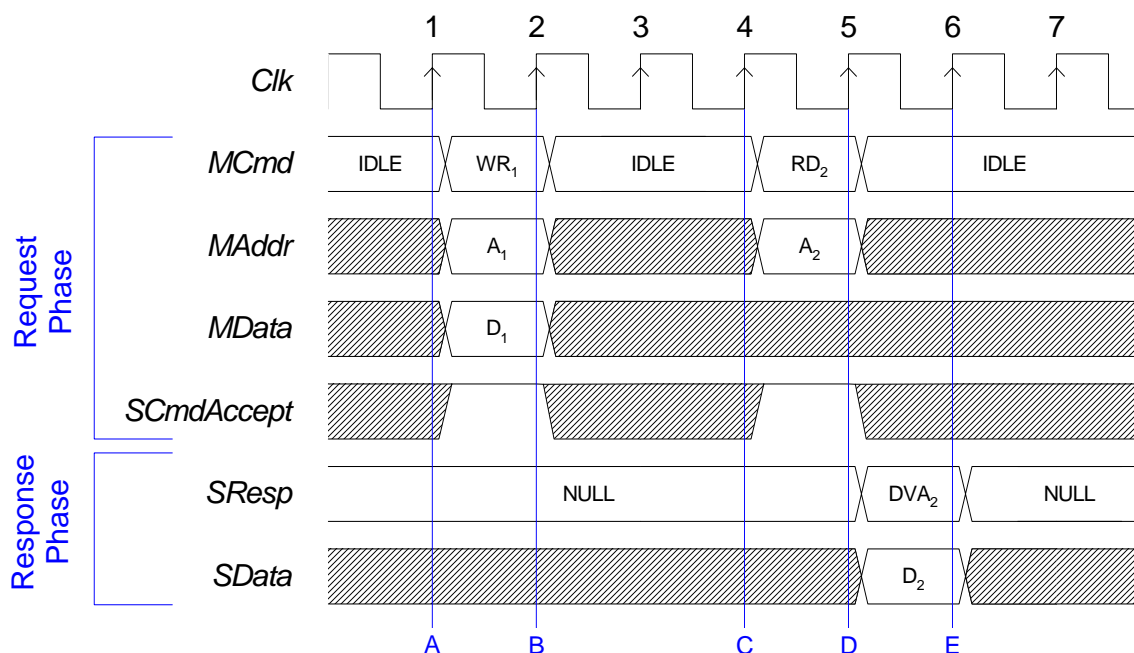
Subscripts on labels in the timing diagrams denote transfer numbers that can be helpful in tracking a transfer across protocol phases.

For a description of timing diagram mnemonics, see Tables 2 and 3 on page 16.

Simple Write and Read Transfer

Figure 14 illustrates a simple write and a read transfer on a basic OCP interface. This diagram shows a write with no response enabled on the write, which is typical behavior for a synchronous SRAM or a register bank.

Figure 14 Simple Write and Read Transfer



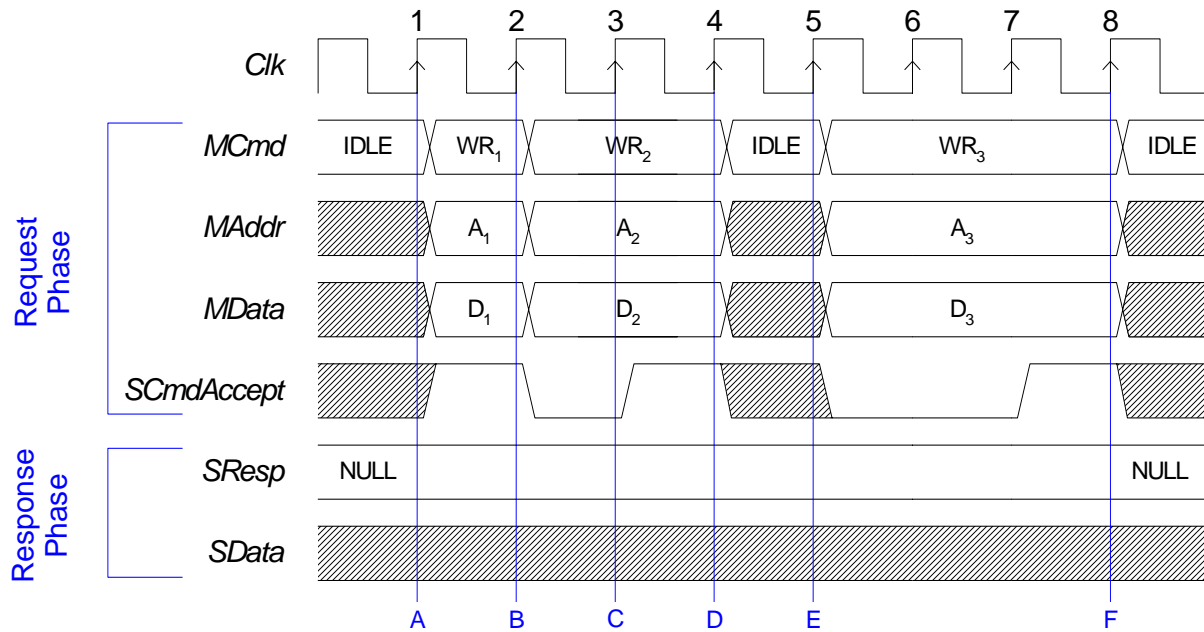
Sequence

- A. The master starts a request phase on clock 1 by switching the MCmd field from IDLE to WR_1 . At the same time, it presents a valid address (A_1) on MAddr and valid data (D_1) on MData. The slave asserts SCmdAccept in the same cycle, making this a 0-latency transfer.
- B. The slave captures the values from MAddr and MData and uses them internally to perform the write. Since SCmdAccept is asserted, the request phase ends.
- C. The master starts a read request by driving RD_2 on MCmd. At the same time, it presents a valid address on MAddr. The slave asserts SCmdAccept in the same cycle for a request accept latency of 0.
- D. The slave captures the value from MAddr and uses it internally to determine what data to present. The slave starts the response phase by switching SResp from NULL to DVA_2 . The slave also drives the selected data on SData. Since SCmdAccept is asserted, the request phase ends.
- E. The master recognizes that SResp indicates data valid and captures the read data from SData, completing the response phase. This transfer has a request-to-response latency of 1.

Request Handshake

Figure 15 illustrates the basic flow-control mechanism for the request phase using SCmdAccept. There are three writes with no responses enabled, each with a different request accept latency.

Figure 15 Request Handshake



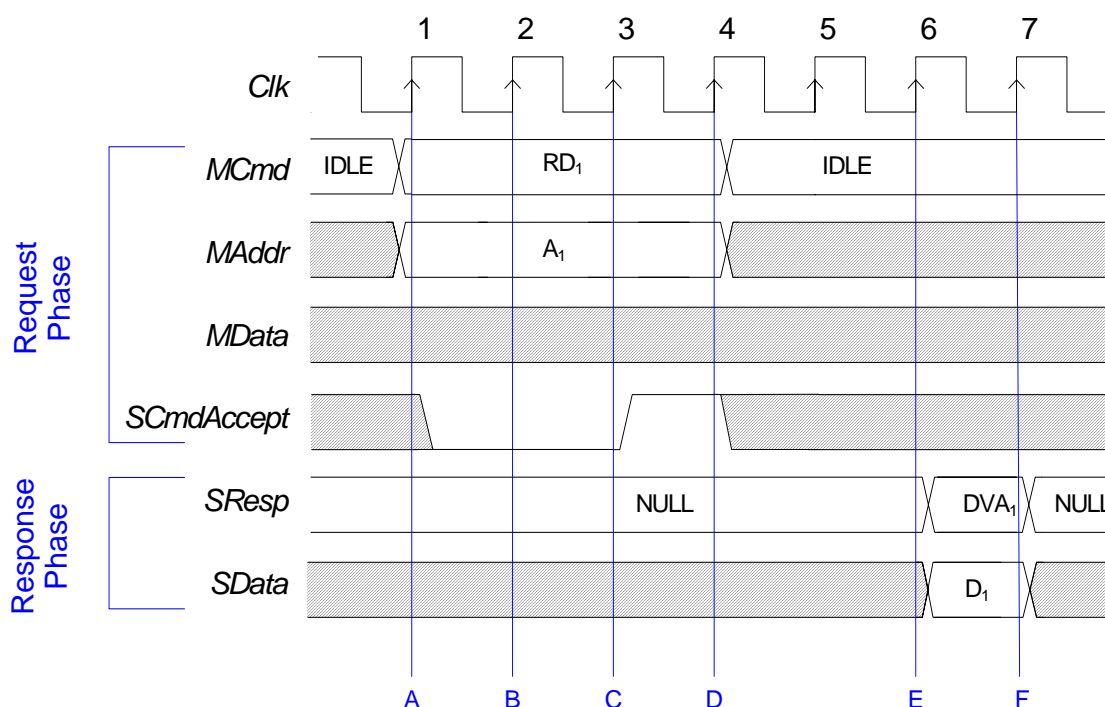
Sequence

- The master starts a write request by driving WR on MCmd and valid address and data on MAddr and MData, respectively. The slave asserts SCmdAccept in the same cycle, for a request accept latency of 0.
- The master starts a new transfer in the next cycle. The slave captures the write address and data. It deasserts SCmdAccept, indicating that it is not yet ready for a new request.
- Recognizing that SCmdAccept is not asserted, the master holds all request phase signals (MCmd, MAddr, and MData). The slave asserts SCmdAccept in the next cycle, for a request accept latency of 1.
- The slave captures the write address and data.
- After 1 idle cycle, the master starts a new write request. The slave deasserts SCmdAccept.
- Since SCmdAccept is asserted, the request phase ends. SCmdAccept was low for 2 cycles, so the request accept latency for this transfer is 2. The slave captures the write address and data.

Request Handshake and Separate Response

Figure 16 illustrates a single read transfer in which a slave introduces delays in the request and response phases. The request accept latency 2, corresponds to the number of clock cycles that SCmdAccept was deasserted. The request to response latency 3, corresponds to the number of clock cycles from the end of the request phase (D) to the end of the response phase (F).

Figure 16 Request Handshake and Separate Response



Sequence

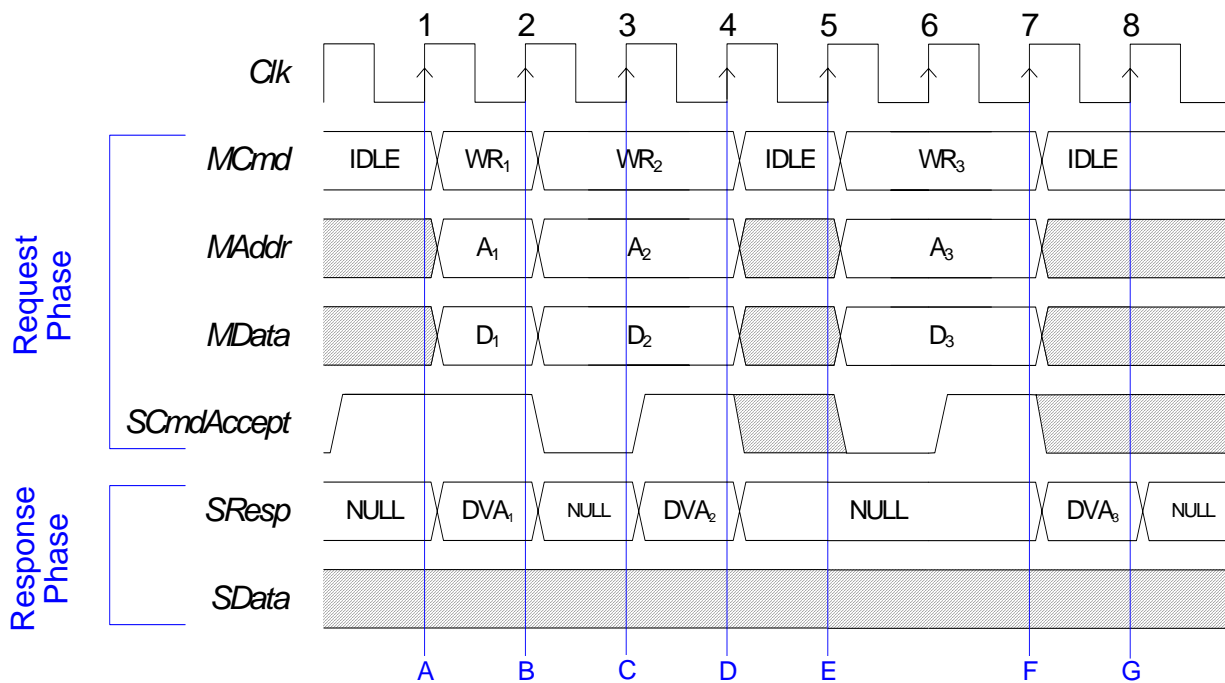
- The master starts a request phase by issuing the RD command on the MCmd field. At the same time, it presents a valid address on MAddr. The slave is not ready to accept the command yet, so it deasserts SCmdAccept.
- The master sees that SCmdAccept is not asserted, so it keeps all request phase signals steady. The slave may be using this information for a long decode operation, and it expects the master to hold everything steady until it asserts SCmdAccept.
- The slave asserts SCmdAccept. The master continues to hold the request phase signals.
- Since SCmdAccept is asserted, the request phase ends. The slave captures the address, and although the request phase is complete, it is not ready to provide the response, so it continues to drive NULL on the SResp field. For example, the slave may be waiting for data to come back from an off-chip memory device.

- E. The slave is ready to present the response, so it issues DVA on the SResp field, and drives the read data on SData.
- F. The master sees the DVA response, captures the read data, and the response phase ends.

Write with Response

Figure 17 is the same example as the waveform on page 107 but with response on write enabled. The response is typically provided to the master in the same cycle as SCmdAccept, but could be delayed (if required to perform an error check for instance). On the first write transaction, the slave uses a default accept scheme, resulting in a 0-wait state write transaction. Using fully-synchronous handshake, this condition is only possible when the slave's ability to accept a command depends solely on its internal state: any command issued by the master can be accepted. Same-cycle SCmdAccept could also be achieved using combinational logic.

Figure 17 Write with Response



Sequence

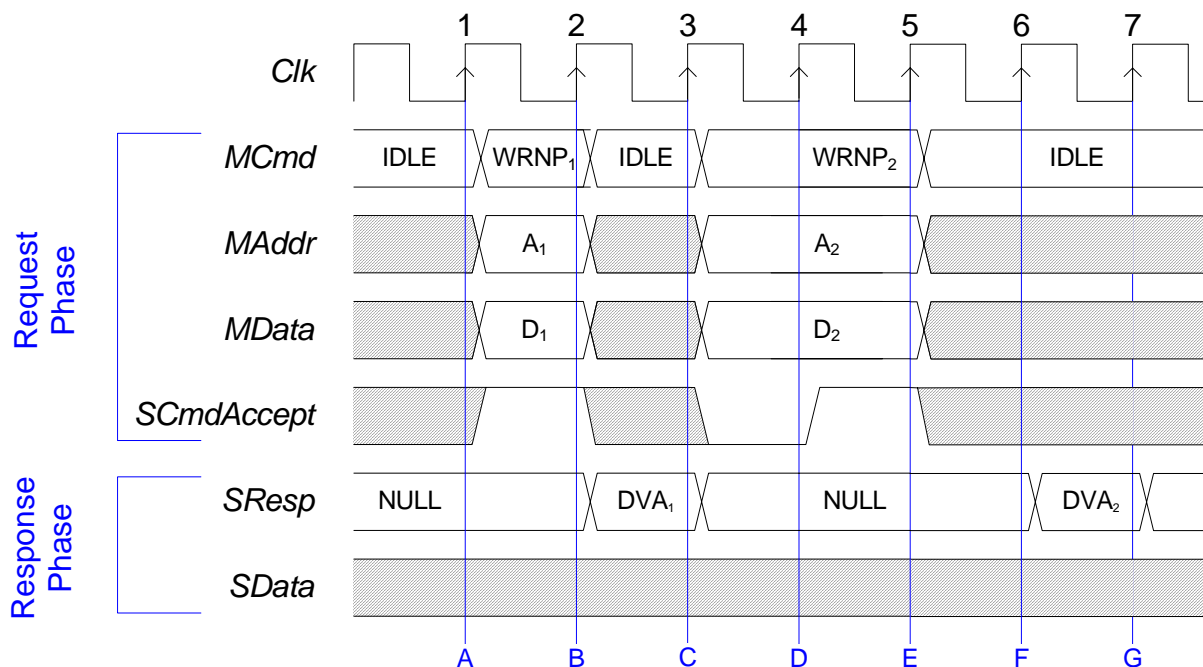
- A. The master starts a write request by driving WR on MCmd and valid address and data on MAddr and MData, respectively. The slave having already asserted SCmdAccept for a request accept latency of 0, drives DVA on SResp to indicate a successful transaction.

- B. The master starts a new transfer in the next cycle. The slave captures the write address and data and deasserts SCmdAccept, indicating that it is not ready for a new request.
- C. With SCmdAccept not asserted, the master holds all request phase signals (MCmd, MAddr, and MData). The slave asserts SCmdAccept in the next cycle, for a request accept latency of 1 and drives DVA on SResp to indicate a successful transaction.
- D. The slave captures the write address and data.
- E. After 1 idle cycle, the master starts a new write request. The slave deasserts SCmdAccept.
- F. Since SCmdAccept is asserted, the request phase ends. SCmdAccept was low for 2 cycles, so the request accept latency for this transfer is 2. The slave captures the write address and data. The slave drives DVA on SResp to indicate a successful transaction.
- G. The master samples the response.

Non-Posted Write

Figure 18 repeats the previous example for a non-posted write transaction. In this case the response must be returned to the master once the write operation occurs. There is no difference in the command acceptance, but the response may be significantly delayed. If this scheme is used for all posting-sensitive transactions, the result is decreased data throughput but higher system reliability.

Figure 18 Non-posted Write



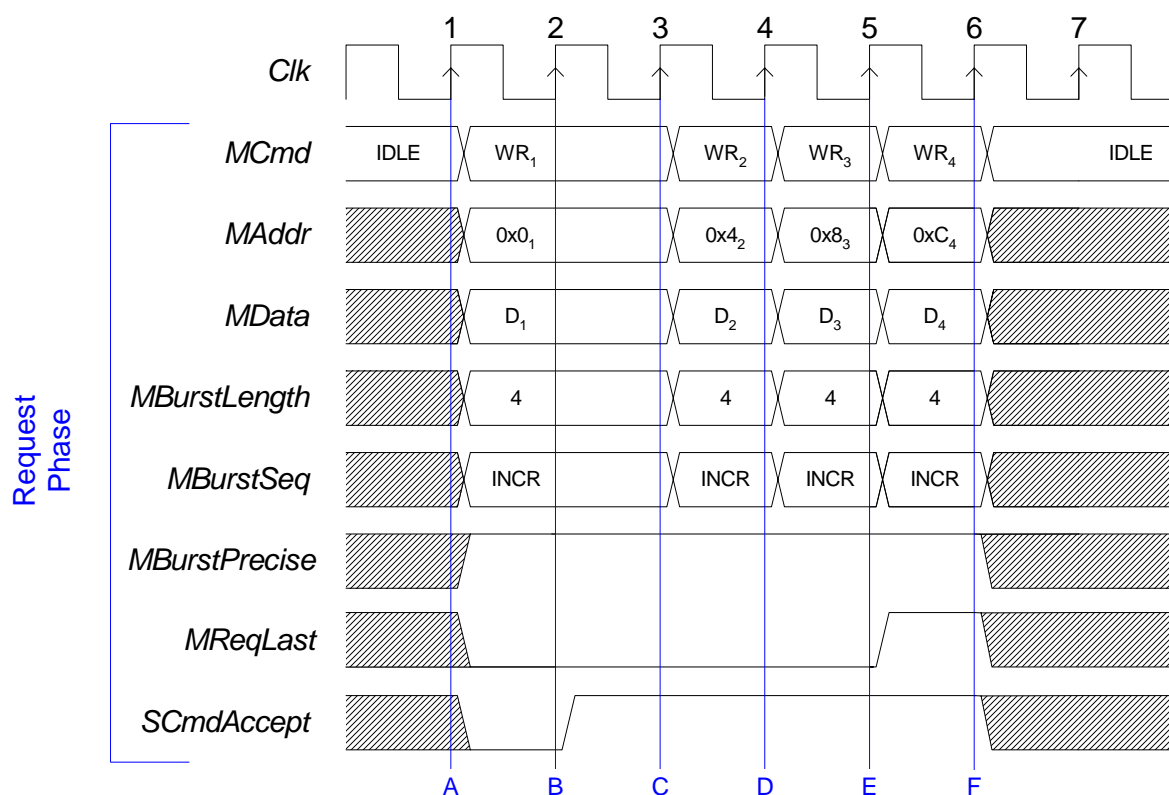
Sequence

- A. The master starts a non-posted write request by driving WRNP on MCmd and valid address and data on MAddr and MData, respectively. The slave asserts SCmdAccept combinationally, for a request accept latency of 0.
- B. The slave drives DVA on SResp to indicate a successful first transaction.
- C. The master starts a new transfer. The slave deasserts the SCmdAccept, indicating it is not yet ready to accept a new request. The master samples DVA on SResp and the first response phase ends.
- D. The slave asserts SCmdAccept for a request accept latency of 1.
- E. The slave captures the write address and data.
- F. The slave drives DVA on SResp to indicate a successful second transaction.
- G. The master samples DVA on SResp and the second response phase ends.

Burst Write

Figure 19 illustrates a burst of four 32-bit words, incrementing precise burst write, with optional burst framing information (MReqLast). As the burst is precise (with no response on write), the MBurstLength signal is constant during the whole burst. MReqLast flags the last request of the burst, and SRespLast flags the last response of the burst. The slave may either count requests or monitor MReqLast for the end of burst.

Figure 19 Burst Write



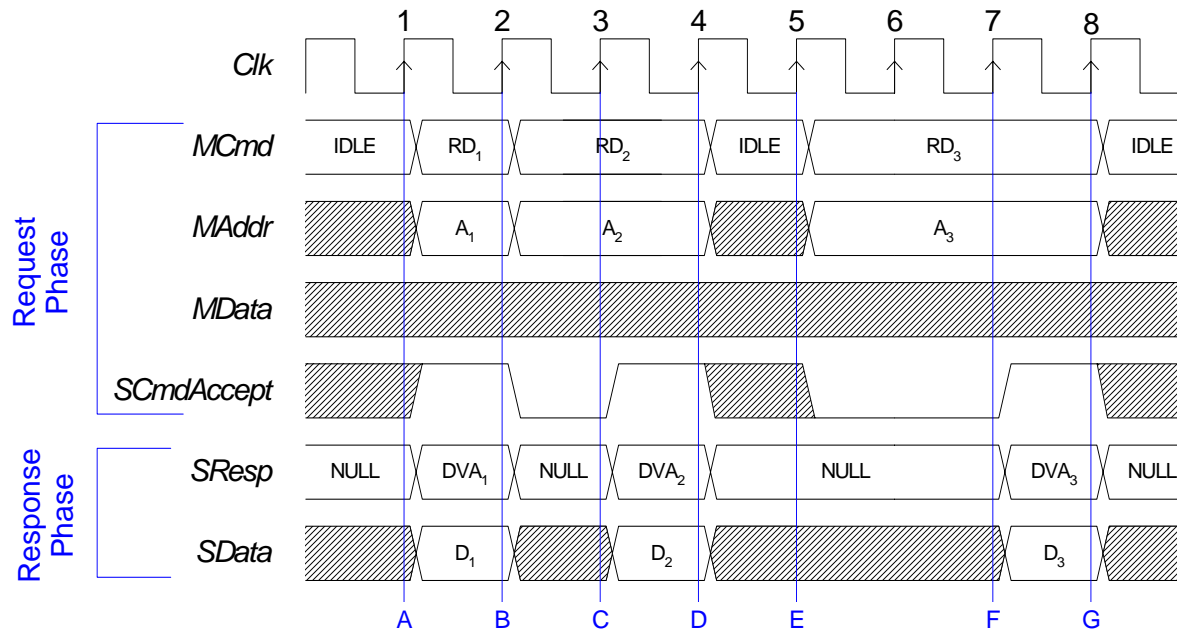
Sequence

- A. The master starts the burst write by driving WR on MCmd, the first address of the burst on MAddr, valid data on MData, a burst length of four on MBurstLength, the burst code INCR on MBurstSeq, and asserts MBurstPrecise. MReqLast must be deasserted until the last request in the burst. The burst signals indicate that this is an incrementing burst of precisely four transfers. The slave is not ready for anything, so it deasserts SCmdAccept.
- B. The slave asserts SCmdAccept for a request accept latency of 1.
- C. The master issues the next write in the burst. MAddr is set to the next word-aligned address. For 32-bit words, the address is incremented by 4. The slave captures the data and address of the first request.
- D. The master issues the next write in the burst, incrementing MAddr. The slave captures the data and address of the second request.
- E. The master issues the final write in the burst, incrementing MAddr, and asserting MBurstLast. The slave captures the data and address of the third request.
- F. The slave captures the data and address of the last request.

Non-Pipelined Read

Figure 20 shows three read transfers to a slave that cannot pipeline responses after requests. This is the typical behavior of legacy computer bus protocols with a single WAIT or ACK signal. In each transfer, SCmdAccept is asserted in the same cycle that SResp is DVA. Therefore, the request-to-response latency is always 0, but the request accept latency varies from 0 to 2.

Figure 20 Non-Pipelined Read



Sequence

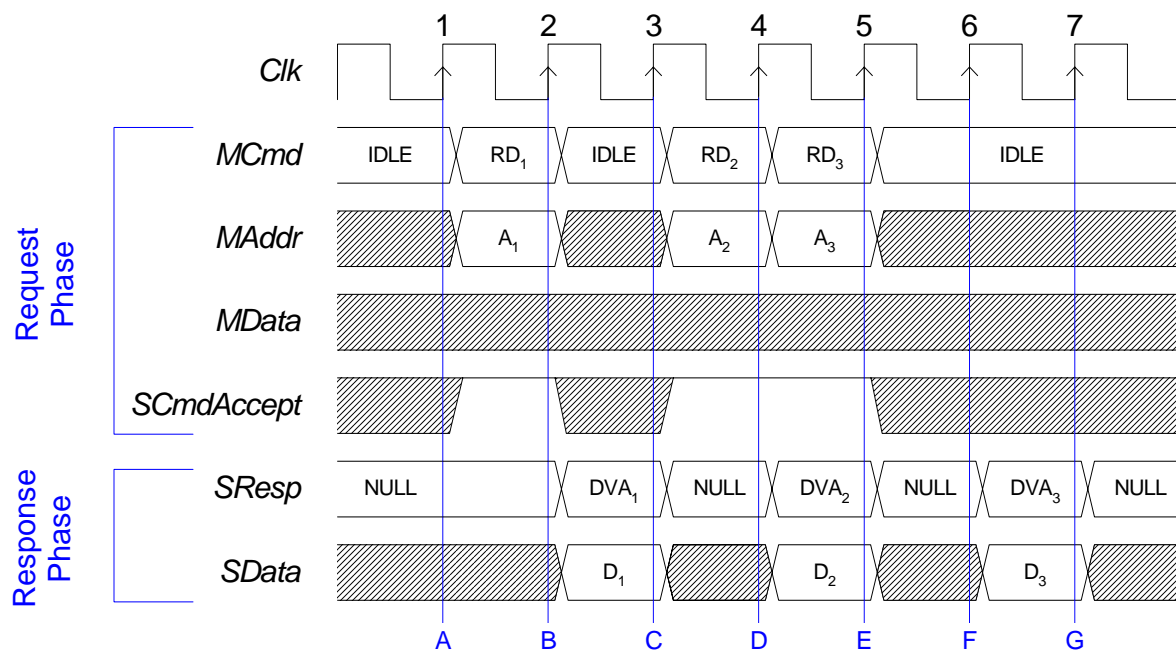
- The master starts the first read request, driving RD on MCmd and a valid address on MAddr. The slave asserts SCmdAccept, for a request accept latency of 0. When the slave sees the read command, it responds with DVA on SResp and valid data on SData. (This requires a combinational path in the slave from MCmd, and possibly other request phase fields, to SResp, and possibly other response phase fields.)
- The master launches another read request. It also sees that SResp is DVA and captures the read data from SData. The slave is not ready to respond to the new request, so it deasserts SCmdAccept.
- The master sees that SCmdAccept is low and extends the request phase. The slave is now ready to respond in the next cycle, so it simultaneously asserts SCmdAccept and drives DVA on SResp and the selected data on SData. The request accept latency is 1.
- Since SCmdAccept is asserted, the request phase ends. The master sees that SResp is now DVA and captures the data.
- The master launches a third read request. The slave deasserts SCmdAccept.

- F. The slave asserts SCmdAccept after 2 cycles, so the request accept latency is 2. It also drives DVA on SResp and the read data on SData.
- G. The master sees that SCmdAccept is asserted, ending the request phase. It also sees that SResp is now DVA and captures the data.

Pipelined Request and Response

Figure 21 shows three read transfers using pipelined request and response semantics. In each case, the request is accepted immediately, while the response is returned in the same or a later cycle.

Figure 21 Pipelined Request and Response



Sequence

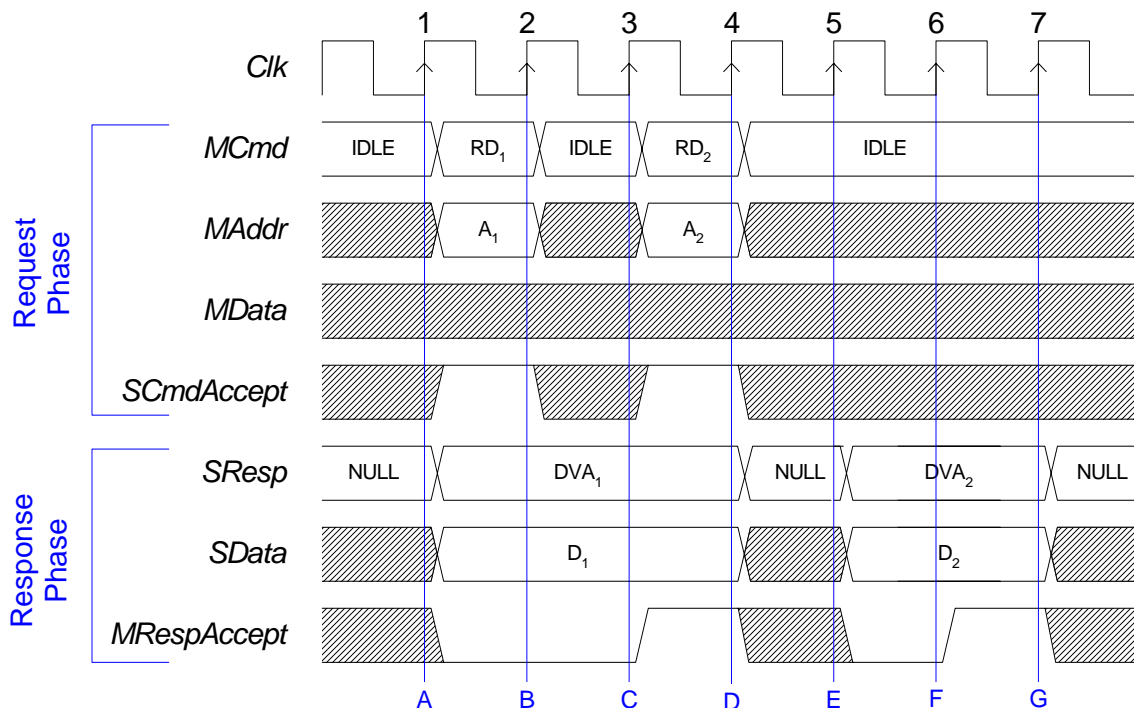
- A. The master starts the first read request, driving RD on MCmd and a valid address on MAddr. The slave asserts SCmdAccept, for a request accept latency of 0.
- B. Since SCmdAccept is asserted, the request phase ends. The slave responds to the first request with DVA on SResp and valid data on SData.
- C. The master launches a read request and the slave asserts SCmdAccept. The master sees that SResp is DVA and captures the read data from SData. The slave drives NULL on SResp, completing the first response phase.

- D. The master sees that SCmdAccept is asserted, so it can launch a third read even though the response to the previous read has not been received. The slave captures the address of the second read and begins driving DVA on SResp and the read data on SData.
- E. Since SCmdAccept is asserted, the third request ends. The master sees that the slave has produced a valid response to the second read and captures the data from SData. The request-to-response latency for this transfer is 1.
- F. The slave has the data for the third read, so it drives DVA on SResp and the data on SData.
- G. The master captures the data for the third read from SData. The request-to-response latency for this transfer is 2.

Response Accept

Figure 22 shows examples of the response accept extension used with two read transfers. An additional field, MRespAccept, is added to the response phase. This signal may be used by the master to flow-control the response phase.

Figure 22 Response Accept



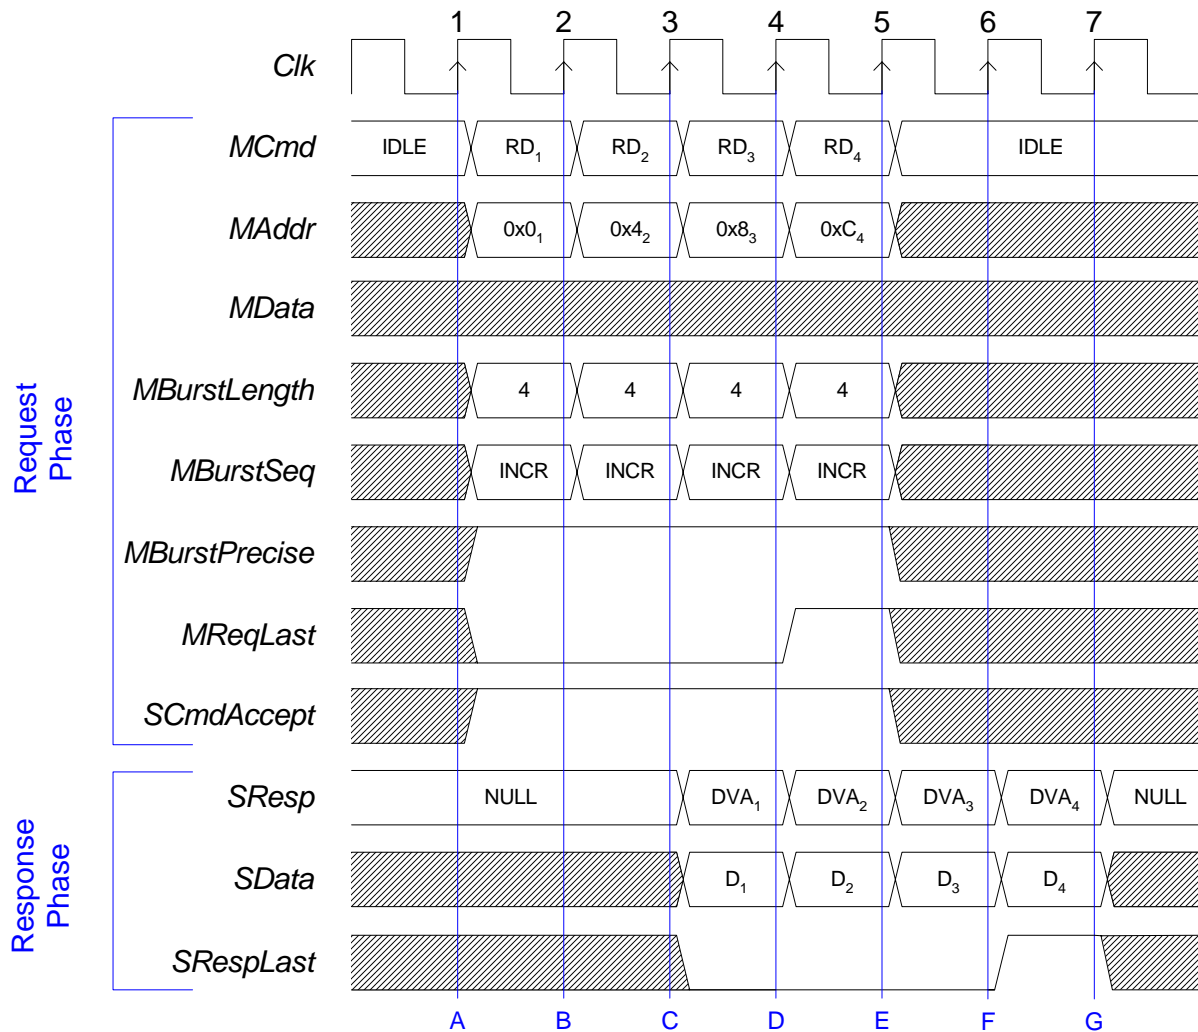
Sequence

- A. The master starts a read request by driving RD on MCmd and a valid address on MAddr. The slave asserts SCmdAccept immediately, and it drives DVA on SResp and the read data on SData as soon as it sees the read request. The master is not ready to receive the response for the request it just issued, so it deasserts MRespAccept.
- B. Since SCmdAccept is asserted, the request phase ends. The master continues to deassert MRespAccept, however. The slave holds SResp and SData steady.
- C. The master starts a second read request and is ready for the response from its first request, so it asserts MRespAccept. This corresponds to a response accept latency of 2.
- D. Since SCmdAccept is asserted, the request phase ends. The master captures the data for the first read from the slave. Since MRespAccept is asserted, the response phase ends. The slave is not ready to respond to the second read, so it drives NULL on SResp.
- E. The slave responds to the second read by driving DVA on SResp and the read data on SData. The master is not ready for the response, however, so it deasserts MRespAccept.
- F. The master asserts MRespAccept, for a response accept latency of 1.
- G. The master captures the data for the second read from the slave. Since MRespAccept is asserted, the response phase ends.

Incrementing Precise Burst Read

Figure 23 illustrates a burst of four 32-bit words, incrementing precise burst read, with burst framing information (MReqLast/SRespLast). Since the burst is precise, the MBurstLength signal is constant during the whole burst. MReqLast flags the last request of the burst, and SRespLast flags the last response of the burst.

Figure 23 Incrementing Precise Burst Read



Sequence

- A. The master starts a read request by driving RD on *MCmd*, a valid address on *MAddr*, four on *MBurstLength*, INCR on *MBurstSeq*, and asserts *MBurstPrecise*. *MBurstLength*, *MBurstSeq* and *MBurstPrecise* must be kept constant during the burst. *MReqLast* must be deasserted until the last request in the burst. The slave is ready to accept any request, so it asserts *SCmdAccept*.
- B. The master issues the next read in the burst. *MAddr* is set to the next word-aligned address (incremented by 4 in this case). The slave captures the address of the first request and keeps *SCmdAccept* asserted.
- C. The master issues the next read in the burst. *MAddr* is set to the next word-aligned address (incremented by 4 in this case). The slave captures the address of the second request and keeps *SCmdAccept* asserted. The slave responds to the first read by driving DVA on *SResp* and the read data on *SData*.

- D. The master issues the last request of the burst, incrementing MAddr and asserting MReqLast. The master also captures the data for the first read from the slave. The slave responds to the second request, and captures the address of the third request.
- E. The master captures the data for the second read from the slave. The slave responds to the third request and captures the address of the fourth.
- F. The master captures the data for the third read from the slave. The slave responds to the fourth request and asserts SRespLast to indicate the last response of the burst.
- G. The master captures the data for the last read from the slave, ending the last response phase.

Incrementing Imprecise Burst Read

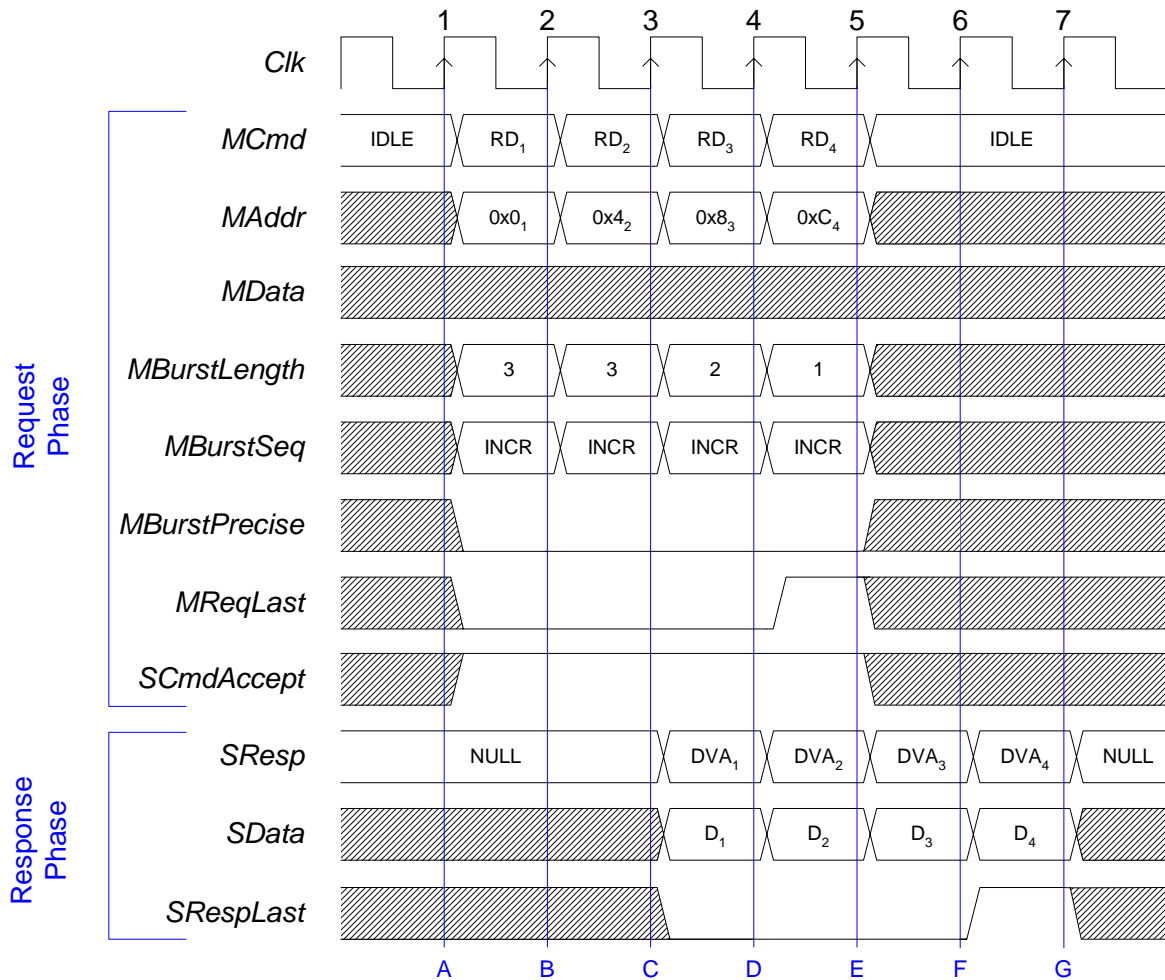
Figure 24 illustrates a burst of four 32-bit words, incrementing imprecise burst read, with burst framing information (MReqLast/SRespLast). MReqLast flags the last request of the burst and SRespLast flags the last response of the burst. The last burst request is signaled primarily by driving the value 1 on MBurstLength.

The burst length sequence (3,3,2,1) is chosen arbitrarily for illustration purposes. The protocol requires that the burst length of the last transfer of the burst be equal to one.

Sequence

- A. The master starts a read request by driving RD on MCmd, a valid address on MAddr, three on MBurstLength, INCR on MBurstSeq, and asserts MBurstPrecise. The burst length is the best guess of the master at this point. MBurstSeq and MBurstPrecise are kept constant during the burst. MReqLast must be deasserted until the last request in the burst. The slave is ready to accept any request, so it asserts SCmdAccept.
- B. The master issues the next read in the burst. MAddr is set to the next word-aligned address (incremented by 4 in this case). The MBurstLength is set to three, since the master knows the burst is longer than it originally thought. The slave captures the address of the first request and keeps SCmdAccept asserted.
- C. The master issues the next read in the burst. MAddr is set to the next word-aligned address (incremented by 4 in this case). The MBurstLength is set to two. The slave captures the address of the second request, and keeps SCmdAccept asserted. The slave responds to the first read by driving DVA on SResp and the read data on SData.

Figure 24 Incrementing Imprecise Burst Read

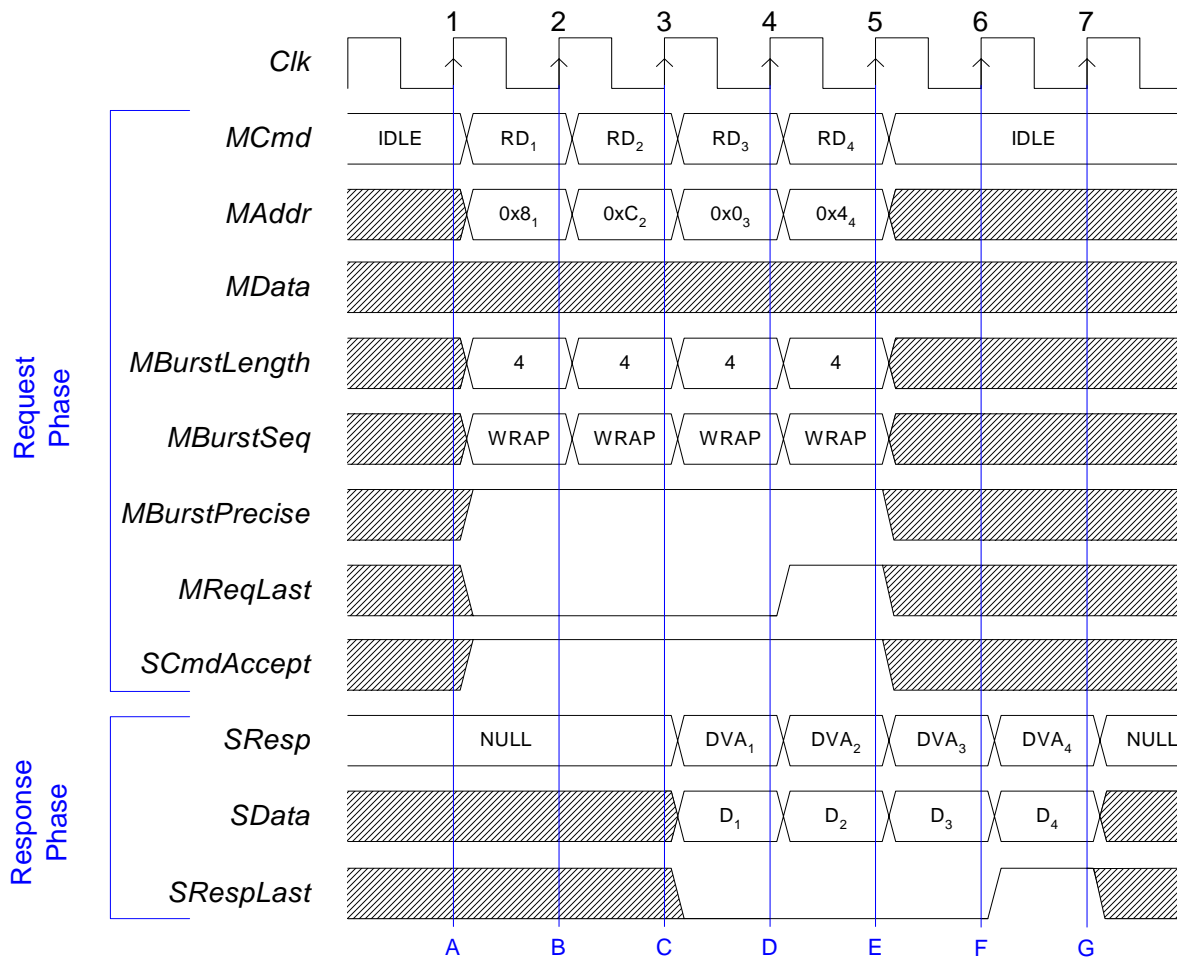


- D. The master issues the last request of the burst, incrementing MAddr, setting MBurstLength to one, and asserting MReqLast. The master also captures the data for the first read from the slave. The slave responds to the second request and captures the address of the last request.
- E. The master captures the data for the second read from the slave. The slave responds to the third request.
- F. The master captures the data for the third read from the slave. The slave responds to the fourth request and asserts SRespLast to indicate the last response of the burst.
- G. The master captures the data for the last read from the slave, ending the last response phase.

Wrapping Burst Read

Figure 25 illustrates a burst of four 32-bit words, wrapping burst read, with optional burst framing information (MReqLast/SRespLast). MReqLast flags the last request of the burst and SRespLast flags the last response of the burst. As a wrapping burst is precise, the MBurstLength signal is constant during the whole burst, and must be power of two. The wrapping burst address must be aligned to boundary MBurstLength times the OCP word size in bytes.

Figure 25 Wrapping Burst Read



Sequence

- A. The master starts a read request by driving RD on MCmd, a valid address on MAddr, four on MBurstLength, WRAP on MBurstSeq, and asserts MBurstPrecise. MBurstLength, MBurstSeq, and MBurstPrecise must be kept constant during the burst. MReqLast must be deasserted until the last request in the burst. The slave is ready to accept any request, so it asserts SCmdAccept.

- B. The master issues the next read in the burst. MAddr is set to the next word-aligned address (incremented by 4 in this case). The slave captures the address of the first request, and keeps SCmdAccept asserted.
- C. If incremented, the next address would be 0x10. Since the first transfer was from address 0x8 and the burst length is 4, the addresses must be between 0 and 0xF. The master wraps the MAddr to 0, which is the closest alignment boundary. (If the first address were 0x14, the address would wrap to 0x10, rather than 0x20.) The slave captures the address of the second request, and keeps SCmdAccept asserted. The slave responds to the first read by driving DVA on SResp and valid data on SData.
- D. The master issues the last request of the burst, incrementing MAddr and asserting MReqLast. The master also captures the data for the first read from the slave. The slave responds to the second request and captures the address of the third.
- E. The master captures the data for the second read from the slave. The slave responds to the third request and captures the address of the fourth.
- F. The master captures the data for the third read from the slave. The slave responds to the fourth request and asserts SRespLast to indicate the last response of the burst.
- G. The master captures the data for the last read from the slave, ending the last response phase.

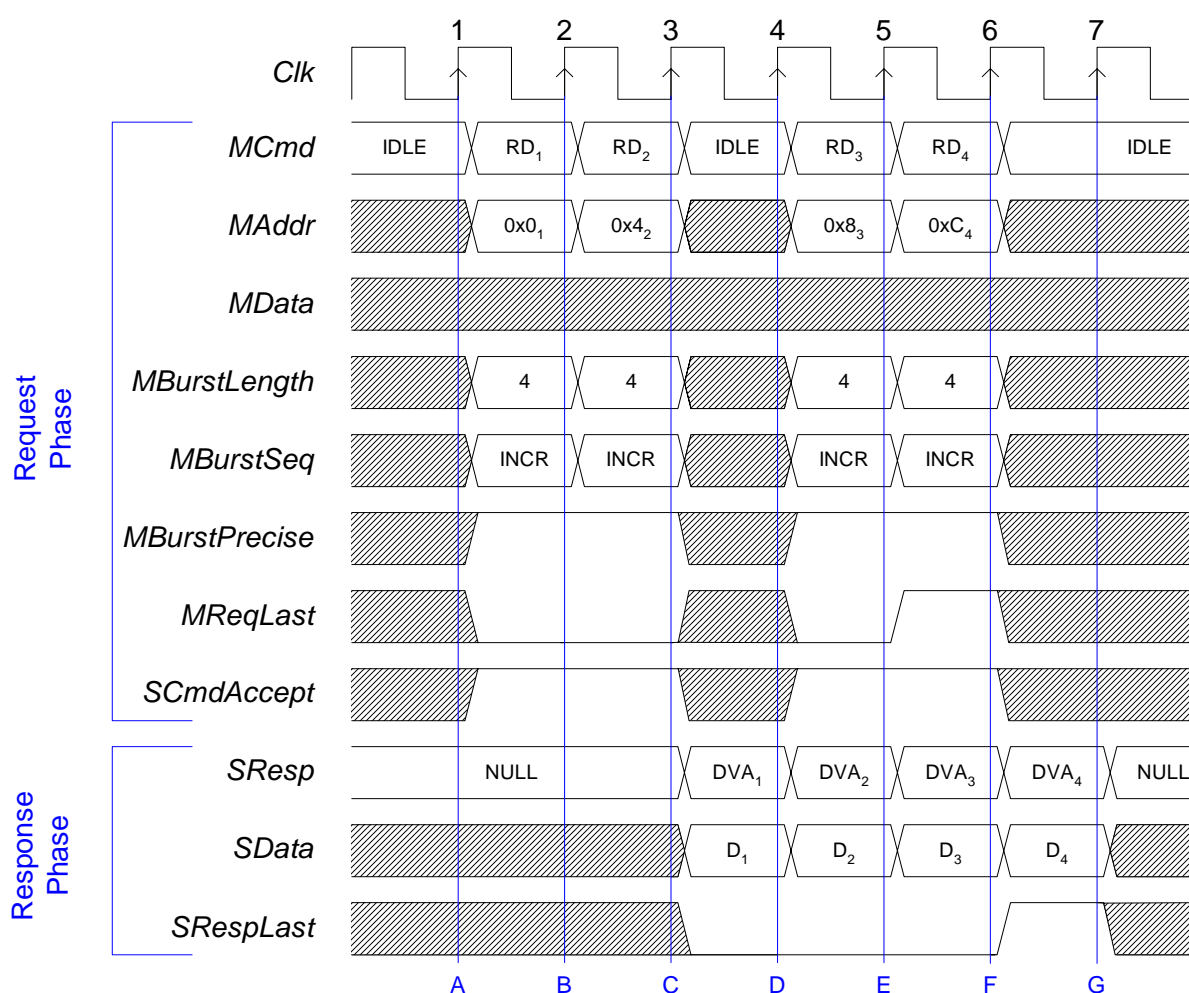
Incrementing Burst Read with IDLE Request Cycle

Figure 26 illustrates a burst of four 32-bit words, incrementing precise burst read, with an IDLE cycle inserted in the middle. The master may insert IDLE requests in any burst type.

Sequence

- A. The master starts a read request by driving RD on MCmd, a valid address on MAddr, four on MBurstLength, INCR on MBurstSeq, and asserts MBurstPrecise. MBurstLength, MBurstSeq, and MBurstPrecise must be kept constant during the burst. MReqLast must be deasserted until the last request in the burst. The slave is ready to accept any request, so it asserts SCmdAccept.
- B. The master issues the next read in the burst. MAddr is set to the next word-aligned address (incremented by 4 in this case). The slave captures the address of the first request and keeps SCmdAccept asserted.
- C. The master inserts an IDLE request in the middle of the burst. The slave does not have to deassert SCmdAccept, anticipating more burst requests to come. The slave captures the address of the second request. The slave responds to the first read by driving DVA on SResp and the read data on SData. The slave must keep SRespLast deasserted until the last response.

Figure 26 Incrementing Precise Burst Read with IDLE Cycle

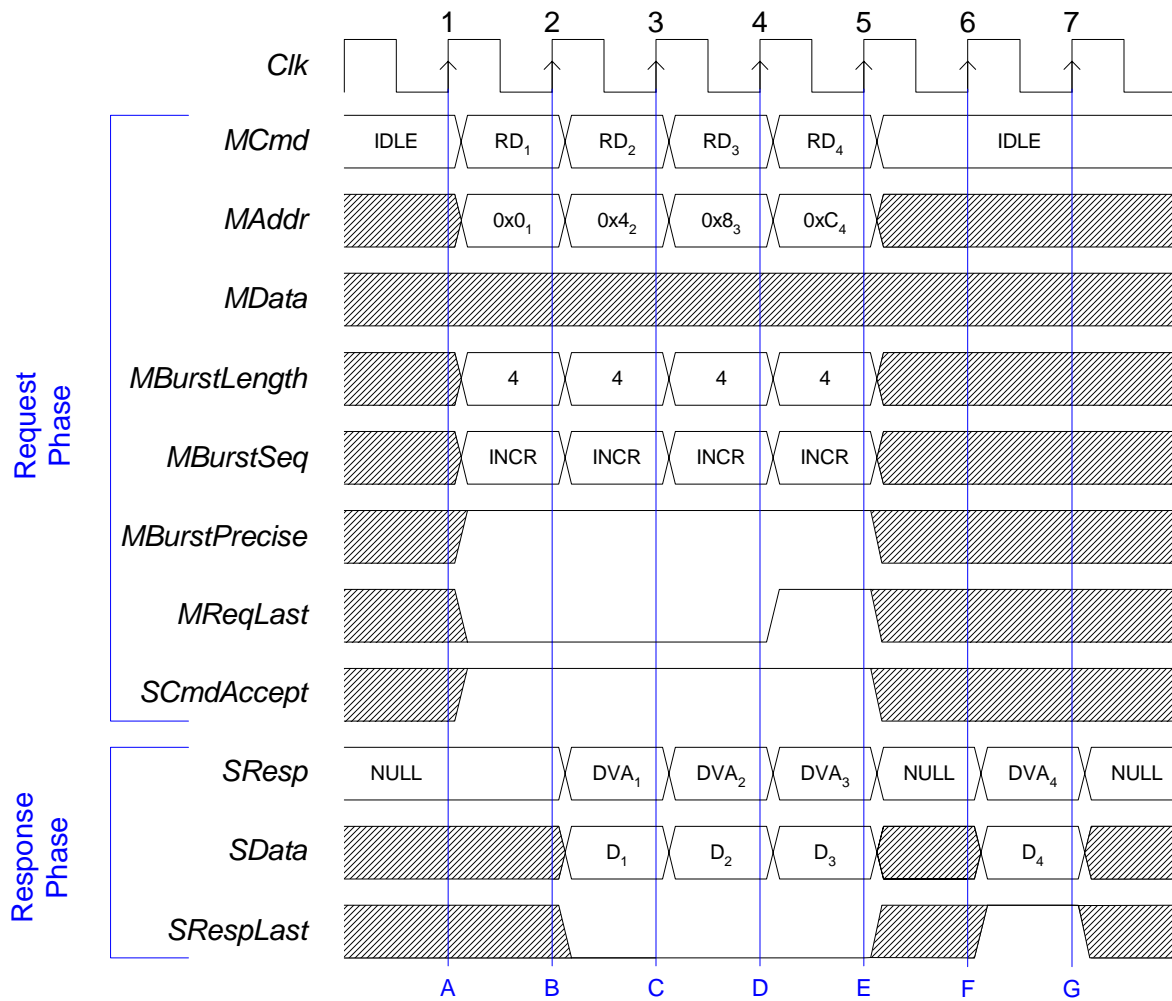


- D. The master issues the next read in the burst. *MAddr* is set to the next word-aligned address (incremented by 4 in this case). The master also captures the data for the first read from the slave. The slave responds to the second read by driving *DVA* on *SResp* and the read data on *SData*. If it has the data available for response, the slave does not have to insert a *NULL* response cycle.
- E. The master issues the last request of the burst, incrementing *MAddr*, and asserting *MReqLast*. The master also captures the data for the second read from the slave. The slave captures the address of the third request and responds to the third request.
- F. The master captures the data for the third read from the slave. The slave captures the address of the fourth request. The slave responds to the fourth request, and asserts *SRespLast* to indicate the end of the slave burst.
- G. The master captures the data for the last read from the slave, ending the last response phase.

Incrementing Burst Read with NULL Response Cycle

Figure 27 illustrates a burst of four 32-bit words, incrementing precise burst read, with a NULL response cycle (wait state) inserted by the slave. Null cycles can be inserted into any burst type by the slave.

Figure 27 Incrementing Burst Read with Null Cycle



Sequence

- The master starts a read request by driving RD on MCmd, a valid address on MAddr, four on MBurstLength, INCR on MBurstSeq, and asserts MBurstPrecise. MBurstLength, MBurstSeq and MBurstPrecise must be kept constant during the burst. MReqLast must be deasserted until the last request in the burst. The slave is ready to accept any request, so it asserts SCmdAccept.

- B. The master issues the next read in the burst. MAddr is set to the next word-aligned address (incremented by 4 in this case). The slave captures the address of the first request and keeps SCmdAccept asserted. The slave responds to the first request by driving DVA on SResp and the read data on SData. The slave must keep SRespLast deasserted until the last response.
- C. The master issues the next read in the burst. MAddr is set to the next word-aligned address (incremented by 4 in this case). The master also captures the data for the first read from the slave. The slave captures the address of the second request and keeps SCmdAccept asserted. The slave responds to the second request.
- D. The master issues the last request of the burst, incrementing MAddr and asserting MReqLast. The master also captures the data for the second read from the slave. The slave captures the address of the third request and keeps SCmdAccept asserted. The slave responds to the third request.
- E. The master captures the data for the third read from the slave. The slave captures the address of the fourth request and keeps SCmdAccept asserted. The slave cannot respond to the fourth request, so it inserts NULL to SResp.
- F. The slave responds to the fourth request and asserts SRespLast to indicate the last response of the burst.
- G. The master captures the data for the last read from the slave, ending the last response phase.

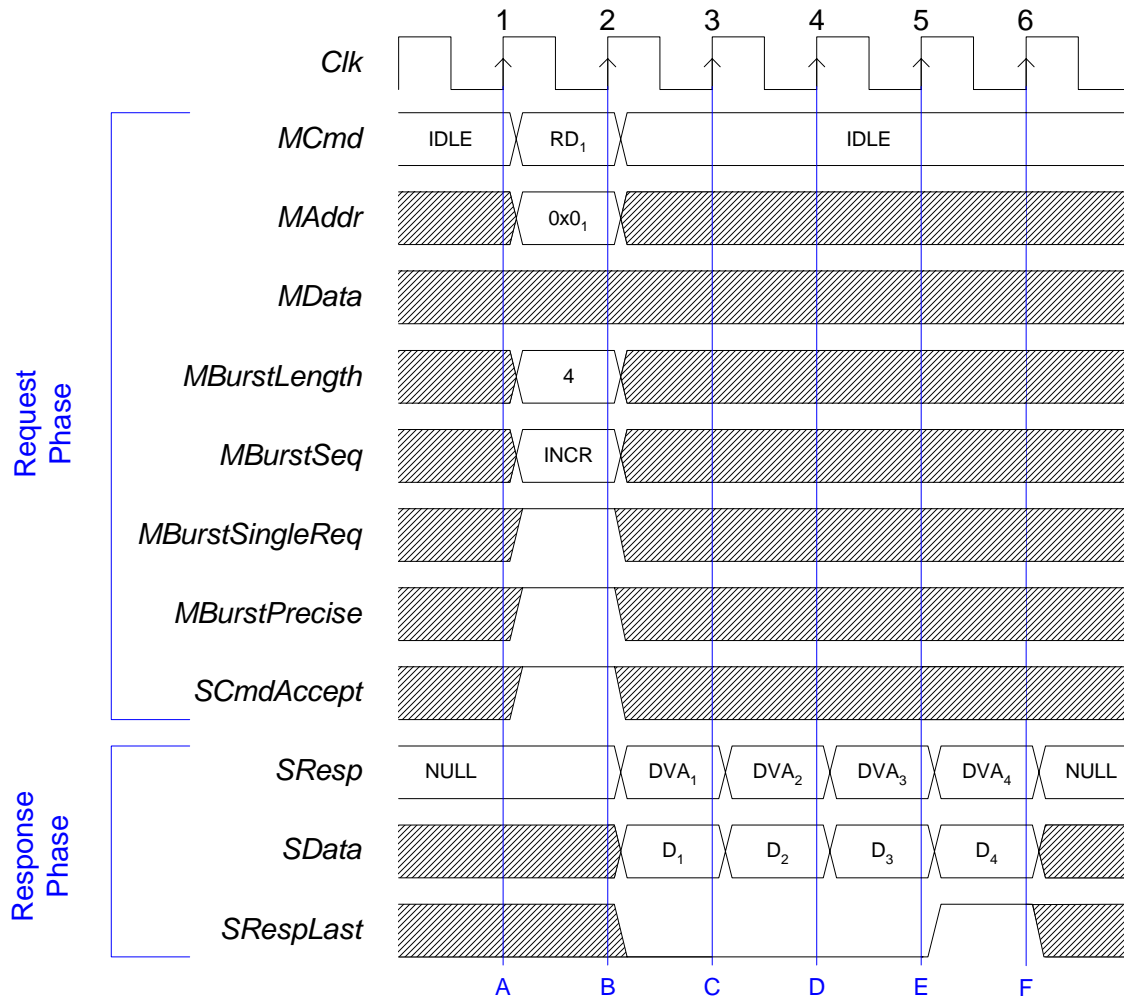
Single Request Burst Read

Figure 28 illustrates a single request, multiple data burst read. The master provides the burst length, start address, and burst sequence, and identifies the burst as a single request with the MBurstSingleReq signal. A single request burst is always precise.

Sequence

- A. The master starts a read request by driving RD on MCmd, a valid address on MAddr, four on MBurstLength, INCR on MBurstSeq, and asserts MBurstPrecise, and MBurstSingleReq. The MBurstPrecise and MBurstSingleReq signals would normally be tied off to logic 1, which is not supplied by the master. The slave is ready to accept any request, so it asserts SCmdAccept.
- B. The master completes the request cycles. The slave captures the address of the request. The slave responds to the request by driving DVA on SResp and the first response data on SData. The slave must keep SRespLast deasserted until the last response.

Figure 28 Single Request Burst Read

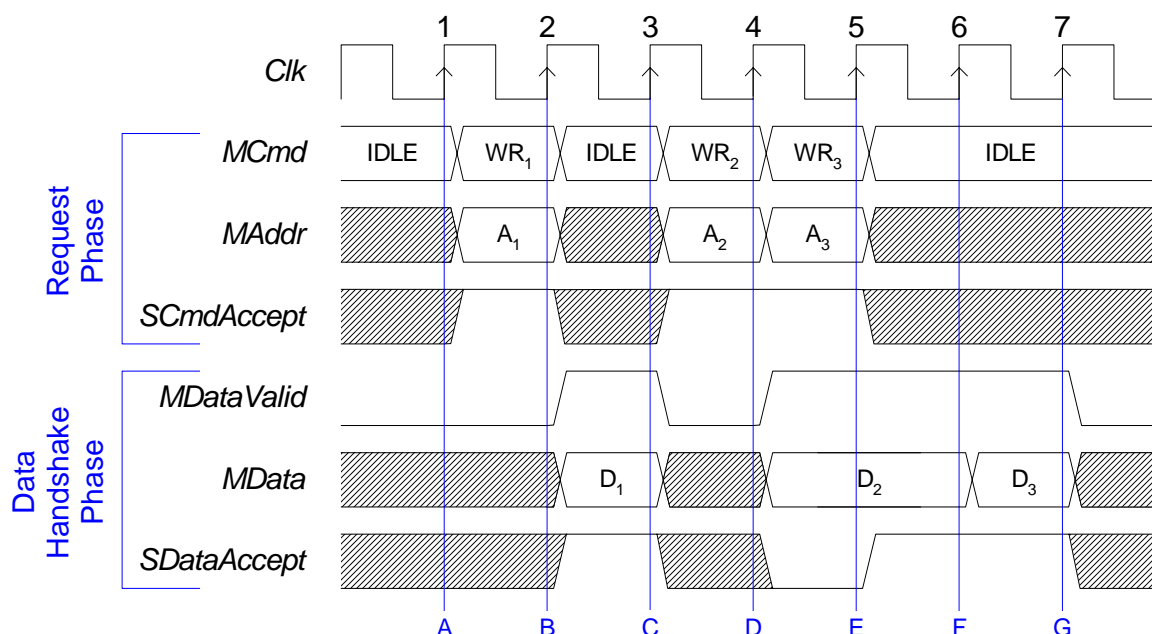


- C. The master captures the first response data. The slave issues the second response.
- D. The master captures the second response data. The slave issues the third response.
- E. The master captures the third response data. The slave issues the fourth response, and asserts SRespLast to indicate the last response of the burst.
- F. The master captures the last response data.

Datahandshake Extension

Figure 29 shows three writes with no responses using the datahandshake extension. This extension adds the datahandshake phase, which is completely independent of the request and response phases. Two signals, MDataValid and SDataAccept, are added, and MData is moved from the request phase to the datahandshake phase.

Figure 29 Datahandshake Extension



Sequence

- A. The master starts a write request by driving WR on MCmd and a valid address on MAddr. It does not yet have the write data, however, so it deasserts MDataValid. The slave asserts SCmdAccept. It does not need to assert or deassert SDataAccept yet, because MDataValid is still deasserted.
- B. The slave captures the write address from the master. The master is now ready to transfer the write data, so it asserts MDataValid and drives the data on MData, starting the datahandshake phase. The slave is ready to accept the data immediately, so it asserts SDataAccept. This corresponds to a data accept latency of 0.
- C. The master deasserts MDataValid since it has no more data to transfer. (Like MCmd and SResp, MDataValid must always be in a valid, specified state.) The slave captures the write data from MData, completing the transfer. The master starts a second write request by driving WR on MCmd and a valid address on MAddr.

- D. Since SCmdAccept is asserted, the master immediately starts a third write request. It also asserts MDataValid and presents the write data of the second write on MData. The slave is not ready for the data yet, so it deasserts SDataAccept.
- E. The master sees that SDataAccept is deasserted, so it holds the values of MDataValid and MData. The slave asserts SDataAccept, for a data accept latency of 1.
- F. Since SDataAccept is asserted, the datahandshake phase ends. The master is ready to deliver the write data for the third request, so it keeps MDataValid asserted and presents the data on MData. The slave captures the data for the second write from MData, and keeps SDataAccept asserted, for a data accept latency of 0 for the third write.
- G. Since SDataAccept is asserted, the datahandshake phase ends. The slave captures the data for the third write from MData.

Burst Write with Combined Request and Data

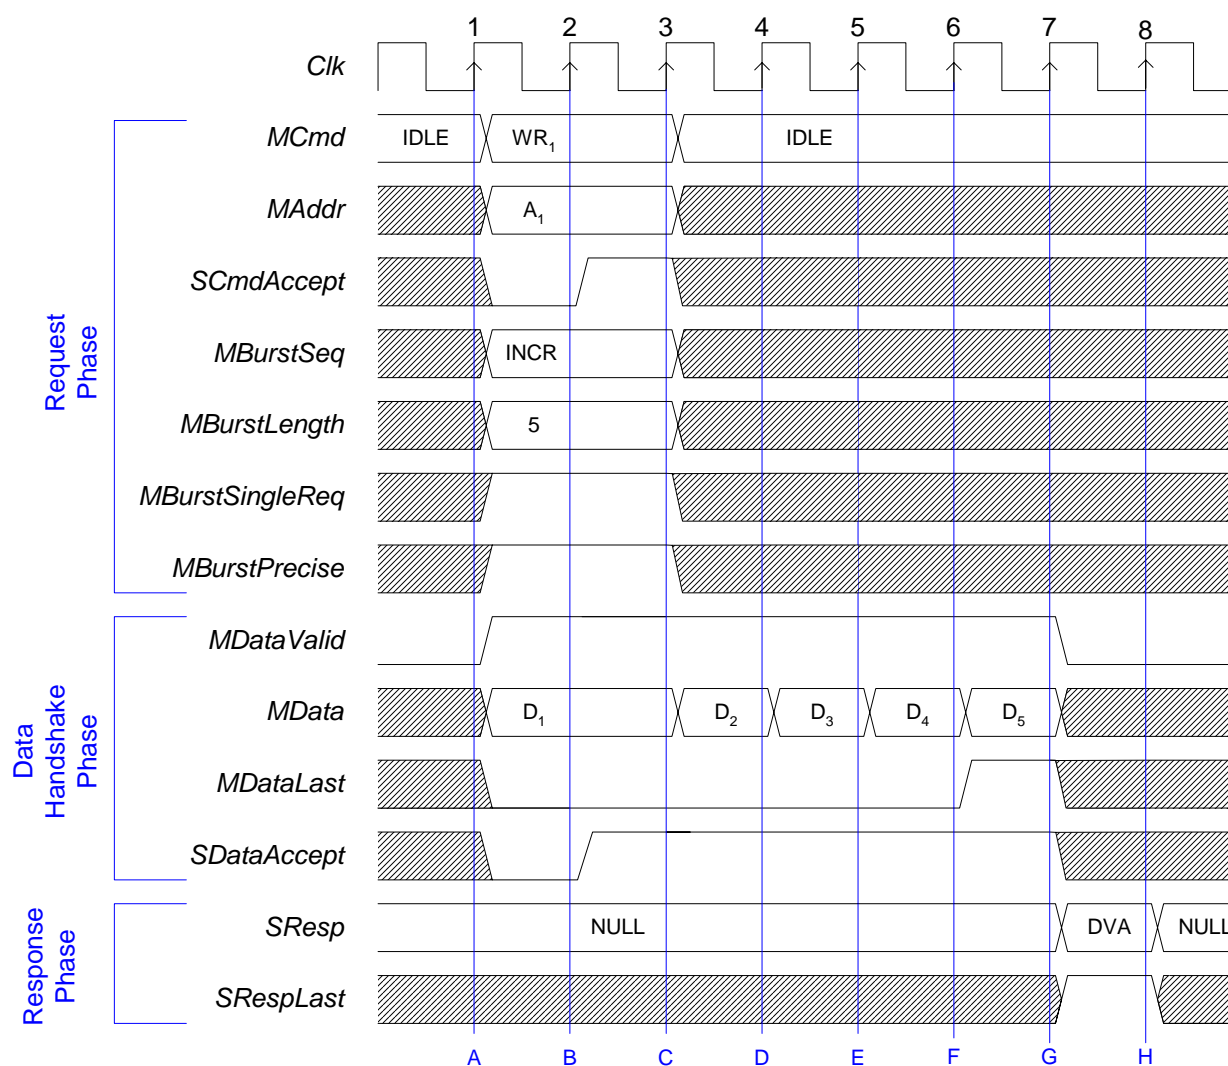
Figure 30 illustrates a single request, multiple data burst write, with datahandshake signaling. Through the request handshake, the master provides the burst length, the start address, and burst sequence, and identifies the burst as a single request with the MBurstSingleReq signal.

The write data is transferred with a datahandshake extension (see Figure 29). The parameter `reqdata_together` forces the first data phase to start with the request, making the design of a slave state machine easier, since it only needs to track one request handshake during the burst. Without this parameter, the MDataValid signal could be asserted any time after the first request. If datahandshake is not used, a single-request write burst is not possible; instead a request is required for each burst word.

Sequence

- A. The master starts a write request by driving WR on MCmd, a valid address on MAddr, INCR on MBurstSeq, five on MBurstLength, and asserts the MBurstPrecise and MBurstSingleReq signals. The master also asserts the MDataValid signal, drives valid data on MData, and deasserts MDataLast. The MDataLast signal must remain deasserted until the last data cycle.
- B. Since it has not received SCmdAccept or SDataAccept, the master holds the request phase signals, keeps MDataValid asserted, and MData steady. The slave asserts SCmdAccept and SDataAccept to indicate it is ready to accept the request and the first data phase.
- C. The master completes the request phase, asserts MDataValid and drives new data to MData. The slave captures the initial data and keeps SDataAccept asserted to indicate it is ready to accept more data.
- D. The master asserts MDataValid and drives new data to MData. The slave captures the second data phase and keeps SDataAccept asserted to indicate it is ready to accept more data.

Figure 30 Burst Write with Combined Request and Data



E. The master asserts MDataValid and drives new data to MData. The slave captures the third data phase and keeps SDataAccept asserted to indicate it is ready to accept more data.

F. The master asserts MDataValid, drives new data to MData, and asserts MDataLast to identify the last data in the burst. The slave captures the fourth data phase and keeps SDataAccept asserted to indicate it is ready to accept more data.

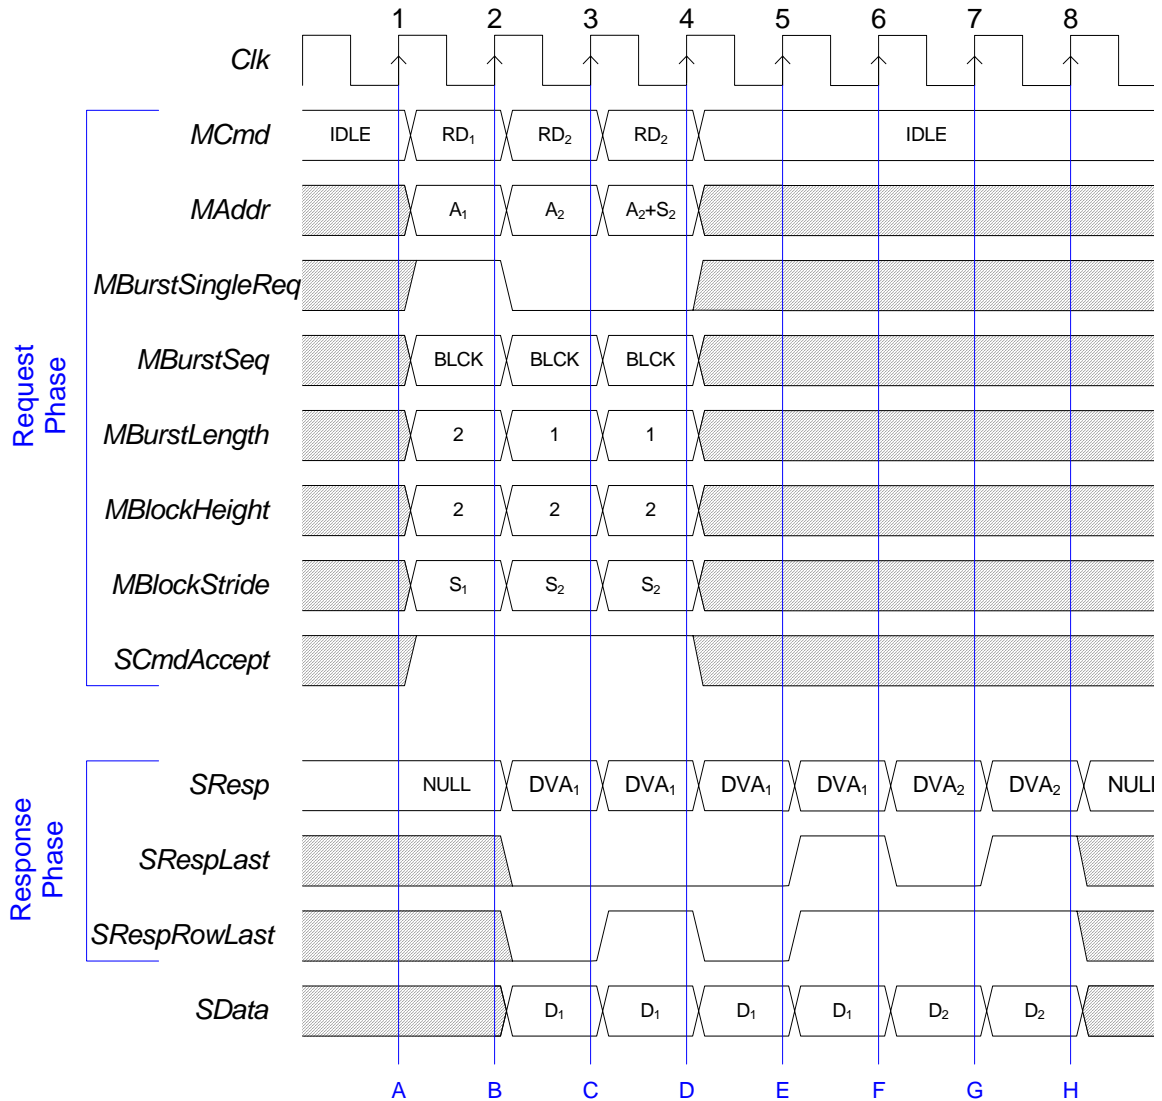
G. The slave captures the last data phase and address.

This example also shows how the slave issues SResp at the end of a burst (when the optional write response is configured in the interface). For single request / multiple data bursts there is only a single response, and it can be issued after the last data has been detected by the slave. The SResp is NULL until point G. in the diagram. The slave may use code DVA to indicate a successful burst, or ERR for an unsuccessful one.

2-Dimensional Block Read

Figure 31 illustrates two read bursts with a 2-dimensional block burst address sequence and optional response phase end-of-row (SRespRowLast) and end-of-burst (SRespLast) framing information. The first transaction is a single-request, multiple-data style block burst of two rows by two words per row, with an address stride of S_1 bytes. The second transaction is a multiple-request, multiple-data style block burst of two rows by one word per row, with an address stride of S_2 bytes. Block bursts are always precise.

Figure 31 2 Dimensional Block Read



Sequence

- A. The master begins the first block read by asserting RD on MCmd, a valid address (A₁) on MAddr, BLCK on MBurstSeq, 2 words per row on MBurstLength, 2 rows on MBlockHeight, and the row-to-row spacing (S₁)

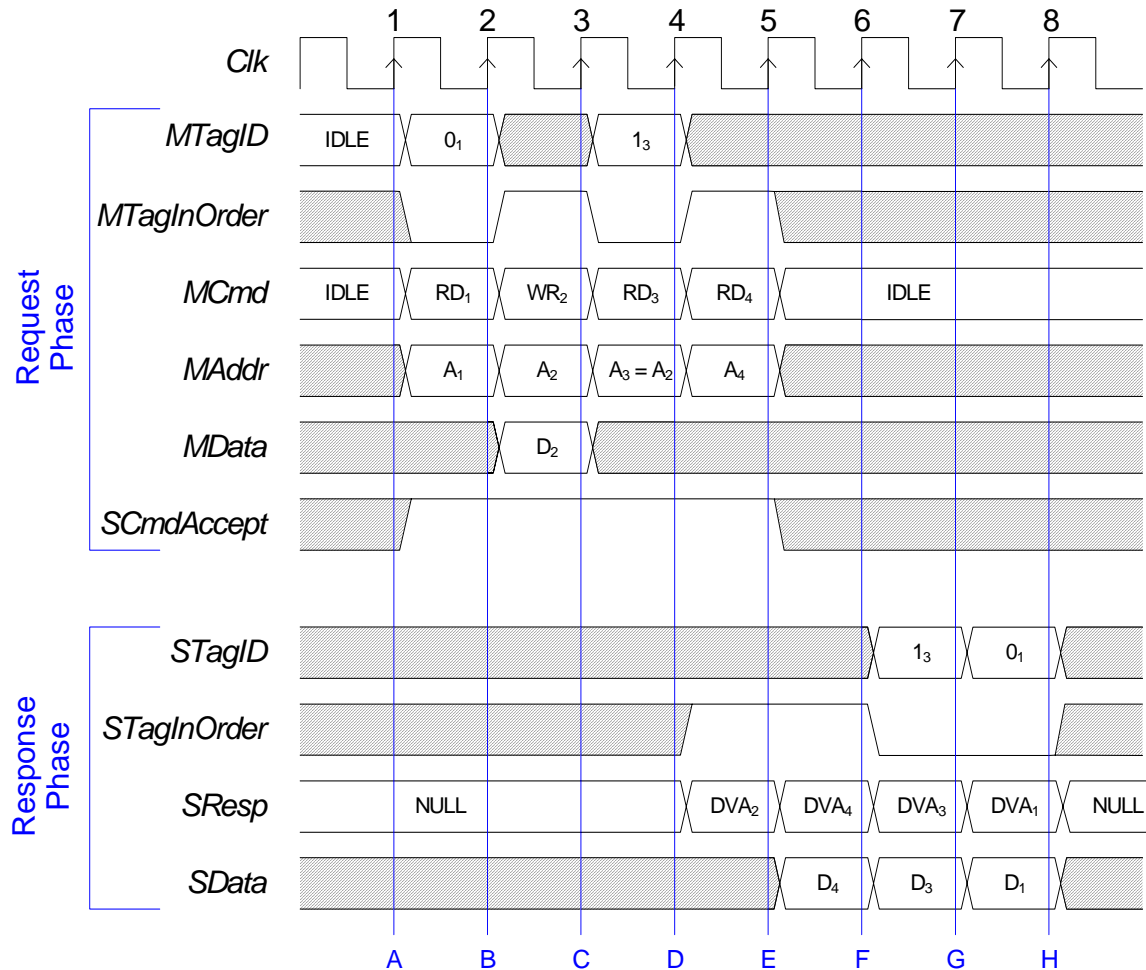
on MBlockStride. The master identifies this as the only request for the read burst by asserting MBurstSingleReq. The slave asserts SCmdAccept signifying that it is ready to accept the request.

- B. The rising edge of the OCP clock ends the first request phase as the slave captures the request. The master starts the second block read at address A_2 , with only a single word per row, and requests 2 rows at a spacing of S_2 . The master deasserts MBurstSingleReq, indicating that there will be one request phase for each data phase. The slave keeps SCmdAccept asserted. The slave also returns a response to the original block burst, including the first word of data. Since there are more data words remaining in the first row of this burst, the slave deasserts SRespRowLast and SRespLast.
- C. The slave captures the first request for the second transaction, and keeps SCmdAccept asserted for the next cycle. The master presents the second (and last) request in the second block burst. The master sets MAddr to the starting address for the second row ($A_2 + S_2$). The master accepts the first response for the first burst. The slave returns the second data word for the first burst, which ends the first row, so the slave also asserts SRespRowLast.
- D. The slave accepts the second request for the second transaction. The master accepts the second response for the first burst. The slave returns the third data word for the first burst, which is the first word of the second row, so the slave deasserts SRespRowLast.
- E. The master accepts the third response for the first burst. The slave returns the fourth (and final) data word for the first burst, and asserts SRespLast and SRespRowLast.
- F. The master accepts the last response for the first burst. The slave returns the first data word for the second burst, which ends the first row, so the slave keeps SRespRowLast asserted and deasserts SRespLast.
- G. The master accepts the first response for the second burst. The slave returns the second (and final) data word for the second burst, and asserts SRespLast.
- H. The master accepts the last response for the second burst.

Tagged Reads

Figure 32 illustrates out-of-order completion of read transfers using the OCP tag extension. The tag IDs, MTagID and STagID, have been added, along with the MTagInOrder and STagInOrder signals. Writes are configured to have responses. There is significant reordering of responses, together with in-order responses forced by both MTagInOrder and address overlap.

Figure 32 Tagged Reads



Sequence

- The master starts the first read request, driving a RD on MCmd and a valid address on MAddr. The master drives a 0 on MTagID, indicating that this read request is for tag 0. The master also deasserts MTagInOrder, allowing the slave to reorder the responses.
- Since SCmdAccept is sampled asserted, the request phase ends with a request accept latency of 0. The master begins a write request to a second address, providing the write data on MData. The master asserts MTagInOrder, indicating that the slave may not reorder this request with respect to other in-order transactions and that MTagID is a “don’t care.”
- When SCmdAccept is sampled asserted, the second request phase ends. The master launches a third request, which is a read to an address that matches the previous write. MTagInOrder is deasserted, enabling reordering, and the assigned tag value is 1.
- Since SCmdAccept is sampled asserted, the third request phase ends. The master launches a fourth request, which is a read. MTagInOrder is asserted, forcing ordering with respect to the earlier in-order write. The

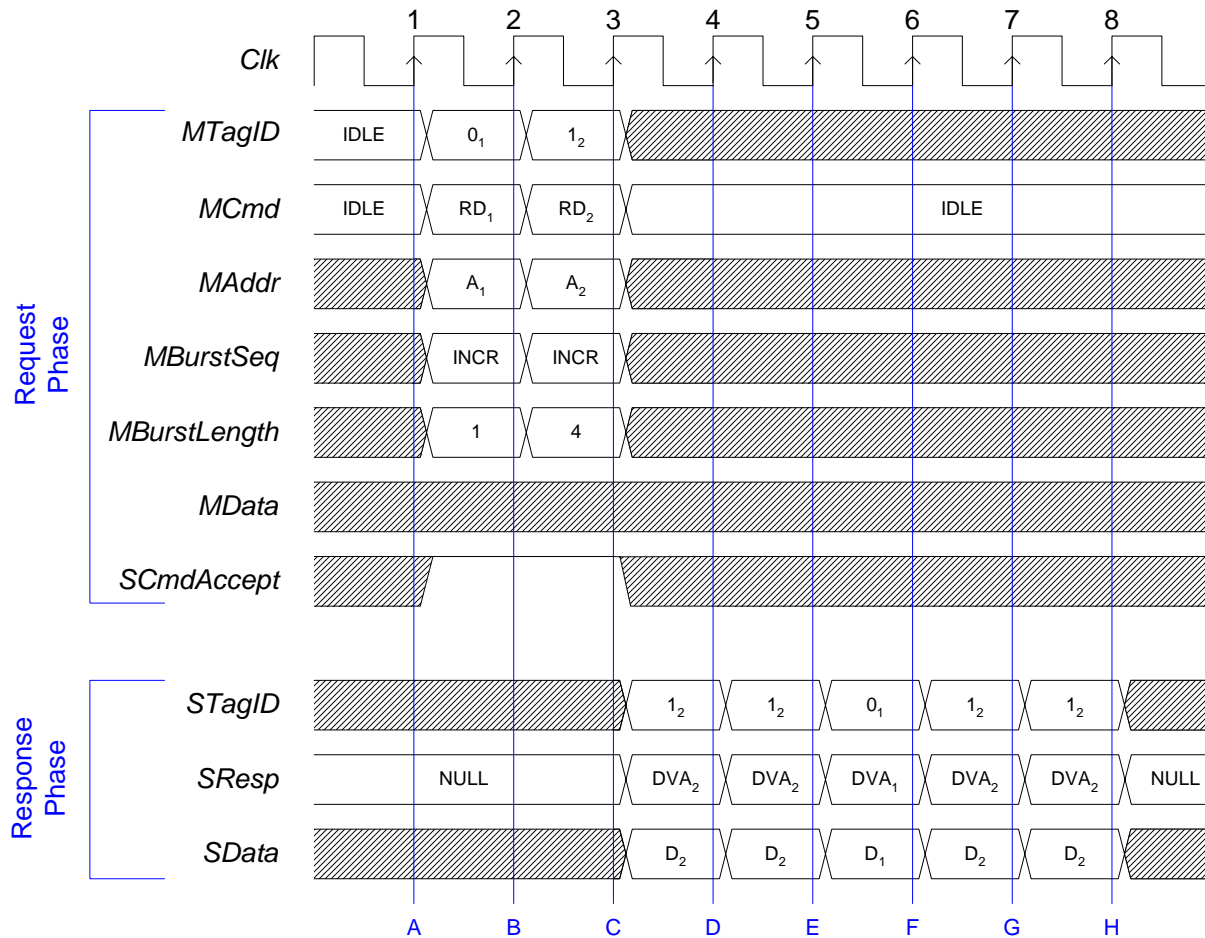
slave responds to the second request (the in-order write presented at B) by driving DVA on SResp. Since the transaction is in-order, STagInOrder is asserted and STagID is a “don’t care.”

- E. SCmdAccept is sampled asserted so the fourth request phase ends. Since the response phase is sampled asserted, the response to the second request ends with a request-to-response latency of 2 cycles. The slave responds to the fourth request (D) by driving DVA on SResp and read data on SData. STagInOrder is asserted to match the associated request. This response was reordered with respect to the first (A) and third (C) requests, which allow reordering.
- F. The response phase is sampled asserted so the response to the fourth request ends with a request-to-response latency of 1 cycle. The slave responds to the third request (C) by driving DVA on SResp, read data on SData, and a 1 on STagID. STagInOrder is deasserted, indicating that reordering is allowed. This response is reordered with respect to the first request (A), but must occur after the second request (B), which has a matching address.
- G. Since the response phase is sampled asserted, the response to the third request ends with a request-to-response latency of 3 cycles. The slave responds to the first request (A) by driving DVA on SResp, read data on SData, and a 0 on STagID. STagInOrder is deasserted.
- H. When the response phase is sampled asserted, the response to the first request ends with a request-to-response latency of 6 cycles.

Tagged Bursts

Figure 33 illustrates out-of-order completion of packing single-request/multiple data read transactions using the OCP tag extension. With the `burstprecise` parameter set to 0, and the `MBurstPrecise` signal tied-off to the default, all bursts are precise. The `burstsinglereq` parameter is 0, and the `MBurstSingleReq` signal is tied-off to 1 (not the default), so all requests have a single request phase. The `taginorder` parameter is set to 0, allowing all transactions to be reordered, subject to the tagging rules. The `tag_interleave_size` parameter is set to 2, so packing bursts must not interleave at a granularity finer than 2 words. Note that the first two words of the second read return before the only word associated with the first read.

Figure 33 Tagged Burst Transactions



Sequence

- The master starts the first read request, driving RD on MCmd and a valid address on MAddr. The request is for a single-word incrementing burst, as driven on MBurstLength and MBurstSeq, respectively. The master also drives a 0 on MTagID, indicating that this read request is for tag 0.
- Once SCmdAccept is sampled asserted, the request phase ends with a request accept latency of 0. The master begins a four-word read request to a second, non-conflicting address, on tag 1.
- SCmdAccept is sampled asserted ending the second request phase. The slave responds to the second request by driving DVA on SResp together with 1 on STagID. The slave provides the first data word from the burst on SData.
- When the response phase is sampled asserted, the first response to the second request ends with a request-to-response latency of 1 cycle. The slave provides the second word of read data for tag 1 on SData, together

with a DVA response. Because `tag_interleave_size` is 2 and the read burst sequence is packing, the slave was forced to return the second word of tag 1's data before responding to tag 0.

- E. The response phase is sampled asserted, terminating the second response to the second request with a request-to-response latency of 2 cycles. The slave responds to the first request by providing the read data for tag 0 together with a DVA response.
- F. When the response phase is sampled asserted, the response to the first request ends with a request-to-response latency of 4 cycles. Since the burst length of the first request is 1, the transaction on tag 0 is complete. The slave provides the third word of read data for tag 1 on `SData`, together with a DVA response.
- G. The response phase is sampled asserted so the third response to the second request ends with a request-to-response latency of 4 cycles. The slave provides the fourth word of read data for tag 1 on `SData`, together with a DVA response.
- H. When the response phase is sampled asserted, the fourth and final response to the second request ends with a request-to-response latency of 5 cycles.

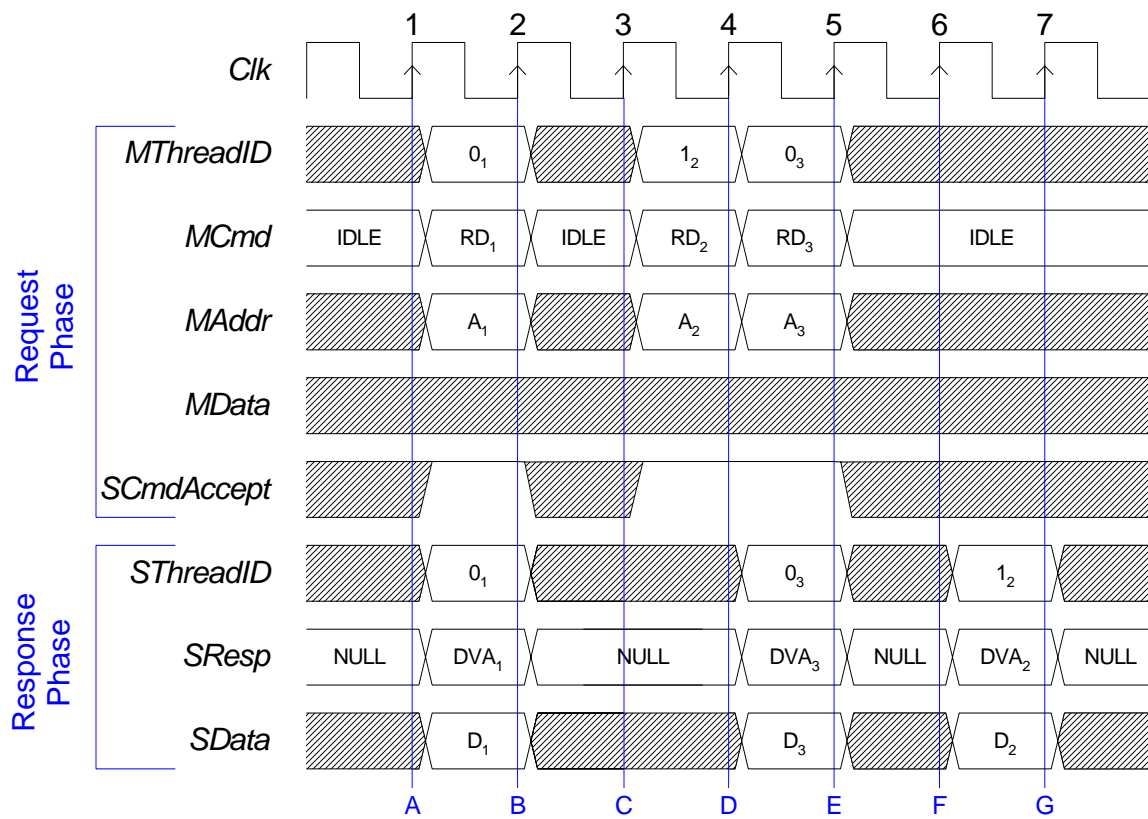
Threaded Read

Figure 34 illustrates out-of-order completion of read transfers using the OCP thread extension. This diagram is developed from Figure 21 on page 114. The thread IDs, `MThreadID` and `SThreadID`, have been added, and the order of two of the responses has been changed.

Sequence

- A. The master starts the first read request, driving `RD` on `MCmd` and a valid address on `MAddr`. The master also drives a 0 on `MThreadID`, indicating that this read request is for thread 0. The slave asserts `SCmdAccept`, for a request accept latency of 0. When the slave sees the read command, it responds with DVA on `SResp` and valid data on `SData`. The slave also drives a 0 on `SThreadID`, indicating that this response is for thread 0.
- B. Since `SCmdAccept` is asserted, the request phase ends. The master sees that `SResp` is DVA and captures the read data from `SData`. Because the request was accepted and the response was presented in the same cycle, the request-to-response latency is 0.
- C. The master launches a new read request, but this time it is for thread 1. The slave asserts `SCmdAccept`, however, it is not ready to respond.

Figure 34 Threaded Read



- D. Since SCmdAccept is asserted, the master can launch another read request. This request is for thread 0, so MThreadID is switched back to 0. The slave captures the address of the second read for thread 1, but it begins driving DVA on SResp, data on SData, and a 0 on SThreadID. This means that it is responding to the third read, before the second read.
- E. Since SCmdAccept is asserted, the third request ends. The master sees that the slave has produced a valid response to the third read and captures the data from SData. The request-to-response latency for this transfer is 0.
- F. The slave has the data for the second read, so it drives DVA on SResp, data on SData, and a 1 on SThreadID.
- G. The master captures the data for the second read from SData. The request-to-response latency for this transfer is 3.

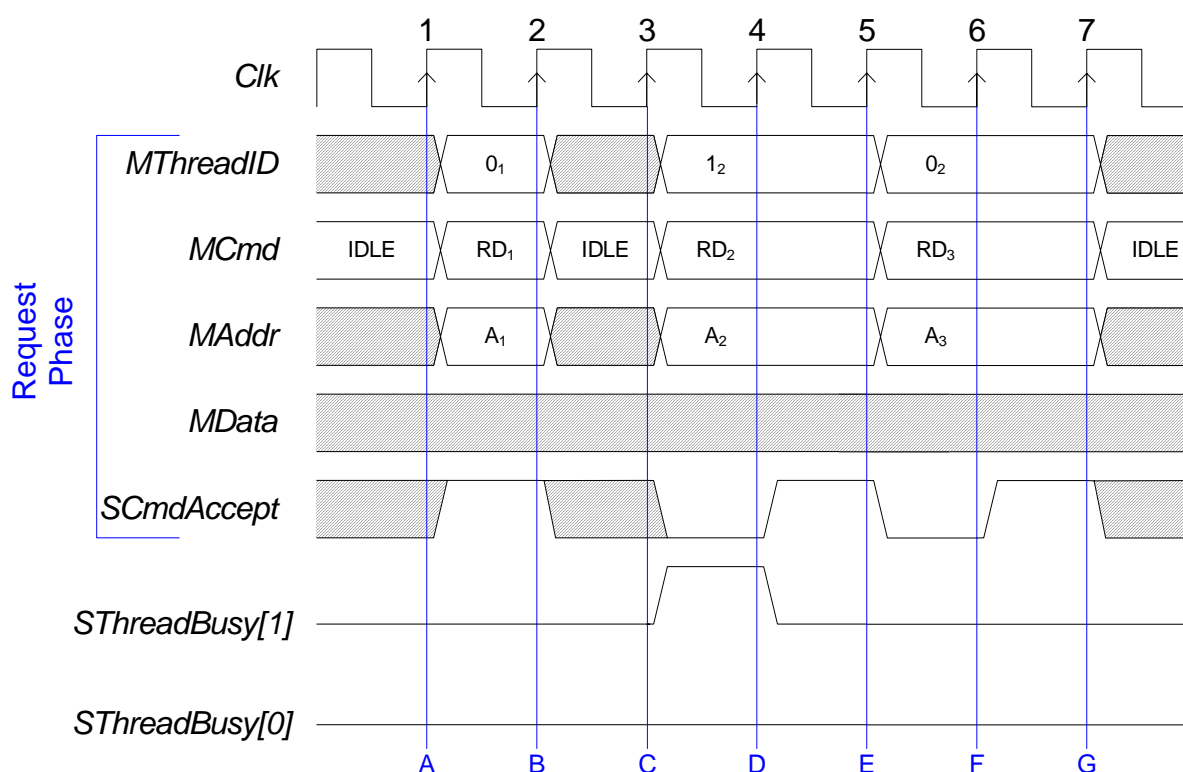
Threaded Read with Thread Busy

Figure 35 illustrates the out-of-order completion of read transfers using the OCP thread extension. The change to Figure 34 is the addition of thread busy signals. In this example, the thread busy is only a hint, since the `sthreadbusy_exact` parameter is not set. In this case the master may ignore the `SThreadBusy` signals, and the slave does not have to accept requests even when it is not busy.

When thread busy is treated as a hint and a request or thread is not accepted, the interface may block for all threads. Blocking of this type can be avoided by treating thread busy as an exact signal using the `sthreadbusy_exact` parameter. For an example, see “Threaded Read with Thread Busy Exact”.

This example shows only the request part of the read transfers. The response part can use a similar mechanism for thread busy.

Figure 35 Threaded Read with Thread Busy



Sequence

- A. The master starts the first read request, driving RD on MCmd and a valid address on MAddr. The master also drives a 0 on MThreadID, associating this read request with thread 0. The slave asserts SCmdAccept for a request accept latency of 0.
- B. Since SCmdAccept is asserted, the request phase ends.

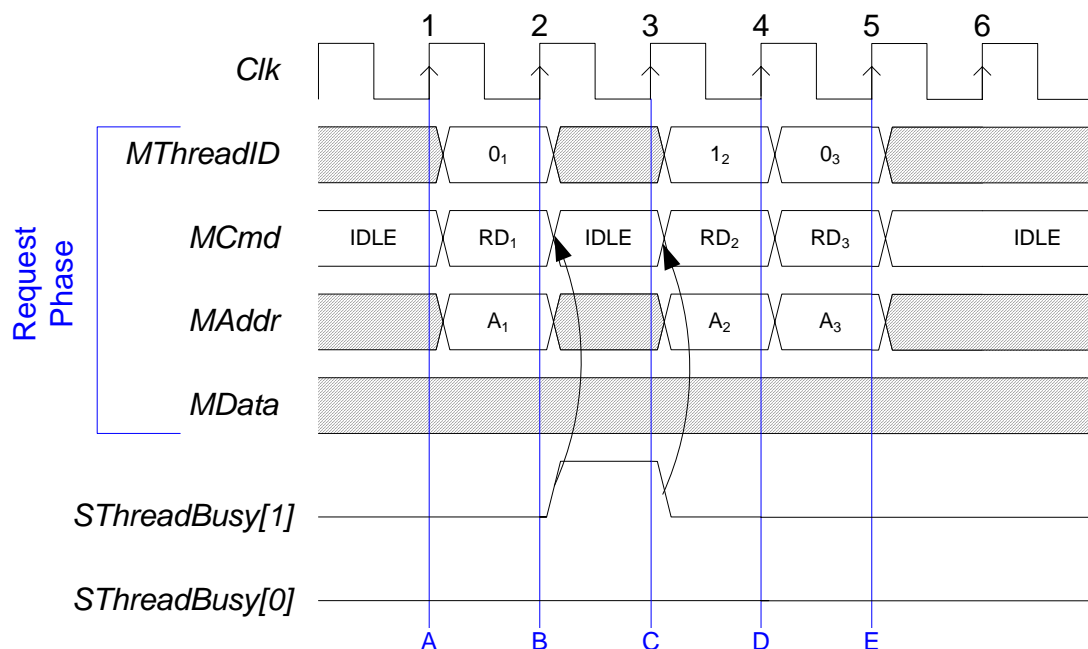
- C. The slave asserts `SThreadBusy[1]` since it is not ready to accept requests on thread 1. The master ignores the hint, and launches a new read request for thread 1. The master can issue a request even though the slave asserts `SThreadbusy` (see transfer 2). All threads are now blocked.
- D. The slave deasserts `SThreadBusy[1]` and asserts `SCmdAccept` to complete the request for thread 1.
- E. Since `SCmdAccept` is asserted, the second request ends. The master issues a new request to thread 0. The slave is not ready to accept the request, and indicates this condition by keeping `SCmdAccept` deasserted. It chooses not to assert `SThreadBusy[0]`. The slave does not have to assert `SCmdAccept` for a request, even if it did not assert `SThreadbusy` (see transfer 3).
- F. The slave asserts the `SCmdAccept` to complete the request on thread 0.
- G. The master captures the `SCmdAccept` to complete the requests.

Threaded Read with Thread Busy Exact

Figure 36 illustrates the out-of-order completion of read transfers using the OCP thread extension. Because the `sthreadbusy_exact` parameter is set, the master may not ignore the `SThreadBusy` signals. The master is using `SThreadBusy` to control thread arbitration, so it cannot present a command on Thread 1 as the slave asserts `SThreadbusy[1]`.

The diagram only shows the request part of the read transfers. The response part can use a similar mechanism for thread busy.

Figure 36 Threaded Read with Thread Busy Exact



Sequence

- A. The master starts the first read request, driving RD on MCmd and a valid address on MAddr. The master also drives a 0 on MThreadID, indicating that this read request is for thread 0.
- B. Since SThreadBusy[0] is not asserted, the request phase ends. The slave samples the data and address and asserts SThreadBusy[1] since it is unready to accept requests on thread 1. The master is prevented from sending a request on thread 1, but it can send a request on another thread.
- C. The slave deasserts SThreadBusy[1] and the master can send the request on thread 1.
- D. Since SThreadBusy[1] is not asserted, the request phase ends and the slave must sample the data and address. The master can send a request on thread 0 (or thread 1).
- E. Since SThreadBusy[0] is not asserted, the request phase ends and the slave must sample the data and address.

Threaded Read with Pipelined Thread Busy

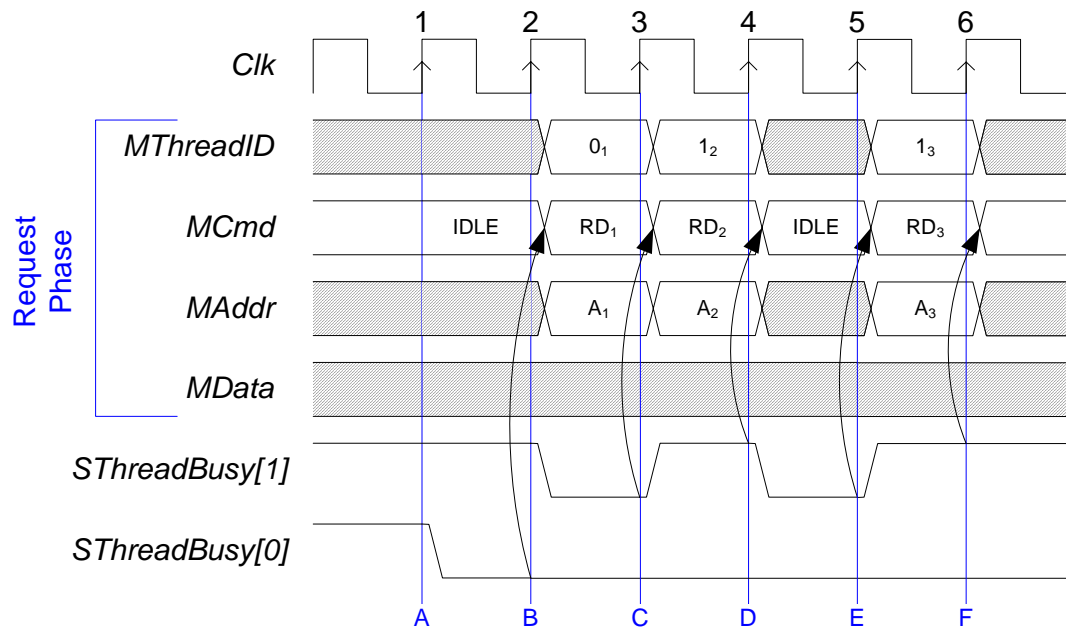
Figure 37 illustrates a set of threaded read requests on an interface where the `sthreadbusy_pipelined` parameter is set. Because pipelining a phase's ThreadBusy signals also forces exact flow control (`sthreadbusy_exact` must be set), the master must obey the SThreadBusy signals.

In this example, the master asserts a single read request phase on thread 0, and multiple requests on thread 1. The slave's SThreadBusy assertions control when the master may assert request phases on each thread. The diagram only shows the request part of the read transfers. The response part uses a similar mechanism for thread busy.

Sequence

- A. Because both SThreadBusy signals were sampled asserted on this rising edge of the OCP clock, the master may not present requests on either thread. The slave indicates its readiness to accept a request on thread 0 in the next cycle by de-asserting SThreadBusy[0].
- B. After sampling SThreadBusy[0] deasserted, the master asserts the first read request on thread 0 by driving a 0 on MThreadID, RD on MCmd and a valid address on MAddr. The slave indicates that it can accept requests on both threads in the next cycle by de-asserting SThreadBusy[1] and leaving SThreadBusy[0] deasserted.

Figure 37 Threaded Read with Pipelined Thread Busy




- C. The master's first request is sampled by the slave and the request phase ends. The master samples SThreadBusy[1] deasserted and uses the information to assert a second read request, this time on thread 1. The slave asserts SThreadBusy[1] since it cannot guarantee that it can accept another request on thread 1 in the next cycle.
- D. The master's second request is sampled by the slave and the request phase ends. The master samples SThreadBusy[1] asserted, and is forced to drive MCmd to IDLE, since it has no more requests for thread 0 and the slave cannot accept a request on thread 1. The slave signals that it will be ready to accept requests on both threads in the next cycle by de-asserting SThreadBusy[1] and leaving SThreadBusy[0] deasserted.
- E. The master samples SThreadBusy[1] deasserted, and uses this information to assert a third read request on thread 1. The slave asserts SThreadBusy[1] since it cannot guarantee that it can accept another request on thread 1 in the next cycle.
- F. The master's third request is sampled by the slave and the request phase ends. The master samples SThreadBusy[1] asserted, and is forced to drive MCmd to IDLE.

Reset

Figure 38 shows the timing of the reset sequence with MReset_n driven from the master to the slave. MReset_n must be asserted for at least 16 cycles of the OCP clock to ensure that the master and slave reach a consistent internal state. Because the interface does not include the EnableClk signal, the OCP clock is simply Clk.

Figure 38 Reset Sequence



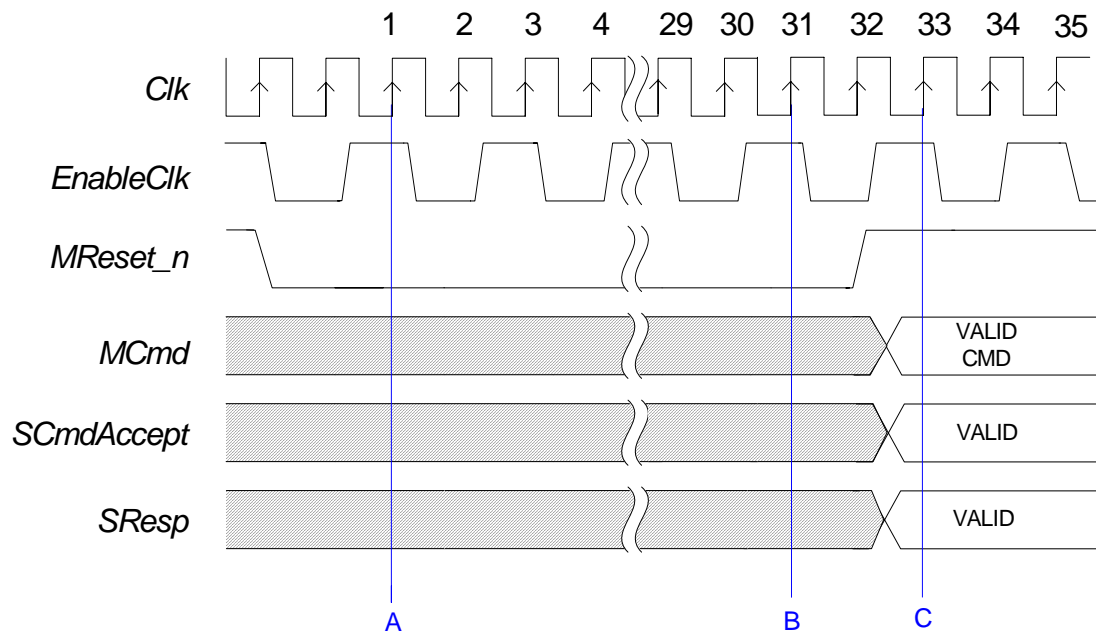
Sequence

- A. MReset_n is sampled active on this clock edge. Master and slave now ignore all other OCP signals. In the first cycle a response to a previously issued request is presented by the slave and ready to be received by the master. Since the master is asserting MReset_n, the response is not received. The associated transaction is terminated by OCP reset so the response is withdrawn by the slave.
- B. MReset_n is asserted for at least 16 Clk cycles.
- C. A new transfer may begin on the same clock edge that MReset_n is sampled deasserted.

Reset with Clock Enable

Figure 39 shows the timing of the reset signal with the EnableClk signal enabled on the interface. In this figure, the EnableClk signal is asserted on every other rising edge of Clk, delivering an OCP clock that is one-half the frequency of Clk. The MReset_n signal is driven from the master to the slave. However, the presence of EnableClk means that MReset_n must be asserted for 16 cycles of the OCP clock (that is, when the rising edge of Clk samples EnableClk asserted), which will require 31 cycles of Clk.

Figure 39 Reset with Clock Enable



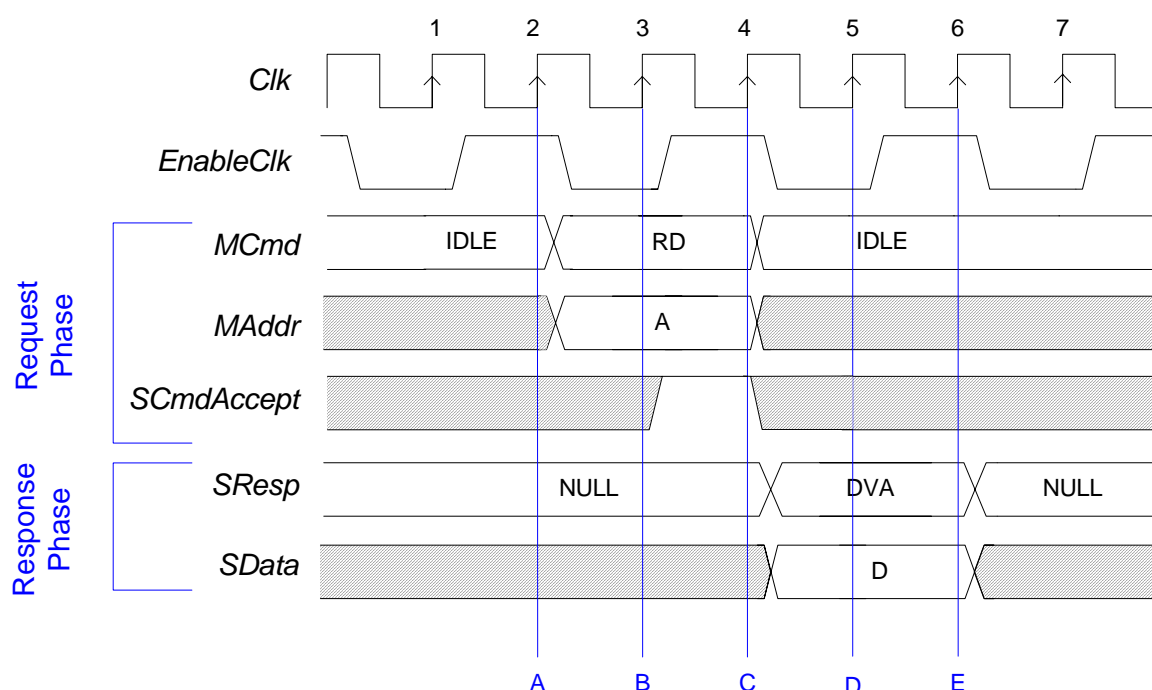
Sequence

- A. MReset_n and EnableClk are sampled active on this clock edge. Master and slave now ignore all other OCP signals.
- B. MReset_n is asserted for at least 16 Clk cycles with EnableClk sampled high.
- C. A new transfer may begin on the same clock edge that MReset_n is sampled deasserted and EnableClk sampled high.

Basic Read with Clock Enable

Figure 40 illustrates a simple read transaction of length one with the EnableClk signal enabled on the interface. In this figure, the EnableClk signal is asserted on every other rising edge of Clk, delivering an OCP clock that is one-half the frequency of Clk. As is shown, interface state only advances on rising edges of Clk that coincide with EnableClk being asserted.

Figure 40 Basic Read with Clock Enable Signal



Sequence

- The master starts a read request by driving *RD* on *MCmd*. At the same time, it presents a valid address on *MAddr*.
- This clock edge is not valid since *EnableClk* is sampled low. The slave asserts *SCmdAccept* in this cycle for a request accept latency of 0.
- The slave uses the value from *MAddr* to determine what data to return. The slave starts the response phase by switching *SResp* from *NULL* to *DVA*. The slave also drives the selected data on *SData*. Since *SCmdAccept* is asserted, the request phase ends.
- Invalid clock edge.
- The master recognizes that *SResp* indicates data valid and captures the read data from *SData*, completing the response phase. This transfer has a request-to-response latency of 1.

9 *Developers Guidelines*

This chapter collects examples and implementation tips that can help you make effective use of the Open Core Protocol and does not provide any additional specification material. This chapter groups together a variety of topics including discussions of:

1. The basic OCP with an emphasis on signal timing, state machines and OCP subsets
2. Simple extensions that cover byte enables, multiple address spaces and in-band information
3. An overview of burst capabilities
4. The concepts of threading, tagging extensions, and connections
5. OCP features addressing write semantics, synchronization issues, and endianness
6. Sideband signals with an emphasis on reset management
7. A description of the debug and test interface

Basic OCP

This section considers the different OCP phases, their relationships, and identifies sensitive timing-related areas and begins with a discussion of support for variable-rate divided clocks. The section includes descriptions of OCP compliant state machines, and also discusses the OCP parameters needed to define simple OCP interfaces.

Divided Clocks

The EnableClk signal allows OCP to provide flexible support for multi-rate systems. When set with the `enableclk` parameter, the EnableClk signal provides a sampling signal that specifies which rising edges of the Clk signal are rising edges of the OCP clock. By driving the appropriate waveforms on EnableClk, the system can control the effective clocking rate of the interface, and frequently the attached cores, without introducing extra outputs from PLLs, or requiring delay matching across multiple clock distribution networks.

When EnableClk is on, the interfaces behave as if the EnableClk signal is not present. All rising edges of Clk are treated as rising edges of the OCP clock allowing the OCP to operate at the Clk frequency. If EnableClk is off, no rising edges of the OCP clock occur, and the OCP clock is effectively stopped.

This feature can be used to reduce dynamic power by idling the attached cores, although the Clk signal may still be active. In normal operation the system drives EnableClk with a periodic signal. For instance, asserting EnableClk for every third Clk cycle causes OCP to operate at one third of the Clk frequency. The system can modify the frequency by changing the repeating pattern on EnableClk.

OCP is fully synchronous (with the exception of reset assertion). All timing paths traversing OCP close in a single OCP clock period. If EnableClk has a maximum duty cycle less than 100%, these timing paths may be constrained as multi-cycle timing paths of the underlying clock domain.

OCP Clock Shape

The *OCP Specification* defines synchronous signals with respect to the rising edge of the OCP clock and makes no assertions about the duty cycle of the OCP clock. Since most designs use the rising-edge clocked flip-flops as the storage element in synchronous designs this is usually not an issue. The OCP Clk signal is frequently the output of a PLL or DLL, which tend to output clock signals with near 50% duty cycles.

An OCP interface with a repeating pattern on EnableClk tends to produce pulsed OCP clock waveforms. For instance, with EnableClk asserted every third Clk cycle, the rising edge of the OCP clock is coincident with the rising edge of Clk that samples EnableClk asserted. For most implementations that use this sampling function, the falling edge of the effective (internal) OCP clock is coincident with the next falling edge of Clk. The effective OCP clock is high for one-half of a Clk period every third Clk cycle, yielding an effective duty cycle of 16.7%.

Divided Clock Timing

Most design flows treat EnableClk as a standard synchronous signal that could have any value for a cycle. If EnableClk is asserted on consecutive cycles the OCP operates at the full Clk frequency, requiring internal and external timing paths to meet the maximum Clk frequency.

Relaxing the internal and external timing by recognizing that the EnableClk signal permits a restricted duty cycle (for instance, only high for every third Clk cycle). Taking advantage of this extra timing margin requires careful control over the timing flow, which may include definition and analysis of multi-cycle paths and other challenges. The design flow must assure that the system-side logic that generates EnableClk does not violate the duty cycle assumption. Finally, the timing flow must ensure that no timing issues arise due to the low duty cycle of the effective OCP clock.

Signal Timing

The Open Core Protocol data transfer model allows many different types of existing legacy IP cores to be bridged to the OCP without adding expensive glue logic structures that include address or data storage. As such, it is possible to draw many state machine diagrams that are compliant with the protocol. This section describes some common state machine models that can be used with the OCP, together with guidance on the use of those models.

Two-way handshaking is the general principle of the dataflow signals in the OCP interface. A group of signals is asserted and must be held steady until the corresponding accept signal is asserted. This allows the receiver of a signal to force the sender to hold the signals steady until it has completely processed them. This principle produces implementations with fewer latches for temporary storage.

OCP principles are built around three fundamental decoupled phases: the request phase, the response phase, and the datahandshake phase.

Request Phase

Request flow control relies on standard request/accept handshaking signals: MCmd and SCmdAccept. Note that in version 2.0 of this specification, SCmdAccept becomes an optional signal, enabled by the cmdaccept parameter. When the signal is not physically present on the interface, it naturally defaults to 1, meaning that a request phase in that case lasts exactly one clock cycle.

The request phase begins when the master drives MCmd to a value other than Idle. When MCmd != Idle, MCmd is referred to as asserted. All of the other request phase outputs of the master must become valid during the same clock cycle as MCmd asserted, and be held steady until the request phase ends. The request phase ends when SCmdAccept is sampled asserted (true) by the rising edge of the OCP clock. The slave can assert SCmdAccept in the same cycle that MCmd is asserted, or stay negated for several OCP clock cycles. The latter choice allows the slave to force the master to hold its request phase outputs until the slave can accomplish its access without latching address or data signals.

The slave designer chooses the delay between MCmd asserted and SCmdAccept asserted based on the desired area, timing, and throughput characteristics of the slave.

As the request phase does not begin until MCmd is asserted, SCmdAccept is a “don’t care” while MCmd is not asserted so SCmdAccept can be asserted before MCmd. This allows some area-sensitive, low frequency slaves to tie SCmdAccept asserted, as long as they can always complete their transfer responsibilities in the same cycle that MCmd is asserted. Since an MCmd value of Idle specifies the absence of a valid command, the master can assert MCmd independently of the current setting of SCmdAccept.

The highest throughput that can be achieved with the OCP is one data transfer per OCP clock cycle. High-throughput slaves can approach this rate by providing sufficient internal resources to end most request phases in the same OCP clock cycle that they start. This implies a combinational path from the master’s MCmd output into slave logic, then back out the slaves SCmdAccept output and back into a state machine in the master. If the master has additional requests to present, it can start a new request phase on the next OCP clock cycle. Achieving such high throughput in high-frequency systems requires careful design including cycle time budgeting as described in “Level2 Timing” on page 182.

Response Phase

The response phase begins when the slave drives SResp to a value other than NULL. When SResp != NULL, SResp is referred to as asserted. All of the other response phase outputs of the slave must become valid during the same OCP clock cycle as SResp asserted, and be held steady until the response phase ends. The response phase ends when MRespAccept is sampled asserted (true) by the rising edge of the OCP clock; if MRespAccept is not configured into a particular OCP, MRespAccept is assumed to be always asserted (that is, the response phase always ends in the same cycle it begins). If present, the master can assert MRespAccept in the same cycle that MResp is asserted, or it may stay negated for several OCP clock cycles. The latter choice allows the master to force the slave to hold its response phase outputs so the master can finish the transfer without latching the data signals.

Since the response phase does not begin until SResp is asserted, MRespAccept is a “don’t care” while SResp is not asserted so MRespAccept can be asserted before SResp. Since an SResp value of NULL specifies the absence of a valid response, the slave can assert SResp independently of the current setting of MRespAccept.

In high-throughput systems, careful use of MRespAccept can result in significant area savings. To maintain high throughput, systems traditionally introduce *pipelining*, where later requests begin before earlier requests have finished. Pipelining is particularly important to optimize Read accesses to main memory.

The OCP supports pipelining with its basic request/response protocol, since a master is free to start the second request phase as soon as the first has finished (before the first response phase, in many cases). However, without MRespAccept, the master must have sufficient storage resources to receive all of the data it has requested. This is not an issue for some masters, but can be expensive when the master is part of a bridge between subsystems such as computer buses. While the original system initiator may have enough storage, the intermediate bridge may not. If the slave has storage resources

(or the ability to flow control data that it is requesting), then allowing the master to de-assert MRespAccept enables the system to operate at high throughput without duplicating worst-case storage requirements across the die.

If a target core natively includes buffering resources that can be used for response flow control at little cost, using MRespAccept can reduce the response buffering requirement in a complex SOC interconnect.

Most simple or low-throughput slave IP cores need not implement MRespAccept. Misuse of MRespAccept makes the slave's job more difficult, because it adds extra conditions (and states) to the slave's logic.

Datahandshake Phase

The datahandshake extension allows the de-coupling of a write address from write data. The extension is typically only useful for master and slave devices that require the throughput advantages available through transfer pipelining (particularly memory). When the datahandshake phase is not present in a configured OCP, MData becomes a request phase signal.

The datahandshake phase begins when the master asserts MDataValid. The other datahandshake phase outputs of the master must become valid during the same OCP clock cycle while MDataValid is asserted, and be held steady until the datahandshake phase ends. The datahandshake phase ends when SDataAccept is sampled asserted (true) by the rising edge of the OCP clock. The slave can assert SDataAccept in the same cycle that MDataValid is asserted, or it can stay negated for several OCP clock cycles. The latter choice allows the slave to force the master to hold its datahandshake phase outputs so the slave can accomplish its access without latching data signals.

The datahandshake phase does not begin until MDataValid is asserted. While MDataValid is not asserted, SDataAccept is a “don't care”. SDataAccept can be asserted before MDataValid. Since MDataValid not being asserted specifies the absence of valid data, the master can assert MDataValid independently of the current setting of SDataAccept.

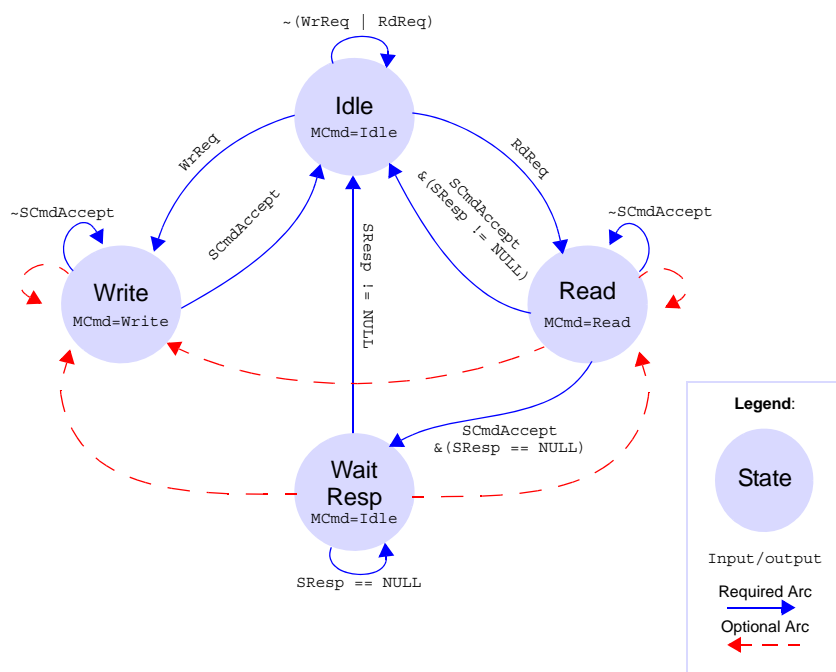
State Machine Examples

The sample state machine implementations in this section use only the features of the basic OCP, request and response phases (the datahandshake phase is not discussed here but can be derived). The examples highlight the flexibility of the basic OCP.

Sequential Master

The first example is a medium-throughput, high-frequency master design. To achieve high frequency, the implementation is a completely sequential (that is, Moore state machine) design. Figure 41 shows the state machine associated with the master's OCP.

Figure 41 Sequential Master



Not shown is the internal circuitry of the master. It is assumed that the master provides the state machine with two control wire inputs, WrReq and RdReq, which ask the state machine to initiate a write transfer and a read transfer, respectively. The state machine indicates back to the master the completion of a transfer as it transitions to its Idle state.

Since this is a Moore state machine, the outputs are only a function of the current state. The master cannot begin a request phase by asserting MCmd until it has entered a requesting state (either write or read), based upon the WrReq and RdReq inputs. In the requesting states, the master begins a request phase that continues until the slave asserts SCmdAccept. At this point (this example assumes write posting with no response on writes), a Write command is complete, so the master transitions back to the idle state.

In case of a Read command, the next state is dependent upon whether the slave has begun the response phase or not. Since MRespAccept is not enabled in this example, the response phase always ends in the cycle it begins, so the master may transition back to the idle state if SResp is asserted. If the response phase has not begun, then the next state is wait resp, where the master waits until the response phase begins.

The maximum throughput of this design is one transfer every other cycle, since each transfer ends with at least one cycle of idle. The designer could improve the throughput (given a cooperative slave) by adding the state transitions marked with dashed lines. This would skip the idle state when there are more pending transfers by initiating a new request phase on the cycle after the previous request or response phase. Also, the Moore state

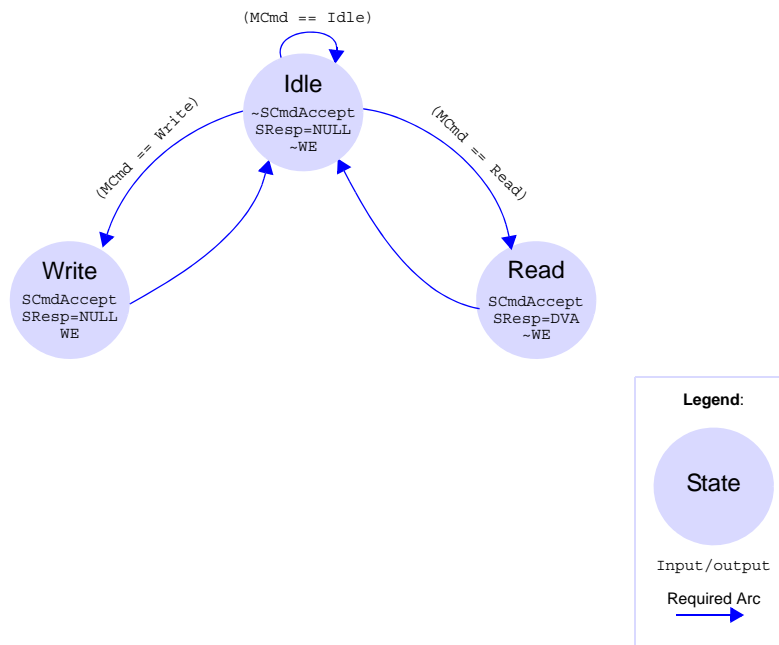
machine adds up to a cycle of latency onto the idle to request transition, depending on the arrival time of WrReq and RdReq. This cost is addressed in “Combinational Master” on page 150.

The benefits of this design style include very simple timing, since the master request phase outputs deliver a full cycle of setup time, and minimal logic depth associated with SResp.

Sequential Slave

An analogous design point on the slave side is shown in Figure 42. This slave’s OCP logic is a Moore state machine. The slave is capable of servicing an OCP read with one OCP clock cycle of latency. On an OCP write, the slave needs the master to hold MData and the associated control fields steady for a complete cycle so the slave’s write pulse generator will store the desired data into the desired location. The state machine reacts only to the OCP (the internal operation of the slave never prevents it from servicing a request), and the only non-OCP output of the state machine is the enable (WE) for the write pulse generator.

Figure 42 Sequential OCP Slave



The state machine begins in an idle state, where it de-asserts SCmdAccept and SResp. When it detects the start of a request phase, it transitions to either a read or a write state, based upon MCmd. Since the slave will always complete its task in one cycle, both active states end the request phase (by asserting SCmdAccept), and the read state also begins the response phase. Since MRespAccept is not enabled in this example, the response phase will end in the same cycle it begins. Writes without responses are assumed so SResp is NULL during the write state. Finally, the state machine triggers the write pulse generator in its write state, since the request phase outputs of the master will be held steady until the state machine transitions back to idle.

As is the case for the sequential master shown in Figure 41 on page 148, this state machine limits the maximum throughput of the OCP to one transfer every other cycle. There is no simple way to modify this design to achieve one transfer per cycle, since the underlying slave is only capable of one write every other cycle. With a Moore machine representation, the only way to achieve one transfer per cycle is to assert SCmdAccept unconditionally (since it cannot react to the current request phase signals until the next OCP clock cycle). Solving this performance issue requires a combinational state machine.

Since the outputs depend upon the state machine, the sequential OCP slave has attractive timing properties. It will operate at very high frequencies (providing the internal logic of the slave can run that quickly).

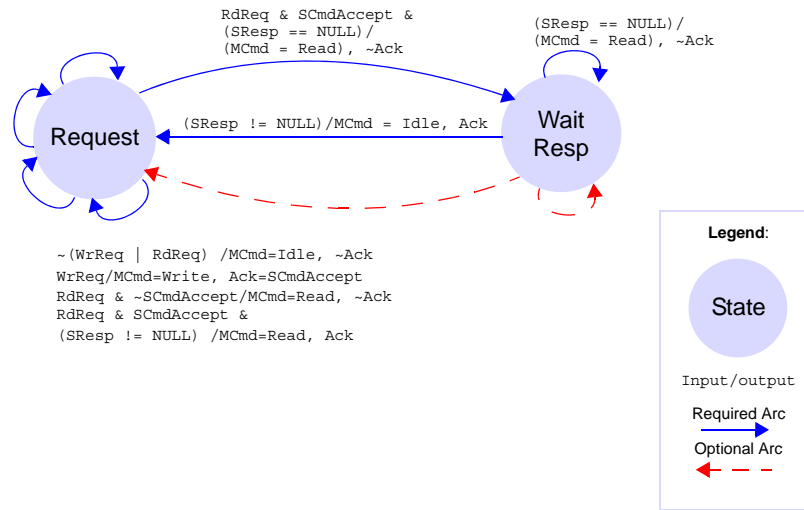
This state machine can be extended to accommodate slaves with internal latency of more than one cycle by adding a counting state between idle and one or both of the active states.

Combinational Master

“Sequential Master” on page 147 describes the transfer latency penalty associated with a Moore state machine implementation of an OCP master. An attractive approach to improving overall performance while reducing circuit area is to consider a combinational Mealy state machine representation. Assuming that the internal master logic is clocked from the OCP clock, it is acceptable for the master's outputs to be dependent on both the current state, the internal RdReq and WrReq signals, and the slave's outputs, since all of these are synchronous to the OCP clock. Figure 43 shows a Mealy state machine for the OCP master. The assumptions about the internal master logic are the same as in “Sequential Master” on page 147, except that there is an additional acknowledge (Ack) signal output from the state machine to the internal master logic to indicate the completion of a transfer.

This state machine asserts MCmd in the same cycle that the request arrives from the internal master logic, so transfer latency is improved. In addition, the state machine is simpler than the Moore machine, requiring only two states instead of four. The request state is responsible for beginning and waiting for the end of the request phase. The wait resp state is only used on Read commands where the slave does not assert SResp in the same cycle it asserts SCmdAccept. The arcs described by dashed lines are optional features that allow a transition directly from the end of the response phase into the beginning of the request phase, which can reduce the turn-around delay on multi-cycle Read commands.

Figure 43 Combinational OCP Master



The cost of this approach is in timing. Since the master request phase outputs become valid a combinational logic delay after RdReq and WrReq, there is less setup time available to the slave. Furthermore, if the slave is capable of asserting SCmdAccept on the first cycle of the request phase, then the total path is:

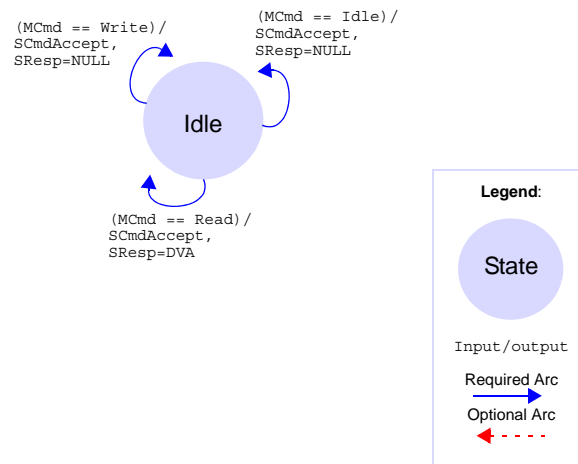
Clk -> (RdReq | WrReq) -> MCmd -> SCmdAccept -> Clk.

To successfully implement this path at high frequency requires careful analysis. The effort is appropriate for highly latency-sensitive masters such as CPU cores. At much lower frequencies, where area is often at a premium, the Mealy OCP master is attractive because it has fewer states and the timing constraints are much simpler to meet. This style of master design is appropriate for both the highest-performance and lowest-performance ends of the spectrum. A Moore state machine implementation may be more appropriate at medium performance.

Combinational Slave

Achieving peak OCP data throughput of one transfer per cycle is most commonly implemented using a combinational Mealy state machine implementation. If a slave can satisfy the request phase in the cycle it begins and deliver read data in the same cycle, the Mealy state machine representation is degenerate - there is only one state in the machine. The state machine always asserts SCmdAccept in the first request phase cycle, and asserts SResp in the same cycle for Read commands (assuming no response on writes as in the write posting model).

Figure 44 Combinational OCP Slave



The implementation shown in Figure 44, offers the ideal throughput of one transfer per cycle. This approach typically works best for low-speed I/O devices with FIFOs, medium-frequency but low-latency asynchronous SRAM controllers, and fast register files. This is because the timing path looks like:

Clk -> (master logic) -> MCmd -> (access internal slave resource) -> SResp -> Clk

This path is simplest to make when:

- OCP clock frequency is low
- Internal slave access time is small
- SResp can be determined based only on MCmd assertion (and not other request phase fields nor internal slave conditions)

To satisfy the access time and operating frequency constraints of higher-performance slaves such as main memory controllers, the OCP supports transfer pipelining. From the state machine perspective, pipelining splits the slave state machine into two loosely-coupled machines: one that accepts requests, and one that produces responses. Such machines are particularly useful with the burst extensions to the OCP.

OCP Subsets

It is possible to define simple interfaces - OCP subsets that are frequently required in complex SOC designs. The subsets provide simple interfaces for HW blocks, typically with one-directional, non-addressed, or odd data size capabilities. Since most of the OCP signals can be individually enabled or disabled, a variety of subsets can be defined. For the command set, any OCP command needs to be explicitly declared as supported by the core with at least one command enabled in a subset.

Some sample interfaces are listed in Table 27. For each example the non-default parameter settings are provided. The list of the corresponding OCP signals is provided for reference. Subset variants can further be derived from these examples by enabling various OCP extensions. For guidelines on suggested OCP feature combinations, see “OCP Profiles” on page 185.

Table 27 OCP Subsets

Usage	Purpose	Non-default parameters	Signals
Handshake-only OCP	Simple request/acknowledge handshake, that can be used to synchronize two processing modules. Using OCP handshake signals with well-defined timing and semantics allows routing this synchronization process through an interconnect. The OCP command WR is used for requests, other commands are disabled.	read_enable=0, addr=0, mdata=0, sdata=0, resp=0	Clk, MCmd, SCmdAccept
Write-only OCP	Interface for cores that only need to support writes.	read_enable=0, sdata=0, resp=0	Clk, MAddr, MCmd, MData, SCmdAccept
Read-only OCP	Interface for cores that only need to support reads.	write_enable=0, mdata=0	Clk, MAddr, MCmd, SCmdAccept, SData, SResp
FIFO Write-only OCP	Interface to FIFO input.	read_enable=0, addr=0, sdata=0, resp=0	Clk, MCmd, MData, SCmdAccept
FIFO Read-only OCP	Interface to FIFO output.	write_enable=0, addr=0, mdata=0	Clk, MCmd, SCmdAccept, SData, SResp
FIFO OCP	Read and write interface to FIFO.	addr=0	Clk, MCmd, MData, SCmdAccept, SData, SResp

Simple OCP Extensions

The simple extensions to the OCP signals add support for higher-performance master and slave devices. Extensions include byte enable capability, multiple address spaces, and the addition of in-band socket-specific information to any of the three OCP phases (request, response, and datahandshake).

Byte Enables

Byte enable signals can be driven during the request phase for read or write operations, providing byte addressing capability, or partial OCP word transfer. This capability is enabled by setting the `byteen` parameter to 1.

Even for simpler OCP cores, it is good practice to implement the byte enable extension, making byte addressing available at the chip level with no restrictions for the host processors.

When a datahandshake phase is used (typically for a single request-multiple data burst), bursts must have the same byte enable pattern on all write data words. It is often necessary to send or receive write byte enables with the write data. To provide full byte addressing capability, the MDataByteEn field can be added to the datahandshake phase. This field indicates which bytes within the OCP data write word are part of the current write transfer.

For example, on its master OCP port, a 2D-graphics accelerator can use variable byte enable patterns to achieve good transparent block transfer performance. Any pixel during the memory copy operation that matches the color key value is discarded in the write by de-asserting the corresponding byte enable in the OCP word. Another example is a DRAM controller that, when connected to a x16-DDR device, needs to use the memory data mask lines to perform byte or 16-bit writes. The data mask lines are directly derived from the byte enable pattern.

Unpacking operations inside an interconnect can generate variable byte enable patterns across a burst on the narrower OCP side, even if the pattern is constant on the wider OCP side. Such unpacking operations may also result in a byte enable pattern of all zeros. Therefore, it is mandatory that slave cores fully support 0 as a legal pattern.

An OCP interface can be configured to include partial word transfers by using either the MByteEn field, or the MDataByteEn field, or both.

- If only MByteEn is present, the partial word is specified by this field for both read and write type transfers as part of the request phase. This is the most common case.
- If only MDataByteEn is present, the partial word is specified by this field for write type transfers as part of the datahandshake phase, and partial word reads are not supported.
- If both MByteEn and MDataByteEn are present, MByteEn specifies partial words for read transfers as part of the request phase, and is don't care for write type transfers. MDataByteEn specifies partial words for write transfers as part of the datahandshake phase, and is don't care for read type transfers.

Multiple Address Spaces

Logically separate memory regions with unique properties or behavior are often scattered in the system address map. The MAddrSpace signal permits explicit selection of these separate address spaces.

Address spaces typically differentiate a memory space within a core from the configuration register space within that same core, or differentiate several cores into an OCP subsystem including multiple OCP cores that can be mapped at non-contiguous addresses, from the top level system perspective.

Another example of the usage of the addressspace extension is the case of an OCP-to-PCI bridge, since PCI natively supports address spaces for configuration registers, memory spaces and i/o spaces.

In-Band Information

OCP can be extended to communicate information that is not assigned semantics by the OCP protocol. This is true for out-of-band information (flag, control/status signals) and also for in-band information. The designer or the chip level architect can define in-band extensions for the OCP phases.

The fields provided for that purpose are MReqInfo for the request phase, SRespInfo for the response phase, MDataInfo for the request phase or the datahandshake phase, and SDataInfo for the response phase. The presence and width of these fields can be controlled individually.

MReqInfo

Uses for MReqInfo can include user/supervisor storage attributes, cacheable storage attributes, data versus program access, emulation versus application access or any other access-related information, such as dynamic endianness qualification or access permission information.

MReqInfo bits have no assigned meanings but have behavior restrictions. MReqInfo is part of the request phase, so when MCmd is Idle, MReqInfo is a "don't care". When MCmd is asserted, MReqInfo must be held steady for the entire request phase. MReqInfo must be constant across an entire transaction, so the value may not change during a burst. This facilitates simple packing and unpacking of data at mismatched master/slave data widths, eliminating the transformation of information.

SRespInfo

Uses for SRespInfo can include error or status information, such as FIFO full or empty indications, or data response endianness information.

SRespInfo bits have no assigned meaning, but have behavior restrictions. SRespInfo is part of the response phase, so when SResp is NULL, SRespInfo is a "don't care". When SResp is asserted, SRespInfo must be held steady for the entire response phase. Whenever possible, slaves that generate SRespInfo values should hold them constant for the duration of the transaction, and choose semantics that favor sticky status bits that stay asserted across transactions. This simplifies the design of interconnects and bridges that span OCP interfaces with different configurations. Holding SRespInfo constant improves simple packing and unpacking of data at mismatched data widths. The spanning element may need to break a single transaction into multiple smaller transactions, and so manage the combination of multiple SRespInfo values when the original transaction has fewer responses than the converted ones.

If you implement SRespInfo as specified, your implementation should work in future versions of the specification. If the current implementation does not meet your needs, please contact techsupport@ocpip.org so that the Specification Working Group can investigate how to satisfy your requirements.

MDataInfo and SDataInfo

MDataInfo and SDataInfo have slightly different semantics. While they have no OCP-defined meaning, they may have packing/unpacking implications. MDataInfo and SDataInfo are only valid when their associated phase is asserted (request or datahandshake phase for MDataInfo, response phase for SDataInfo).

Uses for the MDataInfo and SDataInfo fields might include byte data parity in the low-order bits and/or data ECC values in the high-order (non-packable) bits.

The low-order `mdatainfobyte_wdth` bits of MDataInfo are associated with MData[7:0], and so forth for each higher-numbered byte within MData, so that the low-order `mdatainfobyte_wdth*(data_wdth/8)` bits of MDataInfo are associated with individual data bytes. Any remaining (upper) bits of MDataInfo cannot be packed or unpacked without further specification, although such bits may be used in cases with matched data width, where no transformation is required.

The difference between MReqInfo and the upper bits of MDataInfo is that only MDataInfo is allowed to change during a transaction. Use SDataInfo for information that may change during a transaction.

A slave should be operable when all bits of MReqInfo and MDataInfo are negated; in other words, any MReqInfo or MDataInfo signals defined by an OCP slave, but not present in the master will normally be negated (driven to logic 0) in the tie off rules. A master should be operable when all bits of SRespInfo and SDataInfo are negated.

Burst Extensions

A burst is basically a set of related OCP words. Burst framing signals provide a method for linking together otherwise-independent OCP transfers. This mechanism allows various parts of a system to optimize transfer performance using such techniques as SDRAM page-mode operation, burst transfers, and pre-fetching.

Burst support is a key enabler of SOC performance. The burst extension is frequently used in conjunction with pipelined master and slave devices. For a pipelined OCP device, the request phase is *de-coupled* from the response phase - that is, the request phase may begin and end several cycles before the associated response phase begins and ends. As such, it is useful to think of separate, loosely-coupled state machines to support either the master or the slave. Decoupling for pipeline efficiency remains true even if the OCP includes a separate datahandshake phase.

OCP Burst Capabilities

The OCP burst model includes a variety of options permitting close matching of core design requirements.

Exact Burst Lengths and Imprecise Burst Lengths

A burst can be either precise, of known length when issued by the initiator, or imprecise, the burst length is not specified at the start of the burst.

For precise bursts, MBurstLength is driven to the same value throughout the burst, but is really meaningful only during the first request phase. Precise bursts with a length that is a power-of-two can make use of the WRAP and XOR address sequence types.

For imprecise bursts, MBurstLength can assume different values for words in the burst, reflecting a best guess of the burst length. MBurstLength is a hint. Imprecise bursts are completed by a request with an MBurstLength=1, and cannot make use of the WRAP and XOR address sequence types.

Use the precise burst model whenever possible:

- It is compatible with the single request-multiple data model that provides advantages to the SOC in terms of performance and power.
- Since it is deterministic, it simplifies burst conversion. Restricting burst lengths to power-of-two values and using aligned incrementing bursts (by employing the `burst_aligned` parameter) also reduces the interconnect complexity needed to maintain interoperability between cores.

Address Sequences

Using the MBurstSeq field, the OCP burst model supports commonly-used address sequences. Benefits include:

- A simple incrementing scheme for regular memory type accesses
- A constant addressing mode for FIFO oriented targets (typically peripherals)
- Wrapping on power-of-two boundaries
- XOR for processor cache line fill
- A block transfer scheme for 2-dimensional data stored in memory

User-defined sequences can also be defined. They must be carefully documented in the core specification, particularly the rules to be applied when packing or unpacking. The address behavior for the different sequence types is:

INCR

Each address is incremented by the OCP word size. Used for regular memory type accesses, SDRAM, SRAM, and burst Flash.

STRM

The address is constant during the burst. Used for streaming data to or from a target, typically a peripheral device including a FIFO interface that is mapped at a constant address within the system.

DFLT1

User-specified address sequence. Maximum packing is required.

DFLT2

User-specified address sequence. Packing is not allowed.

WRAP

Similar to INCR, except that the address wraps at aligned MBurstLength * OCP word size. This address sequence is typically used for processor cache line fill. Burst length is necessarily a power-of-two, and the burst is aligned on its size.

XOR

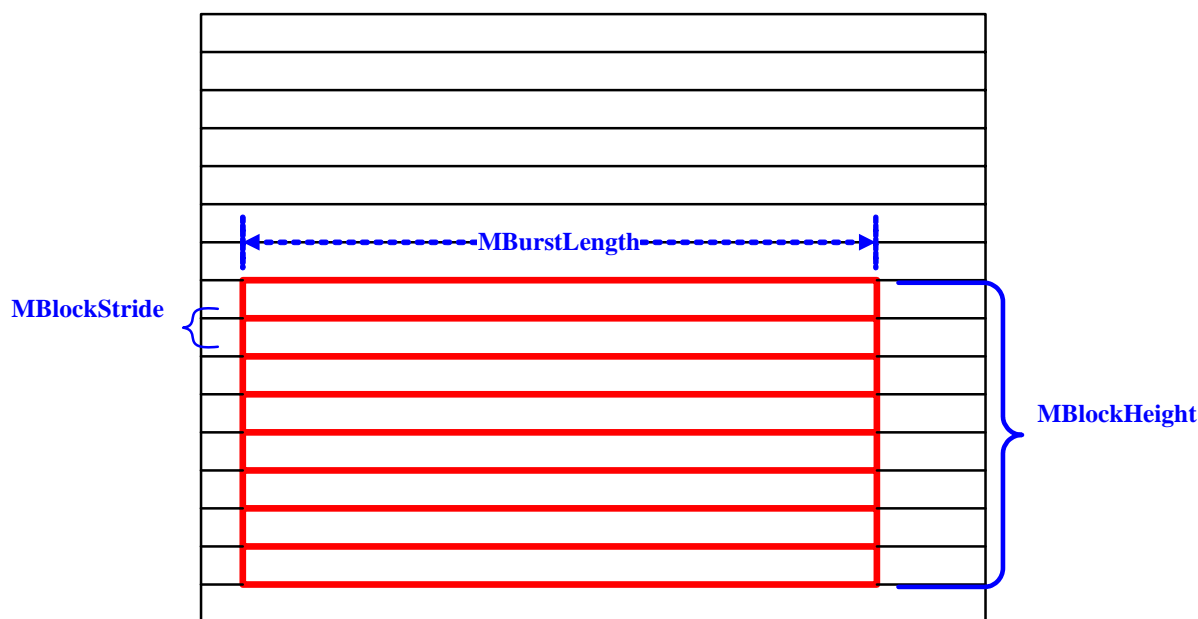
$\text{Addr} = \text{BurstBaseAddress} + (\text{index of first request in burst}) \wedge (\text{current word number})$. XOR is used by some processors for critical-word first cache line fill from wide and slow memory systems.

While it does not always deliver the next sequential words as quickly as WRAP, the XOR sequence maps directly into the interleaved burst type supported by many DRAM devices. The XOR sequence is convenient when there are width differences between OCP interfaces, since the sequence is chosen to successively fill power-of-two sized and aligned words of greater width until the burst length is reached.

BLCK

Describes a sequence of MBlockHeight row transfers, with the starting address MAddr, row-to-row offset MBlockStride (measured from the start of one row to the start of the next row), and rows that are MBurstLength words long in an incrementing row address per word. MBlockHeight and MBlockStride can be considered as don't care for burst sequences other than BLCK. Figure 45 depicts a block transaction representing a 2-dimensional region out of a larger buffer, showing the various parameters.

Figure 45 BLCK Address Sequence



The following example shows how to decompose a BLCK burst into a sequence of INCR bursts.

```
addr = MAddr;
for (row = 0; row < MBlockHeight; row++, addr += MBlockStride) {
    issue INCR using all provided fields (including
    MBurstLength), except use "addr" for MAddr
}
```

UNKN

Indicates that there is no specified relationship between the addresses of different words in the burst. Used to group requests within a burst container, when the address sequence does not match the pre-defined sequences. For example, an initiator can group requests to non-consecutive addresses on the same SDRAM page, so that the target memory bandwidth can be increased.

For targets that have support for some burst sequence, adding support for the UNKN burst sequence can improve the chances of interoperability with other cores and can ease verification since it removes all requirements from the address sequence within a burst.

For single requests, MBurstSeq is allowed to be any value that is legal for that particular OCP interface configuration.

The BLCK, INCR, DFLT1, WRAP, and XOR burst sequences are always considered packing, whereas STRM and DFLT2 sequences are non-packing. Transfers in a packing burst sequence are aggregated / split when translating between OCP interfaces of different widths while transfers in a non-packing sequence are filled / truncated.

The packing behavior of a bridge or interconnect for an UNKN burst sequence is system-dependent. A common policy is to treat an UNKN sequence as packing in a wide-to-narrow OCP request width converter, and as non-packing in a narrow-to-wide OCP request width converter.

Single Request, Multiple Data Bursts for Reads and Writes

A burst model of this type can reduce power consumption, bandwidth congestion on the request path, and buffering requirement at various locations in the system. This model is only applicable for precise bursts, and assumes that the target core can reconstruct the full address sequence using the code provided in the MBurstSeq field.

While the model assumes that the datahandshake extension is on, for those cores that cannot accept the first-data word without the corresponding request, datahandshake can increase design and verification complexity.

For such cores, use the OCP parameter `reqdata_together` to specify the fixed timing relationship between the request and datahandshake phases. When `reqdata_together` is set, each request phase for write-type bursts must be presented and accepted together with the corresponding datahandshake phase. For single request / multiple data write bursts, the request must be presented with the first write data. For multiple request / multiple data write

bursts, the datahandshake phase is locked to the request phase, so this interface configuration only makes sense when both single and multiple request bursts are mixed (that is, `burstsinglereq` is set).

Unit of Atomicity

Use this option when there is a requirement to limit burst interleaving between several threads. Specifying the atomicity allows the master to define a transfer unit that is guaranteed to be handled as an atomic group of requests within the burst, regardless of the total burst length. The master indicates the size of the atomic unit using the `MAtomicLength` field.

Burst Framing with all Transfer Phases

Without burst framing information, cores and interconnects incorporate counters in their control logic: To limit this extra gate count and complexity, enable end-of-burst information for each phase. Use `MReqLast` to specify the last request within a burst, `SRespLast` to specify the last response within a burst, and `MDataLast` to specify the last write data during the datahandshake phase.

If `BLCK` burst sequences are enabled, additional framing information can be provided to eliminate additional counters associated with the `INCR` subsequences that comprise a `BLCK` burst. To limit extra gate count and complexity, enable end-of-row information for each phase using:

- `MReqRowLast` to specify the last request in each row for multiple request/multiple data bursts
- `SRespRowLast` to specify the last response in each row
- `MDataRowLast` to specify the last write data in each row

Compatibility with the OCP 1.0 Burst Model

The OCP 2.0 burst model replaces the OCP 1.0 model, providing a super set in terms of available functionality. To maintain interoperability between cores using the OCP 1.0 burst and cores using the OCP 2.0 bursts requires a thin adaptation layer. Guidelines for the wrapping logic are described in this section.

1.0 Master to 2.0 Slave

For converting an OCP 1.0 burst into an OCP 2.0 burst the suggested mapping is:

- `MBurstPrecise` is available only when the OCP 1.0 `burst_aligned` parameter is set. When set, all incrementing bursts once converted to OCP 2.0 stay precise. Any other OCP 1.0 burst type is mapped to an imprecise burst. When `burst_aligned` is not set, `MBurstPrecise` is tied off to 0, so all bursts are imprecise.
- `MBurstSeq` is derived from `MBurst` as follows:

MBurstSeq = INCR for MBurst {CONT, TWO, FOUR, EIGHT}
 STRM for MBurst {STRM}
 DFLT1 for MBurst {DFLT1}
 DFLT2 for MBurst {DFLT2}

The logic must guarantee that MBurstSeq is constant during the whole burst and must continue driving that MBurstSeq when MBurst=LAST is detected.

- The value of MBurstLength is derived as follows:

MBurstLength = 8 for MBurst {EIGHT}
 4 for MBurst {FOUR}
 2 for MBurst {TWO, CONT, DFLT1, DFLT2, STRM}
 1 for MBurst {LAST}

For precise bursts, MBurstLength is held constant for the entire burst. For imprecise bursts, a new MBurstLength can be derived for each transfer.

- MReqLast is derived from MBurst - it is set when MBurst is LAST.
- SRespLast has no equivalent in OCP1.0, and is discarded by the wrapping logic.
- If required, MDataLast must be generated from a counter or from a queue updated during the request phase.

2.0 Master to 1.0 Slave

For converting an OCP 2.0 burst into an OCP 1.0 burst the suggested mapping is:

- MBurst is derived from MBurstPrecise, MBurstSeq, and MBurstLength, as follows:

```

MBurst =
If MBurstPrecise
if MBurstSeq {INCR}
    EIGHT if MBurstLength >= 8 at start of burst
    FOUR if MBurstLength >= 4 at start of burst
    TWO if MBurstLength >= 2 at start of burst
    load counter with MBurstLength at start of burst,
    decrement counter after every transfer
    subsequent MBurst are generated from counter logic
    LAST when counter==1
else if MBurstSeq {DFLT1, DFLT2, STRM}
    same as MBurstSeq, except when counter==1, must be LAST
else if MBurstSeq {WRAP, XOR, UNKN}
    LAST always: map to consecutive non-burst single transactions

Else if not MBurstPrecise
if MBurstSeq {INCR}
    EIGHT if MBurstLength >= 8
    FOUR if MBurstLength >= 4
  
```

```
TWO if MBurstLength >= 2
LAST if MBurstLength == 1
else if MBurstSeq {DFLT1, DFLT2, STRM}
  LAST if MBurstLength == 1
  same as MBurstSeq if MBurstLength != 1
else if MBurstSeq {WRAP, XOR, UNKN}
  LAST always (map to non-burst)
```

- MAtomicLength, MReqLast, and MDataLast have no equivalents in OCP 1.0, and are discarded by the wrapping logic.
- SRespLast must be generated from counter logic.

The logic described above is not suitable if the OCP 2.0 master generates single request / multiple data bursts. In that case, more complex conversion logic is required.

Tags

Tags are labels that associate requests with responses in order to enable out-of-order return of responses. In the face of different latencies for different requests (for instance, DRAM controllers, or multiple heterogeneous targets), allowing the out-of-order delivery of responses can enable higher performance. Responses are returned in the order they are produced rather than having to buffer and re-order them to satisfy strict response order requirements. As is the case for threads, to make use of tags, the master will normally need a buffer.

The tag value generally only has meaning to the master as it is made up by the master and often corresponds to a buffer ID. The tag is carried by the slave and returned with the response. In the case of datahandshake, the master also tags the datahandshake phase with the tag of the corresponding request.

Out-of-order request and response delivery can also be enabled using multiple threads. The major differences between threads and tags are that threads can have independent flow control per thread and have no ordering rules for transfers on different threads. Tags, on the other hand, exist within a single thread so are restricted to a single flow control for all tags. Also, transfers within the same thread still have some (albeit looser) ordering rules when tags are in use. The need for independent flow control requires independent buffering per thread, leading to more complex implementations. Tags enable lower overhead implementations for out-of-order return of responses.

Tags are local to a single OCP interface. In a system with multiple cores connected to a bridge or interconnect, it is the responsibility of the interconnect to translate the tags from one interface to the other so that a target sees different tags for requests issued from different initiators. Target core implementation can facilitate the job of the bridge or interconnect by supporting a large set of tags.

The MTagInOrder and STagInOrder signals allow both tagged and non-tagged (in-order) initiators to talk to a tagged target. Requests issued with MTagInOrder asserted must be completed in-order with respect to other in-order transactions, so an in-order initiator can guarantee that its responses are returned in request order. To retain compatibility with non-tagged initiators, targets that support tags should also support MTagInOrder and STagInOrder.

When MTagInOrder is asserted any MTagID and MDataTagID values are “don’t care”. Similarly, the STagID value is “don’t care” when STagInOrder is asserted. Nonetheless, it is suggested that the slave return whatever tag value the master provided.

Multi-threaded OCP interfaces can also have tags. Each thread’s tags are independent of the other threads’ tags and apply only to the ordering of transfers within that thread. There are no ordering restrictions for transfers on different threads. The number of tags supported by all threads must be uniform, but a master need not make use of all tags on all threads.

Threads and Connections

Thread extensions add support for concurrency. Without these extensions, there is no way to apply flow control to one set of transfers while allowing another set to proceed. With threading, each transfer is associated with a thread, and independent flow control can be applied to each thread. Additionally, there are no ordering restrictions between transfers associated with different threads. Without threads, ordering is either strict or (if tags are used) somewhat looser.

Threads

The thread capability relies on a thread ID to identify and separate independent transfer streams (threads). The master labels each request with the thread ID that it has assigned to the thread. The thread ID is passed to the slave on MThreadID together with the request (MCmd). When the slave returns a response, it also provides the thread ID (on SThreadID) so the master knows which request is now complete.

The transfers in each thread must remain in-order with respect to each other (as in the basic OCP) or must follow the ordering rules for tagging (if tags are in use), but the order between threads can change between request and response.

The thread capability allows a slave device to optimize its operations. For instance, a DRAM controller could respond to a second read request from a higher-priority initiator before servicing a first request from a lower-priority initiator on a different thread.

As routing congestion and physical effects become increasingly difficult at the back-end stage of the ASIC process, multithreading offers a powerful method of reducing wires. Many functional connections between initiator and target pairs do not require the full bandwidth of an OCP link, so sharing the same

wires between several connections, based on functional requirements and floor planning data, is an attractive mechanism to perform gate count versus performance versus wire density trade-offs.

Multi-threaded behavior is most frequently implemented using one state machine per thread. The only added complexity is the arbitration scheme between threads. This is unavoidable, since the entire purpose for building a multi-threaded OCP is to support concurrency, which directly implies contention for any shared resources.

The MDataThreadID signal simplifies the implementation of the datahandshake extension along with threading, by providing the thread ID associated with the current write data transfer. When datahandshake is enabled, but sdatathreadbusy is disabled, the ordering of the datahandshake phases must exactly match the ordering of the request phases.

The thread busy signals provide status information that allows the master to determine which threads will not accept requests. That information also allows the slave to determine which threads will not accept responses. These signals provide for cooperation between the master and the slave to ensure that requests are not presented on busy threads.

While multithreading support has a cost in terms of gate count (buffers are required on a thread-per-thread basis for maximum efficiency), the protocol can ensure that the multi-threaded interface is non-blocking.

Blocked OCP interfaces introduce a thread dependency. If thread X cannot proceed because the OCP interface is blocked by another thread, Y that is dependent on something downstream that cannot make progress until thread X makes progress, there is a classic circular wait condition that can lead to deadlock.

In the *OCP1.0 Specification*, the semantics of SThreadBusy and MThreadBusy allow these signals to be treated as hints. To guarantee that a multi-threaded interface does not block, both master and slave need to be held to tighter semantics.

OCP 2.2 allows cores to follow exact thread busy semantics. This process enables tighter protocol checking at the interface and guarantees that a multi-threaded OCP interface is non-blocking. Parameters to enable these extensions are sthreadbusy_exact, sdatathreadbusy_exact, and mthreadbusy_exact. There is one parameter for each of the OCP phases, request, datahandshake (assuming separate datahandshake) and response (assuming response flow control). The following conditions are true:

- On an OCP interface that satisfies sthreadbusy_exact semantics, the master is not allowed to issue a command to a busy thread.
- On an OCP interface that complies with sdatathreadbusy_exact semantics, the master is not allowed to issue write data to a busy thread.
- On an OCP interface that complies with mthreadbusy_exact semantics, the slave is not allowed to issue a response to a busy thread.

These rules allow the phase accept signals (SCmdAccept, SDataAccept or MRespAccept) to be tied off to 1 on multi-threaded interfaces for which the corresponding phase handshake satisfies exact thread busy semantics. By eliminating an additional combinational dependency between master and slave, an exact thread busy based handshake can be considered as a substitute for the standard request/accept protocol handshake. For more information, see “Multi-Threaded OCP Implementation” on page 167.

Intra-Phase Signal Relationships on a Multithreaded OCP

This section extends the timing discussion of the “Basic OCP” section, to a multithreaded interface. The ordering and timing relationships between the signals within an OCP phase are designed to be flexible. As described in “Request Phase”, it is legal for SCmdAccept to be driven either combinatorially, dependent upon the current cycle’s MCmd or independently from MCmd, based on the characteristics of the OCP slave. Some restrictions are required to ensure that independently-created OCP masters and slaves will work together. For instance, the MCmd cannot respond to the current state of SCmdAccept; otherwise, a combinational cycle could occur.

Request Phase

If enabled, a slave’s SThreadBusy request phase output should not depend upon the current state of any other OCP signal. SThreadBusy should be stable early enough in the cycle so that the master can factor the current SThreadBusy into the decision of which thread to present a request; that is, all of the master’s request phase outputs may depend upon the current SThreadBusy. SThreadBusy is a hint so the master is not required to include a combinational path from SThreadBusy into MCmd, but such paths become unavoidable if the exact semantics apply (`sthradbusy_exact = 1`). In that case the slave must guarantee that SThreadBusy becomes stable early in the OCP cycle to achieve good frequency performance. A common goal is that SThreadBusy be driven directly from a flip-flop in the slave.

A master’s request phase outputs should not depend upon any current slave output other than SThreadBusy. This ensures that there is no combinational loop in the case where the slave’s SCmdAccept depends upon the current MCmd.

If a slave’s SCmdAccept request phase output is based upon the master’s request phase outputs from the current cycle, there is a combinational path from MCmd to SCmdAccept. Otherwise, SCmdAccept may be driven directly from a flip-flop, or based upon some other OCP signals. It is legal for SCmdAccept to be derived from MRespAccept. This case arises when the slave delays SCmdAccept to force the master to hold the request fields for a multi-cycle access. Once read data is available, the slave attempts to return it by asserting SResp. If the OCP has MRespAccept enabled, the slave then must wait for MRespAccept before negating SResp, so it may need to continue to hold off SCmdAccept until it sees MRespAccept asserted.

While the phase relationships of the OCP specification do not allow the response phase to end before the request phase, it is legal for both phases to complete in the same OCP cycle.

The worst-case combinational path for the request phase could be:

```
Clk -> SThreadBusy -> MCmd -> SResp -> MRespAccept -> SCmdAccept -> Clk
```

The preceding path has too much latency at typical clock frequencies, so must be avoided. Fortunately, a multi-threaded slave (with SThreadBusy enabled) is not likely to exhibit non-pipelined read behavior, so this path is unlikely to prove useful. Slave designers need to limit the combinational paths visible at the OCP. By pipelining the read request, the previous path could be:

```
Clk -> SThreadBusy -> MCmd -> Clk
Clk -> SCmdAccept -> Clk      # Slave accepts if pipeline reg empty
Clk -> SResp -> Clk
Clk -> MRespAccept -> Clk     # Master accepts independent of SResp
```

Response Phase

If enabled, a master's MThreadBusy response phase output should not be dependent upon the current state of any other OCP signal. From the perspective of the OCP, MThreadBusy should become stable early enough in the cycle that the slave can factor the current MThreadBusy into the decision on which thread to present a response; that is, all of the slave's response phase outputs may depend upon the current MThreadBusy. If MThreadBusy is simply a hint (in other words `mthreadbusy_exact = 0`) the slave is not required to include a combinational path from MThreadBusy into SResp, but such paths become unavoidable if the exact semantics apply (`mthreadbusy_exact = 1`). In that case the master must guarantee that MThreadBusy becomes stable early in the OCP cycle to achieve good frequency performance. A common goal is that MThreadBusy be driven directly from a flip-flop in the master.

The slave's response phase outputs should not depend upon any current master output other than MThreadBusy. This ensures that there is no combinational loop in the case where the master's MRespAccept depends upon the current SResp.

The master's MRespAccept response phase output may be based upon the slave's response phase outputs from the current cycle or not. If this is true, there is a combinational path from SResp to MRespAccept. Otherwise, MRespAccept can be driven directly from a flip-flop; MRespAccept should not be dependent upon other master outputs.

Datahandshake Phase

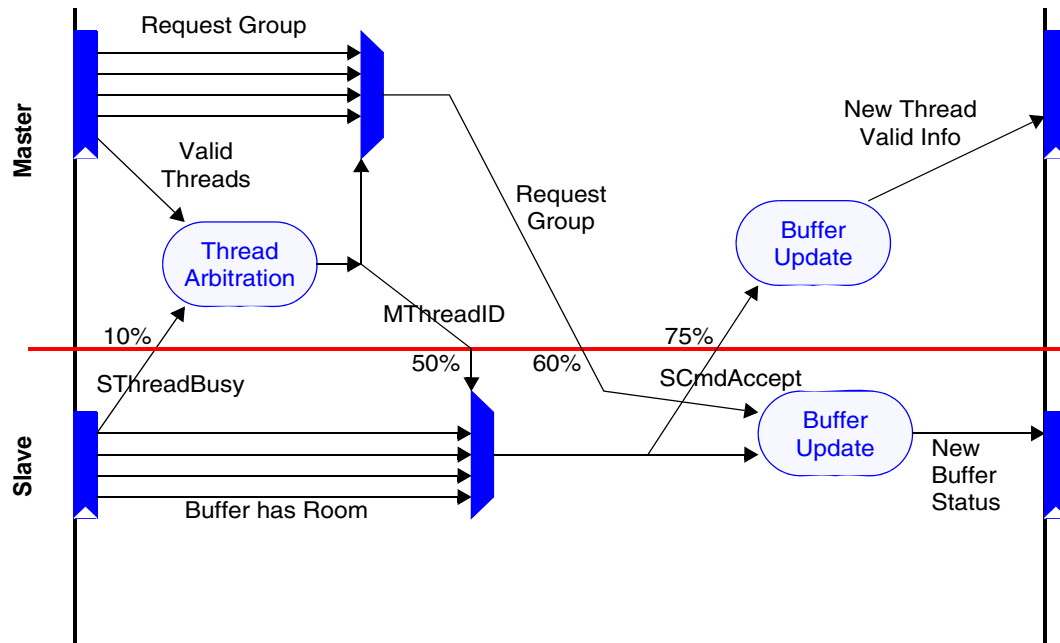
If enabled, a slave's SDataThreadBusy datahandshake phase output should not depend upon the current state of any other OCP signal. SDataThreadBusy should be stable early enough in the cycle so that the master can factor the current SDataThreadBusy into the decision of which thread to present a data; that is, all of the master's data phase outputs may depend upon the current SDataThreadBusy. If SDataThreadBusy is simply a hint (in other words `sdatathreadbusy_exact = 0`) the master is not required to include a combinational path from SDataThreadBusy into MDataValid, but such path becomes unavoidable if the exact semantics apply (`sdatathreadbusy_exact = 1`). In that case, the slave must guarantee that SDataThreadBusy becomes stable early in the OCP cycle to achieve good frequency performance. A common goal is that SDataThreadBusy be driven directly from a flip-flop in the slave.

The master's datahandshake phase outputs should not depend upon any current slave output other than *SThreadBusy*. This ensures that there is no combinational loop in the case where the slave's *SDataAccept* depends upon the current *MDataValid*. The slave's *SDataAccept* output may or may not be based upon the master's datahandshake phase outputs from the current cycle. In the former case, there is a combinational path from *MDataValid* to *SDataAccept*. In the latter case, *SDataAccept* should be driven directly from a flip-flop; *SDataAccept* should not be dependent upon other master outputs.

Multi-Threaded OCP Implementation

Figure 46 on page 167 shows the typical implementation of the combinational paths required to make a multi-threaded OCP work within the framework set by Level-2 timing. While the figure shows a request phase, similar logic can be used for the response and datahandshake phases. The top half of the figure shows logic in the master; the bottom half shows logic in the slave. The width of the figure represents a single OCP cycle.

Figure 46 Multithreaded OCP Interface Implementation



Slave

Information about space available on the per-port buffers comes out of a latch and is used to generate *SThreadBusy* information, which must be generated within the initial 10% of the OCP cycle (as described in “Level2 Timing” on page 182). These signals are also used to generate *SCmdAccept*: if a particular port has room, a command on the corresponding thread is accepted. The correct port information is selected through a multiplexer driven by *MThreadID* at 50% of the clock cycle, making it easy to produce *SCmdAccept*.

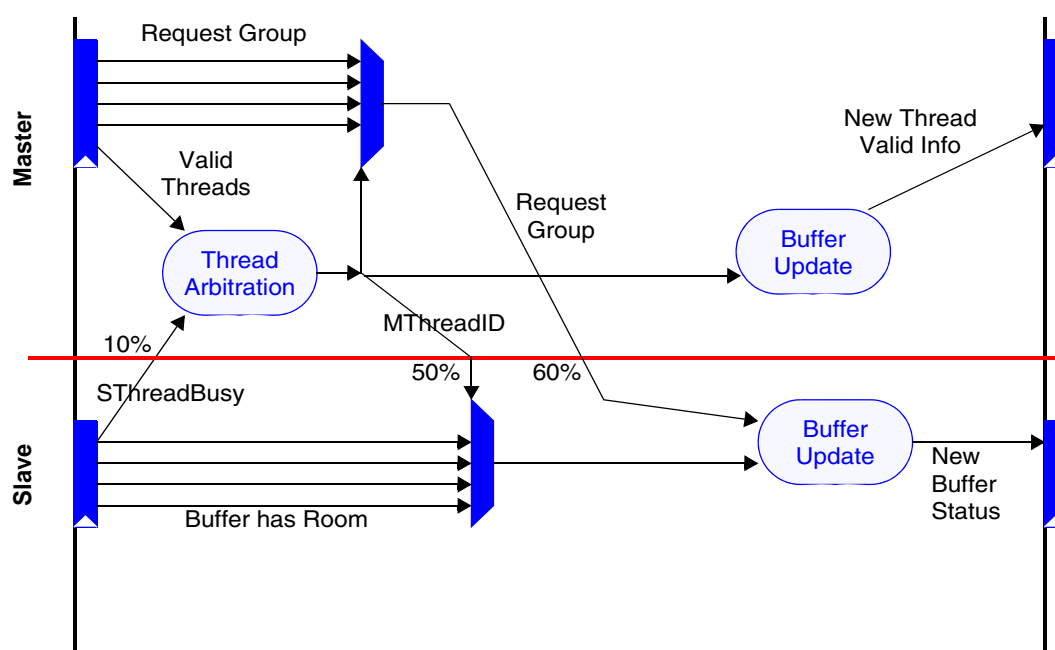
by 75% of the OCP cycle. When the request group arrives at 60% of the OCP cycle, it is used to update the buffer status, which in turn becomes the SThreadBusy information for the next cycle.

Master

The master keeps information on what threads have commands ready to be presented (thread valid bits). When SThreadBusy arrives at 10% of the OCP clock, it is used to mask off requests, that is any thread that has its SThreadBusy signal set is not allowed to participate in arbitration for the OCP. The remaining thread valid bits are fed to thread arbitration, the result is the winning thread identifier, MThreadId. This is passed to the slave at 50% of the OCP clock period. It is also used to select the winning thread's request group, which is then passed to the slave at 60% of the clock period. When the SCmdAccept signal arrives from the slave, it is used to compute the new thread valid bits for the next cycle.

The request phase in Figure 46 assumes a non-exact thread busy model. The exact model shown in Figure 47 is similar, but SCmdAccept is tied off to 1, so any request issued to a non-busy thread is accepted in the same cycle by the slave.

Figure 47 Multithreaded OCP Interface with threadbusy_exact



Connections

In multi-threaded, multi-initiator systems, it is frequently useful to associate a transfer request with a thread operating on a particular initiator. Initiator identification can enable a system to restrict access to shared resources, or

foster an error logging mechanism to identify an initiator whose request has created an error in the system. For devices where these concerns are important, the OCP extensions support connections.

Connections are closely related to threads, but can have end-to-end meaning in the system, rather than the local meaning (that is, master to slave) of a thread.

The connection ID and thread ID seem to provide similar functionality, so it is useful to consider why the OCP needs both. A *thread ID* is an identifier of local scope that simply identifies transfers between the master and slave. In contrast, the *connection ID* is an identifier of global scope that identifies transfers between a system initiator and a system target. A thread ID must be small enough (that is, a few bits) to efficiently index tables or state machines within the master and slave. There are usually more connection IDs in the system than any one slave is prepared to simultaneously accept. Using a connection ID in place of a thread ID requires expensive matching logic in the master to associate the returned connection ID (from the slave) with specific requests or buffer entries.

Using a networking analogy, the thread ID is a level-2 (data link layer) concept, whereas the connection ID is more like a level-3 (transport/session layer) concept. Some OCP slaves only operate at level-2, so it doesn't make sense to burden them or their masters with the expense of dealing with level-3 resources. Alternatively, some slaves need the features of level-3 connections, so in this case it makes sense to pass the connection ID through to them.

A connection ID is not usually provided by an initiator core on its OCP interface but is allocated to that particular initiator in the interconnect logic of the system. The connection ID is system-specific, not core-specific; only the system integrator has the global knowledge of the number of initiators instantiated in the application, and what the requirements are in terms of differentiation.

As an exception to that rule, if the global interconnect consists of multiple hierarchical structures, a complete subsystem can be integrated (including another interconnect with multiple embedded initiators). In that case, the OCP interface between the two interconnects should implement the connid extension, so that the end-to-end meaning of that OCP field can be preserved at the system level.

For a target core, the connid extension is included when such features as access control, error logging or similar initiator-related features require initiator identification.

OCP Specific Features

Write Semantics

OCP semantics specify whether a write is posted or not. A non-posted write model is preferred whenever the originator of the request must be aware of the completion of its write command. An example is clearing an interrupt in a peripheral module using a write command. In that case the processor must be sure that the interrupt line has been effectively released before it can acknowledge the interrupt service in the chip-level interrupt controller.

The concept of the posting semantics diverges from the concept of responses on writes in the following ways:

- A write with a response could have posted semantics in a system (so that a response is returned immediately) or it could have non-posted semantics (so that a response is returned only after the write is completed at the final target).
- A write without a response normally has posted semantics and carries forward the *OCP 1.0 Specification* for backward compatibility.
- A write without a response can be assigned non-posted semantics by not accepting the command until the write has completed, but this is not recommended since it de-pipelines the OCP interface. Since posting makes sense at a system level, adopting a delayed-SCmdAccept scheme can only be efficient locally, with no guarantee of the non-posting semantics at the system level.

The `writeresp_enable` parameter controls whether writes have responses. `writenonpost_enable` controls whether the interface supports the WriteN-onPost command, giving the initiator core the ability to launch two different types of writes. Table 28 summarizes the choices.

Table 28 Write Parameters

writenonpost_enable		
writeresp_enable	0	1
0	Simple posted write model, corresponds to OCP 1.0 Spec	Illegal
1	Initiator core has one flavor of writes (WR), system integrator decides where to post the write requests	Initiator core has two flavors of writes (WR and WRNP), system integrator decides where to post the two different write requests

By separating whether writes have responses (`writeresp_enable`) from whether the core has control over where the responses are generated (`writenonpost_enable`), the OCP specification provides the following features:

- The simple, posted model remains intact. The simplest cores only implement WR, and need not worry about write responses.
- Cores that can generate or use write responses should enable write responses, providing support for in-band error reporting on write commands. The read and write state machines are duplicated from the standpoint of flow control, producing a simpler design. Such cores would normally only implement the WR command. In this case, the system integrator is in control of where in the write path the write response is generated, allowing a choice of the level of posting based upon performance and coherence trade-offs.
- Cores that can distinguish between performance and coherence (really only CPUs and bridges) can enable WRNP to implement dynamic choice between WR and WRNP. The additional signaling gives the system integrator the dynamic information needed to choose the posting point as the CPU requests. The only practical difference between WR and WRNP at the protocol level is the expected latency between request and response. This permits some embedded CPUs to achieve high performance – particularly as interconnects become pipelined and posting buffers are needed.

When the `writeresp_enable` parameter is enabled, responses are required for any write command issued on the OCP, including WR, WRNP, but also BCST and WRC.

Use of the Broadcast command must be limited to specific category of designs (some interconnect designs may benefit from simultaneous update through distributed registers). It is not expected that standard cores support the command.

Lazy Synchronization

Most processors support semaphores through a read-modify-write type of instruction and swap, test-and-set, etc. Using an OCP interconnect, these instructions are mapped onto a pair of OCP commands. A RDEX command sets a lock to the memory location, followed by a WR (or WRNP) command to release the lock. The system must ensure that no other thread will be granted access to that memory location between the RDEX and the unlocking WR.

Because the Write that clears the lock must immediately follow the ReadEx (on the same thread), only a limited number of operations can be performed by a processor between RDEX and WR. Since competing requests to the locked location are also blocked from proceeding until the lock is released, you could lock part of the interconnect for the duration of the RDEX-WR or WRNP pair. Such a mechanism, often referred to as *locked synchronization*, is efficient for handling exclusive accesses to a shared resource, but can result in a significant performance loss when used extensively.

For these reasons, some processors use non-blocking instructions for semaphore handling, breaking the atomicity of the exclusive read/write pair. For the processor, this allows other instructions to be executed by the processor between the read and write accesses. For the system interconnect, it allows requests from other threads to be inserted between the read and

write commands. Referred to as lazy synchronization, this mechanism requires read and write semantics, commonly known as LL/SC semantics, for Load-Linked and Store-Conditional.

OCP's support for lazy synchronization uses the ReadLinked, and WriteConditional commands. A single OCP parameter `rdlwrc_enable` is set to 1, to enable the commands. Because some processors might use both semantics (locked and lazy), the OCP interface supports RDEX, RDL, WR(WRNP) and WRC.

The system relies upon the existence of monitor logic, that can be located either in the system interconnect, or in the memory controller. The ReadLinked command sets a reservation in the logic, associating the accessing thread with a particular address. The WriteConditional command, being transmitted on the same thread, is locally transformed into a memory write access only if the reservation is still set when the command is received. As the tagged address is not locked, the tag can be reset by competing traffic directed to the same location from other threads between RDL and WRC.

Consequently, the WRC command expects a response from the monitor logic, reflecting whether the write operation has been performed. To answer that requirement, OCP provides the value, FAIL for the SResp field (meaning that `writeresp_enable` is on if `rdlwrc_enable` is on). WRC is the only OCP command that makes use of the FAIL code, though new commands in future revisions may. FAIL responses are frequently received in a system using lazy synchronization that operates normally. Do not confuse FAIL with SResp=ERR, which effectively signals a system interconnect error or a target error.

Both RDL and WRC commands assume a single transaction model and cannot be used in a burst.


The semantics of lazy synchronization are defined on the previous page. Some specific sequences resulting from the usage of the RDL and WRC semantics are:

- A thread can issue more than one RDL before issuing a WRC, or issue more than one RDL without issuing WRC. Whether the subsequent RDL clears the reservation or sets a new one is implementation-specific, depending on the number of hardware monitors. At least one monitor per thread is required.
- If a thread issues a WR or WRNP command to an address it previously tagged with a RDL command, the write access clears all reservations from other threads for the same address (but not its own reservation).
- If a thread issues a WRC without having issued a RDL, the WRC will fail.
- If a thread issues a RDEX between the RDL and WRC, the RDEX is executed, sets the lock and waits for the corresponding write to clear the lock. RDL-WRC reservations will not be affected by the RDEX. The WR or WRNP that clears the lock, also clears any reservation set by other initiators for the same address (with the same MAddr, MByteEn and MAddrSpace if applicable).

Because competing requests of any type from other threads to a locked (by RDEX) location are blocked from proceeding until the lock is released, a RDEX command being issued between RDL and WRC commands, also blocks the WRC until the WR or WRNP command clearing the lock is issued. This favours RDEX-WR or WRNP sequence over RDL-WRC, in the sense that competing RDEX-WR or WRNP and RDL-WRC sequences will always result in having the RDEX-WR or WRNP sequence win.

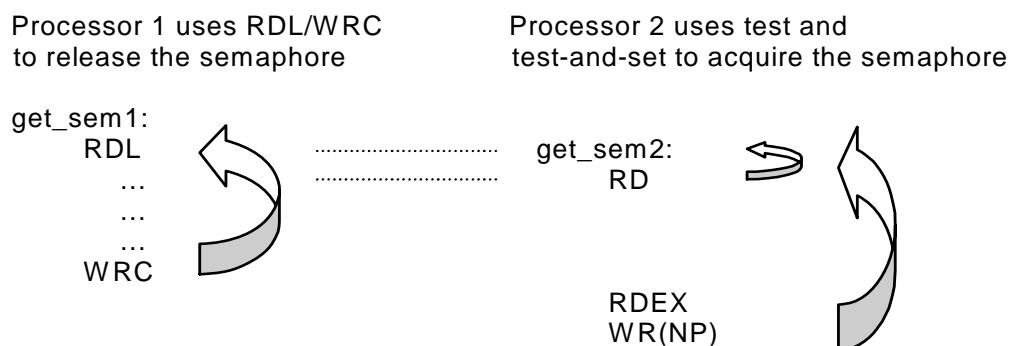
Incorrect use of the two synchronization mechanisms can result in deadlock so for example, the sequence of commands shown in Figure 48 might result in a deadlock. In this example Processor 1 tries to release the semaphore using RDL-WRC commands, Processor 2 tries to acquire the semaphore using RDEX-WR or WRNP commands. The RDEX-WR or WRNP sequence always occurs between the RDL and WRC. Because the WR or WRNP clearing the lock in Processor2 will also clear the reservation for Processor 1, the RDL-WRC sequence will never succeed. Processor 1 will never be able to release the lock or Processor2 to acquire it.

Figure 48 Synchronization Deadlock



The deadlock depicted in Figure 48 is a result of bad programming in Processor 2, and is very unlikely to happen in a real application environment. As shown in Figure 49, to achieve forward progress, Processor 2 should read the semaphore value and wait for the semaphore to be free before trying to retrieve it by issuing a RDEX-WR or WRNP.

Figure 49 Correct Synchronization Sequence



OCP and Endianness

As described in “Endianness” on page 47, OCP is nearly endian-neutral. While OCP specifies a byte address on MAddr, the address must be aligned to the data width of the interface. Sub-word quantities are specified using one bit for each enabled byte in the transfer on MByteEn or MDataByteEn.

While the bit ordering of OCP fields is consistently described in a little-endian fashion, this is conventional, where even big-endian systems tend to number their bits little-endian. Similarly, the MByteEn numbering seems to imply a little-endian byte ordering, but is simply intended to maintain consistency. For example, MByteEn[m] refers to the byte transferred on MData/SData[(8m+7):8m] (provided $m < \text{data width}/8$), regardless of the effective transfer endianness attributes.

If the master OCP and the slave OCP are the same data width, endianness does not matter. Addresses, data, and byte enables must remain consistent across both interfaces. (There are exceptions, since packed sub-word data objects should be swapped if the endianness does not match. OCP does not carry the required signaling to determine sub-word sizes, so full-word transfers must be assumed.)

Endianness problems arise as soon as one looks to connect a master and slave with different data widths. The narrow side has extra (non-zero) address bits, since its word-aligned addresses do not force as many bits to be zero. The wide side has extra byte lanes to carry its wider words. The association of the extra address bits (narrow side) with the extra byte lanes (wide side) specifies an endianness.

To bridge interfaces that suffer from mismatched data widths, packing and unpacking is required. Data width conversion must make some assumptions about the correspondence between the MAddr least-significant bits and the MByteEn field.

If the association maps the low-order byte lanes to lower addresses, the data width conversion is performed in a little-endian manner. If the association maps the high-order byte lanes to lower addresses, the data width conversion is performed in a big-endian manner. This operation is absolutely not an endianness conversion, but rather an endianness-aware packing or unpacking operation, so that the transaction endianness is preserved across the data width converter.

There is no attempt to perform any endian conversion in hardware. Rather, the goal is to enable interconnects that are essentially endian-neutral, but become endian-adaptive to match the endianness of the attached entities. This implies that the native endianness of an OCP core must be specified. OCP captures that property using the `endian` parameter, which can take four values:

LITTLE

Qualifies little-endian only cores

BIG

Qualifies big-endian only cores

BOTH

Qualifies cores that can change endianness:

- Based upon an external input such as a CPU that statically selects its endianness at boot time
- Based upon an internal configuration register such as a DMA engine, that generates OCP read and write requests in accordance with the endianness of the target, as stated by the DMA programmer
- Cores that support dynamic endianness

NEUTRAL

Qualifies cores that have no inherent endianness. Examples are simple memory devices that only work with full OCP-word quantities, or peripheral devices, the endianness of which can be controlled by the software device driver.

While not supported by the standard set of OCP features, it is possible to define a dynamic, endian-aware interconnect using in-band information. By specifying the parameters `reqinfo` (for request packing / unpacking control), `mdatainfo` (for data packing / unpacking control when `datahandshake` is enabled), and `respinfo` (for response packing / unpacking control), the definition of all these qualifiers becomes platform-specific.

Security

To protect against software and some selective hardware attacks use the OCP interface to create a secure domain across the SOC. The domain might include cpu, memory, I/O etc. that need to be secured using a collection of hardware and software features such as secured interrupts, and memory, or special instructions to access the secure mode of the processor.

The master drives the security level of the request using `MReqInfo` as a subnet. The master provides initiator identification using `MConnID`. Table 29 summarizes the relevant parameters.

Table 29 Security Parameters

Parameter	Value	Notes
<code>reqinfo</code>	1	<code>MReqInfo</code> is required
<code>reqinfo_width</code>	Varies	Minimum width is 1
<code>connid</code>	1	To differentiate initiators
<code>connid_width</code>	Varies	Minimum width is 1

The security request is defined as a named subnet `MSecure` within `MReqInfo`, for example:

subnet `MReqInfo M:N MSecure`, where $M \geq N$.

MSecure Bit Codes

With the exception of bit 0, bits are optional and the encoding is user-defined. Bit 0 of the MSecure field is required. The suggested encoding for the MSecure bits is:

Bit	Value 0	Value 1
0	non-secure	secure
1	user mode	privileged mode
2	data request	instruction request
3	user mode	supervisor mode
4	non-host	host
5	functional	debug

A special error response is not specified. A security error can be signaled with response code ERR.

Sideband Signals

The sideband signals provide a means of transmitting control-oriented information. Since the signals are rarely performance sensitive, drive all sideband signals stable early in the OCP clock cycle by making the sideband outputs come directly out of core flip-flops. To allow sideband inputs to arrive late in the OCP clock cycle register the inputs immediately on the receiving core.

Cores that fail to implement this conservative timing may require modification to achieve timing convergence.

Reset Handling

Some anomalous events can result from OCP resets. Among the situations to be aware of are the following:

Power-on reset

At power-on or assertion of any hardware reset, an OCP reset may be asynchronously asserted. Accepting the use of asynchronous resets helps describe the interface behavior.

Asynchronous assertion of a synchronous reset

Asynchronous assertion of resets in OCP may result in an asynchronous reset being fed to a module expecting a synchronous reset. From a design point of view this discrepancy could lead to a setup or hold violation. The required 16 clock cycles of reset guarantees enough time for recovery on the interface receiver side, allowing it to fall back to a safe functional state.

For simulation and verification, such timing violations represent a major hurdle since they may lead to inconsistent states in the master, slave and monitor interfaces. To address this problem, whenever possible, generate resets in a synchronous manner even if the protocol allows for their asynchronous assertion.

Use of OCP resets as asynchronous

OCP reset requires that a reset signal observe the setup and hold times as defined in the core's timing guidelines for at least 16 rising OCP clock edges after reset assertion, making the signal effectively synchronous except at assertion time. To satisfy this requirement do not connect an input OCP reset signal to the asynchronous clear/set pin of a D-flip-flop involved in an OCP signal logic cone. Failure to comply with this rule may violate the OCP protocol. For instance, a glitch on an OCP reset signal that would be sampled as deasserted could be interpreted as asserted, inadvertently causing the receiver to cancel pending transactions and hang the interface.

Dual Reset Signals

Many systems are fully satisfied with a single reset signal applied to both the master and the slave on an OCP interface. Either the master or the slave can drive the reset, or a third entity, such as a chip level reset manager, can provide it to both master and slave.

In some situations, it is more convenient for the master and slave to employ their own reset domains and communicate these internal resets to one another. The OCP interface is unable to communicate until both sides are out of reset since the side still in reset may be driving undetermined values (X) on their OCP outputs and cause problems for the side that is already out of reset. Examples of cases where this might arise are:

- A core with multiple OCP interfaces that are connected to different interconnects, which are each in different reset domains, plus the core has its own internal reset domain.
- Two connected interconnects that both act as initiators of transfers and each with their own reset domain.

Adding a second reset signal to the interface allows each master and slave to have both a reset output and input. The composite reset state for the OCP interface is established as the combination of the two resets, so that either side (or both) asserting reset causes the interface to be in reset. While in reset, the existing rules about the interface state and signal values apply.

Either MReset_n or SReset_n must be present on any OCP interface. Compatibility between different reset configurations of master and slave interfaces is shown in Table 30.

Table 30 Reset Configurations

Master Slave			
	sreset=1, mreset=0	sreset=0, mreset=1	sreset=1, mreset=1
sreset=0, mreset=1	Dual resets driven by the same 3 rd party	Single reset driven by master	Single reset driven by master (SReset_n input tied off to 1)
sreset=1, mreset=0	Single reset driven by slave	Incompatible	Incompatible
sreset=1, mreset=1	Single reset driven by slave (MReset_n input tied off to 1)	Incompatible	Dual resets

The rules describing this table can be stated as follows.

- Either `mreset` or `sreset` or both must be set to 1 for each core.
- The default (and only) tie-off value for `MReset_n` and `SReset_n` is 1.
- If `mreset` is set to 1 for the master and `mreset` is set to 0 for the slave, the reset configurations are incompatible.
- If `sreset` is set to 1 for the slave and `sreset` is set to 0 for the master, the reset configurations are incompatible.

Cores with a reset input are always interoperable with any other core. Add a reset output if it is needed by the core or subsystem to assure proper operation. Typically this is because both sides need to know about the reset state of the other side, or because the overall system does not function properly if the core or subsystem is in reset, while the OCP interface is not in reset.

Compatibility with OCP 1.0

OCP 1.0 cores that have a reset input or output can be converted to OCP 2 cores by renaming the `Reset_n` pin in the core's RTL conf file without touching the actual HDL source of the core. The new name depends on whether the reset is an input or output and whether the core is a master or slave.

In the very unlikely situation of an OCP 1.0 core lacking a reset input or output, the conversion to OCP 2 is achieved by the addition of a dummy reset input pin that is not used inside the core.

Debug and Test Interface

There are three debug and test interface extensions predefined for the OCP: scan, clock control, and IEEE 1149. The scan extension enables internal scan techniques, either in a pre-designed hard-core or end user inserted into a soft-core. Clock control extensions assist in scan testing and debug when the IP core has at least one other clock domain that is not derived from the OCP

clock. The IEEE 1149 extension is for interfacing to cores that have an IEEE 1149.1 test access port built-in and accessible. This is primarily the case with cores, such as microprocessors, that were derived from standalone products.

These three extensions along with sideband signals (flags) can yield a highly debuggable and testable IP core and device.

Scan Control

The width and meaning of the Scanctrl field is user-defined. At a minimum this field carries a signal to specify when the device is in scan chain shifting mode. The signal can be used for the scan clock if scan-clock style flip-flops are being used. When this is a multi-bit field, another common signal to carry would be one specifying the scan mode. This signal can be used to put the IP core into any special test mode that is necessary before scanning and application of ATPG vectors can begin.

Clock Control

The clock control test extensions are included to ease the integration of IP cores into full or partial scan test environments and support of debug scan operations in designs that use clock sources other than the OCP clock.

When an external clock source exists (for example, non-Clk derived clock), the ClkByp signal specifies a bypass of the external clock. In that case the TestClk signal usually becomes the clock source. The TestClk toggles in the correct sequence for applying ATPG vectors, stopping the internal clocks, and doing scan dumps as required by the user.

10 *Timing Guidelines*

To provide core timing information to system designers, characterize each core into one of the following timing categories:

- Level0 identifies the core interface as having been designed without adhering to any specific timing guidelines.
- Level1 timing represents conservative interface timing.
- Level2 represents high performance interface timing.

One category is not necessarily better than another. The timing categories are an indication of the timing characteristics of the core that allow core designers to communicate at a very high level about the interface timing of the core. Table 31 represents the inter-operability of two OCP interfaces.

Table 31 Core Interface Compatibility

	Level0	Level1	Level2
Level0	X	X	X
Level1	X	V	V
Level2	X	V	V*

X no guarantee

V guaranteed inter-operability with possible performance loss (extra latency)

V* high performance inter-operability but some minor changes may be required

The timing guidelines apply to dataflow and sideband signals only. There is no timing guideline for the scan and test related signals.

Timing numbers are specified as a percentage of the minimum supported clock-cycle (at maximum operating frequency). If a core is specified at 100MHz and the `c2qtime` is given as 30%, the actual `c2qtime` is 3ns.

Level0 Timing

Level0 timing indicates that the core developer has not followed any specific guideline in designing the core interface. There is no guarantee that the interface can operate with any other core interface. Inter-operability for the core will need to be determined by comparing timing specifications for two interfaces on a per-signal basis.

Level1 Timing

Level1 timing indicates that a core has been developed for minimum timing work during system integration. The core uses no more than 25% of the clock period for any of its signals, either at the input (`setuptime`) or at the output (`outputtime`). A core interface in this category must not use any of the combinational paths allowed in the OCP interface.

Since inputs and outputs each only use 25%, 50% of the cycle remains available. This means that a Level1 core can always connect to other Level1 and Level2 cores without requiring any additional modification.

Level2 Timing

Level2 timing indicates that a core interface has been developed for high performance timing. A Level2 compliant core provides or uses signals according to the timing values shown in Table 32. There are separate values for single-threaded and multi-threaded OCP interfaces. The number for each signal indicates the percentage of the minimum cycle time at which the signal is available, that is the `outputtime` at the output. `setuptime` at the input is calculated by subtracting the number given from the minimum cycle time. For example, a time of 30% indicates that the `outputtime` is 30% and the `setuptime` is 70% of the minimum clock period.

In addition to meeting the timing indicated in Table 32, a Level2 compliant core must not use any combinational paths other than the preferred paths listed in Table 33.

There is no margin between `outputtime` and `setuptime`. When using Level2 cores, extra work may be required during the physical design phase of the chip to meet timing requirements for a given technology/library.

Table 32 Level2 Signal Timing

Signal	Single-threaded Interface %	Multi-threaded Interface %	Multi-threaded Pipelined
Control, Status	25	25	25
ControlBusy, StatusBusy	10	10	10
ControlWr, StatusRd	25	25	25
Datahandshake Group (excluding MDataThreadID)	30	60	30
EnableClk	20	20	20
MDataThreadID	n/a	50	30
MRespAccept	50	75	50
MThreadBusy	10	10	50
MThreadID	n/a	50	30
Request Group (excluding MThreadID)	30	60	30
MReset_n, SReset_n	10	10	10
Response Group (excluding SThreadID)	30	60	30
SCmdAccept	50	75	50
SDataAccept	50	75	50
SDataThreadBusy	10	10	50
SError, SFlag, SInterrupt, MFlag, MError	40	40	40
SThreadBusy	10	10	50
SThreadID	n/a	50	30

Table 33 Allowed Combinational Paths for Level2 Timing (No Pipelining)

Core	From	To
Master	SThreadBusy	Request Group
	SThreadBusy SDataThreadBusy	Datahandshake Group
	Response Group	MRespAccept
Slave	MThreadBusy	Response Group
	Request Group	SCmdAccept and SDataAccept
	Datahandshake Group	SCmdAccept and SDataAccept

Table 34 Allowed Combinational Paths for Level2 Timing (Pipelined)

Core	From	To
Master	Request Group	SThreadBusy
	Datahandshake Group	SDataThreadBusy
	Response Group	MRespAccept
Slave	Response Group	MThreadBusy
	Request Group	SCmdAccept and SDataAccept
	Datahandshake Group	SCmdAccept and SDataAccept

11 *OCP Profiles*

A strength of OCP is the ability to configure an interface to match a core's communication requirements. Developing an OCP configuration can be challenging for a new user. This chapter provides profiles that capture the OCP features associated with standard communication functions. The pre-defined profiles help map the requirements of a new core to OCP configuration guidelines. The expected benefits include:

- Reduced risk of incompatibility when integrating OCP based cores originating from different providers
- Reduced learning curve in applying OCP for standard purposes
- Simplified circuitry needed to bridge an OCP based core to another communication interface standard
- Improved core maintenance
- Simplified creation of reusable core test benches

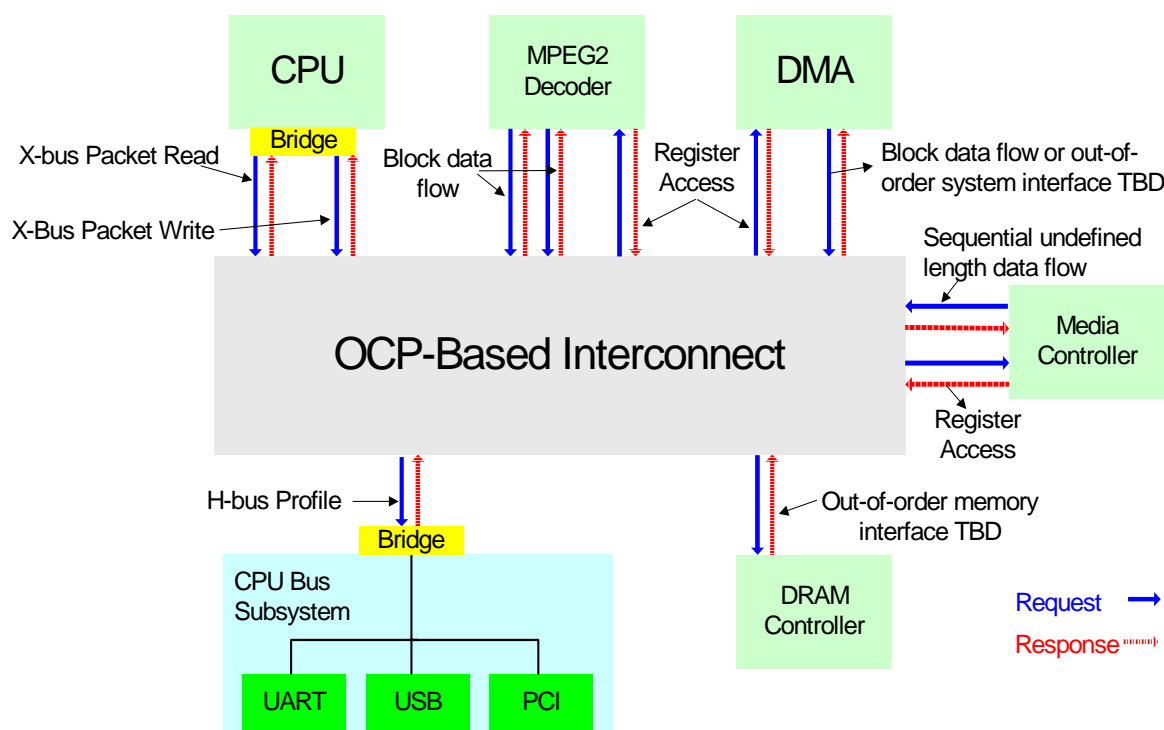
Profiles address only the OCP interface, with each profile consisting of OCP interface signals, specific protocol features, and application guidelines. For cores natively equipped with OCP interfaces, profiles minimize the number of interface and protocol options that need to be considered.

Two sets of OCP profiles are provided: profiles for new IP cores implementing native OCP interfaces and profiles that are targeted at designers of bridges between OCP and other bus protocols.

Since the other bus protocols may have several implementation flavors that require custom OCP parameter sets, the bridging profiles are incomplete. The bridging profiles can be used with OCP serving as either a master or a slave.

Figure 50 contains an SOC created using a mix of profiles. The diagram shows cores connected using two different interconnects. One interconnect is OCP-based, connecting a variety of cores using differing profiles. The other interconnect reflects a legacy on-chip bus that is populated largely with cores compatible with that bus protocol.

Figure 50 OCP Profiles



The figure shows the two styles of profiles. The CPU and CPU bus subsystem use the bridging profiles, while the rest of the IP cores use the native OCP profiles.

The CPU connects to the OCP-based interconnect twice; using the X-bus packet read and write profiles. These interfaces support cacheable and non-cacheable instruction and data traffic between the CPU and the memories and register interfaces of other targets. These profiles might be used with a CPU core that for example, internally used the AMBA AXI protocols, and were externally bridged to OCP.

The CPU bus subsystem connects to the OCP-based interconnect using the H-bus profile, through an external bridge. This profile might be used if the CPU bus subsystem was based on the AMBA AHB protocol.

The MPEG2 decoder core has multiple OCP interfaces. The core has two OCP master interfaces that use the block data flow profile, which is particularly effective for managing pipelined access of defined-length traffic (for example, MPEG macroblocks) to and from memory. Two interfaces separate independent traffic, enabling additional parallelism and higher peak

performance without using OCP threads. The core also has an OCP slave interface that uses the register access profile, providing the control processor (in this case, the CPU) with the ability to program the operation of the decoder.

The DMA and media controllers also include register access, profile-based slave interfaces, enabling CPU control over their operation.

The principle data flow interface of the DMA controller is through its OCP master interface. Depending upon the degree of parallelism needed by the controller, it could either follow the block data flow profile, or exploit thread-level parallelism using the out-of-order system interface profile (which is not included in this revision of the Specification).

The OCP master interface of the media controller uses the sequential undefined length profile to communicate a data stream with a memory based buffer. Because of the sequential undefined length characteristics of a data stream, this profile is a good fit for the media controller. In this example, packetization for transport at the OCP based interconnect level would be left to the interconnect.

A shared SDRAM controller core is attached to the OCP-based interconnect using a slave OCP interface. As this controller optimizes bank and page accesses to SDRAM, it can minimize local buffering, maximize throughput, and minimize latency by re-ordering requests. It therefore uses the out-of-order memory interface profile (which is not included in this revision of the Specification) to leverage both OCP threads and tags to fully exploit transaction parallelism.

Profile Types

The profiles cover the most common core communication requirements. Each OCP profile defines one or more applications. The available profiles are:

- Block data flow
- Sequential undefined length data flow
- Register access
- Simple H-bus
- X-bus packet write
- X-bus packet read
- Security

Each profile addresses distinct OCP interfaces, though most systems constructed using OCP will have a mix of functional elements. As different cores and subsystems have diverse communication characteristics and constraints, various profiles will prove useful at different interfaces within the system.

Native OCP Profiles

The native OCP profiles are designed for new IP cores implementing native OCP interfaces.

Block Data Flow Profile

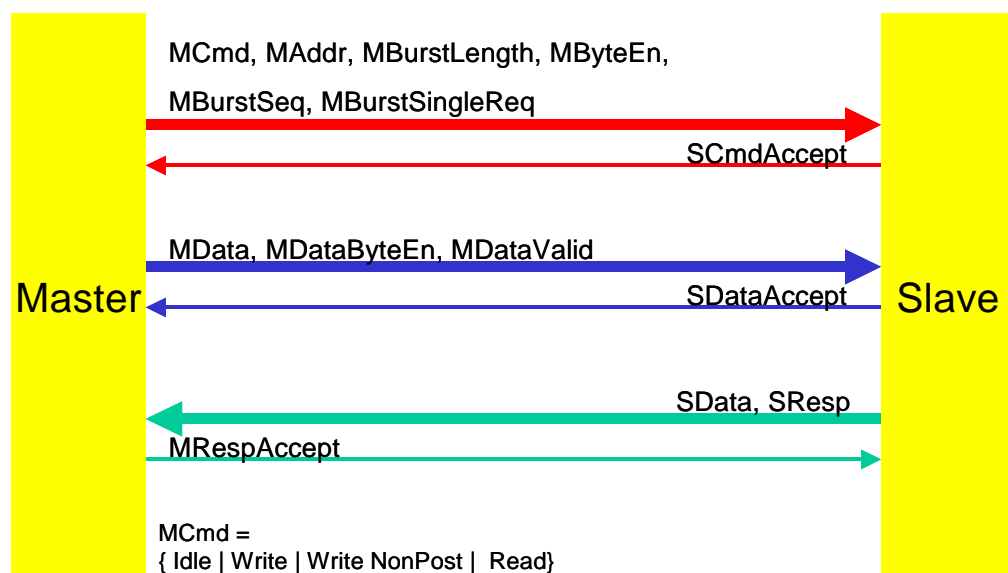
The block data flow profile is designed for master type (read/write, read-only, or write-only) interfaces of cores exchanging data blocks with memory. This profile is particularly effective for managing pipelined access of defined-length traffic (for example, MPEG macroblocks) to and from memory.

Core Characteristics

Cores with the following characteristics would benefit from using this profile:

- Block-based communication (includes a single data element block)
- A single request/multiple data burst model using incremental or a stream burst sequence
- De-coupled, pipelined command and data flows
- 32 bit address
- Natural data width and block size
- Use of byte enables
- Single threaded
- Support for producer/consumer synchronization through non-posted writes

Figure 51 Block Data Flow Signal Processing



Interface Configuration

Table 35 lists the OCP configuration parameters that need to be set along with the recommended values. For default values refer to Table 25, “Configuration Parameter Defaults,” on page 64.

Table 35 Block Data Flow Parameter Settings

Parameter	Value	Notes
addr_width	32	
burstlength	1	
burstlength_width	Core specific	Choose a burst length that is natural to the operation of the core
burstseq	1	If the core is STRM burst capable
burstseq_strm_enable	1	If the core is STRM burst capable
burstsinglereq	1	
byteen	1	For interfaces that are capable of partial access
data_width	Core specific	Choose a data width that is natural to the operation of the core
dataaccept	1	
datahandshake	1	
mdatabyteenable	1	
read_enable	1	For read capable interface only
respaccept	1	
write_enable	1	For write capable interface only
writenonpost_enable	1	See note on synchronization below
writeresp_enable	1	

Implementation Notes

When implementing this profile, consider the following suggestions:

- Start read transactions as early as possible to minimize read latency behind ongoing transactions.
- To implement a synchronization scheme (typically the case when data written by the IP core to shared memory is read by another core in the system), the IP core should issue a synchronization request (for instance through an OCP flag interface) only after receiving notification that the last write transaction is complete. To accomplish this step, perform all write transactions up to the last one as posted writes. Make the final write transaction a non-posted write. This will lead to reception of a response once the non-posted write transaction has completed.
- Error responses should lead to an interrupt.

Sequential Undefined Length Data Flow Profile

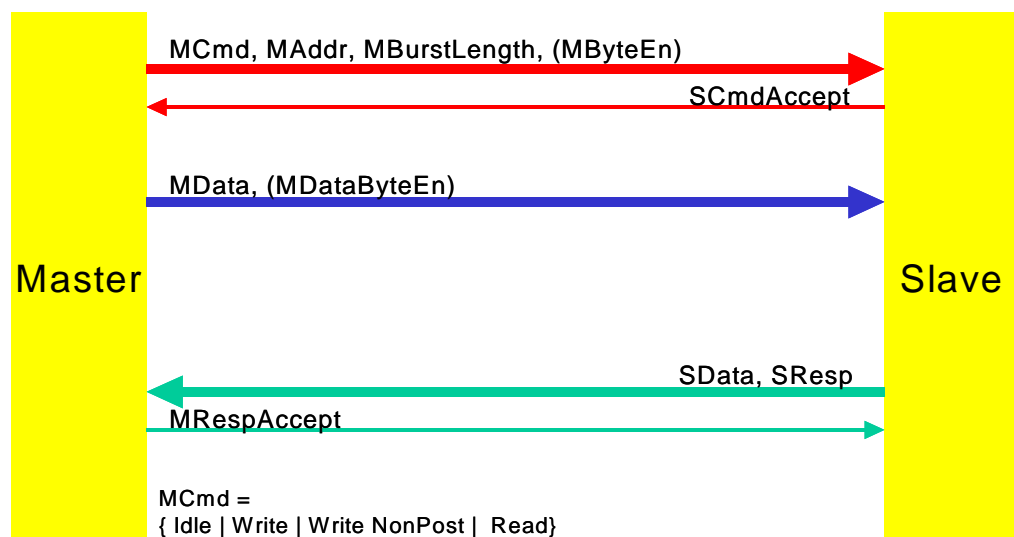
This profile is a master type (read/write or read-only, or write-only) interface for cores that communicate data streams with memory.

Core Characteristics

Cores with the following characteristics would benefit from using this profile:

- Communication of an undefined amount of data elements to consecutive addresses
- Imprecise burst model
- Aligned command/write data flows, decoupled read data flow
- 32 bit address
- Natural data width
- Optional use of byte enables
- Single thread
- Support for producer/consumer synchronization

Figure 52 Sequential Undefined Length Data Flow Signals Processing



Interface Configuration

Table 36 lists the OCP configuration parameters that need to be set along with the recommended values. For default values refer to Table 25, “Configuration Parameter Defaults,” on page 64.

Table 36 Sequential Undefined Length Data Flow Parameter Settings

Parameter	Value	Notes
addr_width	32	
burstlength	1	
burstlength_width	2	
burstseq	1	
byteen	1 (optional)	For interfaces that are capable of partial access
data_width	core specific	Choose a data width that is natural to the operation of the core
read_enable	1 (optional)	For read capable interface only
respaccept	1	
write_enable	1 (optional)	For write capable interface only
writenonpost_enable	1	
writeresp_enable	1	

Implementation Notes

When implementing this profile, consider the following suggestions:

- The core streams data to and from a memory-based buffer at sequential addresses. When hitting the boundaries of that buffer a non-sequential address is occasionally given.

MBurstLength equals 2 while the data stream proceeds to sequential addresses; MBurstLength equals 1 to indicate that a non-sequential address follows, or that the stream terminates.

- Start read transactions as early as possible to hide read latency behind ongoing transactions.
- To implement a producer/consumer synchronization scheme (typically the case when data written by the IP core to shared memory is read by another core in the system), the IP core should issue a synchronization request (for instance through an OCP flag interface) only after receiving notification that the last write transaction is complete. To accomplish this step, perform all write transactions up to the last one as posted writes. Make the final write transaction a non-posted write. This will lead to reception of a response once the non-posted write transaction has completed at the final destination.
- Error response should lead to an interrupt.

Register Access Profile

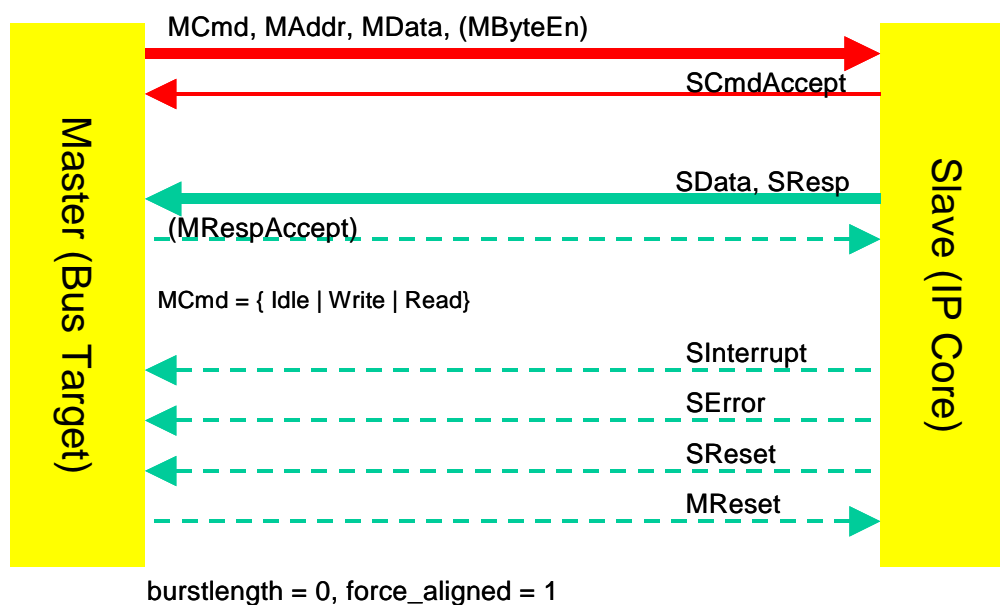
The register access profile offers a control processor the ability to program the operation of an attached core. This profile supports programmable register interfaces across a wide range of IP cores, such as simple peripherals, DMA engines, or register-controlled processing engines. The IP core would be an OCP slave on this interface, connected to a master that is a target addressed by a control processor.

Core Characteristics

Cores with the following characteristics would benefit from using this profile:

- Address mapped communication, but target decoding is handled upstream
- Natural data width (unconstrained, but 32 bits is most common)
- Natural address width (based on the number of internal registers X data width)
- No bursting
- Precise write responses indicate completion of write side-effects
- Single threaded
- Optional aligned byte enables for sub-word CPU access
- Optional response flow control
- Optional use of side-band signals for interrupt, error, and DMA ready signaling

Figure 53 Register Access Signals Processing



Interface Configuration

Table 37 lists the OCP configuration parameters that need to be set along with the recommended values. For default values refer to Table 25, “Configuration Parameter Defaults,” on page 64.

Table 37 Register Access Parameter Settings

Parameter	Value	Notes
addr	1	
addr_width	Varies	Num_regs * data_width/8
byteen	Varies	Not suggested for new designs
cmdaccept	1	
data_width	Varies	8, 16, 32 and 64 bits; 32 is preferred
force_aligned	1	Normally read/written by aligned CPU
interrupt	Varies	For cores with multiple interrupt lines use SFlag
mdatainfo	0	
mreset	Varies	
respaccept	1	Include when possible!
error	Varies	If core has internally-generated errors
sreset	Varies	If core receives own (non-interface) reset
writenonpost_enable	Varies	Not usually needed
writeresp_enable	1	Precise write responses needed.

Implementation Notes

When implementing this profile, consider the following suggestions:

- Choose a data width based on the internal register width of the core.
- Choose an address width based on the data width and number of registers.
- Design new cores so that all read/write side-effects can be managed without using byte enables.
- Design new cores so that registers are at least 32 bit aligned.
- Use force_aligned byte enables for cores with side effects in multiple-use registers.
- Implement response flow control if convenient.
- Implement sideband signaling towards CPU (interrupts, sideband errors, etc.) on this interface, since it is largely controlled by the CPU.
- Select reset options based on the expected use of the core in a larger system.

- Use regular Write commands for precise completion, unless the core is capable of per-transfer posting decisions, where mixing Write and WriteNonPost helps.

Bridging Profiles

These profiles are designed to simplify or automate the creation of bridges to other interface protocols. The bridge can have an OCP master or slave port. There are two types:

- The simple H-bus profile is intended to provide a connection through an external bridge, for example to a CPU with an AMBA AHB protocol.
- The X-bus interfaces support cacheable and non-cacheable instruction and data traffic between a CPU and the memories and register interfaces of other targets. The X-bus profiles might be used with a CPU core that internally uses the AMBA AXI protocols, and is externally bridged to OCP.

Simple H-bus Profile

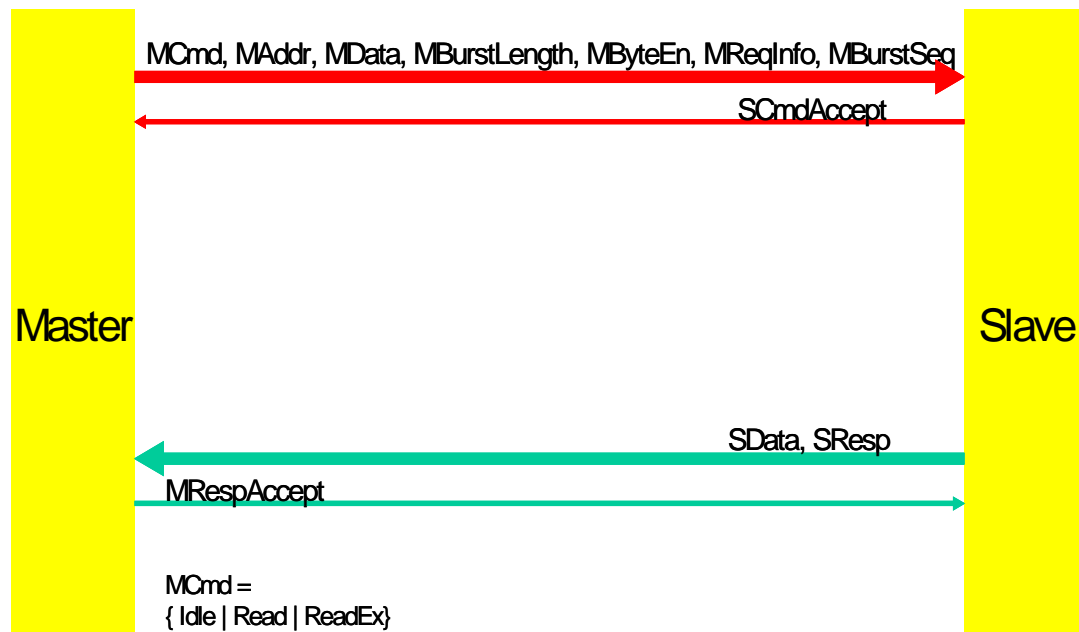
This profile allows you to create OCP master wrappers to native interfaces of simple CPU type initiators with multiple-request/multiple-data, read and write transactions.

Core Characteristics

Cores with the following characteristics would benefit from using this profile:

- Address mapped communication
- Natural address width
- Byte enable
- Natural data width
- Constrained burst size
- Single thread
- Caching and similar extensions mapped to MReqInfo or corresponding OCP commands

Figure 54 Simple H-Bus Signal Processing



Interface Configuration

Table 38 lists the OCP configuration parameters that need to be set along with the recommended values. For default values refer to Table 25, “Configuration Parameter Defaults,” on page 64.

Table 38 Simple H-Bus Parameter Settings

Parameter	Value	Notes
addr_width	Varies	Use native address width
burstlength	1	
burstlength_width	5	Only short burst supported
burstprecise	0	MBurstPrecise signal is not part of the interface. All bursts are precise
burstseq	1	Subset of burst codes common with OCP and CPU
burstseq_wrap_enable	1	
byteen	1	
data_width	Varies	8, 16, 32 and 64 bits
force_aligned	1	
mreset	1	
readex_enable	1	For CPUs with locked access
reqinfo	1	
reqinfo_width	Varies	Map CPU-specific inband info here

Parameter	Value	Notes
reqlast	1	Can be created of burst length
respaccept	1	
sreset	1	
writeresp_enable	1	

Implementation Notes

When implementing this profile, consider the following suggestions:

- If the CPU's burst parameters such as data alignment are not supported in OCP, such bursts are broken into single transactions.
- All requests have a response.
- If the CPU address and write data are pipelined, the OCP bridge will align them.
- CPU specific information is mapped to the MReqInfo field, however, since this field is often used for proprietary bits by OCP users, a precise bit-to-bit mapping is not provided. For this field concatenate bit fields starting from bit 0. If the inband field contains control information that has equivalent native OCP functionality, map the information to the corresponding OCP request. For example, a bit that can be buffered can be mapped to OCP posted or non-posted write types.

X-Bus Packet Write Profile

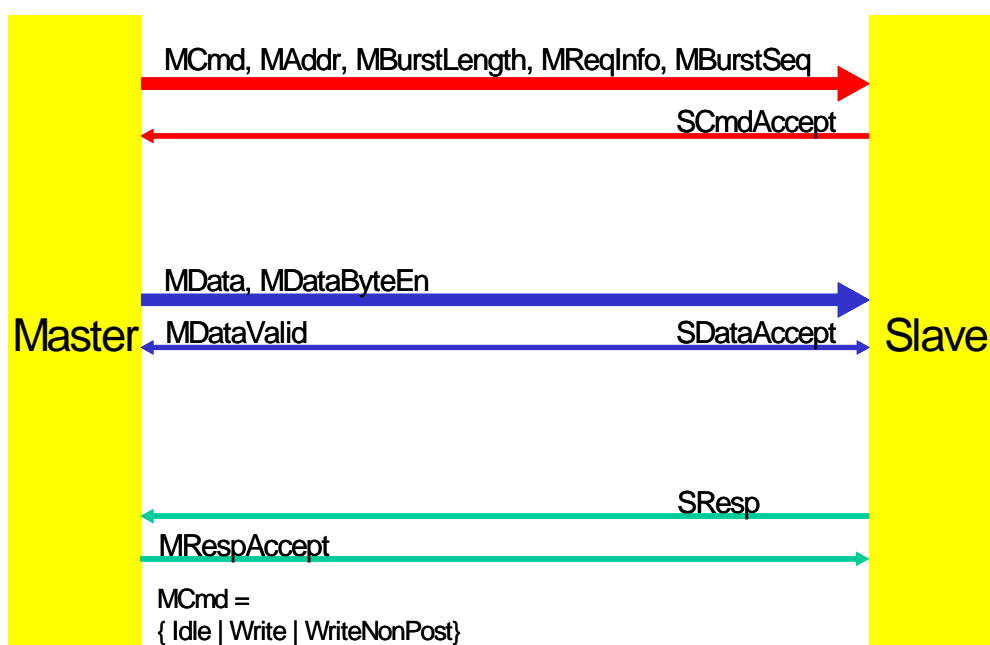
This profile is designed to create OCP master wrappers to native interfaces of CPU type initiators with single-request/multiple-data, write-only transactions.

Core Characteristics

Cores with the following characteristics would benefit from using this profile:

- Packet type communication
- Natural address width
- Separate command/write data handshake
- Byte enable
- Natural data width
- Multi-thread (blocking)
- Caching and similar extensions mapped to MReqInfo

Figure 55 X-bus Packet Write Signal Processing



Interface Configuration

Table 39 lists the OCP configuration parameters that need to be set along with the recommended values. For default values refer to Table 25, “Configuration Parameter Defaults,” on page 64.

Table 39 X-bus Packet Write Parameter Settings

Parameter	Value	Notes
addr_width	Varies	Use native address width
burstlength	1	
burstlength_width	5	Only short burst support
burstprecise	0	MBurstPrecise signal is not part of the interface. All bursts are precise
burstseq	1	Subset of burst codes common with OCP and CPU
burstseq_strm_enable	1	
burstseq_wrap_enable	1	
byteen	0	Only databyteen needed for write-only
data_width	Varies	8, 16, 32 and 64 bits
dataaccept	1	
datahandshake	1	
datalast	1	Can be created of burst size

Parameter	Value	Notes
interrupt	0	Interrupts are not part of this interface
mdatabyteen	1	
mreset	1	
read_enable	0	Write only
reqdata_together	Varies	A simpler bridge can often be made if 1, some performance loss possible.
reqinfo	1	
reqinfo_width	Varies	Map CPU-specific inband info here
reqlast	1	Superfluous for single request, datalast suffices
respaccept	1	
resplast	1	
sdata	0	
sreset	1	
sthreadbusy	0	Blocking threads only
threads	Varies	Natural number of threads
writenonpost_enable	1	
writeresp_enable	1	

Implementation Notes

When implementing this profile, consider the following suggestions:

- Data ordering among read and write port transactions is the responsibility of the CPU. The bridge must consider the read and write ports as single master with regard to exclusive access.
- Only precise bursts are supported. If CPU burst parameters do not map to OCP, break the burst into single accesses.
- CPU specific information is mapped as in the H-bus profile.

X-Bus Packet Read Profile

This profile helps you create OCP master wrappers for native interfaces of CPU type initiators with single-request multiple-data read-only transactions.

Core Characteristics

Cores with the following characteristics would benefit from using this profile:

- Packet type communication
- Natural address width

- Single-request multiple-data read
- Byte enable
- Natural data width
- Multi-thread (blocking)
- Caching and similar extensions mapped to MReqInfo
- Write support for ReadEx/Write synchronization

Figure 56 X-bus Packet Read Signal Processing



Interface Configuration

Table 40 lists the OCP configuration parameters that need to be set along with the recommended values. For default values refer to Table 25, “Configuration Parameter Defaults,” on page 64.

Table 40 X-bus Packet Read Parameter Settings

Parameter	Value	Notes
addr_width	Varies	Use native address width
burstlength	1	
burstlength_width	5	Only short burst support
burstprecise	0	MBurstPrecise signal is not part of the interface. All bursts are precise
burstseq	1	Subset of burst codes common with OCP and CPU

Parameter	Value	Notes
burstseq_strm_enable	1	
burstseq_wrap_enable	1	
burstsinglereq	1	
byteen	1	
data_width	Varies	8, 16, 32 and 64 bits
force_aligned	0	If CPU alignment is not supported in OCP, such bursts are broken into single transactions.
interrupt	0	Interrupts are not part of this interface
mreset	1	
readex_enable	1	
reqinfo	1	
reqinfo_width	Varies	Map CPU-specific inband info here
respaccept	1	
resplast	1	
sreset	1	
sthreadbusy	0	Blocking threads only
threads	Varies	Natural number of threads
write_enable	1	To support ReadEx

Implementation Notes

When implementing this profile, consider the following suggestions:

- Data integrity among read and write port transactions is the responsibility of the CPU. The bridge must consider the read and write ports as single master with regard to exclusive access. Both transfers in the ReadEx/Write pair should be issued on the X-bus packet read interface.
- Only precise bursts are supported. If the CPU burst parameters do not map to OCP, break the burst into single accesses.
- CPU specific information is mapped as in the H-bus profile.

Layered Profiles


Layered profiles extend the OCP interface as an add-on to any other profile, when additional features are required.

Security

To protect against software and some selective hardware attacks use the OCP interface to create a secure domain across the SOC. The domain might include cpu, memory, I/O etc. that need to be secured using a collection of hardware and software features such as secured interrupts, and memory, or special instructions to access the secure mode of the processor.

The master drives the security level of the request using MReqInfo as a subnet. The master provides initiator identification using MConnID.

Figure 57 Security Signal Processing



Interface Configuration

Table 41 lists the OCP configuration parameters that need to be set along with the recommended values. For default values refer to Table 25, “Configuration Parameter Defaults,” on page 64.

Table 41 Security Parameters

Parameter	Value	Notes
reqinfo	1	MReqInfo is required
reqinfo_width	Varies	Minimum width is 1
connid	1	To differentiate initiators, if required
connid_width	Varies	Minimum width is 1

Implementation Notes

When implementing this profile, consider the following suggestions:

- Define the security request as a named subnet MSecure within MReqInfo, for example: subnet MReqInfo M:N MSecure, where M is \geq N.

With the exception of bit 0, other bits are optional and the encoding is user-defined. Bit 0 of the MSecure field is required and must use the specified value. The suggested encoding for the MSecure bits is:

Bit	Value 0	Value 1
0	non-secure	secure
1	user mode	privileged mode
2	data request	instruction request
3	user mode	supervisor mode
4	non-host	host
5	functional	debug

- A special error response is not specified. A security error can be signaled with response code ERR.

12 *Core Performance*

To make it easier for the system integrator to choose cores and architect the system, an IP core provider should document a core's performance characteristics. This chapter supplies a template for a core performance report on page 208, and directions on how to fill out the template.

Report Instructions

To document the core, you will need to provide the following information:

1. Core name. Identify the core by the name you assigned.
2. Core ID. Specify the precise identification of the core inside the system-on-chip. The information consists of the vendor code, core code, and revision code.
3. Core is/is not process dependent. Specify whether the core is process-dependent or not. This is important for the frequency, area, and power estimates that follow.

If multiple processes are supported, name them here and specify corresponding frequency/area/power numbers separately for each core if they are known.

4. Frequency range for this core. Specify the frequency range that the core can run at. If there are conditions attached, state them clearly.
5. Area. Specify the area that the core occupies. State how the number was derived and be precise about the units used.
6. Power estimate. Specify an estimate of the power that the core consumes. This naturally depends on many factors, including the operations being processed by the core. State all those conditions clearly, and if possible, supply a file of vectors that was used to stimulate the core when the power estimate was made.

7. Special reset requirements. If the core needs MReset_n/SReset_n asserted for more than the default (16 OCP clock cycles) list the requirement.
8. Number of interfaces.
9. Interface information. For each OCP interface that the core provides, list the name and type.

The remaining sections focus on the characteristics and performance of these OCP interfaces.

For master OCP interfaces:

- a. Issue rate (per OCP cycle for sequences of reads, writes, and interleaved reads/writes). State the maximum issue rate. Specify issue rates for sequences of reads, writes, and interleaved reads and writes.
- b. Maximum number of operations outstanding (pipelining support). State the number of outstanding operations that the core can support; is there support for pipelining.
- c. If the core has burst support, state how it makes use of bursts, and how the use of bursts affects the issue rates.
- d. High level flow-control. If the core makes use of high-level flow control, such as full/empty bits, state what these mechanisms are and how they affect performance.
- e. If multiple threads are present, explain the use of threads.
- f. Connection ID support. Explain the use and meaning of connection information.
- g. Use of side-band signals. For each sideband signal (such as SInterrupt, MFlag) explain the use of the signal.
- h. If the OCP interface has any implementation restrictions, they need to be clearly documented.

For slave OCP interfaces:

- a. Unloaded latency for each operation (in OCP cycles). Describe the unloaded latency of each type of operation.
- b. Throughput of operations (per OCP cycle for sequences of reads, writes, and interleaved reads/writes). State the maximum throughput of the operations for sequences of reads, writes, and interleaved reads and writes.
- c. Maximum number of operations outstanding (pipelining support). State the number of outstanding operations that the core can support, i.e. is there support for pipelining.

- d. Burst support and effect on latency and throughput numbers. If the core has burst support, state how it makes use of bursts, and how the use of bursts affects the latency and throughput numbers stated above.
- e. High level flow-control. If the core makes use of high-level flow control, such as full/empty bits, state what these mechanisms are and how they affect performance.
- f. If multiple threads are present, explain the use of threads.
- g. Connection ID support. Explain the use and meaning of connection information.
- h. Use of side-band signals. For each sideband signal (such as SInterrupt, MFlag) explain the use of the signal.
- i. If the OCP interface has any implementation restrictions, they need to be clearly documented.

For every non-OCP interface, you will need to provide all of the same information as for OCP interfaces wherever it is applicable.

Sample Report

1. Core name	flashctrl
2. Core identity	
Vendor code	0x50c5
Core code	0x002
Revision code	0x1
3. Core is/is not process dependent	Not
4. Frequency range for this core	≤100Mhz with NECCBC9-VX library
5. Area	4400 gates 2input NAND equivalent gates
6. Power estimate	not available
7. Special reset requirements	
8. Number of interfaces	2
9. Interface information:	
Name	ip
Type	slave
For master OCP interfaces:	
a. Issue rate (per OCP cycle for sequences of reads, writes, and interleaved reads/writes)	
b. Maximum number of operations outstanding (pipelining support)	
c. Effect of burst support on issue rates	
d. High level flow-control	
e. Use of threads (if any)	
f. Use of connection information	
g. Use of side-band signals	
h. Implementation restrictions	

For slave OCP interfaces:	
a. Unloaded latency for each operation (in OCP cycles)	Register read or write: 1 cycle. The flash read takes SBFL_TAA (read access time). Can be changed by writing corresponding register field of emem configuration register. The flash write operation takes about 2000 cycles since it has to go through the sequence of operations - writing command register, reading the status register twice.
i. Throughput of operations (per OCP cycle for sequences of reads, writes, and interleaved reads/writes)	No overlap of operations therefore reciprocal of latency.
j. Maximum number of operations outstanding (pipelining support)	No pipelining support.
k. Effect of burst support on latency and throughput numbers	No burst support.
l. High level flow-control	No high-level flow-control support.
m. Use of threads (if any)	No thread support.
n. Use of connection information	No connection information support.
o. Use of side-band signals	Reset_n, Control, SError. Control is used to provide additional write protection to critical blocks of flash memory. SError is used when an illegal width of write is performed. Only 16 bit writes are allowed to flash memory.
p. Implementation restrictions	
For every non-OCP interface Provide all of the same information as for OCP interfaces wherever it is applicable.	Hitachi flash card HN29WT800 Only 1 flash ROM part is supported, therefore the CE_N is hardwired on the board. The ready signal RDY_N, is not used since not all parts support it. For the BYTE_N signal, only 16-bit word transfers are supported

Performance Report Template

Use the following template to document a core.

1. Core name	
2. Core identity Vendor code Core code Revision code	
3. Core is/is not process dependent	
4. Frequency range for this core	
5. Area	
6. Power estimate	
7. Special reset requirements	
8. Number of interfaces	
9. Interface information: Name Type	
For master OCP interfaces:	
a. Issue rate (per OCP cycle for sequences of reads, writes, and interleaved reads/writes)	
b. Maximum number of operations outstanding (pipelining support)	
c. Effect of burst support on latency and throughput numbers	
d. High level flow-control	
e. Use of threads (if any)	
f. Use of connection information	
g. Use of side-band signals	
h. Implementation restrictions	

For slave OCP interfaces:	
a. Unloaded latency for each operation (in OCP cycles)	
i. Throughput of operations (per OCP cycle for sequences of reads, writes, and interleaved reads/writes)	
j. Maximum number of operations outstanding (pipelining support)	
k. Effect of burst support on latency and throughput numbers	
l. High level flow-control	
m. Use of threads (if any)	
n. Use of connection information	
o. Use of side-band signals	
p. Implementation restrictions	
For every non-OCP interface Provide all of the same information as for OCP interfaces wherever it is applicable.	

Part III *Protocol Compliance*

13 *Compliance*

This section contains the OCP compliance checks that can help you create checking solutions in the language and tool of your choice.

The guidelines listed in this section are based on the "Specification" and "Guidelines" parts of this document and allow you to verify an IP/VIP for OCP compliance. In all cases, "Part I, Specification" is the definitive reference. Any references made to "Part II, Guidelines" are not definitive as Part I supersedes the guidelines.

For a core to be considered OCP compliant it must satisfy the compliance definition as described in "Compliance" on page 3.

Configuration Compliance

Interface Configuration

The main challenge in developing an OCP VIP lies in accounting for the high degree of configurability of OCP. Figure 58 shows the different inputs that can affect OCP configurability. To properly define the OCP interfaces of an IP/VIP, consider the following contexts.

Open System Context

For an open system, it must be possible to setup all of the OCP interfaces with a file using the <core>_rtl.conf syntax, which is required for OCP compliance. Fixed configuration IP/VIP must be delivered with a core_rtl.conf file describing the configuration. The metadata properties for this are described in Chapter 6.

Configurable IP/VIP supporting multiple OCP configurations must support the setup of any configuration using a <core>_rtl.conf file. The mechanisms used to fix the configuration must provide a method for generating a core_rtl.conf file that represents the fixed configuration and can be used to configure the IP/VIP directly.

For IP providers <core>_rtl.conf generation likely occurs during the IP generation step. When the IP code is generated based on configuration, and other settings in the GUI, the <core>_rtl.conf file is generated along with the IP.

Closed System Context

In a closed system, the verification of an IP/VIP with one or more OCP interfaces may be driven from a <core>_rtl.conf file. A vendor is free to implement any other solution. For example, a VERA verification environment could use a VERA object to control the OCP stimuli generators instead of a <core>_rtl.conf file.

If the IP/VIP is being developed in a closed system for delivery in an open system context, then the verification must include the <core>_rtl.conf files and any applicable <core>_rtl.conf generators that are delivered with the IP/VIP.

Configuration Parameter Extraction

Depending on the system context, the VIP must extract the OCP configuration parameters from the <core>_rtl.conf file (open) or from any alternate solution (closed). Parameters with indeterminate values must be retrieved using the configuration parameter defaults summarized in Table 25, “Configuration Parameter Defaults,” on page 64 (Table 22 of the *OCP 2.0 Specification*). Some parameters are required in certain configurations, and for those, no default is specified. For example: `addr_width` must always be specified if `addr == 1`.

Protocol Compliance

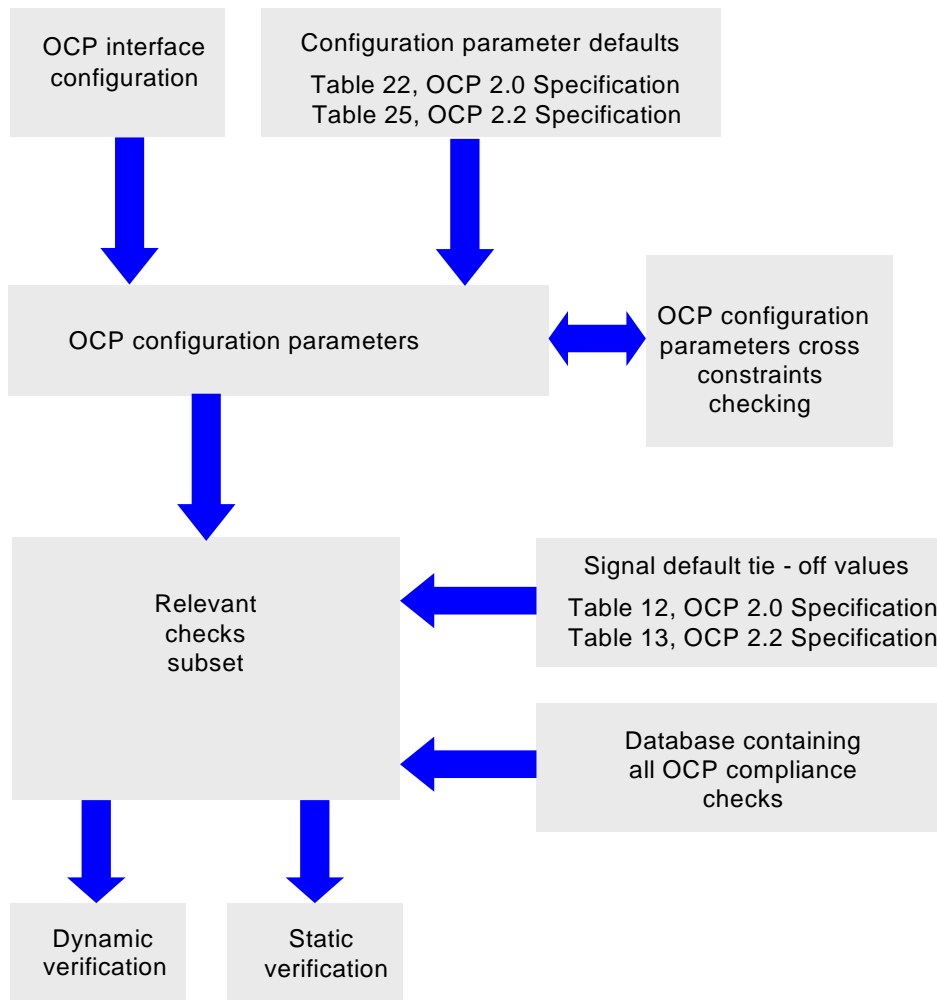
Once all the OCP configuration parameters are known, illegal OCP configurations must be flagged. Chapter 15 contains compliance checks for the configuration parameters. Chapter 4 contains most of the cross-constraints. For example: if `readex_enable` is set to 1, `write_enable` or `writenonpost_enable` must be set to 1.

Select the Relevant Checks

Based on the OCP configuration parameters, select a subset of the checks in the VIP OCP library. This subset is used for the actual verification. If a signal used by a check is not configured in the OCP interface and if no other tie-off value is specified, Table 13, “OCP Signal Configuration Parameters,” on page 30 (Table 12 of the *OCP 2.0 Specification*) specifies the inferred default tie-off values. For example, the `MBurstPrecise` default tie-off value is 1 or precise.

Check the compliance of the DUT OCP interfaces using static or dynamic verification techniques described in the next section

Figure 58 OCP Configurability



Verification Techniques

The verification guidelines are valid for developers using static or dynamic verification methods. This section provides an overview of static and dynamic verification methods, along with guidance on how these checks can be used to support these verification efforts.

Dynamic Verification

Dynamic verification methodology consists of:

- Driving a set of stimuli through the OCP interface into the DUT.

- Using a protocol checker on the VIP monitor traces of OCP interface activity to make sure that the protocol is not violated.
- Assessing the quality of the stimuli using functional and code coverage.

The OCP configuration parameters determine which protocol checks must be active and how the OCP functional coverage is defined.

Stimuli

Because of the degree of difficulty of defining a golden set of stimuli for any OCP interface configuration, you will need to implement a smart and efficient set of stimuli. This may be accomplished using constraint-driven random stimuli generation. The quality of the stimuli must be assessed using both functional and code coverage as described below.

Protocol Checks

The protocol checker is a passive component that monitors a specific set of OCP configuration parameters to determine whether the OCP protocol is violated. The protocol checker can be written in a variety of languages including HDL, PSL, SVA, E, NSCa, or VERA. The protocol checker must be instantiated on each OCP interface of the DUT.

To allow this document to be easily referenced, the names of the protocol checks must match the names given to the compliance checks described in this document.

Functional Coverage

Measure the quality of the applied stimuli. The target is 100% functional coverage. Run code coverage on the RTL to determine whether there are any verification holes such as uncovered FSM states or missed branches. Based on the code coverage analysis, additional coverage metrics may be required.

The guidelines for OCP functional coverage are provided in Chapter 16. Any additional coverage metrics, based on code-coverage analysis, are design dependent and are out of the scope of this document.

Static Verification

The static verification approach is also referred to as formal verification and relies on the following key elements:

- OCP protocol assertions (or protocol checkers)
- OCP protocol constraints (optional)
- OCP functional coverage

This approach uses a formal tool to prove that, given the stimulus limits defined by the OCP protocol constraints, the OCP interface of the DUT never violates OCP protocol assertions. The formal proof may be exhaustive (the assertions are never violated) or bounded (up to a certain depth of the state

space the assertions are never violated). Stimuli are not needed; instead the tool relies on the 'all acceptable stimuli' definitions provided by the OCP protocol constraints.

The OCP configuration parameters determine:

- Which protocol assertions must be active.
- How the functional coverage must be defined.
- Which protocol constraints must be active.

Protocol Assertions

Formal verification revolves around taking the protocol assertions and attempting to prove that they are never violated. If a violation is found, the formal tool provides a test sequence that illustrates the violation on the design. The assertions can be written in different languages such as HDL, PSL or SVA.

To allow this document to be easily referenced, the names of the assertions must match the names given to the compliance checks described in this document.

Protocol Constraints

To place bounds on the stimuli that a formal engine must consider, the design must be connected to protocol constraints or some form of generator description. Constraints can be specified using the same language as is used for protocol assertions, typically in HDL, PSL, SVA, or OVA.

Protocol constraints are not provided and must be obtained or created for use with formal tools. Constraints must be specific enough to prevent invalid test sequences that can lead to false negative test results (as indicated by protocol assertion failures) but not so specific that they prevent valid test sequences. The latter situation can lead to false positive test results implied by protocol assertion success over an incomplete set of test sequences.

Using functional coverage, false positive results can be checked, however, the false negatives cannot be checked as easily.

Functional Coverage

You must insure that all of the protocol assertions are verified to a reasonable extent, and that the protocol constraints are sound. To accomplish this, some functional coverage must be added as a function of the OCP configuration parameters. The functional coverage definition should cover:

- Which assertions warrant exhaustive proofs
- Which assertions are ok with just bounded proofs, at what depth
- Detect and correct an over-constrained environment

The guidelines for the OCP functional coverage are described in Chapter 16.

14 *Protocol Compliance Checks*

The compliance checks listed in this chapter are extracted from the *OCP Specification* and are intended to serve as guidelines to verify an IP for OCP compliance. In all cases "Part I, Specification" is the definitive reference.

The compliance check names have been created using the following template:

<hierarchy>_<check type>_<critical signal>_<extra details>

In which:

<hierarchy>	signal, request, datahs, response, burst, transfer, rdex
<check type>	valid, hold, value, exact, phase_order, lock_release, sequence, order, reorder
<critical signal>	(optional) any OCP signal name that is impacted by the compliance check

Activation Tables

Tables 42 - 47 list the parameters needed for each check to be initiated. The following assumptions are made with respect to these tables:

- An understanding of how a combination of these parameters can lead to an illegal configuration.
- The tables only shows the minimum parameters needed for a check to be fired. Each configuration needs the parameters defined for each check plus the parameters needed to make it a legal configuration. For example, a check for INCR bursts would need some command (read_enable, write_enable, etc.) parameters defined to test the check.

Table 42 Dataflow Signal Checks

Name	Activation Parameters
signal_valid_<signal>_when_reset_inactive	
MCmd	-
MDataValid	datahandshake
MThreadBusy	mthreadbusy
SDataThreadBusy	sdatathreadbusy
SResp	resp
SThreadBusy	sthreadbusy
request_valid_<signal>	
MAddr	addr
MAddrSpace	addrspace
MAtomicLength	atomiclength
MBlockHeight	blockheight
MBlockStride	blockstride
MBurstLength	burstlength
MBurstPrecise	burstprecise
MBurstSeq	burstseq
MBurstSingleReq	burstsinglereq
MByteEn	byteen
MConnID	connid
MReqLast	reqlast
MThreadID	threads > 1
SCmdAccept	cmdaccept
datahs_valid_<signal>	
MDataByteEn	mdatabyteen
MDataLast	datalast
MDataThreadID	datahandshake & threads > 1
SDataAccept	dataaccept
response_valid_<signal>	
MRespAccept	respaccept
SRespLast	resplast
SThreadID	resp & threads > 1
request_valid_MTagInOrder	Taginorder
response_valid_STagInOrder	resp & taginorder

Name	Activation Parameters
request_valid_MTagID_when_MTagInOrder_zero	tags > 1
datahs_valid_MTagID_when_MTagInOrder_zero	datahandshake & tags > 1
response_valid_STagID_when_STagInOrder_zero	resp & tags > 1

Table 43 Dataflow Phase Checks

Name	Activation Parameters
request_exact_SThreadBusy	sthradbusy & sthradbusy_exact & ~sthradbusy_pipelined
request_pipelined_SThreadBusy	sthradbusy & sthradbusy_exact & sthradbusy_pipelined
request_hold_<signal> MAddr MAddrSpace MAAtomicLength MBlockHeight MBlockStride MBurstLength MBurstPrecise MBurstSeq MBurstSingleReq MByteEn MCmd MConnID MData MDataInfo MReqInfo MReqLast MThreadID	cmdaccept & addr cmdaccept & addrspace cmdaccept & atomiclength cmdaccept & blockheight cmdaccept & blockstride cmdaccept & burstlength cmdaccept & burstprecise cmdaccept & burstseq cmdaccept & burstsinglereq cmdaccept & byteen cmdaccept cmdaccept & connid cmdaccept & mdata & !datahandshake cmdaccept & mdatainfo & !datahandshake cmdaccept & reqinfo cmdaccept & reqlast cmdaccept & threads > 1
request_value_MCmd_<command> BCST RDL WRC RD RDEX WR WRNP	!broadcast_enable !rdlwr_enable !rdlwr_enable !read_enable !readex_enable !write_enable !writenonpost_enable
request_value_MAddr_word_aligned	addr
request_value_<signal>_0x0 MAAtomicLength MBurstLength MBlockHeight MBlockStride	atomiclength burstlength blockheight blockstride

Name	Activation Parameters
request_value_MBurstSeq_<sequence> BLCK DLFT1 DLFT2 INCR STRM UNKN WRAP XOR	burstlength & !burstseq_blk_enable burstlength & !burstseq_dflt1_enable burstlength & !burstseq_dflt2_enable burstlength & !burstseq_incr_enable burstlength & !burstseq_strm_enable burstlength & !burstseq_unkn_enable burstlength & !burstseq_wrap_enable burstlength & !burstseq_xor_enable
request_value_MByteEn_force_aligned	byteen & force_aligned
request_value_MThreadID	threads > 1
datahs_exact_SDataThreadBusy	datahandshake & sdatathreadbusy & sdatathreadbusy_exact & ~sdatathreadbusy_pipelined
datahs_pipelined_SDataThreadBusy	datahandshake & sdatathreadbusy & sdatathreadbusy_exact & sdatathreadbusy_pipelined
datahs_hold_<signal> MData MDataByteEn MDataInfo MDataThreadID MDataValid MDataLast	dataaccept & mdata & datahandshake mdatabyteen & dataaccept & datahandshake dataaccept & mdatainfo & datahandshake datahandshake & threads > 1 & dataaccept dataaccept & datahandshake datalast & dataaccept & datahandshake
datahs_value_MDataByteEn_force_aligned	Mdatabyteen & datahandshake & force_aligned
datahs_value_MDataThreadID	datahandshake & threads > 1
response_exact_MThreadBusy	resp & mthreadbusy & mthreadbusy_exact & ~mthreadbusy_pipelined
response_pipelined_MThreadBusy	resp & mthreadbusy & mthreadbusy_exact & mthreadbusy_pipelined
response_hold_<signal> SData SDataInfo SResp SRespInfo SRespLast SThreadID	respaccept & sdata & resp & (read_enable readex_enable rdlwrc_enable) respaccept & sdatainfo respaccept & resp respaccept & resp & respinfo respaccept & resp & resplast respaccept & resp & threads > 1
response_value_SResp_FAIL_without_WRC	resp & rdlwrc_enable
response_value_SThreadID	resp & threads > 1
request_hold_MTagInOrder	cmdaccept & taginorder & tags > 1
response_hold_STagInOrder	resp & respaccept & taginorder & tags > 1
request_hold_MTagID_when_MTagInOrder_zero	cmdaccept & tags > 1
datahs_hold_MDataTagID_when_MTagInOrder_zero	datahandshake & dataaccept & tags > 1
response_hold_STagID_when_STagInOrder_zero	resp & respaccept & tags > 1

Name	Activation Parameters
request_value_MTagID_when_MTagInOrder_zero	tags > 1
datahs_value_MTagID_when_MTagInOrder_zero	datahandshake & tags > 1
response_value_STagID_when_STagInOrder_zero	resp & tags > 1
datahs_order_MDataTagID_when_MTagInOrder_zero	burstlength & datahandshake & tags > 1
response_reorder_STagID_tag_interleave_size	burstsinglereq & resp & tags > 1
response_reorder_STagID_overlapping_addresses	resp & tags > 1 & (addr addrspace byteen)

Table 44 Dataflow Burst Checks

Name	Activation Parameters
burst_hold_MBurstLength_precise	burstlength
burst_hold_<signal> MAddrSpace MAtomicLength MBurstPrecise MBurstSeq MBurstSingleReq MCmd MConnID MReqInfo SRespInfo	burstlength & addrspace burstlength & atomiclength burstlength & burstprecise burstseq & burstlength burstlength & burstsinglereq burstlength burstlength & connid burstlength & reqinfo burstlength & respinfo
burst_hold_<signal>_BLCK MBlockHeight MBlockStride	burstlength & burstseq_blk_enable & burstseq & blockheight burstlength & burstseq_blk_enable & burstseq & blockstride
burst_hold_<signal>_STRM MByteEn MDataByteEn	burstlength & burstseq_strm_enable & byteen & burstseq burstlength & burstseq_strm_enable & datahandshake & mdatabyteen & burstseq
burst__phase_order_reqdata_together	reqdata_together & datahandshake
burst_sequence_MAddr_BLCK	burstlength & addr & burstseq_blk_enable & burstseq
burst_sequence_MAddr_INCR	burstlength & addr & burstseq_incr_enable & burstseq
burst_sequence_MAddr_STRM	burstlength & addr & burstseq_strm_enable & burstseq
burst_sequence_MAddr_WRAP	burstlength & burstseq_wrap_enable & addr & burstseq
burst_sequence_MAddr_XOR	burstlength & burstseq_xor_enable & addr & burstseq

Name	Activation Parameters
burst_value_<signal>_<sequence> MByteEn STRM	burstlength & burstseq_strm_enable & byteen & burstseq
MDataByteEn STRM	burstseq & burstseq_strm_enable & mdatabyteen & datahandshake
MByteEn DFLT2	burstlength & burstseq_dflt2_enable & byteen & burstseq
MDataByteEn DFLT2	burstseq & burstseq_dflt2_enable & mdatabyteen & datahandshake
burst_value_MAddr_INCR_burst_aligned	burstlength & burstseq_incr_enable & burst_aligned & burstseq
burst_value_MAddr_<sequence>_no_wrap INCR	burstlength & burstseq_incr_enable & addr
BLCK	burstlength & burstseq_blk_enable & addr
burst_value_MBurstLength_<sequence> WRAP	burstlength & burstseq & burstseq_wrap_enable
XOR	burstlength & burstseq & burstseq_xor_enable
burst_value_MBurstLength_INCR_burst_aligned	burstlength & burstseq_incr_enable & burst_aligned & burstseq
burst_value_MBurstPrecise_<sequence> WRAP	burstprecise & burstseq_wrap_enable & burstlength & burstseq
XOR	burstprecise & burstseq_xor_enable & burstlength & burstseq
BLCK	burstprecise & burstseq_blk_enable & burstlength & burstseq
burst_value_MBurstPrecise_INCR_burst_aligned	burstaligned & burstprecise & burstseq_incr_enable & burstseq
burst_value_MBurstPrecise_SRMD	burstprecise & burstsinglereq
burst_value_MBurstSeq_UNKN_SRMD	burstsinglereq & burstreq & burstseq_unkn_enable
burst_value_MCmd_<command> RDEX	burstlength & readex_enable
RDL	burstlength & rdlwrc_enable
WRC	burstlength & rdlwrc_enable
burst_value_MReqLast_MRMD	reqlast
burst_value_MReqLast_SRMD	Reqlast & burstsinglereq
burst_value_MReqRowLast_MRMD	mreqlast & mreqrowlast
burst_value_MReqRowLast_SRMD	mreqlast & mreqrowlast
burst_value_MDataLast_MRMD	datalast & mdata
burst_value_MDataLast_SRMD	datalast & mdata
burst_value_MDataRowLast_MRMD	mdatalast & mdatarowlast
burst_value_MDataRowLast_SRMD	mdatalast & mdatarowlast
burst_value_SRespLast_MRMD	resplast & resp
burst_value_SRespLast_SRMD	resplast & resp & burstsinglereq

Name	Activation Parameters
burst_value_SRespRowLast_MRMD	srespplast & sresprowlast
burst_value_SRespRowLast_SRMD	srespplast & sresprowlast
burst_hold_MTagID_when_MTagInOrder_zero	burtstlength & tags > 1
burst_hold_MTagInOrder	burstlength & tags > 1 & taginorder

Table 45 Dataflow Transfer Checks

Name	Activation Parameters
transfer_phase_order_datahs_before_request_begin	datahandshake
transfer_phase_order_datahs_before_request_end	datahandshake
transfer_phase_order_response_before_request_begin	resp
transfer_phase_order_response_before_request_end	resp
transfer_phase_order_response_before_datahs_begin	resp & datahandshake
transfer_phase_order_response_before_datahs_end	resp & datahandshake
transfer_phase_order_response_before_last_datahs_begin_SRMD_wr	resp & datahandshake & burstsinglereq
transfer_phase_order_response_before_last_datahs_end_SRMD_wr	resp & datahandshake & burstsinglereq

Table 46 Dataflow ReadEx Checks

Name	Activation Parameters
rdex_hold_<signal>	
MAddr	readex_enable & addr
MAddrSpace	readex_enable & addrspace
MByteEn	readex_enable & byteen
MDataByteEn)	readex_enable & mdatabyteen & datahandshake
rdex_lock_release_no_WR/WRNP	readex_enable
rdex_lock_release_no_burst_allowed	burstlength & readex_enable

Table 47 Sideband Checks

Name	Activation Parameters
signal_valid_<signal> MReset_n SReset_n	mreset sreset
signal_valid_<signal> when_reset_inactive ControlBusy ControlWr MError SError SInterrupt StatusBusy StatusRd	controlbusy controlwr merror serror interrupt statusbusy statusrd
signal_hold_<signal>_16_cycles MReset_n SReset_n	mreset sreset
signal_hold_Control_after_reset	control
signal_hold_Control_2_cycles	control
signal_hold_Control_ControlBusy_active	controlbusy
signal_hold_ControlWr_after_reset	controlwr
signal_value_ControlWr_Control_transitioned	control & controlwr
signal_value_ControlWr_ControlBusy_active	controlbusy & controlwr
signal_hold_ControlWr_2_cycle	controlwr
signal_value_ControlBusy	controlwr & controlbusy
signal_hold_StatusRd_2_cycles	statusrd
signal_value_StatusRd_StatusBusy_active	statusrd & statusbusy

Compliance Checks

Dataflow Signals Checks

1.1.1 signal_valid_<signal>_when_reset_inactive

When reset is inactive, the following signals should never have an X or Z value on the rising edge of the OCP clock:

MCmd	MDataValid	MThreadBusy
SDataThreadBusy	SResp	SThreadBusy

Protocol hierarchy	Reset activity
Signal group	Dataflow
Critical signals	MCmd, MDataValid, MThreadBusy, SDataThreadBusy, SResp, SThreadBusy
Assertion type	X, Z
Reference	“Reset” on page 44

1.1.2 request_valid_<signal>

The following signals should never have an X or Z value on the rising edge of the OCP clock during a request phase:

MAddr	MAddrSpace	MAtomicLength
MBurstLength	MBurstPrecise	MBurstSeq
MBurstSingleReq	MByteEn	MConnID
MReqLast	MThreadID	SCmdAccept
MBlockHeight	MBlockStride	MReqRowLast

If datahandshake=1 and mdatabyteen=1 then MByteEn can be invalid for write accesses during the request phase

Protocol hierarchy	Request phase
Signal group	Dataflow
Critical signals	MAddr, MAddrSpace, MAtomicLength, MBurstLength, MBurstPrecise, MBurstSeq, MBurstSingleReq, MByteEn, MConnID, MReqLast, MThreadID, SCmdAccept, MBlockHeight, MBlockStride, MReqRowLast
Assertion type	X, Z
References	“Reset” on page 44 “Request Phase” on page 145

1.1.3 datahs_valid_<signal>

The following signals should never have an X or Z value on the rising edge of the OCP clock during a datahandshake phase:

MDataByteEn MDataLast MDataThreadID SDataAccept

Protocol hierarchy	Datahandshake
Signal group	Dataflow
Critical signals	MDataByteEn, MDataLast, MDataThreadID, SDataAccept
Assertion type	X, Z
Reference	“Reset” on page 44 “Datahandshake Phase” on page 147

1.1.4 response_valid_<signal>

The following signals should never have an X or Z value on the rising edge of the OCP clock during a response phase:

MRespAccept SRespLast SThreadID

Protocol hierarchy	Response
Signal group	Dataflow
Critical signals	MRespAccept, SRespLast, SThreadID
Assertion type	X, Z
Reference	“Reset” on page 44 “Response Phase” on page 146

1.1.5 request_valid_MTagInOrder

MTagInOrder should not be X/Z during the request phase.

Protocol hierarchy	Request
Signal group	Dataflow - tag extensions
Critical signals	MTagInOrder
Assertion type	X, Z
Reference	“Tags” on page 162

1.1.6 response_valid_STagInOrder

STagInOrder should not be X/Z during the response phase.

Protocol hierarchy	Response
Signal group	Dataflow - tag extensions
Critical signals	STagInOrder
Assertion type	X, Z
Reference	"Tags" on page 162

1.1.7 request_valid_MTagID_when_MTagInOrder_zero

If MTagInOrder is 0 during the request phase, MTagID should not be X/Z during the request phase.

Protocol hierarchy	Response
Signal group	Dataflow - tag extensions
Critical signals	MTagID, MTagInOrder
Assertion type	X, Z
Reference	"Tags" on page 162

1.1.8 datahs_valid_MTagID_when_MTagInOrder_zero

If datahandshake is active and MTagInOrder is 0 (during the request phase), MDataTagID should not be X/Z during the datahandshake phase.

Protocol hierarchy	Datahandshake
Signal group	Dataflow - tag extensions
Critical signals	MDataTagID, MTagInOrder
Assertion type	X, Z
Reference	"Tags" on page 162

1.1.9 response_valid_STagID_when_STagInOrder_zero

If STagInOrder is 0 during the request phase, STagID should not be X/Z during the response phase.

Protocol hierarchy	Response
Signal group	Dataflow - tag extensions
Critical signals	STagID, STagInOrder
Assertion type	X, Z
Reference	"Tags" on page 162

DataFlow Phase Checks

1.2.1 request_exact_SThreadBusy

If `sthreadbusy_exact = 1` and `sthreadbusy_pipelined = 0`, when a given slave thread is busy, the master must stay idle on this thread.

Protocol hierarchy	Request
Signal group	Dataflow - thread extensions
Critical signals	MCmd
Assertion type	Value
References	“Ungrouped Signals” on page 42 “Threads” on page 163

1.2.2 request_pipelined_SThreadBusy

If `sthreadbusy_exact = 1` and `sthreadbusy_pipelined = 1`, and an SThreadbusy bit was set to 1 in the prior cycle, the master cannot present a request on a thread in the current cycle.

Protocol hierarchy	Request
Signal group	Dataflow - thread extensions
Critical signals	MCmd
Assertion type	Value
Reference	“Ungrouped Signals” on page 42

1.2.3 request_hold_<signal>

Once a request phase has begun, the following signals may not change their value until the OCP slave has accepted the request.

Basic Signals	Burst Extensions
MAddr	MAtomicLength
MCmd	MBurstLength
MData	MBurstPrecise
	MBurstSeq
	MBurstSingleReq
Simple Extensions	MReqLast
MAddrSpace	Thread Extensions
MByteEn	MConnID
MDataInfo	MThreadID
MReqInfo	MBlockHeight
	MBlockStride
	MReqRowLast

The following exceptions apply:

1. If datahandshake=1 and mdatabyteen=1 then MByteEn can change for write accesses during the request phase.
2. For read requests the MData and MDataInfo fields can change during the request phase.
3. For write requests the SData and SDataInfo fields can change during the response phase.
4. Non-enabled data bytes in MData and bits in MDataInfo fields can change during the request and datahandshake phases.
5. Non-enabled data bytes in SData and bits in SDataInfo fields can change during the response phase.
6. MDataByteEn can change during read-type transfers.
7. MTagID can change if MTagInOrder is asserted, and MDataTagID can change for the corresponding datahandshake phase.
8. STagID can change if STagInOrder is asserted.

Protocol hierarchy	Request
Signal group	Dataflow
Critical signals	MAddr, MCmd, MData, MAddrSpace, MByteEn, MDataInfo, MReqInfo, MAtomicLength, MBurstLength, MBurstPrecise, MBurstSeq, MBurstSingleReq, MReqLast, MConnID, MThreadID, MBlockHeight, MBlockStride, MReqRowLast
Assertion type	Hold
References	“Request Phase” on page 145

1.2.4 request_value_MCmd_<command>

The following <Command> is illegal if the corresponding <Parameter> is set to 0.

Command	Parameter
BCST	broadcast_enable
RD	read_enable
RDEX	readex_enable
RDL	rdlwrc_enable
WR	write_enable
WRC	rdlwrc_enable
WRNP	writenonpost_enable

Protocol hierarchy	Request
Signal group	Dataflow - basic signals
Critical signals	MCmd
Assertion type	Value
Reference	“Optional Commands” on page 55

1.2.5 request_value_MAddr_word_aligned

Signal MAddr must be OCP word aligned as follows:

if data_width = 16 then MAddr[0] = 0
if data_width = 32 then MAddr[1:0] = 0
if data_width = 64 then MAddr[2:0] = 0
if data_width = 128 then MAddr[3:0] = 0

Protocol hierarchy	Request
Signal group	Dataflow - basic signals
Critical signals	MAddr
Assertion type	Value
Reference	“Basic Signals” on page 14

1.2.6 request_value_<signal>_0x0

During a request phase, the <signal> must not be zero.

Protocol hierarchy	Request
Signal group	Dataflow - burst extensions
Critical signals	MAtomicLength, MBurstLength, MBlockHeight, MBlockStride
Assertion type	Value
References	“Burst Extensions” on page 19 Footnotes on page 32

1.2.7 request_value_MBurstSeq_<sequence>

The following <burst type> is illegal if its corresponding <parameter > is set to 0.

Burst type	Parameter
BLCK	burstseq_blk_enable
DLFT1	burstseq_dflt1_enable
DLFT2	burstseq_dflt2_enable
INCR	burstseq_incr_enable
STRM	burstseq_strm_enable

UNKN	burstseq_unkn_enable
WRAP	burstseq_wrap_enable
XOR	burstseq_xor_enable

Protocol hierarchy	Request
Signal group	Dataflow - burst extensions
Critical signals	MBurstSeq
Assertion type	Value
References	“Optional Burst Sequences” on page 55

1.2.8 request_value_MByteEn_force_aligned

If `force_aligned=1`, the byte enable values during a request phase are restricted to the following patterns for `data_width ≥ 32`:

`data_width=32` and MByteEn has one of the following values:

```
0001
0010
0100
1000
0011
1100
1111
0000
```

`data_width=64` and MByteEn has one of the following values:

```
00000001
00000010
00000100
00001000
00010000
00100000
01000000
10000000
00000011
00001100
00110000
11000000
00001111
11110000
11111111
00000000
```

`data_width=128` and MByteEn has one of the following values:

```
0000000000000001
0000000000000010
0000000000000100
0000000000001000
```

```
00000000000010000
00000000000010000
000000000000100000
00000000010000000
00000000100000000
00000001000000000
00000010000000000
00000100000000000
00001000000000000
00010000000000000
00100000000000000
01000000000000000
00000000000000011
00000000000001100
00000000000110000
00000000110000000
00000011000000000
00000110000000000
00110000000000000
11000000000000000
0000000000001111
00000000111100000
00001111000000000
11110000000000000
0000000011111111
11111111000000000
1111111111111111
00000000000000000
```

If `datahandshake=1` and `mdatabyteen=1` then `MByteEn` can change for write accesses during the request phase.

Protocol hierarchy	Request
Signal group	Dataflow - simple extensions
Critical signals	MByteEn
Assertion type	Value
References	“Byte Enable Patterns” on page 56

1.2.9 request_value_MThreadID

MThreadID value is always < threads.

Protocol hierarchy	Request
Signal group	Dataflow - thread extensions
Critical signals	MThreadID
Assertion type	Value
Reference	“Thread Extensions” on page 23

1.2.10 datahs_exact_SDataThreadBusy

If `sdatathreadbusy_exact = 1` and `sdatathreadbusy_pipelined = 0`, when a given slave data thread is busy, the master must not present a data phase on this thread.

Protocol hierarchy	Datahandshake
Signal group	Dataflow - thread extensions
Critical signals	MDataValid
Assertion type	Value
Reference	“Ungrouped Signals” on page 42

1.2.11 datahs_pipelined_SDataThreadBusy

If `sdatathreadbusy_exact = 1` and `sdatathreadbusy_pipelined = 1`, and an `SDataThreadbusy` bit was set to 1 in the prior cycle, the master cannot present a datahandshake on a thread in the current cycle.

Protocol hierarchy	Datahandshake
Signal group	Dataflow - thread extensions
Critical signals	MDataValid
Assertion type	Value
Reference	“Ungrouped Signals” on page 42

1.2.12 datahs_hold_<signal>

Once a datahandshake phase has begun, the following signals may not change their value until the OCP slave has accepted the data.

Basic Signals	Burst Extensions
MData	MDataLast
MDataValid	
Simple Extensions	Thread Extensions
MDataByteEn	MDataThreadID
MDataInfo	

Protocol hierarchy	Datahandshake
Signal group	Dataflow
Critical signals	MData, MDataValid, MDataByteEn, MDataInfo, MDataLast, MDataThreadIDd
Assertion type	Hold
Reference	“Datahandshake Phase” on page 147

1.2.13 datahs_value_MDataByteEn_force_aligned

If force_aligned=1, the data byte enable values during a datahandshake phase are restricted to the following patterns for data_width \geq 32:

data_width=32 and MDataByteEn has one of the following values:

0001
0010
0100
1000
0011
1100
1111
0000

data_width=64 and MDataByteEn has one of the following values:

00000001
00000010
00000100
00001000
00010000
00100000
01000000
10000000
00000011
00001100
00110000
11000000
00001111
11110000
11111111
00000000

data_width=128 and MDataByteEn has one of the following values:

0000000000000001
0000000000000010
0000000000000100
0000000000001000
0000000000010000
0000000001000000
0000000010000000
0000000100000000
0000001000000000
0000010000000000
0000100000000000
0001000000000000
0010000000000000
0100000000000000
0000000000000011
00000000000001100
00000000000110000

```

00000000011000000
00000001100000000
0000110000000000
0011000000000000
1100000000000000
0000000000000111
0000000011110000
0000111100000000
1111000000000000
0000000011111111
1111111100000000
1111111111111111
0000000000000000

```

Protocol hierarchy	Datahandshake
Signal group	Dataflow - simple extensions
Critical signals	MDataByteEn
Assertion type	Value
Reference	“Byte Enable Patterns” on page 56

1.2.14 datahs_value_MDataThreadID

MDataThreadID value must be < threads.

Protocol hierarchy	Datahandshake
Signal group	Dataflow - thread extensions
Critical signals	MDataThreadID
Assertion type	Value
Reference	“Thread Extensions” on page 23

1.2.15 response_exact_MThreadBusy

If mthreadbusy_exact = 1 and mthreadbusy_pipelined = 0, when a given master thread is busy, the slave must not present a response on that thread.

Protocol hierarchy	Response
Signal group	Dataflow - thread extensions
Critical signals	SResp
Assertion type	Value
Reference	“Ungrouped Signals” on page 42

1.2.16 response_pipelined_MThreadBusy

If `sthreadbusy_exact = 1` and `sthreadbusy_pipelined = 1`, and an `MThreadbusy` bit was set to 1 in the prior cycle, the slave cannot present a response on a thread in the current cycle.

Protocol hierarchy	Response
Signal group	Dataflow - thread extensions
Critical signals	SResp
Assertion type	Value
Reference	“Ungrouped Signals” on page 42

1.2.17 response_hold_<signal>

Once a response phase has begun, the following signals may not change their value until the master has accepted the response.

Basic Signals	Burst Extensions
SData	SRespLast
SResp	
Simple Extensions	Thread Extensions
SDataInfo	SThreadID
SRespInfo	

Protocol hierarchy	Response
Signal group	Dataflow
Critical signals	SData, SResp, SDataInfo, SRespInfo, SRespLast, SThreadID
Assertion type	Hold
Reference	“Response Phase” on page 146

1.2.18 response_value_SResp_FAIL_without_WRC

The FAIL response can occur only on a WRC request.

Protocol hierarchy	Response
Signal group	Dataflow - basic signals
Critical signals	SResp
Assertion type	Value
References	“Transfer Effects” on page 45

1.2.19 response_value_SThreadID

SThreadID value must be < threads.

Protocol hierarchy	Response
Signal group	Dataflow - thread extensions
Critical signals	SThreadID
Assertion type	Value
Reference	“Thread Extensions” on page 23

1.2.20 request_hold_MTagInOrder

If taginorder = 1, the MTagInOrder signal cannot change until accepted by the OCP slave (SCmdAccept = 1).

Protocol hierarchy	Request
Signal group	Dataflow - tag extensions
Critical signals	MTagInOrder
Assertion type	Hold
Reference	“Request Phase” on page 145

1.2.21 response_hold_STagInOrder

If taginorder = 1, the STagInOrder signal cannot change until accepted by the master (MRespAccept = 1).

Protocol hierarchy	Response
Signal group	Dataflow - tag extensions
Critical signals	STagInOrder
Assertion type	Hold
Reference	“Response Phase” on page 146

1.2.22 request_hold_MTagID_when_MTagInOrder_zero

If tags > 1, the MTagID signal cannot change until accepted by the OCP slave (SCmdAccept = 1).

Protocol hierarchy	Request
Signal group	Dataflow - tag extensions
Critical signals	MTagID, MTagInOrder
Assertion type	Hold
Reference	“Request Phase” on page 145

1.2.23 datahs_hold_MDataTagID_when_MTagInOrder_zero

When tags > 1, during a datahandshake phase corresponding to a non in-order request phase (MTagInOrder = 0), the MDataTagID signal cannot change value until accepted by the OCP slave (SDataAccept = 1).

Protocol hierarchy	Datahandshake
Signal group	Dataflow - tag extensions
Critical signals	MDataTagID, MTagInOrder
Assertion type	Hold
Reference	“Datahandshake Phase” on page 147

1.2.24 response_hold_STagID_when_STagInOrder_zero

If tags > 1, the STagID signal cannot change until it is accepted by the master (MRespAccept = 1).

Protocol hierarchy	Response
Signal group	Dataflow - tag extensions
Critical signals	STagID, STagInOrder
Assertion type	Hold
Reference	“Response Phase” on page 146

1.2.25 request_value_MTagID_when_MTagInOrder_zero

The MTagID signal must always be < tags.

Protocol hierarchy	Request
Signal group	Dataflow - tag extensions
Critical signals	MTagID, MTagInOrder
Assertion type	Value
Reference	“Tag Extensions” on page 22

1.2.26 datahs_value_MTagID_when_MTagInOrder_zero

The MDataTagID signal must always be < tags.

Protocol hierarchy	Datahandshake
Signal group	Dataflow - tag extensions
Critical signals	MDataTagID, MTagInOrder
Assertion type	Value
Reference	“Tag Extensions” on page 22

1.2.27 response_value_STagID_when_STagInOrder_zero

The STagID signal must always be < tags.

Protocol hierarchy	Response
Signal group	Dataflow - tag extensions
Critical signals	STagID, STagInOrder
Assertion type	Value
Reference	“Tag Extensions” on page 22

1.2.28 datahs_order_MDataTagID_when_MTagInOrder_zero

When datahandshake = 1, for tagged write transactions, the datahandshake phase must observe the same order as the request phase.

Protocol hierarchy	Datahandshake
Signal group	Dataflow - tag extensions
Critical signals	MDataTagID, (MTagInOrder
Assertion type	Data_order
Reference	“Ordering Restrictions” on page 53

1.2.29 response_reorder_STagID_tag_interleave_size

When tags > 1 and tag_interleave_size > 0 the slave must ensure that responses associated with packing burst sequences stay together up to the tag_interleave_size. When tags > 1 and tag_interleave_size == 0 no interleaving of responses between any packing burst sequences with different tags is allowed.

Protocol hierarchy	Response
Signal group	Dataflow - tag extensions
Critical signals	STagID
Assertion type	Reorder
References	“Ordering Restrictions” on page 53 “Burst Interleaving with Tags” on page 58

1.2.30 response_reorder_STagID_overlapping_addresses

Responses to requests that target overlapping addresses (as determined by MAddrSpace, MAddr and MByteEn) must not be re-ordered with respect to another. Note that this property does not need take into account the value of MTagInOrder.

Protocol hierarchy	Response
Signal group	Dataflow - tag extensions
Critical signals	STagID
Assertion type	Reorder
References	“Tags” on page 10 “Ordering Restrictions” on page 53

Dataflow Burst Checks

1.3.1 burst_hold_MBurstLength_precise

For precise bursts, MBurstLength must hold its value during all request phases of the entire burst.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MBurstLength
Assertion type	Hold
References	“Constant Fields in Bursts” on page 51

1.3.2 burst_hold_<signal>

The following signals must hold the same value on all request phases of the entire burst:

MAddrSpace	MBurstSingleReq
MAtomicLength	MCmd
MBurstPrecise	MConnID
MBurstSeq	MReqInfo

The hold requirements for SRespInfo in a burst are different for the 2.0 versus 2.2 specifications.

OCP 2.0 page 44 states that:

SRespInfo must be held steady by the slave for every transfer in a burst.

OCP 2.2 page 51 states that:

If possible, slaves should hold SRespInfo steady for every transfer in a burst

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MAddrSpace, MAtomicLength, MBurstPrecise, MBurstSeq, MBurstSingleReq, MCmd, MConnID, MReqInfo, (SRespInfo)
Assertion type	Hold
Reference	“Constant Fields in Bursts” on page 51

1.3.3 burst_hold_<signal>_BLCK

<signal> must hold for BLCK bursts. Applicable to MBlockHeight and MBlockStride signals.

Protocol hierarchy	Burst
Signal group	Dataflow - simple extensions
Critical signals	MBlockHeight, MBlockStride
Assertion type	Hold
Reference	“Constant Fields in Bursts” on page 51

1.3.4 burst_hold_<signal>_STRM

For STRM bursts, MByteEn / MDataByteEn must hold the same value on all request / datahandshake phases of the entire burst.

Protocol hierarchy	Burst
Signal group	Dataflow - simple extensions
Critical signals	MByteEn, MDataByteEn
Assertion type	Hold
References	“Byte Enable Restrictions” on page 50

1.3.5 burst__phase_order_reqdata_together

For single request multiple data bursts, if reqdata_together = 1, the master must present the request and first write data in the same cycle, and the slave must accept the request and the first write data in the same cycle.

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	MCmd, MDatavalid
Assertion type	Ordering
Reference	“Phase Options” on page 59

1.3.6 burst_sequence_MAddr_BLK

Within a block burst, the address begins with the provided address and proceeds through a set of MBlockHeight subsequences, each of which follows the normal INCR burst sequence for MBurstLength transfers. The starting address of each subsequence should be the starting address of the prior subsequence plus MBurstStride.

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	MBlockHeight, MBurstLength, MBurstStride
Assertion type	Ordering
Reference	“Burst Address Sequences” on page 49

1.3.7 burst_sequence_MAddr_INCR

Within an INCR burst, the address increases for each new master request by the OCP word size.

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	MAddr
Assertion type	Ordering
Reference	Table 19 on page 49

1.3.8 burst_sequence_MAddr_STRM

Within a STRM burst, the address remains constant on all request phases of the burst.

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	MAddr
Assertion type	Ordering
Reference	Table 19 on page 49

1.3.9 burst_sequence_MAddr_WRAP

Within a WRAP burst, the address increases for each new master request by the OCP word size, and wraps on the burst length x OCP word size.

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	MAddr
Assertion type	Ordering
Reference	Table 19 on page 49

1.3.10 burst_sequence_MAddr_XOR

Within an XOR burst, the address increases for each new OCP master request as follows:

BASE

Is the lowest byte address in the burst, which must be aligned with the total burst size.

FIRST_OFFSET

Is the byte offset (from BASE) of the first transfer in the burst.

CURRENT_COUNT

Is the count of current transfer in the burst starting at 0.

WORD_SHIFT

Is the log2 of the OCP word size in bytes.

The current address of the transfer is $\text{BASE} \mid (\text{FIRST_OFFSET} \wedge (\text{CURRENT_COUNT} \ll \text{WORD_SHIFT}))$.

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	MAddr
Assertion type	Ordering
Reference	“Burst Address Sequences” on page 49

1.3.11 burst_value_<signal>_<sequence>

When `mdatabyteen = 0`, during STRM or DFLT2 bursts, `MByteEn` should never take the value 0.

When `mdatabyteen = 1`, during read-type STRM or DFLT2 bursts, `MByteEn` should never take the value 0.

When `mdatabyteen = 1`, during write-type STRM or DFLT2 bursts, `MDataByteEn` should never take value 0.

Protocol hierarchy	Burst
Signal group	Dataflow - simple extensions
Critical signals	MByteEn, MDataByteEn
Assertion type	Value
Reference	“Byte Enable Restrictions” on page 50

1.3.12 burst_value_MAddr_INCR_burst_aligned

When `burst_aligned=1`, the first burst request of an INCR burst must have its address aligned. The equation below indicates which MAddr bits must be 0.

Equation

$$\text{MAddr}[(\text{size}-1)+\text{BL}:0] = 0$$

Where:

$$\begin{array}{ll} \text{size} = \text{ceil}(\log_2(\text{bytes}(\text{data_width}))) & \text{for } \text{data_width} > 1 \text{ byte} \\ \text{BL} = \log_2(\text{MBurstLength}) & \text{for } \text{MBurstLength} \geq 1 \end{array}$$

Example

For an interface with `data_width=32`, `size=2` and:

$$\text{MBurstLength} = 2: \text{MAddr}[2:0] = 0$$

$$\text{MBurstLength} = 4: \text{MAddr}[3:0] = 0$$

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	MAddr
Assertion type	Value
References	“Burst Alignment” on page 56

1.3.13 burst_value_MAddr_<sequence>_no_wrap

An INCR or BLCK burst can never cross the address space boundary.

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	MAddr
Assertion type	Value
Reference	“Burst Address Sequences” on page 49

1.3.14 burst_value_MBurstLength_<sequence>

The length of a WRAP or XOR burst must be a power of two.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MBurstLength
Assertion type	Value
Reference	“Burst Address Sequences” on page 49

1.3.15 burst_value_MBurstLength_INCR_burst_aligned

When `burst_aligned = 1`, the length of an INCR burst must be a power of two.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MBurstLength
Assertion type	Value
References	“Burst Alignment” on page 56

1.3.16 burst_value_MBurstPrecise_<sequence>

BLCK, WRAP and XOR bursts can be issued only as precise bursts.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MBurstPrecise
Assertion type	Value
References	“Burst Address Sequences” on page 49

1.3.17 burst_value_MBurstPrecise_INCR_burst_aligned

When `burst_aligned = 1`, INCR bursts can be issued only as precise bursts.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MBurstPrecise
Assertion type	Value
Reference	“Burst Address Sequences” on page 49

1.3.18 burst_value_MBurstPrecise_SRMD

Single request multiple data transfers can be issued only as precise bursts.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MBurstPrecise
Assertion type	Value
Reference	“Single Request / Multiple Data Bursts (Packets)” on page 51

1.3.19 burst_value_MBurstSeq_UNKN_SRMD

An unknown burst sequence (value UNKN) is illegal during a single request multiple data transfer.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MBurstSeq
Assertion type	Value
Reference	“Phases in a Transfer” on page 40

1.3.20 burst_value_MCcmd_<command>

The RDEX, RDL, and WRC commands cannot be part of a burst.

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	MCmd
Assertion type	Value
References	“Burst Definition” on page 48

1.3.21 burst_value_MReqLast_MRMD

The signal MReqLast must be 0 for all request phases of a MRMD burst, except on the last one when it must be 1. For BLCK bursts the last request phase is the last request phase of the last MBlockHeight subsequence.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MReqLast
Assertion type	Value
References	“MReqLast, MDataLast, SRespLast” on page 52

1.3.22 burst_value_MReqLast_SRMD

The signal MReqLast must be 1 for any single request (SRMD being active or not).

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MReqLast
Assertion type	Value
References	“MReqLast, MDataLast, SRespLast” on page 52

1.3.23 burst_value_MReqRowLast_MRMD

For BLCK bursts the signal MReqRowLast must be 0 for all request phases other than the last phases in each row, when it must be 1. For non-BLCK bursts the signal MReqRowLast must be 0 for all request phases of a MRMD

burst, except on the last one when it must be 1. When `mreqlast` and `mreqrowlast` are both enabled, whenever `MReqLast` is asserted `MReqRowLast` must also be asserted.

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	<code>MReqRowLast</code>
Assertion type	Ordering
Reference	" <code>MReqLast</code> , <code>MDataLast</code> , <code>SRespLast</code> " on page 52

1.3.24 `burst_value_MReqRowLast_SRMD`

The signal `MReqRowLast` must be 1 for any single request (SRMD being active or not).

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	<code>MReqRowLast</code>
Assertion type	Ordering
Reference	" <code>MReqLast</code> , <code>MDataLast</code> , <code>SRespLast</code> " on page 52

1.3.25 `burst_value_MDataLast_MRMD`

The `MDataLast` signal must be 0 for all datahandshake phases in an MRMD burst, except on the last one when it must be 1. For BLCK bursts the last datahandshake phase is the last datahandshake phase of the last MBlock-Height subsequence.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	<code>MDataLast</code>
Assertion type	Value
References	" <code>MReqLast</code> , <code>MDataLast</code> , <code>SRespLast</code> " on page 52

1.3.26 burst_value_MDataLast_SRMD

The MDataLast signal must be 0 for all datahandshake phases of an SRMD burst, except on the last one when it must be 1. For BLCK bursts the last datahandshake phase is the last datahandshake phase of the last MBlock-Height subsequence.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MDataLast
Assertion type	Value
References	"MReqLast, MDataLast, SRespLast" on page 52

1.3.27 burst_value_MDataRowLast_MRMD

For BLCK bursts the signal MDataRowLast must be 0 for all datahandshake phases other than the last phases in each row, when it must be 1. For non-BLCK bursts the signal MDataRowLast must be 0 for all datahandshake phases of a MRMD burst, except on the last one when it must be 1. If mdatalast and mdatarowlast are both enabled, whenever MDataLast is asserted MDataRowLast must also be asserted.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MDataRowLast, MDataLast
Assertion type	Value
References	"MReqLast, MDataLast, SRespLast" on page 52

1.3.28 burst_value_MDataRowLast_SRMD

For BLCK bursts the signal MDataRowLast must be 0 for all datahandshake phases other than the last phases in each row, when it must be 1. For non-BLCK bursts the signal MDataRowLast must be 0 for all datahandshake phases of a SRMD burst, except on the last one when it must be 1. When mdatalast and mdatarowlast are both enabled, whenever MDataLast is asserted MDataRowLast must also be asserted.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MDataLast
Assertion type	Value
References	"MReqLast, MDataLast, SRespLast" on page 52

1.3.29 burst_value_SRespLast_MRMD

The signal SRespLast must be 0 for all response phases of an MRMD burst, except on the last one where it must be 1. For BLCK bursts the last response phase is the last response phase of the last MBlockHeight subsequence.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	SRespLast
Assertion type	Value
Reference	“MReqLast, MDataLast, SRespLast” on page 52

1.3.30 burst_value_SRespLast_SRMD

The signal SRespLast must be 1 for any single response (with SRMD active or not).

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	SRespLast
Assertion type	Value
Reference	“MReqLast, MDataLast, SRespLast” on page 52

1.3.31 burst_value_SRespRowLast_MRMD

For BLCK bursts the signal MRespRowLast must be 0 for all response phases other than the last phases in each row, when it must be 1. For non-BLCK bursts the signal MRespRowLast must be 0 for all response phases of a MRMD burst, except on the last one when it must be 1. If sresplast and sresprowlast are both enabled, whenever SRespLast is asserted SRespRowLast must also be asserted.

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MRespRowLast, SRespLast
Assertion type	Value
Reference	“MReqLast, MDataLast, SRespLast” on page 52

1.3.32 burst_value_SRespRowLast_SRMD

The signal MRespRowLast must be 1 for any single response (SRMD being active or not).

Protocol hierarchy	Burst
Signal group	Dataflow - burst extensions
Critical signals	MRespRowLast
Assertion type	Value
Reference	"MReqLast, MDataLast, SRespLast" on page 52

1.3.33 burst_hold_MTagID_when_MTagInOrder_zero

The MTagID signal must remain constant for all transfers of a burst when MTagInOrder is zero. The master cannot interleave requests (or datahandshake) phases with different tags within a transaction. This check should only focus on the request phase. The datahandshake phase is covered by phase property "datahs_order_MDataTagID_when_MTagInOrder_zero". This last property checks that the datahandshake phase observes the same order as the request phase.

Protocol hierarchy	Burst
Signal group	Dataflow - tag extensions
Critical signals	MTagID, (MTagInOrder)
Assertion type	Hold
Reference	"Ordering Restrictions" on page 53

1.3.34 burst_hold_MTagInOrder

The MTagInOrder signal must remain constant for all transfers of a burst.

Protocol hierarchy	Burst
Signal group	Dataflow - tag extensions
Critical signals	MTagInOrder
Assertion type	Hold
Reference	"Ordering Restrictions" on page 53

DataFlow Transfer Checks

1.4.1 transfer_phase_order_datahs_before_request_begin

For each thread, for each transaction tag, a datahandshake phase cannot begin before the associated request phase begins, but can begin in the same clock cycle.

Protocol hierarchy	Transfer
Signal group	Dataflow - basic signals
Critical signals	MDataValid, MCmd
Assertion type	Ordering
Reference	“Phase Ordering Within a Transfer” on page 41

1.4.2 transfer_phase_order_datahs_before_request_end

For each thread, for each transaction tag, a datahandshake phase cannot end before the associated request phase ends, but can end in the same clock cycle.

Protocol hierarchy	Transfer
Signal group	Dataflow - basic signals
Critical signals	MDataValid, MCmd
Assertion type	Ordering
Reference	“Phase Ordering Within a Transfer” on page 41

1.4.3 transfer_phase_order_response_before_request_begin

For each thread, for each transaction tag, a response phase cannot begin before the associated request phase begins, but can begin in the same clock cycle.

Protocol hierarchy	Transfer
Signal group	Dataflow - basic signals
Critical signals	MCmd, SResp
Assertion type	Ordering
References	“Phase Ordering Within a Transfer” on page 41

1.4.4 transfer_phase_order_response_before_request_end

For each thread, for each transaction tag, a response phase cannot end before the associated request phase ends, but can end in the same clock cycle.

Protocol hierarchy	Transfer
Signal group	Dataflow - basic signals
Critical signals	MCMD, SResp
Assertion type	Ordering
References	“Phase Ordering Within a Transfer” on page 41

1.4.5 transfer_phase_order_response_before_datahs_begin

For each thread, for each transaction tag, when datahandshake = 1, the response phase cannot begin before the associated datahandshake begins, but can begin in the same clock cycle.

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	MDataValid, SResp
Assertion type	Ordering
References	“Phase Ordering Within a Transfer” on page 41

1.4.6 transfer_phase_order_response_before_datahs_end

For each thread, for each transaction tag, when datahandshake = 1, the response phase cannot end before the associated datahandshake ends, but can end in the same clock cycle.

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	MDataValid, SResp
Assertion type	Ordering
References	“Phase Ordering Within a Transfer” on page 41

1.4.7 `transfer_phase_order_response_before_last_datahs_begin_SRMD_wr`

For each thread, for each transaction tag, with a write-type SRMD, the response phase cannot begin before the last datahandshake phase begins, but it can begin in the same clock cycle.

Protocol hierarchy	Transfer
Signal group	Dataflow - basic signals
Critical signals	MDataValid, SResp
Assertion type	Ordering
References	“Phase Ordering Within a Transfer” on page 41

1.4.8 `transfer_phase_order_response_before_last_datahs_end_SRMD_wr`

For each thread, for each transaction tag, with a write-type SRMD, the response phase cannot end before the last datahandshake phase ends, but it can end in the same clock cycle.

Protocol hierarchy	Burst
Signal group	Dataflow - basic signals
Critical signals	MDataValid, SResp
Assertion type	Ordering
Reference	“Phase Ordering Within a Transfer” on page 41

DataFlow ReadEx Checks

1.5.1 `rdex_hold_<signal>`

The unlocking command following a ReadEx must retain the same address and address space values

When `mdatabyteen = 0`, the unlocking command following a ReadEx must retain the same `MByteEn` value.

When `mdatabyteen = 1`, the unlocking command following a `ReadEx` must retain for `MDataByteEn` the value given to `MByteEn` during the `ReadEx` command. If `MByteEn` is absent, `MDataByteEn` must be all 1s.

Protocol hierarchy	<code>ReadEx</code>
Signal group	Dataflow - basic signals, simple extensions
Critical signals	<code>MAddr</code> , <code>MAddrSpace</code> , <code>MByteEn</code> , <code>MDataByteEn</code>
Assertion type	Hold
References	“Transfer Effects” on page 45 “Burst Definition” on page 48

1.5.2 `rdex_lock_release_no_WR/WRNP`

If a `ReadEx` is issued on an address on a particular thread, no other request with the same address can be issued on any other thread until the `ReadEx` is unlocked.

The command following the `ReadEx` on the same thread must be a write command (`WR` or `WRNP`). This command unlocks the `ReadEx`.

Protocol hierarchy	<code>ReadEx</code>
Signal group	Dataflow - basic signals
Critical signals	<code>MCmd</code>
Assertion type	Ordering
References	“Transfer Effects” on page 45 “Burst Definition” on page 48

1.5.3 `rdex_lock_release_no_burst_allowed`

The unlocking command following a `RDEX` must have `MBurstLength = 1`.

Protocol hierarchy	<code>ReadEx</code>
Signal group	Dataflow - basic signals
Critical signals	<code>MBurstLength</code>
Assertion type	Value
Reference	“Transfer Effects” on page 45 “Burst Definition” on page 48

Sideband Checks

1.6.1 signal_valid_<signal>

Signals MReset_n and SReset_n are never X or Z.

Protocol hierarchy	Reset activity
Signal group	Sideband - reset
Critical signals	MReset, SReset
Assertion type	X, Y
Reference	“Reset” on page 44 “Status and Control” on page 45

1.6.2 signal_valid_<signal> when_reset_inactive

When reset is inactive, the following signals should never have an X or Z value on the rising edge of the OCP clock:

ControlBusy	ControlWr	MError
SError	SInterrupt	
StatusBusy	StatusRd	

Protocol hierarchy	Reset activity
Signal group	Sideband - reset
Critical signals	ControlBusy, ControlWr, MError, SError, SInterrupt, StatusBusy, StatusRd
Assertion type	X, Y
Reference	“Status and Control” on page 45

1.6.3 signal_hold_<signal>_16_cycles

If they are active, signals MReset_n and SReset_n must stay active at least 16 consecutive cycles.

Protocol hierarchy	Reset activity
Signal group	Sideband - reset
Critical signals	MReset, SReset
Assertion type	Hold
References	“Reset” on page 44

1.6.4 signal_hold_Control_after_reset

The Control signal must be held steady the first cycle after reset is de-asserted.

Protocol hierarchy	Control
Signal group	Sideband - control
Critical signals	Control
Assertion type	Hold
References	“Status and Control” on page 45

1.6.5 signal_hold_Control_2_cycles

The Control signal must be held steady for a full cycle after the cycle in which it has transitioned.

Protocol hierarchy	Control
Signal group	Sideband - control
Critical signals	Control
Assertion type	Hold
References	“Status and Control” on page 45

1.6.6 signal_hold_Control_ControlBusy_active

If the ControlBusy signal was sampled active at the end of the previous cycle, the Control signal must not transition in the current cycle.

Protocol hierarchy	Control
Signal group	Sideband - control
Critical signals	Control
Assertion type	Value
Reference	“Status and Control” on page 45

1.6.7 signal_hold_ControlWr_after_reset

The ControlWr signal must not be asserted in the cycle following a reset.

Protocol hierarchy	Control
Signal group	Sideband - control
Critical signals	ControlWr
Assertion type	Hold
Reference	“Status and Control” on page 45

1.6.8 signal_value_ControlWr_Control_transitioned

If signal Control transitions in a cycle, signal ControlWr must be driven active on that cycle.

Protocol hierarchy	Control
Signal group	Sideband - control
Critical signals	ControlWr
Assertion type	Hold
References	“Status and Control” on page 45

1.6.9 signal_value_ControlWr_ControlBusy_active

The ControlWr signal must not be asserted if ControlBusy is active.

Protocol hierarchy	Control
Signal group	Sideband - control
Critical signals	ControlWr
Assertion type	Value
Reference	“Status and Control” on page 45

1.6.10 signal_hold_ControlWr_2_cycle

The ControlWr signal must not remain asserted for two consecutive cycles.s

Protocol hierarchy	Control
Signal group	Sideband - control
Critical signals	ControlWr
Assertion type	Hold
Reference	“Status and Control” on page 45

1.6.11 signal_value_ControlBusy

The ControlBusy signal can only be asserted in a cycle after the ControlWr signal is asserted or after the reset transitions to inactive.

Protocol hierarchy	Control
Signal group	Sideband - control
Critical signals	ControlBusy
Assertion type	Value
Reference	“Status and Control” on page 45

1.6.12 signal_hold_StatusRd_2_cycles

If the StatusRd signal was asserted in the previous cycle, it must not be asserted in the current cycle.

Protocol hierarchy	Status
Signal group	Sideband - status
Critical signals	StatusRd
Assertion type	Hold
References	“Status and Control” on page 45

1.6.13 signal_value_StatusRd_StatusBusy_active

The StatusRd signal must not be asserted while StatusBusy is asserted.

Protocol hierarchy	Status
Signal group	Sideband - status
Critical signals	StatusRd, StatusBusy
Assertion type	Value
Reference	“Status and Control” on page 45

15 Configuration Compliance Checks

The configuration checks listed in this chapter are extracted from the OCP Specification and are intended to serve as guidelines to verify an IP for OCP compliance. In all cases "Part I, Specification" is the definitive reference.

The configuration checks listed in this chapter are based on the "Specification" and "Guidelines" parts of this document and allow you to verify an IP/VIP for OCP compliance. In all cases, "Part I, Specification" is the definitive reference. Any references made to "Part II, Guidelines" are not definitive as Part I supersedes the guidelines.

The section describes the configuration checks needed for an OCP port. The names assigned to the configuration compliance checks have been created using the following template:

<hierarchy>_cfg_<critical_param>_<relationship>_<extra_details>

In which:

<hierarchy>: request, datahandshake, response, sideband, test,
 master_slave

<critical_param> : any OCP parameter that is impacted by the configuration
 check

<relationship> : (optional) enable, depends, match

<extra details> : a short additional explanation

The majority of the configuration checks involve an enable relationship. For these enable checks 'paramA_enable_paramB' implies that paramA is somehow enabled by paramB. In these situations the individual check descriptions provide details on the enabling relationship between the parameters.

Request Group

2.1.1 request_cfg_cmd_enable

One of the command enable parameters must be enabled. The critical parameters are: read_enable, readex_enable, write_enable, writenonpost_enable, broadcast_enable, or rdlwrc_enable.

Protocol hierarchy Request

Critical parameters read_enable, readex_enable, write_enable, writenonpost_enable, broadcast_enable, rdlwrc_enable

Reference “Basic Signals” on page 14

2.1.2 request_cfg_readex_enable_write_writenonpost

readex_enable can only be enabled if write_enable or writenonpost_enable is enabled.

Protocol hierarchy Request

Critical parameters readex_enable, write_enable, writenonpost_enable

Reference “Optional Commands” on page 55

2.1.3 request_cfg_addr_width_depends_data_width

data_width defines a minimum addr_width value that is based on the data bus byte width, and is defined as:

$$\text{min_addr_width} = \max(1, \text{ceil}(\log_2(\text{data_width}))-3))$$

Protocol hierarchy Request

Critical parameters addr_width, data_width

Reference MAddr on page 15

2.1.4 request_cfg_byteen_enable_mdata_sdata

byteen can only be enabled when either sdata or mdata is also enabled.

Protocol hierarchy Request

Critical parameters byteen, sdata, mdata

Reference “Simple Extensions” on page 16

2.1.5 request_cfg_byteen_enable_data_width

byteen is only supported when data_width is a multiple of 8.

Protocol hierarchy Request
Critical parameters byteen, data_width
Reference “Simple Extensions” on page 16

2.1.6 request_cfg_sthreadbusy_exact_enable_sthreadBusy

sthreadbusy can only be enabled if sthreadbusy_exact is enabled.

Protocol hierarchy Request
Critical parameters sthreadbusy, sthreadbusy_exact
Reference Table 22 on page 57

2.1.7 request_cfg_sdata_enable_resp

sdata can only be enabled if resp is enabled.

Protocol hierarchy Request
Critical parameters sdata, resp
Reference Table 22 on page 57

2.1.8 request_cfg_sthreadbusy_enable_sthreadbusy_exact_cmdaccept

sthreadbusy can only be enabled if either sthreadbusy_exact or cmdaccept is enabled.

Protocol hierarchy Request
Critical parameters sthreadbusy, sthreadbusy_exact, cmdaccept
Reference Table 22 on page 57

2.1.9 request_cfg_atomiclength_enable_burstlength

atomiclength can only be enabled if burstlength is enabled.

Protocol hierarchy Request
Critical parameters atomiclength, burstlength
Reference Footnotes on page 32

2.1.10 request_cfg_burstprecise_enable_burstlength

burstprecise can only be enabled if burstlength is enabled.

Protocol hierarchy Request
Critical parameters burstprecise, burstlength
Reference Footnotes on page 32

2.1.11 request_cfg_burstseq_enable_burstlength

burstseq can only be enabled if burstlength is enabled.

Protocol hierarchy Request
Critical parameters burstseq, burstlength
Reference Footnotes on page 32

2.1.12 request_cfg_burstsinglereq_enable_burstlength

burstsinglereq can only be enabled if burstlength is enabled.

Protocol hierarchy Request
Critical parameters burstsinglereq, burstlength
Reference Footnotes on page 32

2.1.13 request_cfg_reqlast_enable_burstlength

reqlast can only be enabled if burstlength is enabled.

Protocol hierarchy Request
Critical parameters reqlast, burstlength
Reference Footnotes on page 32

2.1.14 request_cfg_reqrowlast_enable_burstlength

reqrowlast can only be enabled if burstlength is also enabled.

Protocol hierarchy Request
Critical parameters reqrowlast, burstlength
Reference Footnotes on page 32

2.1.15 request_cfg_reqrowlast_enable_reqlast_burstseq_blk_enable

reqrowlast can only be enabled if reqlast and burstseq_blk_enable are enabled.

Protocol hierarchy Request

Critical parameters reqlast, reqrowlast, burstseq_blk_enable

Reference Footnotes on page 32

2.1.16 request_cfg_burstlength_enable_burstseq_enable

burstlength can only be enabled if at least one of the burst sequences is enabled.

Protocol hierarchy Request

Critical parameters burstlength

Reference “Burst Extensions” on page 19

2.1.17 request_cfg_burstseq_enable_burstseq_enable

burstseq can only be enabled if at least one burst sequence is enabled.
burstseq must be enabled if 2 or more burst sequences are enabled.

Protocol hierarchy Request

Critical parameters burstseq, burst_<type>_enable

Reference “Optional Burst Sequences” on page 55

2.1.18 request_cfg_reqdata_together_enable_burstsinglereq

reqdata_together can only be enabled if burstsinglereq is enabled.

Protocol hierarchy Request

Critical parameters reqdata_together, burstsinglereq

Reference “Request and Data Together” on page 59

2.1.19 request_cfg_force_aligned_enable_data_width

force_aligned can only be enabled if data_width is a power of 2

Protocol hierarchy Request

Critical parameters force_aligned, data_width

Reference “Byte Enable Patterns” on page 56

2.1.20 request_cfg_mdainfo_enable_mdata

mdainfo can only be enabled if mdata is enabled.

Protocol hierarchy Request
Critical parameters mdainfo, mdata
Reference MDatInfo on page 18

2.1.21 request_cfg_sdainfo_enable_sdata

sdainfo can only be enabled if sdata is enabled.

Protocol hierarchy Request
Critical parameters sdainfo, sdata
Reference MDatInfo on page 19

2.1.22 request_cfg_atomiclength_width_depends_burstlengthwidth

atomiclength_width must be less than or equal to burstlength_width.

Protocol hierarchy Request
Critical parameters atomiclength_width, burstlength_width
Reference Footnotes on page 32

2.1.23 request_cfg_value_burstlength_width_0x1

The burstlength_width must be 0 if burstlength is disabled.
burstlength_width must be greater than 1 if burstlength is enabled.

Protocol hierarchy Request
Critical parameters burstlength_width
Reference Footnotes on page 32

2.1.24 request_cfg_burst_aligned_enable_burstlength

burst_aligned can only be enabled if burstlength is enabled.

Protocol hierarchy Request
Critical parameters burst_aligned, burstlength
Reference “Burst Alignment” on page 56

2.1.25 request_cfg_burst_aligned_enable_burstseq_enable

burst_aligned can only be enabled if burst_seq_incr_enable is enabled or burstseq is disabled.

Protocol hierarchy Request

Critical parameters burst_seq_incr_enable, burst_aligned, burstseq

Reference “Burst Alignment” on page 56

2.1.26 request_cfg_burstprecise_enable_burstseq_enable

If the only enabled burst sequences are WRAP, XOR, or BLCK, burstprecise must be disabled. If other burst sequences are enabled, and one of WRAP, XOR, or BLCK is enabled, then burstprecise must be enabled.

Protocol hierarchy Request

Critical parameters burstprecise, burst_wrap_enable, burst_xor_enable, burst_blk_enable

Reference Footnotes on page 32

2.1.27 request_cfg_burstseq_enable_addr

burstseq can only be enabled if addr is enabled.

Protocol hierarchy Request

Critical parameters addr, burstseq

Reference MBurstSeq on page 21

2.1.28 request_cfg_burstsinglereq_enable_burstprecise

burstsinglereq can only be enabled if burstprecise is enabled.

Protocol hierarchy Request

Critical parameters burstsinglereq, burstprecise

Reference “Burst Length, Precise and Imprecise Bursts” on page 50

2.1.29 request_cfg_burstsinglereq_enable_burstseq_enable

burstsinglereq must be disabled if burstseq_unkn_enable is the only enabled burst sequence.

Protocol hierarchy Request

Critical parameters burstseq_unkn_enable, burstsinglereq

Reference “Burst Address Sequences” on page 49

2.1.30 request_cfg_taginorder_enable_tags

taginorder is only enabled if tags > 1.

Protocol hierarchy Request

Critical parameters taginorder, tags

Reference Footnotes on page 32

2.1.31 request_cfg_tag_interleave_size_depends_burstlength_width

tag_interleave_size must be 1 if burstlength_width is 0 and otherwise must be 0 or a power-of-two which is less than or equal to $2^{(burstlength_width-1)}$.

Protocol hierarchy Request

Critical parameters tag_interleave_size, burstlength_width

Reference “Burst Interleaving with Tags” on page 58

2.1.32 request_cfg_<block_signal>_enable_burstseq_blk_enable

MBlockHeight and MBlockStride are only enabled when burstseq_blk_enable is also enabled.

Protocol hierarchy Request

Critical parameters MBlockHeight, MBlockStride, burstseq_blk_enable

Reference Footnotes on page 32

2.1.33 request_cfg_value_blockheight_width_0x1

The `blockheight_width` must be 0 if `blockheight` is disabled.
`blockheight_width` must be greater than 1 if `blockheight` is enabled.

Protocol hierarchy Request

Critical parameters `blockheight`, `blockheight_width`

Reference Footnotes on page 32

2.1.34 request_cfg_<threadbusy_pipelined_cfg>_enable_<threadbusy_exact_cfg>

The parameters `mthreadbusy_pipelined`, `sdatathreadbusy_pipelined`, and `stthreadbusy_pipelined` can be enabled to 1 only when the corresponding `_exact` parameter is enabled.

Protocol hierarchy Request

Critical parameters `mthreadbusy_pipelined`, `sdatathreadbusy_pipelined`, and `stthreadbusy_pipelined`

Reference “Ungrouped Signals” on page 42

2.1.35 request_cfg_reqdata_together_enable_cmdaccept_dataaccept

`reqdata_together` can only be enabled if `cmdaccept` and `dataaccept` match.

Protocol hierarchy Request

Critical parameters `reqdata_together`, `cmdaccept`, `dataaccept`

Reference Implicit

Datahandshake Group**2.2.1 datahandshake_cfg_datalast_enable_burstlength**

`datalast` can only be enabled if `burstlength` is enabled.

Protocol hierarchy Request

Critical parameters `datalast`, `burstlength`

Reference Footnotes on page 32

2.2.2 `request_cfg_burstsinglereq_enable_datahandshake_cmd_enable`

`burstsinglereq` can only be enabled if `datahandshake` is enabled or none of the write command types are enabled.

Protocol hierarchy `Datahandshake`

Critical parameters `burstsinglereq`, `datahandshake`, `write_enable`, `writenonpost_enable`, `rdlwrc_enable`

Reference “Single Request / Multiple Data Bursts (Packets)” on page 51

2.2.3 `datahandshake_cfg_datahandshake_enable_mdata`

`datahandshake` can only be enabled if `mdata` is also enabled.

Protocol hierarchy `Datahandshake`

Critical parameters `datahandshake`, `mdata`

Reference “Basic Signals” on page 14

2.2.4 `datahandshake_cfg_datalast_enable_datahandshake`

`datalast` can only be enabled if `datahandshake` is also enabled.

Protocol hierarchy `Datahandshake`

Critical parameters `datalast`, `datahandshake`

Reference “Burst Extensions” on page 19

2.2.5 `datahandshake_cfg_datarowlast_enable_datahandshake`

`datarowlast` can only be enabled if `datahandshake` is also enabled.

Protocol hierarchy `Datahandshake`

Critical parameters `datarowlast`, `datahandshake`

Reference “Burst Extensions” on page 19

2.2.6 `datahandshake_cfg_datrowalast_enable_burstlength`

`datarowlast` can only be enabled if `burstlength` is also enabled.

Protocol hierarchy `Datahandshake`

Critical parameters `datalast`, `datahandshake`

Reference “Burst Extensions” on page 19

2.2.7 datahandshake_cfg_datarowlast_enable_datalast_burstseq_blk_enable

datarowlast can only be enabled if datalast and burstseq_blk_enable are enabled.

Protocol hierarchy Datahandshake

Critical parameters datarowlast, datalast, datahandshake, burstseq_blk_enable

Reference “Burst Extensions” on page 19

2.2.8 datahandshake_cfg_dataaccept_enable_datahandshake

dataaccept can only be enabled if datahandshake is also enabled.

Protocol hierarchy Datahandshake

Critical parameters dataaccept, datahandshake

Reference “Basic Signals” on page 14

2.2.9 datahandshake_cfg_sdatathreadbusy_enable_datahandshake

sdatathreadbusy can only be enabled if datahandshake is also enabled.

Protocol hierarchy Datahandshake

Critical parameters sdatathreadbusy, datahandshake

Reference “Thread Extensions” on page 23

2.2.10 datahandshake_cfg_mdatabyteen_enable_datahandshake

mdatabyteen can only be enabled if datahandshake is also enabled.

Protocol hierarchy Datahandshake

Critical parameters mdatabyteen, datahandshake

Reference “Simple Extensions” on page 16

2.2.11 datahandshake_cfg_mdatabyteen_enable_mdata

mdatabyteen can only be enabled if mdata is also enabled.

Protocol hierarchy Datahandshake

Critical parameters mdatabyteen, mdata

Reference “Simple Extensions” on page 16

2.2.12 datahandshake_cfg_mdatabyteen_depends_data_width

mdatabyteen can only be enabled if data_width is a multiple of 8.

Protocol hierarchy Datahandshake

Critical parameters mdatabyteen, data_width

Reference “Simple Extensions” on page 16

2.2.13 datahandshake_cfg_mdainfo_depends_data_width

mdainfo can only be enabled if data_width is a multiple of 8.

Protocol hierarchy Datahandshake

Critical parameters mdainfo, data_width

Reference “Simple Extensions” on page 16

2.2.14 datahandshake_cfg_mdainfo_width_depends_mdainfobyte_width

mdainfo_width must be greater than or equal to mdainfobyte_width * data_width / 8.

Protocol hierarchy Datahandshake

Critical parameters mdainfo_width, mdainfobyte_width, data_width

Reference “Simple Extensions” on page 16

2.2.15 datahandshake_cfg_sdainfo_depends_data_width

sdainfo can only be enabled if data_width is a multiple of 8.

Protocol hierarchy Datahandshake

Critical parameters sdainfo, data_width

Reference “Simple Extensions” on page 16

2.2.16 datahandshake_cfg_sdainfo_width_depends_sdainfobyte_width

sdainfo_width must be greater than or equal to sdainfobyte_width * data_width / 8.

Protocol hierarchy Datahandshake

Critical parameters sdainfo_width, sdainfobyte_width, data_width

Reference “Simple Extensions” on page 16

2.2.17 datahandshake_cfg_sdatathreadbusy_enable_sdatathreadbusy_exact_dataaccept

sdatathreadbusy can only be enabled if either sdatathreadbusy_exact or dataaccept is enabled.

Protocol hierarchy Datahandshake

Critical parameters dataaccept, sdatathreadbusy, sdatathreadbusy_exact

Reference Table 22 on page 57

2.2.18 datahandshake_cfg_sdatathreadbusy_exact_enable_sdatathreadbusy

sdatathreadbusy can only be enabled if sdatathreadbusy_exact is enabled.

Protocol hierarchy Datahandshake

Critical parameters sthreadbusy_exact, sdatathreadbusy

Reference Table 22 on page 57

2.2.19 datahandshake_cfg_dataaccept_enable_sdatathreadbusy_exact

dataaccept can only be enabled if sdatathreadbusy_exact is not enabled.

Protocol hierarchy Datahandshake

Critical parameters dataaccept, sdatathreadbusy

Reference Table 22 on page 57

2.2.20 datahandshake_cfg_reqdata_together_enable_datahandshake

reqdata_together is only enabled if datahandshake is enabled.

Protocol hierarchy Datahandshake

Critical parameters reqdata_together, datahandshake

Reference Table 20 on page 55

Response Group

2.3.1 response_cfg_resplast_enable_burstlength

resplast can only be enabled if burstlength is enabled.

Protocol hierarchy Response

Critical parameters resplast, burstlength

Reference Footnotes on page 32

2.3.2 response_cfg_respaccept_enable_resp

respaccept can only be enabled if resp is also enabled.

Protocol hierarchy Response

Critical parameters respaccept, resp

References “Simple Extensions” on page 16
Footnotes on page 32

2.3.3 response_cfg_resplast_enable_resp

resplast can only be enabled if resp is also enabled.

Protocol hierarchy Response

Critical parameters resplast, resp

References “Burst Extensions” on page 19
Footnotes on page 32

2.3.4 response_cfg_resprowlast_enable_resp

resprowlast can only be enabled if resp is also enabled.

Protocol hierarchy Response

Critical parameters resprowlast, resp

References “Burst Extensions” on page 19
Footnotes on page 32

2.3.5 response_cfg_resprowlast_enable_burstlength

resprowlast can only be enabled if burstlength is also enabled.

Protocol hierarchy	Response
Critical parameters	resprowlast, burstlength
References	“Burst Extensions” on page 19 Footnotes on page 32

2.3.6 response_cfg_resprowlast_enable_resplast_burstseq_blk_enable

resprowlast can only be enabled if resplast and burstseq_blk_enable are enabled.

Protocol hierarchy	Response
Critical parameters	resprowlast, resplast, burstseq_blk_enable
References	“Burst Extensions” on page 19 Footnotes on page 32

2.3.7 response_cfg_respinfo_enable_resp

respinfo can only be enabled if resp is also enabled.

Protocol hierarchy	Response
Critical parameters	respinfo, resp
References	“Burst Extensions” on page 19 Footnotes on page 32

2.3.8 response_cfg_mthreadbusy_enable_resp

mthreadbusy can only be enabled if resp is enabled.

Protocol hierarchy	Response
Critical parameters	mthreadbusy, resp
References	“Burst Extensions” on page 19 Footnotes on page 32

2.3.9 response_cfg_sdata_enable_resp

sdata can only be enabled if resp is also enabled.

Protocol hierarchy Response

Critical parameters sdata, resp

References “Simple Extensions” on page 16
Footnotes on page 32

2.3.10 response_cfg_sdatainfo_enable_resp

sdatainfo can only be enabled if resp is also enabled.

Protocol hierarchy Response

Critical parameters sdatainfo, resp

References “Simple Extensions” on page 16
Footnotes on page 32

2.3.11 response_cfg_<cmd_enable>_enable_writeresp_enable

writenonpost_enable and rdllwrc_enable are only enabled if writeresp_enable is enabled.

Protocol hierarchy Response

Critical parameters writenonpost_enable, writeresp_enable, rdllwrc_enable

Reference “Optional Commands” on page 55

2.3.12 response_cfg_<cmd_enable>_enable_resp

read_enable and writeresp_enable are only enabled if resp is enabled.

Protocol hierarchy Response

Critical parameters read_enable, writeresp_enable, resp

References “Basic Signals” on page 14

2.3.13 response_cfg_mthreadbusy_exact_enable_mthreadbusy

mthreadbusy_exact can only be enabled if mthreadbusy is enabled.

Protocol hierarchy Response

Critical parameters mthreadbusy, mthreadbusy_exact

Reference “Flow Control Options” on page 57

2.3.14 response_cfg_respaccept_enable_mthreadbusy_exact

respaccept can only be enabled if mthreadbusy_exact is not enabled.

Protocol hierarchy Response

Critical parameters respaccept, mthreadbusy_exact

Reference Table 22 on page 57

2.3.15 response_cfg_mthreadbusy_enable_mthreadbusy_exact_respaccept

mthreadbusy can only be enabled if exactly one of mthreadbusy_exact and respaccept is enabled.

Protocol hierarchy Response

Critical parameters mthreadbusy, mthreadbusy_exact respaccept

Reference Table 22 on page 57

Sideband Group**2.4.1 sideband_cfg_statusbusy_enable_status**

statusbusy can only be enabled if status is enabled.

Protocol hierarchy Request

Critical parameters statusbusy, status

Reference Footnotes on page 32

2.4.2 sideband_cfg_mreset_sreset

Either mreset or sreset must be enabled.

Protocol hierarchy Sideband

Critical parameters mreset, sreset

Reference “Sideband Signals” on page 25

2.4.3 `sideband_cfg_controlwr_enable_control`

`controlwr` can only be enabled if `control` is enabled.

Protocol hierarchy Sideband

Critical parameters `control`, `controlwr`

Reference Footnotes on page 32

2.4.4 `sideband_cfg_controlbusy_enable_control`

`controlbusy` is enabled but `control` is not enabled.

Protocol hierarchy Sideband

Critical parameters `control`, `controlbusy`

Reference Footnotes on page 32

2.4.5 `sideband_cfg_controlbusy_enable_controlwr`

`controlbusy` is enabled but `controlwr` is not enabled.

Protocol hierarchy Sideband

Critical parameters `controlbusy`, `controlwr`

Reference Footnotes on page 32

2.4.6 `sideband_cfg_statusrd_enable_status`

`statusrd` can only be enabled if `status` is enabled.

Protocol hierarchy Sideband

Critical parameters `status`, `statusrd`

Reference Footnotes on page 32

2.4.7 `sideband_cfg_statusbusy_enable_status`

`statusbusy` can only be enabled if `status` is enabled.

Protocol hierarchy Sideband

Critical parameters `status`, `statusbusy`

Reference Footnotes on page 32

Test Group

2.5.1 test_cfg_jtagreset_enable_jtag_enable

jtagtrst_enable can only be enabled if jtag_enable is also enabled.

Protocol hierarchy Test

Critical parameters jtagtrst_enable, jtag_enable

Reference Footnotes on page 32

Interface Interoperability

The checks contained in this section identify configuration checks for connected devices. These checks are written under the assumption that the configurations accurately reflect the enabled protocol features of the individual devices. They do not reflect exceptions that are noted in the specification and that are acceptable when used in conjunction with tie-offs.

2.6.1 master_slave_cfg_read_enable_match

If the slave has read_enable set to 0, the master must have read_enable set to 0.

Protocol hierarchy Request

Critical parameters read_enable

Reference “OCP Interface Interoperability” on page 60

2.6.2 master_slave_cfg_readex_enable_match

If the slave has readex_enable set to 0, the master must have readex_enable set to 0.

Protocol hierarchy Request

Critical parameters readex_enable

Reference “OCP Interface Interoperability” on page 60

2.6.3 master_slave_cfg_rdlwrc_enable_match

If the slave has rdlwrc_enable set to 0, the master must have rdlwrc_enable set to 0.

Protocol hierarchy Request

Critical parameters rdlwrc_enable

Reference “OCP Interface Interoperability” on page 60

2.6.4 master_slave_cfg_write_enable_match

If the slave has `write_enable` set to 0, the master must have `write_enable` set to 0.

Protocol hierarchy Request

Critical parameters `write_enable`

Reference “OCP Interface Interoperability” on page 60

2.6.5 master_slave_cfg_writenonpost_enable_match

If the slave has `writenonpost_enable` set to 0, the master must have `writenonpost_enable` set to 0.

Protocol hierarchy Request

Critical parameters `writenonpost_enable`

Reference “OCP Interface Interoperability” on page 60

2.6.6 master_slave_cfg_broadcast_enable_match

If the slave has `broadcast_enable` set to 0, the master must have `broadcast_enable` set to 0.

Protocol hierarchy Request

Critical parameters `broadcast_enable`

Reference “OCP Interface Interoperability” on page 60

2.6.7 master_slave_cfg_burstseq_blk_enable_match

If the slave has `burstseq_blk_enable` set to 0, the master must have `burstseq_blk_enable` set to 0.

Protocol hierarchy Request

Critical parameters `burstseq_blk_enable`

Reference “OCP Interface Interoperability” on page 60

2.6.8 master_slave_cfg_burstseq_incr_enable_match

If the slave has `burstseq_incr_enable` set to 0, the master must have `burstseq_incr_enable` set to 0.

Protocol hierarchy Request

Critical parameters `burstseq_incr_enable`

Reference “OCP Interface Interoperability” on page 60

2.6.9 master_slave_cfg_burstseq_strm_enable_match

If the slave has `burstseq_strm_enable` set to 0, the master must have `burstseq_strm_enable` set to 0.

Protocol hierarchy Request

Critical parameters `burstseq_strm_enable`

Reference “OCP Interface Interoperability” on page 60

2.6.10 master_slave_cfg_burstseq_dflt1_enable_match

If the slave has `burstseq_dflt1_enable` set to 0, the master must have `burstseq_dflt1_enable` set to 0.

Protocol hierarchy Request

Critical parameters `burstseq_dflt1_enable`

Reference “OCP Interface Interoperability” on page 60

2.6.11 master_slave_cfg_burstseq_dflt2_enable_match

If the slave has `burstseq_dflt2_enable` set to 0, the master must have `burstseq_dflt2_enable` set to 0.

Protocol hierarchy Request

Critical parameters `burstseq_dflt2_enable`

Reference “OCP Interface Interoperability” on page 60

2.6.12 master_slave_cfg_burstseq_wrap_enable_match

If the slave has `burstseq_wrap_enable` set to 0, the master must have `burstseq_wrap_enable` set to 0.

Protocol hierarchy Request

Critical parameters `burstseq_wrap_enable`

Reference “OCP Interface Interoperability” on page 60

2.6.13 master_slave_cfg_burstseq_xor_enable_match

If the slave has `burstseq_xor_enable` set to 0, the master must have `burstseq_xor_enable` set to 0.

Protocol hierarchy Request

Critical parameters `burstseq_xor_enable`

Reference “OCP Interface Interoperability” on page 60

2.6.14 master_slave_cfg_burstseq_unkn_enable_match

If the slave has `burstseq_unkn_enable` set to 0, the master must have `burstseq_unkn_enable` set to 0.

Protocol hierarchy Request

Critical parameters `burstseq_unkn_enable`

Reference “OCP Interface Interoperability” on page 60

2.6.15 master_slave_cfg_force_aligned_match

If the slave has `force_aligned`, the master has `force_aligned` or it must limit itself to aligned byte enable patterns.

Protocol hierarchy Request

Critical parameters `force_aligned`

Reference “OCP Interface Interoperability” on page 60

2.6.16 master_slave_cfg_mdatabyteen_match

Configuration of the mdatabyteen parameter is identical between master and slave.

Protocol hierarchy Request

Critical parameters mdatabyteen

Reference “OCP Interface Interoperability” on page 60

2.6.17 master_slave_cfg_burst_aligned_match

If the slave has burst_aligned, the master has burst_aligned or it must limit itself to issue all INCR bursts using burst_aligned rules.

Protocol hierarchy Request

Critical parameters burst_aligned

Reference “OCP Interface Interoperability” on page 60

2.6.18 master_slave_cfg_<sthreadbusy_param>_match

If the interface includes SThreadBusy, the sthreadbusy_exact and sthreadbusy_pipelined parameters are identical between master and slave.

Protocol hierarchy Response

Critical parameters sthreadbusy_exact, sthreadbusy_pipelined

Reference “OCP Interface Interoperability” on page 60

2.6.19 master_slave_cfg_<mthreadbusy_param>_match

If the interface includes MThreadBusy, the mthreadbusy_exact and mthreadbusy_pipelined parameters are identical between master and slave.

Protocol hierarchy Response

Critical parameters mthreadbusy_exact, mthreadbusy_pipelined

Reference “OCP Interface Interoperability” on page 60

2.6.20 master_slave_cfg_<sdathreadbusy_param>_match

If the interface includes SDataThreadBusy, the sdathreadbusy_exact and sdathreadbusy_pipelined parameters are identical between master and slave.

Protocol hierarchy Response

Critical parameters sdathreadbusy_exact, sdathreadbusy_pipelined

Reference “OCP Interface Interoperability” on page 60

2.6.21 master_slave_cfg_tag_interleave_size_match

If tags > 1, the master's tag_interleave_size is smaller than or equal to the slave's tag_interleave_size.

Protocol hierarchy Request

Critical parameters tag_interleave_size, tag_interleave_size

Reference “OCP Interface Interoperability” on page 60

2.6.22 master_slave_cfg_datahandshake_match

Configuration of the datahandshake parameter is identical between master and slave.

Protocol hierarchy Request

Critical parameters datahandshake

Reference “OCP Interface Interoperability” on page 60

2.6.23 master_slave_cfg_writeresp_enable_match

Configuration of the writeresp_enable parameter is identical between master and slave.

Protocol hierarchy Request

Critical parameters writeresp_enable

Reference “OCP Interface Interoperability” on page 60

2.6.24 master_slave_cfg_reqdata_together_match

Configuration of the `reqdata_together` parameter is identical between master and slave.

Protocol hierarchy Request

Critical parameters `reqdata_together`

Reference “OCP Interface Interoperability” on page 60

2.6.25 master_slave_cfg_mreset_match

If the master has `mreset` enabled to 1, the slave has `mreset` enabled to 1.

Protocol hierarchy Sideband

Critical parameters `mreset`

Reference “OCP Interface Interoperability” on page 60

2.6.26 master_slave_cfg_sreset_match

If the slave has `sreset` enabled to 1, the master has `sreset` enabled to 1.

Protocol hierarchy Sideband

Critical parameters `sreset`

Reference “OCP Interface Interoperability” on page 60

2.6.27 master_slave_cfg_tags_match

The master and slave `tags` values must match.

Protocol hierarchy Request

Critical parameters `tags`

Reference “OCP Interface Interoperability” on page 60

16 *Functional Coverage*

The functional coverage approach described in this chapter is bottom-up, meaning the analysis starts at the signal level and goes up to the transaction level. The transfer level has been skipped for reasons highlighted in “Transfer Level” on page 293. Along this path several coverage types are used. The signal level uses toggle, state, and meta coverages, while the transaction level uses cross and meta coverages.

Toggle coverage

Toggle coverage provides baseline information that a system is connected properly, and that higher level coverage or compliance failures are not simply the result of connectivity issues. Toggle coverage answers the question: Did a bit change from a value of 0 to 1 and back from 1 to 0? This type of coverage does not indicate that every value of a multi-bit vector was seen but measures that all the individual bits of a multi-bit vector did toggle. In certain cases, not all bits can toggle. A system that only supports RD commands, (“010”) for example, will only need toggle coverage on MCmd bit1. MCmd bit0 and bit2 will always be 0. Therefore they must be filtered from the MCmd toggle coverage.

State coverage

State coverage applies to signals that are a minimum of two bits wide. In most cases, the states (also commonly referred to as coverage bins) can be easily identified as all possible combinations of the signal. For example, for the SResp signal, the states could be 00 (IDLE), 01 (DVA), 10 (FAIL) and 11 (ERR). If the state space is too large, an intelligent classification of the states must be made. In the case of the MAddr signal for example, a possible choice of the coverage bins could be one bin to cover the lower address range, one bin to cover the upper address range and one bin to cover all other intermediary addresses.

Meta coverage

Meta coverage is collecting second-order coverage data. Possible meta coverage measurements include accept backpressure delays, threadbusy backpressure delays and inter-phase delays. Meta coverage information is

particularly useful to flag excessive latencies (possibly indicating deadlocks) and to evaluate the OCP backpressure mechanisms (accept / threadbusy).

Cross coverage

Cross coverage measures the activity of one or multiple categories. A category is defined at the transaction level that typically groups multiple OCP signals to form a more abstract, higher-level view of a particular aspect of the OCP protocol. The most pertinent category example is the `transTypes`. This category combines the `MCmd`, `MBurstLength` and `MBurstSingleReq` signals into a higher-level category. Cross coverage on one category, for example the `transTypes` category, indicates which kind of transactions were applied to the system under test (for instance, `MRMD-RD-4`, `SINGLE-WRNP`, etc.). Cross coverage on multiple categories, for example the `transTypes` and `transResults` categories, not only provides information about the transactions applied to the system, but also on their results. In essence, cross coverage measures the types of transactions passing through a system.

Signal Level

Table 48 summarizes the OCP functional coverage approach for the signal level. The table maps all OCP signals (non-sideband) into phase groups (req / datahs / resp) and provides coverage information in the two outermost right columns. Each coverage type field is colored either in green or in yellow. Green fields are mandatory for functional coverage. Yellow fields are optional for functional coverage.

Level 1 - Baseline Coverage

Level 1 - baseline column establishes a solid baseline for the signal level functional coverage, so it contains only mandatory coverage. The coverage type is toggle coverage. Toggle coverage provides a minimum level of confidence to the verification engineer that the device under test is alive and properly connected to the rest of the system. It proves as well that no OCP signals are stuck at 0 or 1. In some cases, filters should be applied to the toggle coverage to exclude coverage of bits that can never toggle (refer to the `MCmd` example on page 289).

Level 2 Coverage

The level 2 coverage type column defines additional coverage. Possible coverage types are state or meta. State coverage defines states (bins) for a multi-bit vector to provide a higher level of abstraction. Meta coverage covers accept / threadbusy backpressure delays.

Mandatory fields must be covered and are with their respective coverage type:

<code>MAddr/MCmd/SResp</code>	:state coverage
<code>MAddrSpace/MByteEn/MDataByteEn</code>	:state coverage
<code>MAtomicLength/MBurstLength/MBurstSeq</code>	:state coverage
<code>MBurstHeight/MBlockStride</code>	:state coverage
<code>MTagID/MDataTagID/STagID</code>	:state coverage
<code>MThreadID/MDataThreadID/SthreadID</code>	:state coverage

Optional column level 2 fields and their coverage types are:

MData/SData :state coverage
 MDataValid :meta coverage
 SCmdAccept/SDataAccept/MRespAccept :meta coverage
 MReqInfo/MDataInfo/SDataInfo/SRespInfo :state coverage
 MConnID :state coverage
 SThreadBusy/SDataThreadBusy/MThreadBusy :meta coverage

Signals that are only one bit wide only have toggle coverage:

MBurstPrecise/MBurstSingleReq
 MReqLast/MDataLast/SRespLast
 MTagInOrder/STagInOrder
 MReqRowLast/MDataRowLast/SRespRowLast

Table 48 Signal Level Functional Coverage

		Phase Groups			Coverage Type	
Signal Group	Signal	Req	Datahs	Resp	Level 1 Baseline	Level 2
Basic	MAddr	X			Toggle	State
	MCmd	X			Toggle	State
	MData	X			Toggle	State
	MDataValid		X		Toggle	Meta
	MRespAccept			X	Toggle	Meta
	SCmdAccept	X			Toggle	Meta
	SData			X	Toggle	State
	SDataAccept		X		Toggle	Meta
	SResp			X	Toggle	State
Simple	MAddrSpace	X			Toggle	State
	MByteEn	X			Toggle	State
	MDataByteEn		X		Toggle	State
	MDataInfo		X		Toggle	State
	MReqInfo	X			Toggle	State
	SDataInfo			X	Toggle	State
	SRespInfo			X	Toggle	State
Burst	MAatomicLength	X			Toggle	State
	MBlockStride	X			Toggle	State
	MBurstHeight	X			Toggle	State
	MBurstLength	X			Toggle	State
	MBurstPrecise	X			Toggle	

Signal Group	Signal	Phase Groups			Coverage Type	
		Req	Datahs	Resp	Level 1 Baseline	Level 2
	MBurstSeq	X			Toggle	State
	MBurstSingleReq	X			Toggle	
	MDataLast		X		Toggle	
	MDataRowLast		X		Toggle	
	MReqLast	X			Toggle	
	MReqRowLast	X			Toggle	
	SRespLast			X	Toggle	
	SRespRowLast			X	Toggle	
Tag	MDataTagID		X		Toggle	State
	MTagID	X			Toggle	State
	MTagInOrder	X			Toggle	
	STagID			X	Toggle	State
	STagInOrder			X	Toggle	
Thread	MConnID	X			Toggle	State
	MDataThreadID		X		Toggle	State
	MThreadBusy			X	Toggle	Meta
	MThreadID	X			Toggle	State
	SDataThreadBusy		X		Toggle	Meta
	SThreadBusy	X			Toggle	Meta
	SThreadID			X	Toggle	State

Table 49 outlines options for signal level meta coverage. For each phase group (req/datahs/resp), two meta coverage types are identified: accept backpressure delay and threadbusy backpressure delay. Other meta coverage types could be identified.

Table 49 Signal Level Meta Coverage Examples

Phase Group	Coverage Types	Details
Request phase	Accept backpressure delay	MCmd – SCmdAccept delay
	Thread busy backpressure delay	SThreadBusy backpressure delay (per bit)

Phase Group	Coverage Types	Details
Datahs phase	Accept backpressure delay	MDataValid – SDataAccept delay
	Thread busy backpressure delay	SDataThreadBusy backpressure delay (per bit)
Response phase	Accept backpressure delay	SResp – MRespAccept delay
	Thread busy backpressure delay	MThreadBusy backpressure delay (per bit)

Transfer Level

The transfer level for functional coverage is being skipped. The underlying reasons are:

- The most obvious order for the OCP functional coverage definition is to follow the OCP hierarchy: signal, phase, transfer, transaction. However, such reasoning does not work well for SRMD bursts. SRMD bursts can be constructed as 1 req + n datahs + 1 resp. As such, the transfer concept does not apply 100% because the number of phases per transfer is not constant. Since it is desirable to have a uniform functional coverage definition, which applies to all OCP transactions (MRMD, SRMD, or SINGLE), it makes sense to skip the transfer level.
- Even if the transfer level was included, there are no valuable coverage points. The combination of phases into transfers is a pure protocol check related matter. Meta coverage to measure inter-phase delays may be useful and is discussed at the transaction level.

Transaction Level

transTypes Concept

Before discussing coverage at the transaction level, clarification is required concerning the process of getting from the signal level to the transaction level. In essence, signals are combined into phases that are then combined into transactions. The unique transaction types represented in this table are referred to as transTypes. Table 50 below summarizes this process and is based on the phases in a transfer described in “Phases in a Transfer” on page 40.

Table 50 *transTypes*

MBurstSingleReq	MCmd	Phases		Enabling Condition		
		Req	Datahs	Resp	Datahandshake	Writeresp_enable
0 MRMD SINGLE transfer	RD/RDEX/RDL	H*L		H*L		
	WR/BCST	H*L			0	0
	WR/BCST/ WRNP/WRC	H*L		H*L	0	1
	WR/BCST	H*L	H*L		1	0
	WR/BCST/ WRNP/WRC	H*L	H*L	H*L	1	1
1 SRMD SINGLE transfer	RD	1		H*L		
	WR/BCST	1	H*L		1	0
	WR/BCST/ WRNP	1	H*L	1	1	1

Notes to Table 50:

1. The table shows how phases are combined into SINGLE transfers, MRMD bursts and SRMD bursts. SINGLE transfers can be de-generated from either MRMD or SRMD bursts.
2. L stands for the MBurstLength (one in the case of a SINGLE transfer) and H stands for MBurstHeight.
3. transTypes are controlled by the signals MBurstSingleReq, MCmd and MBurstHeight (H), MBurstLength (L) and the parameters datahandshake and writeresp_enable
4. RDEX, RDL and WRC commands only apply to SINGLE transfers.
5. RD, RDEX and RDL are not controlled by the datahandshake and writeresp_enable parameters.
6. The possible transTypes are:
 - SINGLE transfers RD, WR, BCST, WRNP, RDEX, RDL, and WRC
 - MRMD bursts RD, WR, BCST, and WRNP
 - SRMD bursts RD, WR, BCST, and WRNP

The following example illustrates this.

If MBurstSingleReq supports values 0 and 1 and
If MCmd only supports RD,WR,WRNP,RDEX and
If datahandshake == 1 and writeresp_enable == 1

then the transTypes will be:

- a) SINGLE transfers, RD/WR/WRNP/RDEX
- b) MRMD bursts, RD/WR/WRNP
- c) SRMD bursts, RD/WR/WRNP

Category Concept

A category groups one or more OCP signals and serves as a building block for cross coverage. A category also represents a higher level view of the OCP protocol, allowing intelligent crosses to be made of one or more categories. Table 51 lists and describes the proposed categories:

Table 51 Categories

Name	Description
transTypes	Category containing the transaction types based on Table 50
transTargets	Category containing the transaction targets (MAddr / MAddrSpace)
transResults	Category containing the transaction results (SResp)
transBurstProps	Category containing the burst properties (AtomicLength / MBurstPrecise / MBurstSeq)
transByteens	Category containing the transaction byte enables (MByteEn / MDataByteEn)
flowThreads	Category containing the flows as function of the ThreadID
flowTagTypes	Category containing the flows as function of the TagInOrder / TagID

Notes to Table 51:

1. The transBurstProps category may be split into multiple categories enabling a higher granularity for cross coverage.
2. The flowTagTypes category combines the TagInOrder and TagID signals. The tag type is encoded as follows:
 - If TagInOrder, then tag type == tag-in-order (1 enumerated value)
 - If not, then tag type == the TagID (multiple enumerated values)

If the TagID range is too large, sub-ranges should be defined. As such, the tag type will have 1 + x enumerated values.

Mapping Signals into Categories

Table 52 shows the OCP signals (non-sideband) mapped into the categories described in the previous section. Only signals that are not optional in the level 2 column of Table 48, and are not MReqLast, MDataLast, or SRespLast, are mapped.

Table 52 Signal Mapping into Categories

Signal	Categories						
	transTypes	transTargets	transResults	transBurstProps	transByteens	Flow Threads	FlowTagTypes
MAddr		X					
MAddrSpace		X					
MAtomicLength				X			
MBurstHeight	X						
MBurstLength	X						
MBurstPrecise				X			
MBurstSeq				X			
MBurstSingleReq	X						
MByteEn					X		
MCmd	X						
MDataByteEn					X		
MDataTagID							X
MDataThreadID						X	
MTagID							X
MTagInOrder							X
MThreadID						X	
SResp			X				
STagID							X
STagInOrder							X
SThreadID						X	

Cross Coverage of One Category

Cross coverage can be applied to just one category. Since this kind of cross coverage only makes sense if a category contains more than one signal, the transResults and transByteens categories are excluded from this type of cross coverage.

Cross coverage of one category can be useful in measuring what kind of transTypes flowed through a design regardless of the signals contained in other categories (for example, the transResults). A more useful coverage result from applying crosses among several categories. Cross coverage of one category is considered optional while cross coverage on multiple categories is considered mandatory.

Cross Coverage on Multiple Categories

Cross coverage can also be applied to combinations of categories. Theoretically, many crosses are possible (128 in total), but only some will make sense for a specific OCP interface configuration and design architecture.

The crosses between the transTypes category and other categories are considered mandatory and establish a solid base for cross coverage at the transaction level. Table 53 shows some of the mandatory crosses that form a sub-set of the theoretical possibilities. It is up to the user to declare additional crosses that exclude the transTypes category, but are important for the system under test. Such crosses are considered optional.

Table 53 Mandatory Crosses of Multiple Categories (Including transType)

Categories						Cross Description
transTypes	transTargets	transResults	transBurstProps	transByteens	Flow Threads	
X	X					Cross all transaction types with all targets
X					X	Cross all transaction types for all threads
X		X				Cross all transaction types with all transaction results
X			X	X		Cross all transaction types for all bursts with all byte enable patterns
X		X	X			Cross all transaction types for all bursts with the transaction results

Meta Coverage

Table 54 outlines possibilities for the transaction level meta coverage. Three meta coverage types are identified: accept backpressure delays relative to the position in a transaction, threadbusy backpressure delays relative to the position in a transaction and several inter-phase delays. Other interesting meta coverage types could be identified.

Table 54 Transaction Level Meta Coverage Examples

Meta Coverage Types	Coverage Details
Inter-phase delays	req to req / datahs to datahs / resp to resp delays
	req to datahs / req to resp / datahs to resp delays
	First req accepted delay / last req accepted delay for MRMD bursts
Accept backpressure delays relative to the position in the transaction	Measure when accept backpressure occurs in a transaction
Threadbusy backpressure delays relative to the position in the transaction	Measure when threadbusy backpressure occurs in a transaction

Sideband Signals Coverage

Toggle coverage

Toggle coverage must be applied to each individual bit of the sideband signals to establish a solid coverage baseline.

State coverage

Sideband signals that consist of multiple bits can have state coverage similar to the dataflow signals.

Meta coverage

Meta coverage can be added for the control and status signals. Some examples of meta coverage that might be added for the control signals handshake are:

- The delay between two ControlWr signal assertions.
- The length of the ControlBusy signal assertion.
- The ControlBusy assertion relative to the previous ControlWr assertion.

Some examples of meta coverage that might be added for the status signals handshake are:

- The delay between two StatusRd signal assertions.
- The length of the StatusBusy signal assertion.

Naming Conventions

This section describes the naming conventions for functional coverage.

Signal Level (Dataflow Signals)

Naming template:

signal_<coverage type>_<signal name | meta name>_<bin>

In which:

<code><coverage type></code>	: toggle state meta
<code><signal name></code>	: OCP signal
<code><meta name></code>	: SCmdAcceptDelay SDataAcceptDelay MRespAcceptDelay SThreadBusyDelay SDataThreadBusyDelay MThreadBusyDelay
<code><bin></code>	: if enumerated types are defined in OCP use them for example: SResp in [ERR,DVA,FAIL,IDLE] else be free to choose a clear name

Examples:

```
signal_toggle_MAddr_bit0_0to1
signal_state_MByteEn_allOnes
signal_meta_SThreadBusyDelay_2
```

Transaction Level (Dataflow Signals)

Naming template:

trans_<coverage type>_<cross name | meta name>[_<bin>]

In which:

<code><coverage type></code>	: cross meta
<code><cross name></code>	: for cross coverage of 1 category: transTypes transTargets transBurstProps flowThreads flowTagTypes for cross coverage of multiple categories: trans_<list of ...>_flow_<list of ...> Types Threads Results TagTypes Targets BurstProps Byteens
<code><meta name></code>	: SCmdAcceptDelay SDataAcceptDelay MRespAcceptDelay SThreadBusyDelay SDataThreadBusyDelay MThreadBusyDelay ReqReqPhaseDelay ReqRespPhaseDelay ...
<code>[<bin>]</code>	: The bin naming is optional for cross coverage.

In most cases the bins will automatically be chosen by the verification tool itself. However if the cross includes signals which have specific OCP enumerated values defined (as DVA for SResp), it's advisable to use them.

Examples:

```
trans_cross_transTypes
trans_cross_trans_Results
trans_cross_flow_ThreadsTagTypes
trans_cross_trans_Results_flow_Threads
trans_meta_ReqReqPhaseDelay_4
```

Sideband Signals

Naming template:

sideband_<coverage type>_<signal name | meta name>_<bin>

In which:

<coverage type>	: toggle state meta
<signal name>	: OCP signal
<meta name>	: ControlWrControlWrDelay ControlBusyDuration ControlWrControlBusyDelay StatusRdStatusRdDelay StatusRdDuration
<bin>	: be free to choose a clear name

Examples:

```
sideband_toggle_ControlWr
sideband_state_Control_3
sideband_meta_StatusRdDuration_5
```

Index

A

- access mode information 18
- addr parameter 65
- addr_base statement 83
- addr_size statement 83
- addr_width parameter 15, 65
- address
 - conflict 45
 - match 45
 - region statement 83
 - sequence
 - BLCK 158
 - burst 49, 157
 - DFLT1 157
 - DFLT2 158
 - INCR 157
 - STRM 157
 - UNKN 159
 - user defined 157
 - WRAP 158
 - XOR 158
 - space 9
 - transfer 17
- addrspace parameter 17, 65
- addrspace_width 17
- addrspace_width parameter 65
- arbitration, shared resource 8
- area savings 146, 150
- asynchronous
 - reset assertion 44, 176
- atomicity requirements 51
- atomiclength parameter 20, 65
- atomiclength_width parameter 20, 65
- ATPG vectors 179

B

- basic OCP signals 14
- BCST 15
- bit
 - naming 82
- BLCK address sequence 158
- block
 - last
 - request 22
 - response 22
 - transfer 22
- block data flow profile 188

- bridging profiles 194
- Broadcast command
 - description 8
 - enabling 55
 - transfer effects 46
- broadcast_enable parameter 64
- bundle
 - characteristics 80
 - defining
 - core 79
 - non-OCP interface 69
 - name 79
 - nets 80
 - port mapping 79
 - signals 71
 - statement 70
- bundle statement 70
- burst
 - address sequence 21, 49
 - address sequences 157
 - alignment 56
 - burst_aligned parameter 56
 - command restrictions 48
 - constant fields 51
 - definition 48
 - exclusive OR 49
 - extension 156
 - framing 116
 - imprecise
 - definition 48
 - MBurstLength 51
 - read 118
 - uses 157
 - INCR 56
 - incrementing precise read 121
 - interleaving 58
 - length 21
 - lengths 157
 - null cycles 123
 - packets 48
 - phases
 - datahandshake 52
 - precise
 - definition 48
 - MBustLength 50
 - read 116
 - precise write 111
 - request, last 22
 - response, last 22
 - sequences 55
 - signals 19
 - single request 21
 - single request/multiple data

- conditions 159
- definition 48
- read 124
- write 127
- state machine 152
- support
 - reporting instructions 204, 205
 - signals 19
- tagged 132
- transfers
 - atomic unit 20
 - total 20
- types 48
- wrapping 120
- write
 - last 21
- burst_aligned parameter 56, 64
- burstlength parameter 21, 65
- burstlength_width parameter 21, 65
- burstprecise parameter 21, 65
- bursts
 - precise
 - uses 157
- burstseq 64
- burstseq parameter 21, 65
- burstseq_dflt1_enable parameter 64
- burstseq_dflt2_enable parameter 64
- burstseq_incr_enable parameter 64
- burstseq_strm_enable parameter 64
- burstseq_unkn_enable parameter 64
- burstseq_wrap_enable parameter 64
- burstseq_xor_enable parameter 64
- burstsinglereq parameter 21, 65
- bus
 - independence 8
 - of signals 71
 - wrapper interface module 2
- byte enable
 - data width conversion 50
 - field 17
 - MByteEn signal 17
 - pattern 52
 - supported patterns 56
 - write 17
- byteen parameter 17, 65

C

- c2qtime
 - port constraints 95
 - timing 88
- c2qtimemin 95
- cacheable storage attributes 18

- capacitance
 - wireload 97
 - wireloaddelay 97
- capacitive load 89
- cell library name 89
- chipparam variable 93
- Clk signal
 - function 14
 - summary 30
- ClkByp signal
 - function 28
 - summary 32
 - test extensions 179
 - timing 45
- clkctrl_enable parameter 28, 67
- clock
 - bypass signal 28
 - control test extensions 179
 - divided
 - enabling 144
 - timing 144
 - gated test 29
 - non-OCP 94
 - portname 95
 - signal 14
 - test 29
- clockName 94
- clockname field 95
- clockperiod variable 92
- cmdaccept parameter 16, 65
- combinational
 - dependencies 37, 183, 184
 - Mealy state machine 150
 - paths 91, 100
 - slave state machine 151
- command
 - encoding 15
 - limiting 15
 - request types 15
- commands
 - basic 8
 - extensions 8
 - mnemonic 15
 - required 60
- concurrency 163
- configurable interfaces 81
- connection
 - description 169
 - identifier
 - definition 169
 - field 24
 - support 204, 205
 - transfer handling 54
 - uses 10

- transfers 54
- connid parameter 24, 65
- connid_width parameter 24, 65
- control
 - event signal 27
 - field 27
 - information
 - specifying 27
 - timing 45
 - parameter 27, 67
 - timing 45
- Control signal
 - function 25
 - summary 31
 - timing 45
- control_width parameter 27, 67
- controlbusy parameter 27, 67
- ControlBusy signal
 - function 25
 - summary 31
 - timing 45
- controlwr parameter 27, 67
- ControlWr signal 45
 - function 25
 - summary 31
- core
 - area 94, 203
 - code 77
 - compliance 3
 - control
 - busy 27
 - event 27
 - information 27
 - documentation 203
 - documentation template 208
 - endianness 58
 - frequency range 203
 - ID 203
 - interconnecting 90
 - interface
 - defining 79
 - endianness 48
 - timing parameters 88
 - interoperability 60
 - name 203
 - power consumption 203
 - process dependent 203
 - revision 77
 - RTL configuration file 75
 - status
 - busy 27
 - event 27
 - information 27
 - synthesis configuration file
 - defining 87
 - tie-off constants 63

- timing 87, 181
- core_id statement 76
- core_name 75
- core_stmt 75

D

- data byte parity 18, 19
- data width
 - conversion 50
 - endianness 174
- data_width parameter 15, 16, 65
- dataaccept parameter 16, 65
- dataflow signals
 - definitions 14
 - naming 14
 - timing 39
- datahandshake
 - extension 126
 - intra-phase output 167
 - parameter 16
 - phase
 - active 40
 - order 41
 - sequence 147
 - signal group 36

- datahandshake parameter 59, 64
- datalast parameter 21, 65
- datarowlast parameter 65
- ddr_space statement 83
- debug and test interface 29, 178
- DFLT1 burst sequence 49, 56
- DFLT2 burst sequence 49, 56
- direction statement 71
- divided clock 144
- driver strength 89
- drivingcellpin parameter
 - timing requirements 89
 - values 95
- DVA 16
- DVA response 16, 45

E

- enableclk parameter 144
- EnableClk signal 144
- endian parameter 58, 64, 174
- endianness
 - attributes 58
 - bit ordering 174
 - concepts 47
 - data width issues 174

- definitions 58
- dynamic interconnect 175
- ERR response 16
- error
 - correction code 18, 19
 - master 26
 - report mechanisms 11
 - signal 25
 - slave 26
- exclusive OR bursts 49
- extended OCP signals 16, 153

F

- FAIL response 16, 45
- false path constraints 91, 100
- falsepath parameter 91, 100
- fanout
 - maximum 96
- FIFO
 - full 19
- flags
 - core-specific 25
 - master 26
 - slave 26
- flow-control 204, 205
- force_aligned parameter 56, 64

H

- H-bus profile 194
- high-frequency design 147
- hold time
 - checking 88
- holdtime
 - description 96
 - timing 88

I

- icon statement 76
- implementation restrictions 204, 205
- imprecise burst 51
- INCR burst sequence 49, 56
- inout ports 97, 98
- input
 - load 89
 - port syntax 97
 - signal timing 88
- instance
 - size 94
- interface

- characteristics 204
- clock control 28
- compatibility 60
- configurable 81
- configuration file 69
- core RTL description 75
- debug and test 29
- endianness 48
- location 82
- multiple 79
- parameters 81
- scan 28
- statement 79
- type statement 80
- types 71
- interface_types statement 71
- interfaceparam variable 93
- internal
 - scan techniques 178
- interoperability rules 60
- interrupt
 - parameter 26
 - processing 10
 - signal 25
 - slave 26
- interrupt parameter 67

J

- jtag_enable parameter 29, 67
- jtagtrst_enable 29
- jtagtrst_enable parameter 67

L

- latency
 - sensitive master 151
- lazy synchronization
 - command sequences 172
 - mechanism 171
- level0 timing 181, 182
- level1 timing 181, 182
- level2 timing 181, 182
- loadcellpin
 - description 96
 - timing 89
- loads parameter
 - description 96
 - timing 89
- location statement 82
- locked synchronization 171
- longest path 95

M

- MAddr signal
 - function 14
 - summary 30
- MAddrSpace signal
 - function 17
 - summary 30
- master
 - error 26
 - flags 26
 - interface documentation 204
 - reset 26, 44
 - response accept 16
 - signal compatibility 60
 - slave interaction 165
 - thread busy 24
- MAtomicLength signal
 - atomicity 51
 - function 20
 - summary 30
- maxdelay parameter
 - description 100
 - timing 91
- maxfanout variable 96
- MBurstLength signal
 - burst lengths 50
 - function 20
 - summary 30
 - values 157
- MBurstPrecise signal
 - function 20
 - summary 30
- MBurstSeq signal
 - address sequences 157
 - encoding 21
 - function 20
 - summary 30
- MBurstSingleReq signal
 - conditions 51
 - function 20
 - summary 30
- MBurstSingleReq signal
 - transfer phases 40
- MByteEn signal
 - function 17
 - summary 30
- MCmd signal
 - function 14
 - summary 30
- MConnID signal
 - function 23
 - summary 31
- mdata parameter 15, 65
- MData signal
 - data valid 16
 - description 14
 - request phase 147
 - summary 30
- mdatabyteen parameter 17, 65
- MDataByteEn signal
 - function 17
 - phases 39
 - summary 30
- mdatainfo parameter 18, 65
- MDataInfo signal
 - function 17
 - summary 30
- mdatainfo_width parameter 18, 66
- mdatainfobyte_width parameter 66
- MDataLast signal
 - function 20
 - phases 52
 - summary 30
- MDataTagID signal
 - flow 53
 - function 22
 - summary 31
- MDataThreadID signal
 - datahandshake 164
 - function 23
 - summary 31
- MDataValid signal
 - datahandshake 147
 - function 14
 - summary 30
 - timing 147
- merror parameter 26, 67
- MError signal
 - summary 31
- mflag parameter 26, 67
- MFlag signal
 - function 25
 - summary 31
- mflag_width parameter 26, 67
- MReqInfo signal
 - function 17
 - summary 30
- MReqLast signal
 - function 20
 - phases 52
 - summary 30
- mreset parameter 26, 67
- MReset_n signal
 - function 25
 - required cycles 44
 - summary 31
 - timing 44

- MRespAccept signal
 - definition 14
 - response phase 146
 - response phase output 166
 - saving area 146
 - summary 30

- MSecure subnet 175

- MTagID signal
 - flow 53
 - function 22
 - summary 31

- MTagInOrder signal 163
 - function 23
 - summary 31

- mthreadbusy parameter 24, 66

- MThreadBusy signal
 - definition 23
 - information 54
 - intra-phase output 166
 - semantics 42
 - summary 31
 - timing cycle 42

- mthreadbusy_exact parameter 42, 57, 64

- mthreadbusy_pipelined parameter 42, 66

- MThreadID signal
 - function 23
 - summary 31

N

- nets
 - bit naming 82
 - characterizing 71
 - redirection 80
 - statement 71

- NULL response 16, 123

O

- out-of-band information 26

- output
 - port syntax 98
 - signal timing 88

P

- packets 48

- param variable 93

- parameter summary 29

- partial word transfer 47

- path
 - longest 95
 - shortest 95

- phase
 - interoperability 62

- intra-phase 165
- options 59
- ordering
 - between transfers 41
 - request 145
 - within transfer 41
- protocol 38
- timing 145
- transfer 40

- physical design parameters 89

- pin
 - level timing 88

- pipeline
 - decoupling request phase 156
 - request/response protocol 146
 - support 204
 - transfer 152
 - without MRespAccept 146
 - write data, slave 16

- pipelined access 188

- point-to-point signals 8

- port
 - constraint variables 95
 - delay 100
 - inout 97, 98
 - input, syntax 97
 - mapping 79
 - module names 81
 - output, syntax 98
 - renaming 80
 - statement 80
 - timing constraints 87

- posted write
 - model 40
 - timing diagram 109

- power
 - consumption estimates 203
 - idling cores 144

- precise burst 50

- precise bursts 157

- prefix command 81

- profiles
 - benefits 185
 - block data flow 188
 - bridging 194
 - H-bus 194
 - register access 192
 - sequential undefined length data flow 190
 - types 187
 - X-bus packet read 198
 - X-bus packet write 196

- programmable register interfaces 192

- proprietary statement 83

- protocol
 - interoperability 60

- phases
 - mapping 39
 - order 41
 - rules 39
- R**
- RDEX 15
- rdlwrc_enable parameter 64
- read
 - burst wrapping 120
 - data field 16
 - imprecise burst 118
 - incrementing precise burst 121
 - information 19
 - non-pipelined timing diagram 113
 - optimizing access 146
 - out-of-order completion 134
 - tagged 130, 132
 - precise burst 116
 - single request/multiple data 124
 - tagged 130, 132
 - threaded 134
 - timing diagram 105
- Read command
 - enabling 55
 - transfer effects 46
- read_enable parameter 64
- ReadEx command
 - burst restrictions 48
 - enabling 55
 - transfer effects 46
- readex_enable parameter 64
- ReadLinked command 172
 - burst restrictions 48
 - enabling 55
 - encoding 15
 - mnemonic 15
 - transfer effects 46
- reference_port statement 80
- register
 - access profile 192
- reqdata_together parameter 59, 64, 127, 159
- reqinfo parameter 18, 66
- reqinfo_width parameter 18, 66
- reqlast parameter 22, 66
- reqrowlast parameter 66
- request
 - delays 108
 - flow-control mechanism 107
 - handshake 107
 - information 18
 - interleaving 51
 - last 52
 - last in a burst 22
- order 23
- phase
 - intra-phase 165
 - order 41
 - outputs 145, 165
 - signal group 39
 - timing 145
 - transfer ordering 41
 - worst-case combinational path 165
- pipelined 114
- signals
 - active 39
 - group 36
- tag identifier 23
- thread identifier 24
- reset
 - asynchronous
 - completed transactions 44
 - asynchronous assertion 176
 - conditions 176
 - domains 177
 - dual signals 177
 - interface compatibility 177
 - master 26
 - phases 44
 - power-on 176
 - signal 25
 - slave 26
 - special requirements 204
 - state 44
 - timing 140
- resistance
 - wireload 97
 - wireloaddelay 97
- resp parameter 16, 66
- respaccept parameter 16, 66
- respinfo parameter 19, 66
- respinfo_width parameter 19, 66
- resplast parameter 22, 66
- response
 - accept 16
 - accept extension 115
 - delays 108
 - encoding 16
 - field 16
 - information 19
 - last in a burst 22
 - mnemonics 16
 - null cycle 123
 - order 23
 - phase
 - active 39
 - intra-phase 166
 - order 41
 - slave 166
 - timing 146

- pipelined 114
- required types 60
- signal group 36
- tag identifier 23
- thread identifier 25

resprowlast parameter 66

revision_code 77

rootclockperiod variable 93

RTL

- proprietary extensions 83

S

scan

- clock 179
- control 179
- data
 - in 28
 - out 28
- interface signals 28
- mode control 28
- Scanctrl signal 28
- Scanin signal 28
- Scanout signal 28
- test environments 179
- test mode 179

Scanctrl signal

- function 28
- summary 32
- uses 179

scanctrl_width parameter 28, 67

Scanin signal

- function 28
- summary 32
- timing 45

Scanout signal

- function 28
- summary 32
- timing 45

scanport parameter 67

scanport_width parameter 28, 67

SCmdAccept signal

- definition 14
- request phase 145
- request phase output 165
- summary 30

sdata parameter 16, 66

SData signal

- function 14
- summary 30

SDataAccept signal

- datahandshake 147
- function 14
- summary 30

sdatainfo parameter 19, 66

SDataInfo signal

- function 17
- summary 30

sdatainfo_width parameter 19, 66

sdatainfobyte_width parameter 66

sdatathreadbusy parameter 24, 66

SDataThreadBusy signal

- function 24
- semantics 42
- summary 31
- timing cycle 42

SDataThreadbusy signal

- information 54

sdatathreadbusy_exact parameter 42, 64

sdatathreadbusy_pipelined parameter 42, 66

security

- level 175, 201
- parameters 175

semaphores 171

sequential undefined length data flow profile 190

serror parameter 26, 67

SError signal

- function 25
- summary 31

setuptime

- description 96
- timing 88

sflag parameter 26, 67

SFlag signal

- function 25
- summary 31

sflag_width parameter 26, 67

shared resource arbitration 8

shortest path 95

sideband signals

- definitions 25
- timing 44, 176

signal

- attribute list 81
- basic OCP 14
- configuration 29
- dataflow 14
- direction 33
- driver strength 89
- extensions
 - simple 16
 - thread 22, 23
- group
 - division 36
 - mapping 39
- interface interoperability 62
- ordering 165
- requirements 60

- sideband 25
- test 27
- tie-off 63, 81
 - rules 62
- tie-off values 29
- timing
 - input 88
 - output 88
 - requirements 38
 - restrictions 165
- ungrouped 42
- width 82
- width mismatch 62
- SInterrupt signal
 - function 25
 - summary 31
- slave
 - combinational paths 166
 - error 26
 - flag
 - description 26
 - interrupt 26
 - optimizing 163
 - pipelined write data 16
 - reset 26, 44
 - response field 16
 - response phase 166
 - signal compatibility 60
 - successful transfer 45
 - thread busy 25
 - transfer accept 16
 - write
 - accept 16
 - thread busy 24
- sreset parameter 26, 67
- SReset_n signal
 - function 25
 - required cycles 44
 - summary 31
 - timing 44
- SResp signal
 - function 14
 - summary 30
- SRespInfo signal
 - function 17
 - summary 30
- SRespLast signal
 - function 20
 - phases 52
 - summary 30
- STagID signal
 - flow 53
 - function 23
 - summary 31
- STagInOrder signal 163
 - function 23
- summary 31
- state machine
 - combinational
 - master 150
 - Mealy 150
 - slave 151
 - diagrams 145
 - multi-threaded behavior 164
 - sequential master 147
 - sequential slave 149
- status
 - busy 27
 - core 27
 - event 27
 - information
 - response 19
 - signals 27
 - parameter 27
 - timing 45
- status parameter 67
- Status signal
 - function 25
 - summary 31
- status_width parameter 27, 67
- statusbusy parameter 27, 67
- StatusBusy signal
 - function 25
 - summary 31
 - timing 45
- statusrd parameter 27, 67
- StatusRd signal
 - function 25
 - summary 31
 - timing 45
- stthreadbusy parameter 25, 66
- SThreadBusy signal
 - function 24
 - information 54
 - semantics 42
 - slave request phase 165
 - summary 31
 - timing cycle 42
- stthreadbusy_exact parameter 42, 57, 64, 136
- stthreadbusy_pipelined parameter 42, 66
- SThreadID signal
 - function 24
 - summary 31
- STRM burst sequence 49, 56
- subnet statement 82
- synchronization
 - deadlock 173
 - lazy 171
 - locked 171
- synchronous

- handshaking signals 3
- interface 8
- system
 - initiator 2
 - target 2

T

- tag
 - burst handling 132
 - burst interleaving 58
 - definition 162
 - flow 53
 - identifier
 - binary-encoded value 23
 - interleaving 53
 - ordering 53
 - read handling 130
 - value 162
- tag_interleave_size parameter 58
- taginorder parameter 23, 66
- tags parameter 66
- TCK signal
 - function 28
 - summary 32
- TDI signal
 - function 28
 - summary 32
- TDO signal
 - function 28
 - summary 32
- test
 - clock 29
 - clock control extensions 178
 - data
 - in 29
 - out 29
 - logic reset 29
 - mode 29
 - signals
 - definitions 27
 - timing 44, 45
- TestClk signal
 - function 28
 - summary 32
 - timing 45
- thread
 - arbitration 137
 - blocking 136
 - busy
 - hint 136
 - master 24
 - signals 54
 - slave 25
 - busy semantics 164
 - busy signals 164

- dependency 164
- description 163
- end-to-end identification 54
- identifier
 - binary-encoded value 24
 - definition 169
 - request 24
 - response 25
 - uses 54
- mapping 54
- multiple
 - concurrent activity 54
 - non-blocking 57
- multi-threaded interface 167
- ordering 10
- signal extensions 23
- state machine implementation 164
- transfer order 163
- valid bits 168

- threads parameter 24, 66
- throughput
 - documenting 204
 - maximum 146
 - peak data 151
- timing
 - categories
 - definitions 181
 - level0 182
 - level1 182
 - level2 182
 - combinational path 38
 - combinational paths 91
 - constraints
 - ports 87
 - core 87
 - connecting 90
 - interface parameters 88
 - dataflow signals 39
 - diagrams 105
 - max delay 91
 - parameters 95
 - pin-level 88
 - pins 89
 - sideband signals 44
 - signals 38, 165
 - test signals 44
- TMS signal
 - function 28
 - summary 32
- transfer
 - accept 16
 - address 15
 - address region 17
 - assigning 54
 - burst
 - linking 48
 - byte enable 17

- command 15
- concurrent activity 54
- connection ID 24
- data widths 9
- effects of commands 45
- efficiency 9, 50
- endianness 48
- order 10, 41
- out-of-order 54
- phases 40
- pipelining 9
- successful 45
- type 15
- TRST_N signal
 - function 28
 - summary 32
- type statement 71

U

- UNKN burst sequence 49, 56

V

- vendor code 77
- version 94
 - statement 70, 76
- VHDL
 - ports 71
 - signals 71
- vhdl_type command 71
- Virtual Socket Interface Alliance 1

W

- width
 - data 9
 - interoperability 62
 - mismatch 62
- wireloadcapacitance
 - description 97
 - See also wireloaddelay
 - timing 89
- wireloaddelay
 - description 97
 - timing 89
- wireloadresistance
 - description 97
 - timing 89
- word
 - corresponding bytes 17
 - packing 50
 - padding 50
 - partial 47
 - power-of-2 15
 - size 9, 15

- stripping 50
- transfer 9, 47
- worstcasedelay 94
- WRAP burst sequence 49, 56
- wrapper interface modules 2
- write
 - byte enables 17
 - completion model 47
 - data
 - burst 21
 - extra information 18
 - master to slave 15
 - slave accept 16
 - tag ID 23
 - thread ID 24
 - valid 16
 - nonpost
 - enabling 59
 - phases 40
 - semantics 47
 - timing diagram 110
 - non-posted
 - semantics 170
 - posted
 - phases 40
 - semantics 47, 170
 - precise burst 111
 - response enabled 109
 - responses 59
 - single request, multiple data 127
 - timing diagram 105
- Write command
 - burst restrictions 48
 - enabling 55
 - transfer effects 46
- write_enable parameter 64
- WriteConditional command 172
 - burst restrictions 48
 - enabling 55
 - encoding 15
 - mnemonic 15
 - transfer effects 46
- WriteNonPost command
 - burst restrictions 48
 - enabling 55
 - encoding 15
 - mnemonic 15
 - posting options 47
 - semantics 8
 - transfer effects 46
- writenonpost_enable parameter 64, 170
- writeresp_enable parameter 59, 64, 170

X

- X-bus packet

read profile 198
write profile 196

XOR address sequence 49
XOR burst sequence 56

OCP International Partnership

3855 SW 153rd Drive
Beaverton, OR 97006

Ph: 503-619-0560

Fax: 503-297-1090

admin@ocpip.org

www.ocpip.org



Part Number: 161-000125-0003